

CSCI 460 Operating Systems

Computer System & Operating System Overview

Professor Travis Peters


Fall 2019

Goals for Today


- Learning Objectives
 - Wrap up high-level computer system & OS concepts (Chapters 1 & 2)
 - Connect some dots and pave the way for the rest of the course
 - *not going deep into history, which was covered (enough) last week...*
 - *gory details in the text, of course...*
- Announcements
 - Note taker?!! ;-)
 - Grades for HW1 coming soon (restructuring things in D2L over the next week or so)
 - More HW soon, I promise ;-)

Questionnaire - *The results are in (mostly...)*


Programming Languages:



Who is in this class:



Sys Programming Skillzzz:



I've heard from software engineers that taking OS is a very helpful class for once you get out into the field

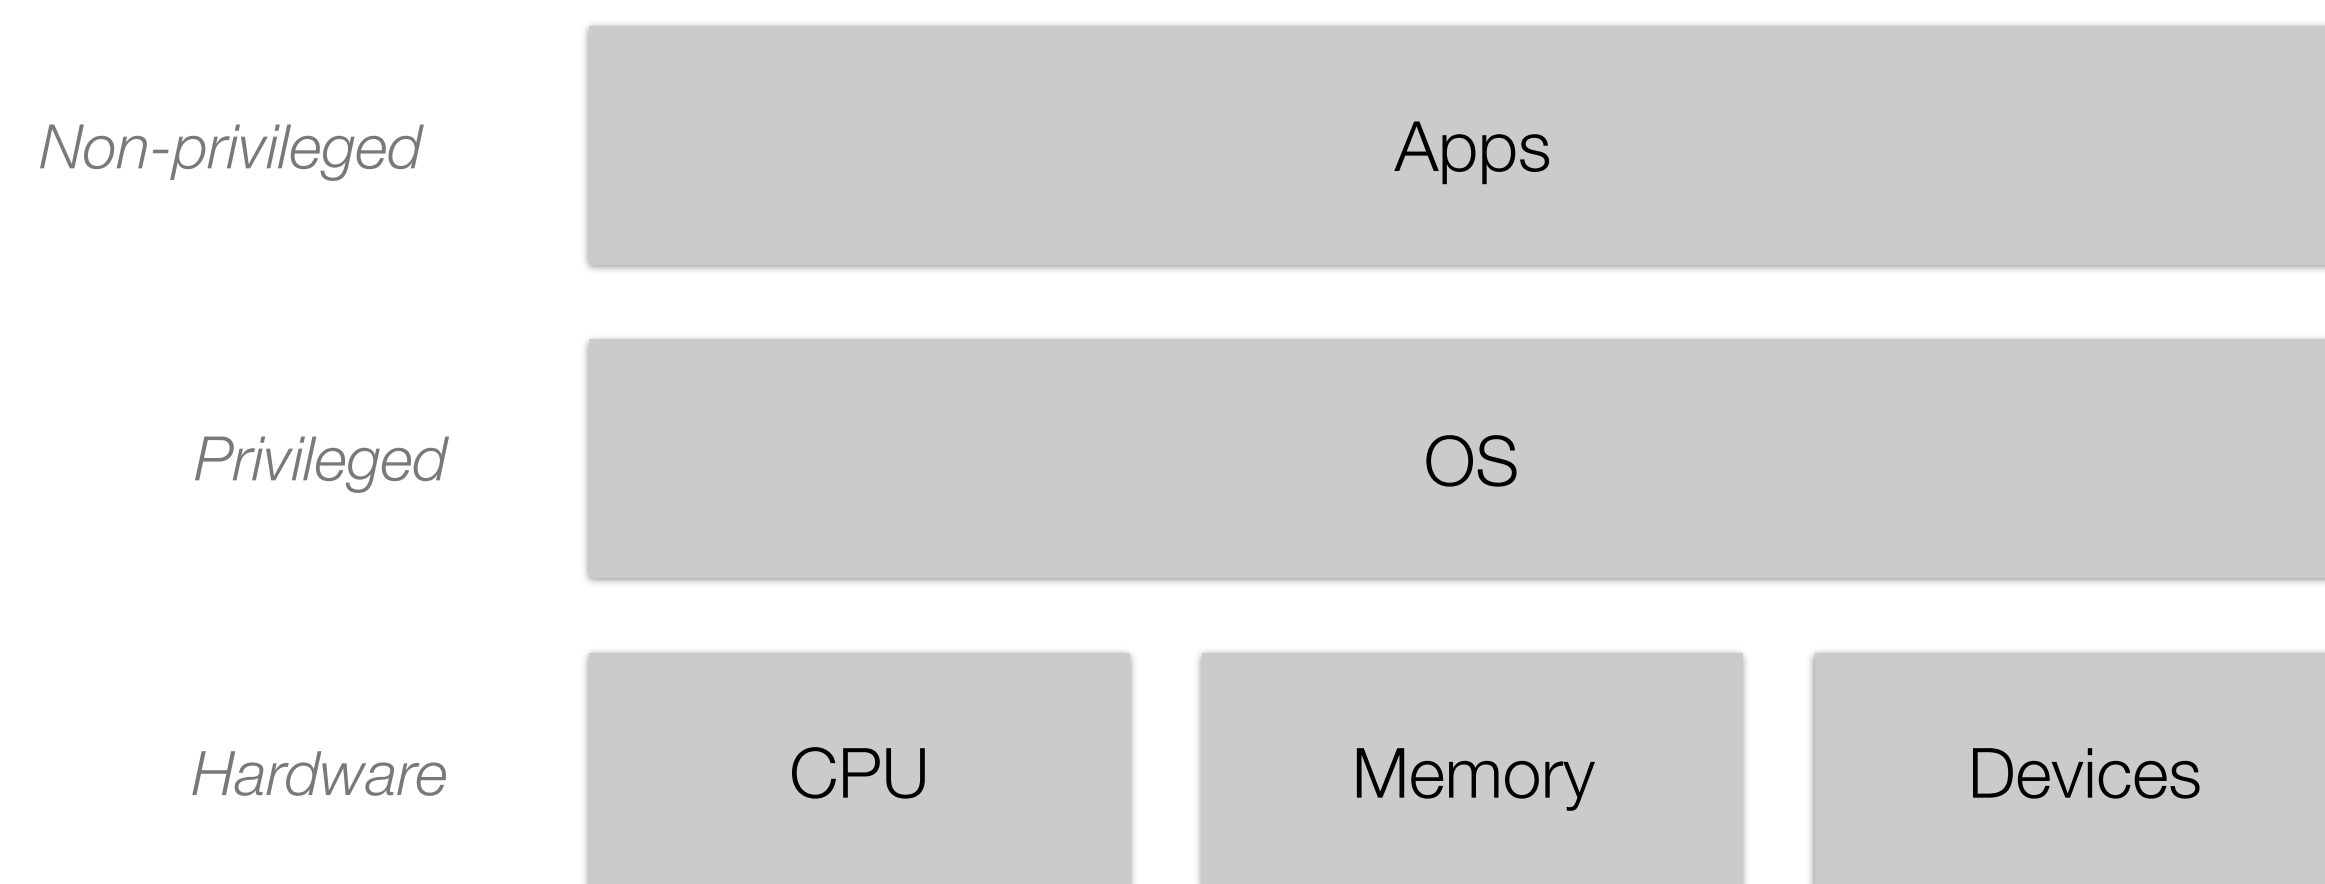
Why do all my profs use Macs?

I did not like the other course offerings


I'm just interested to learn more about the inner workings of operating systems.

Computer System Overview

- Basic Model



Computer System Overview

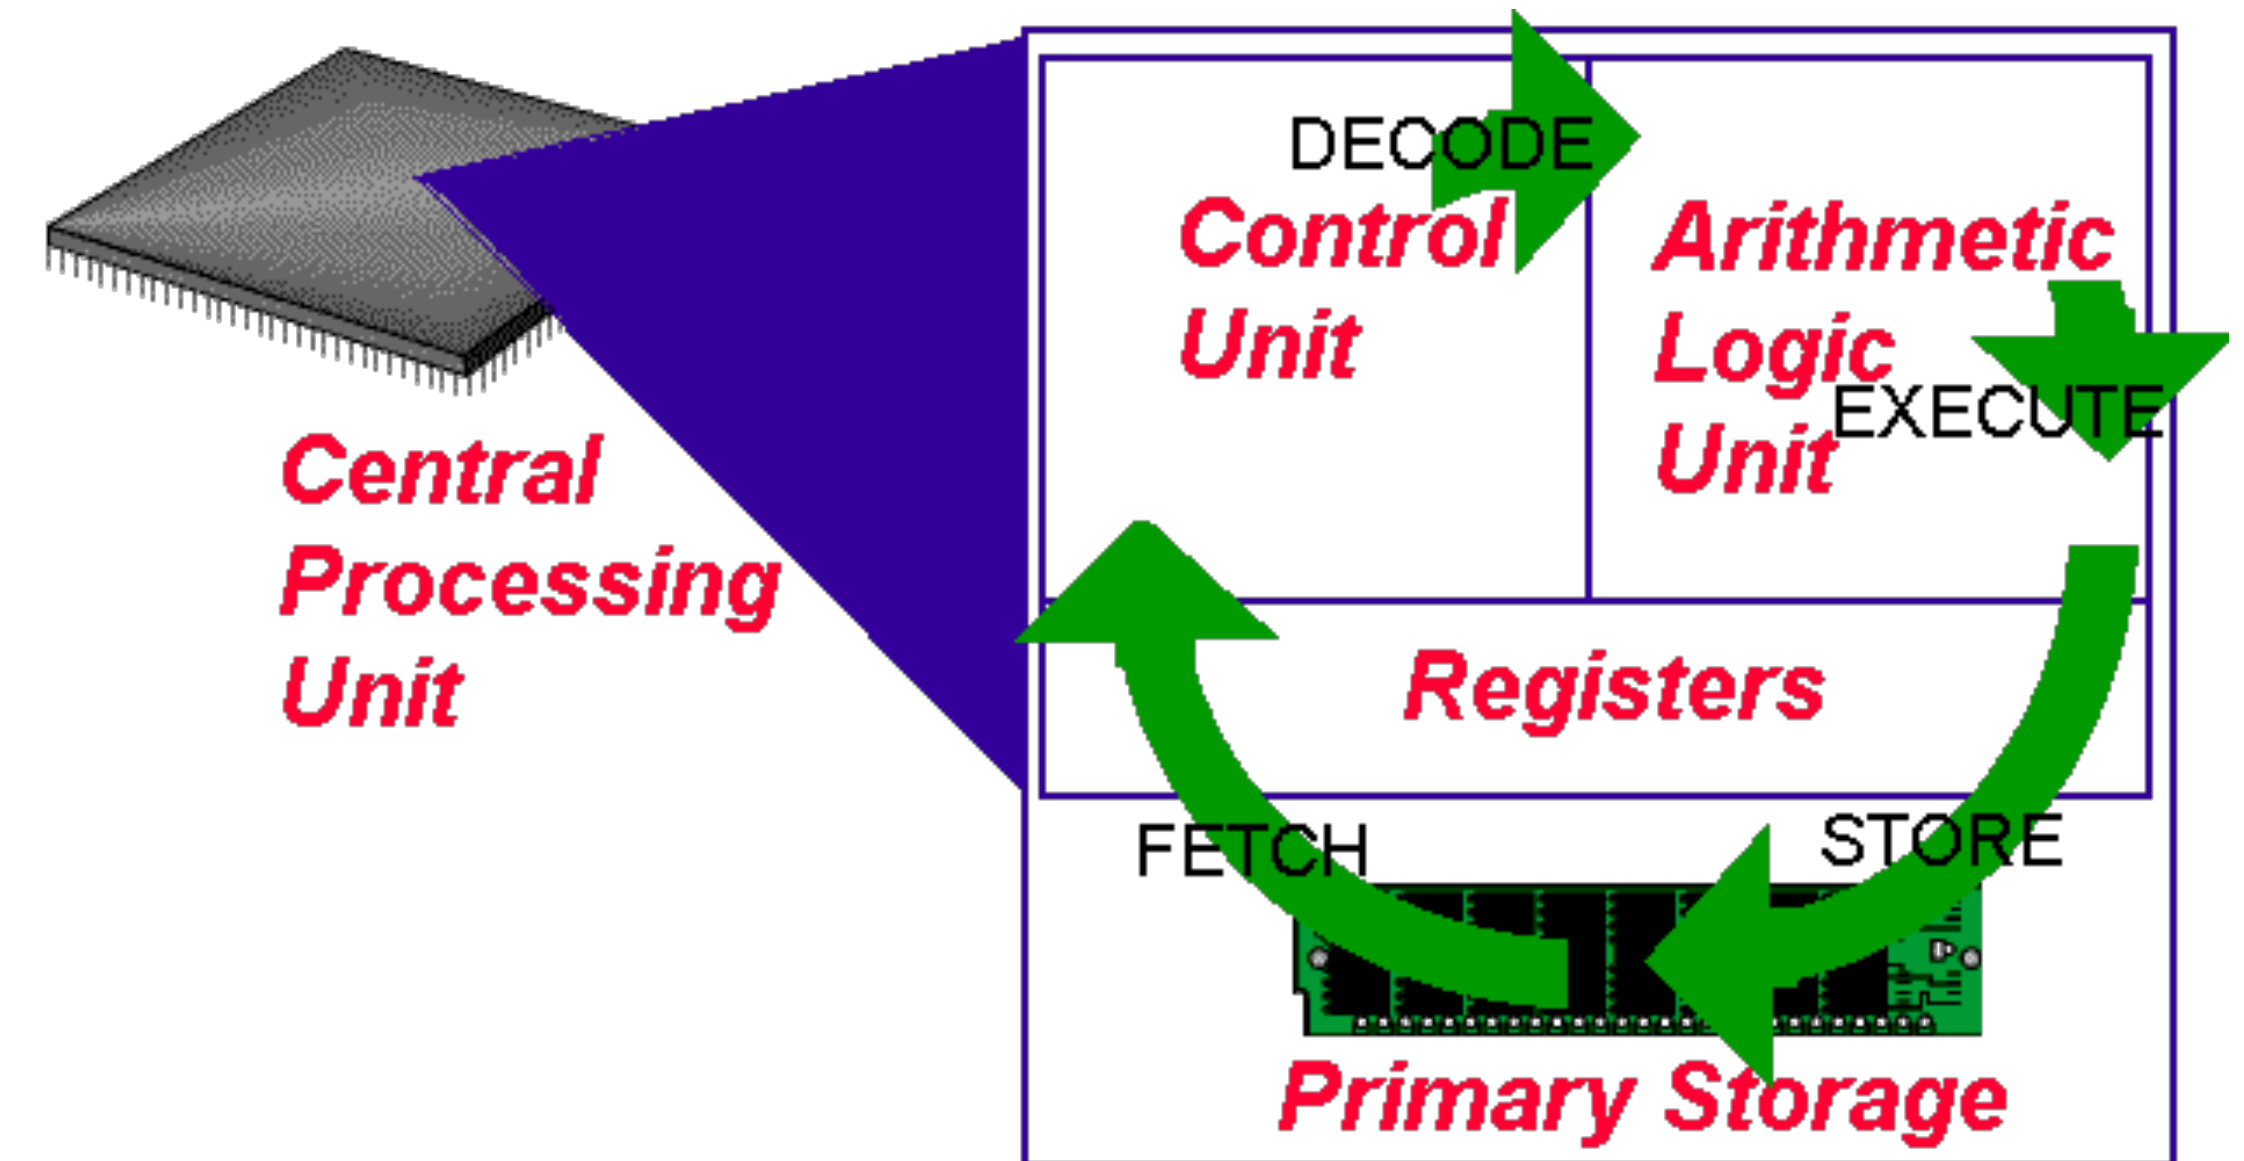
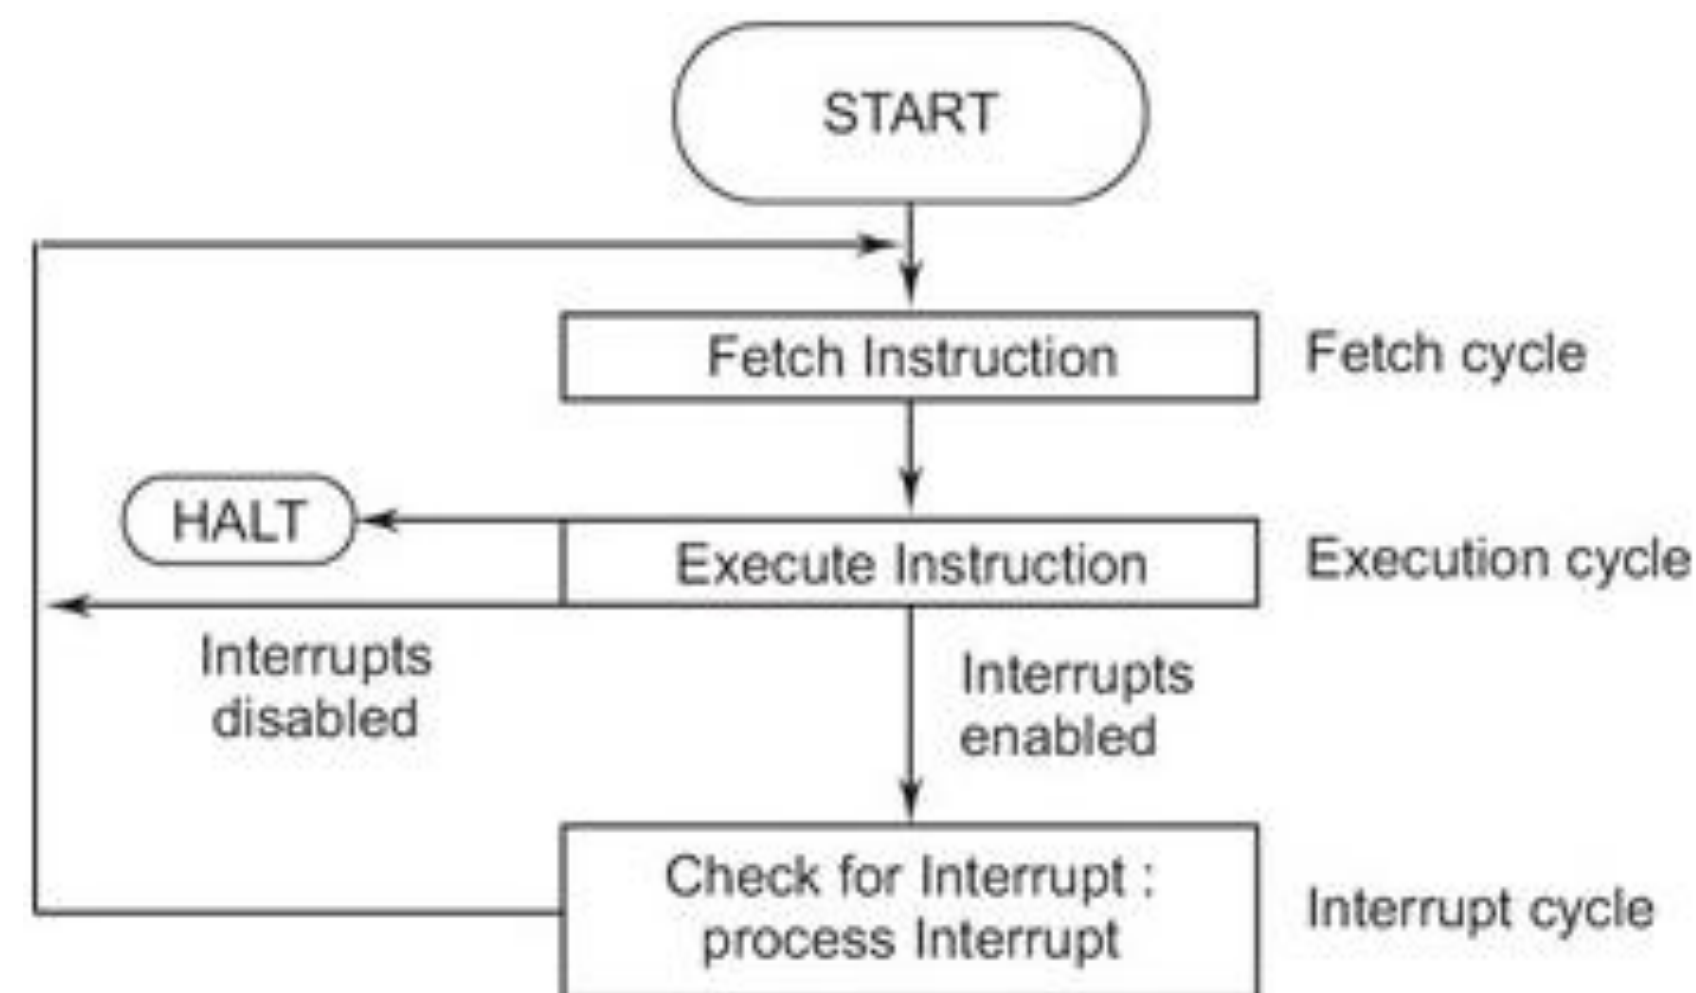


- Instruction Execution / Instruction Cycle


- **FETCH** the next instruction, **DECODE** it, **EXECUTE** it, and **STORE** the result.

- Interrupts (*interrupt the processor!*)

- E.g., program int. (e.g., illegal instruction, out-of-bounds access), timer, I/O, hardware failure
- Run corresponding “handler”



Computer System Overview



- Memory

- The Memory Hierarchy

want **fast access** to a **large amount of memory**

locality of reference: memory access tends to cluster (e.g., loops, subroutines)

- Cache Memory

keep as much as possible, as close as possible...

...but still support large amounts of memory.

Multiple levels of cache... need **cache coherency**.

- Direct Memory Access (DMA)

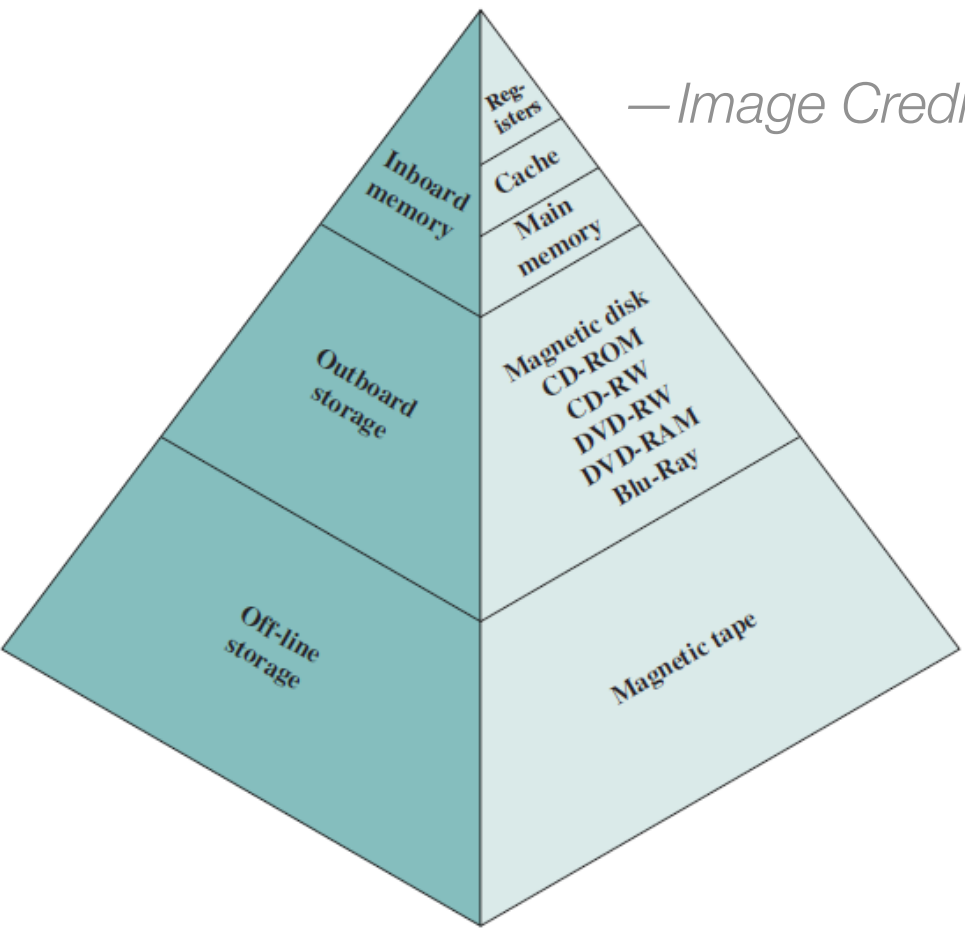
programmed I/O (active; processor polls I/O device)

vs.


interrupt-driven I/O (assist; I/O device does work; interrupt processor to help)

vs.


DMA I/O (delegate; DMA module is given OP, DEV, ADDR, #WORDS; sent INT when complete)




—Image Credit: Stallings




(a) Single cache



—Image Credit: Stallings

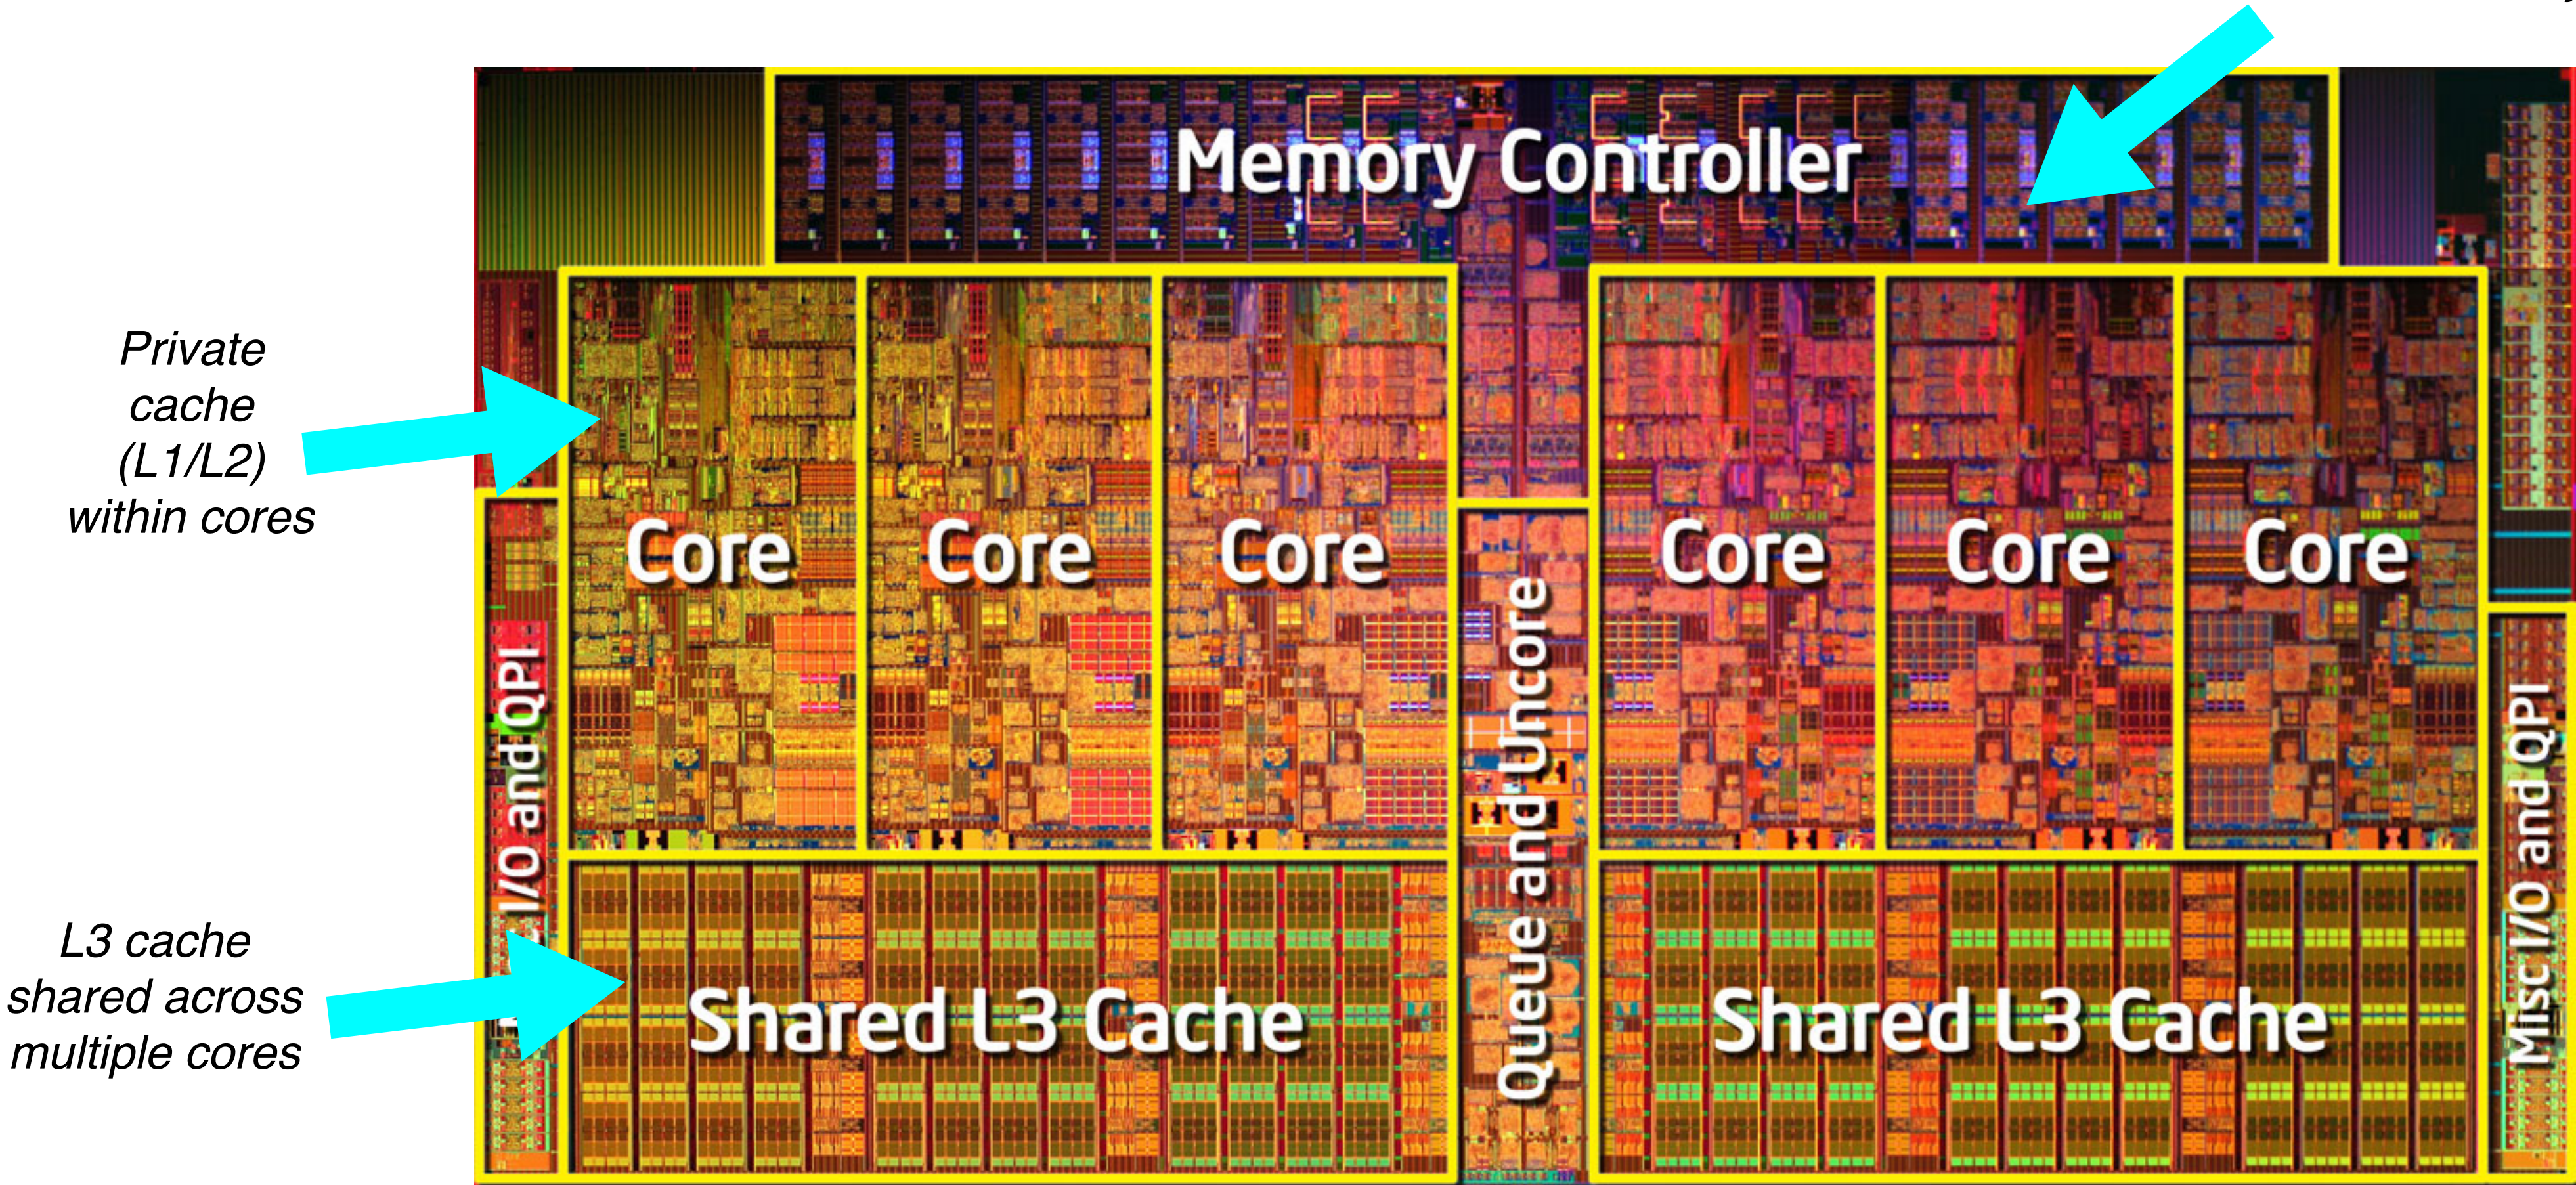


Computer System Overview



- Organization IRL

Access to RAM and Disk via Memory Controller

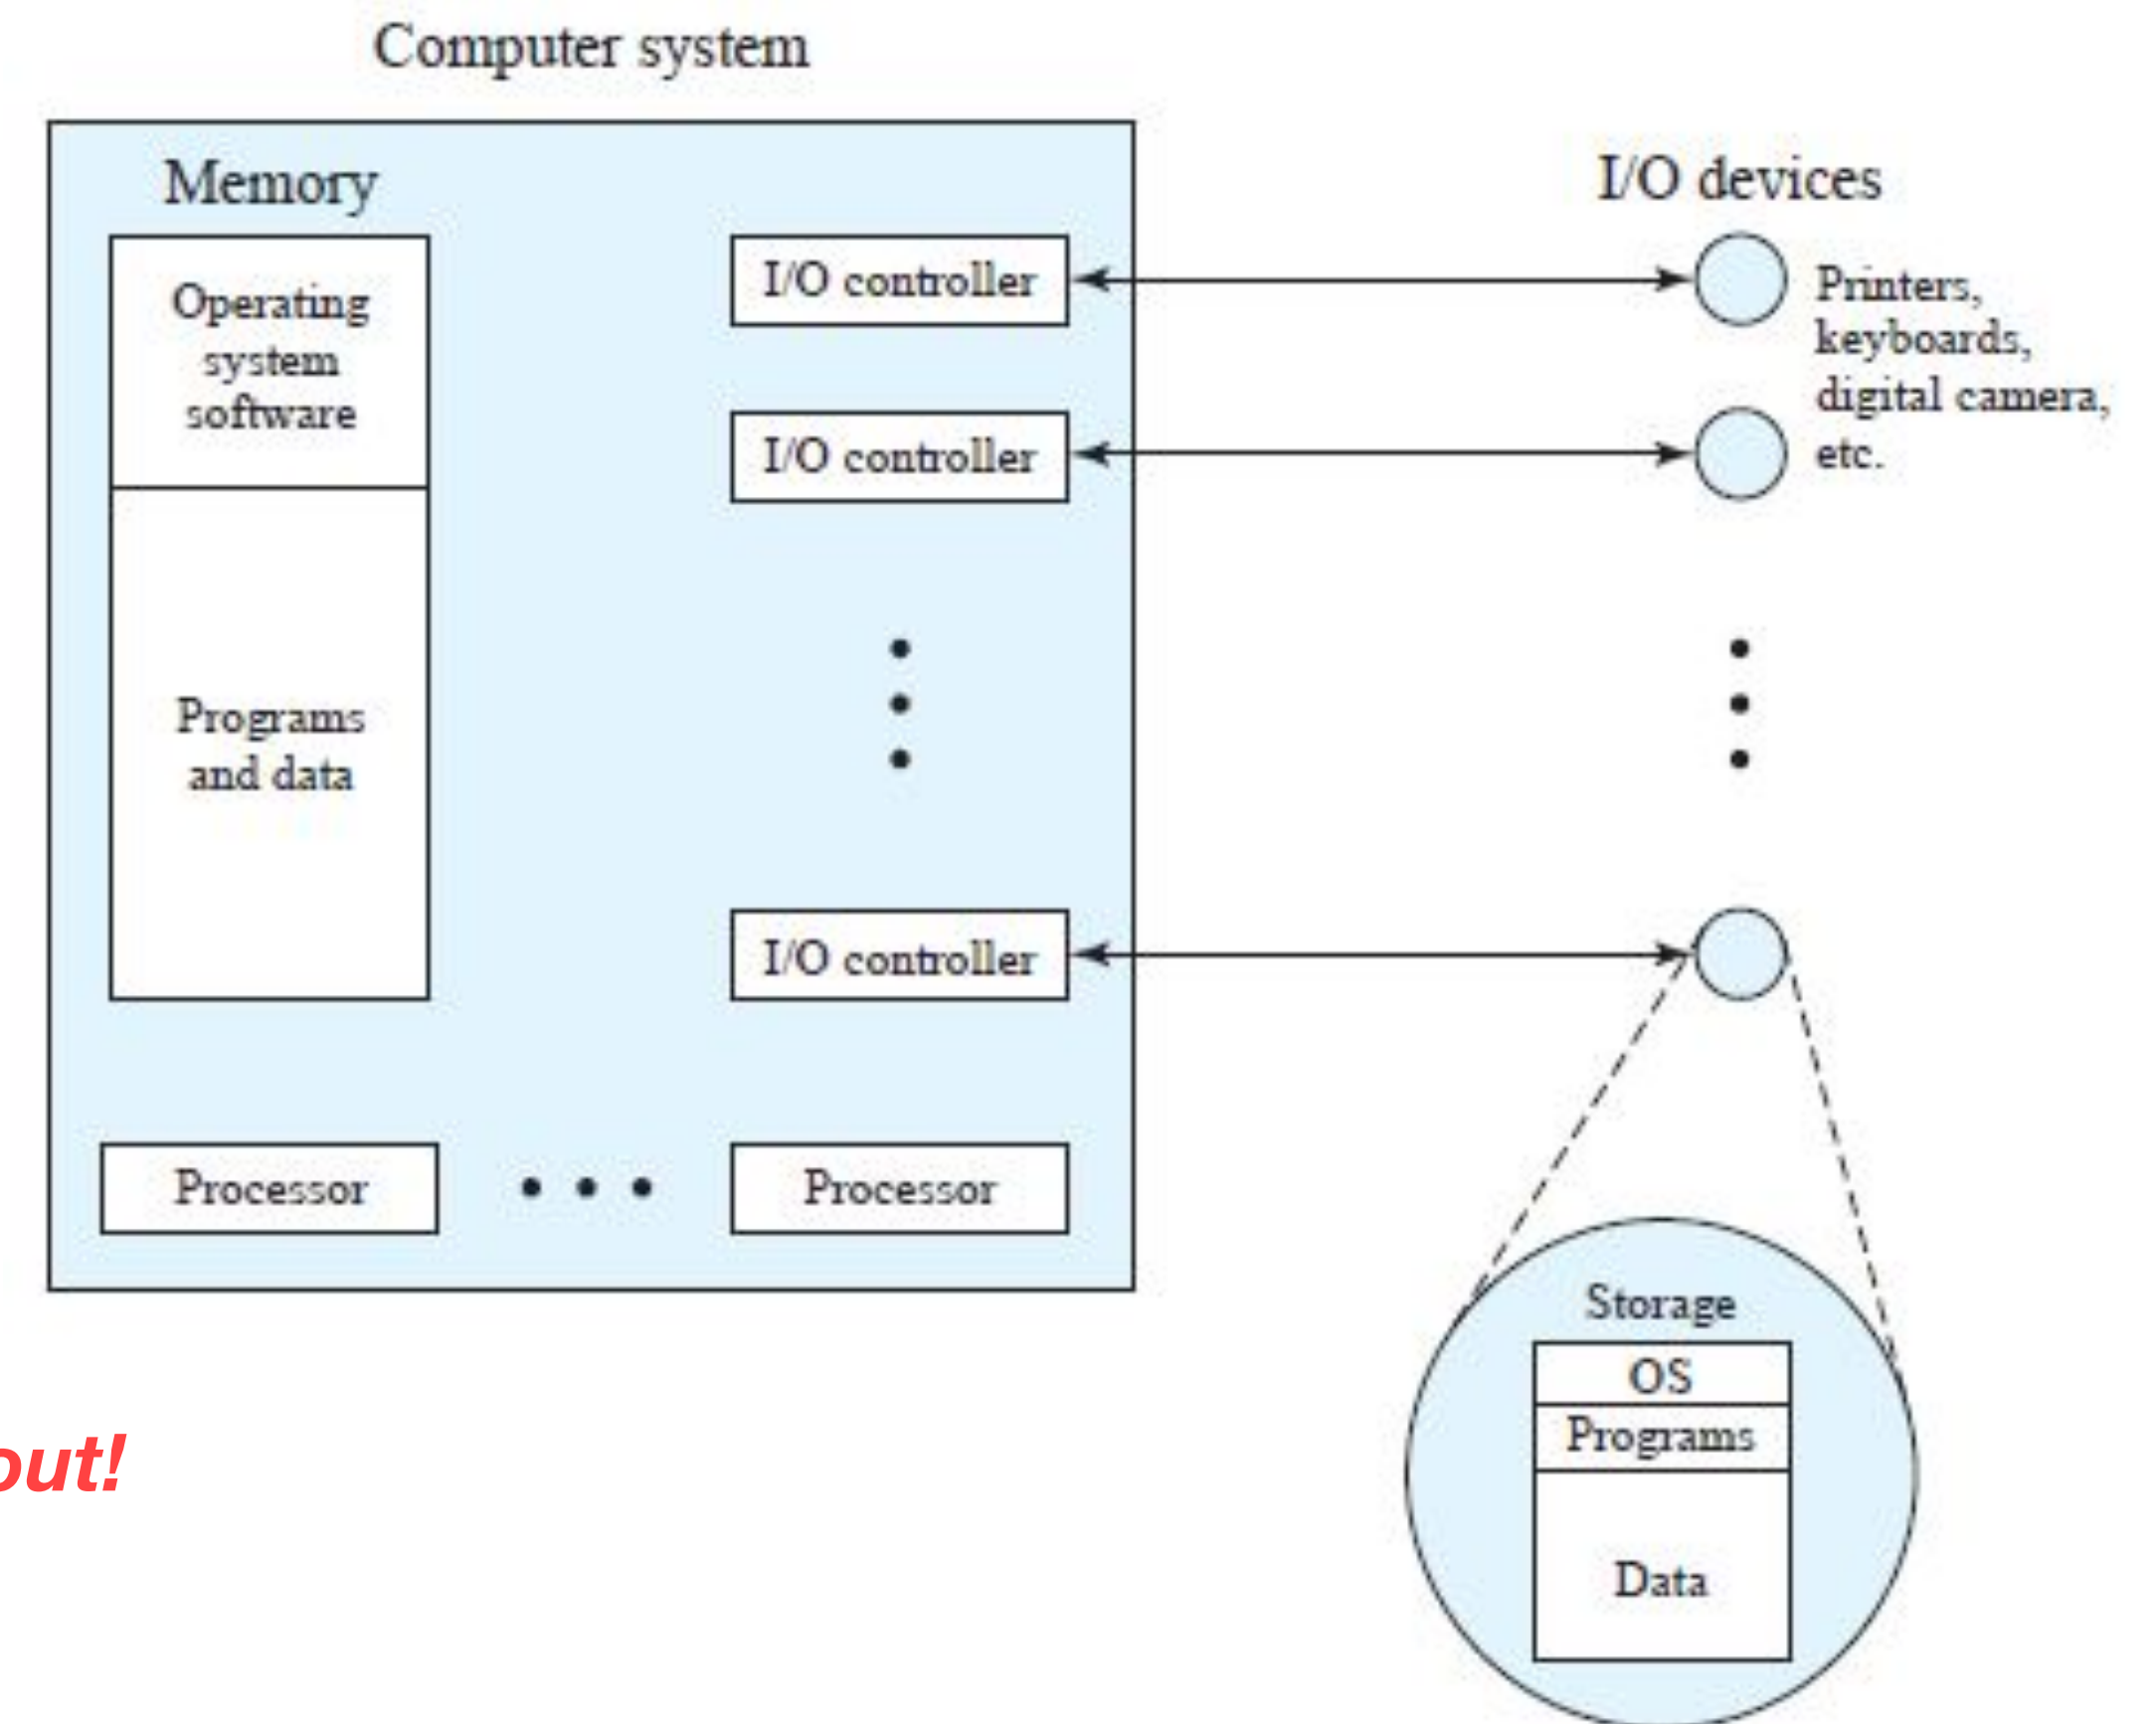


—Image Credit: Stallings

Operating System Overview

- Objectives of an OS
 - Convenience, Efficient, Modular
vs.
Referee, Illusionist, Glue
- Major Achievements
 - The Process
 - Memory Management
 - Security
 - Scheduling & Resource Management

These topics are what most of the rest of the class is about!




—Image Credit: Stallings

OS Objectives/Roles

The textbook offers one way of looking at the role of the OS. Here is another:

- Referee

- manage resources between apps
- isolate apps and users from one another
- facilitate communications between apps/users



- Illusionist

- make apps believe they have the whole machine to themselves
- create appearance of infinite processors and memory
- abstract away complexity of storage, network communications, etc.



- Glue

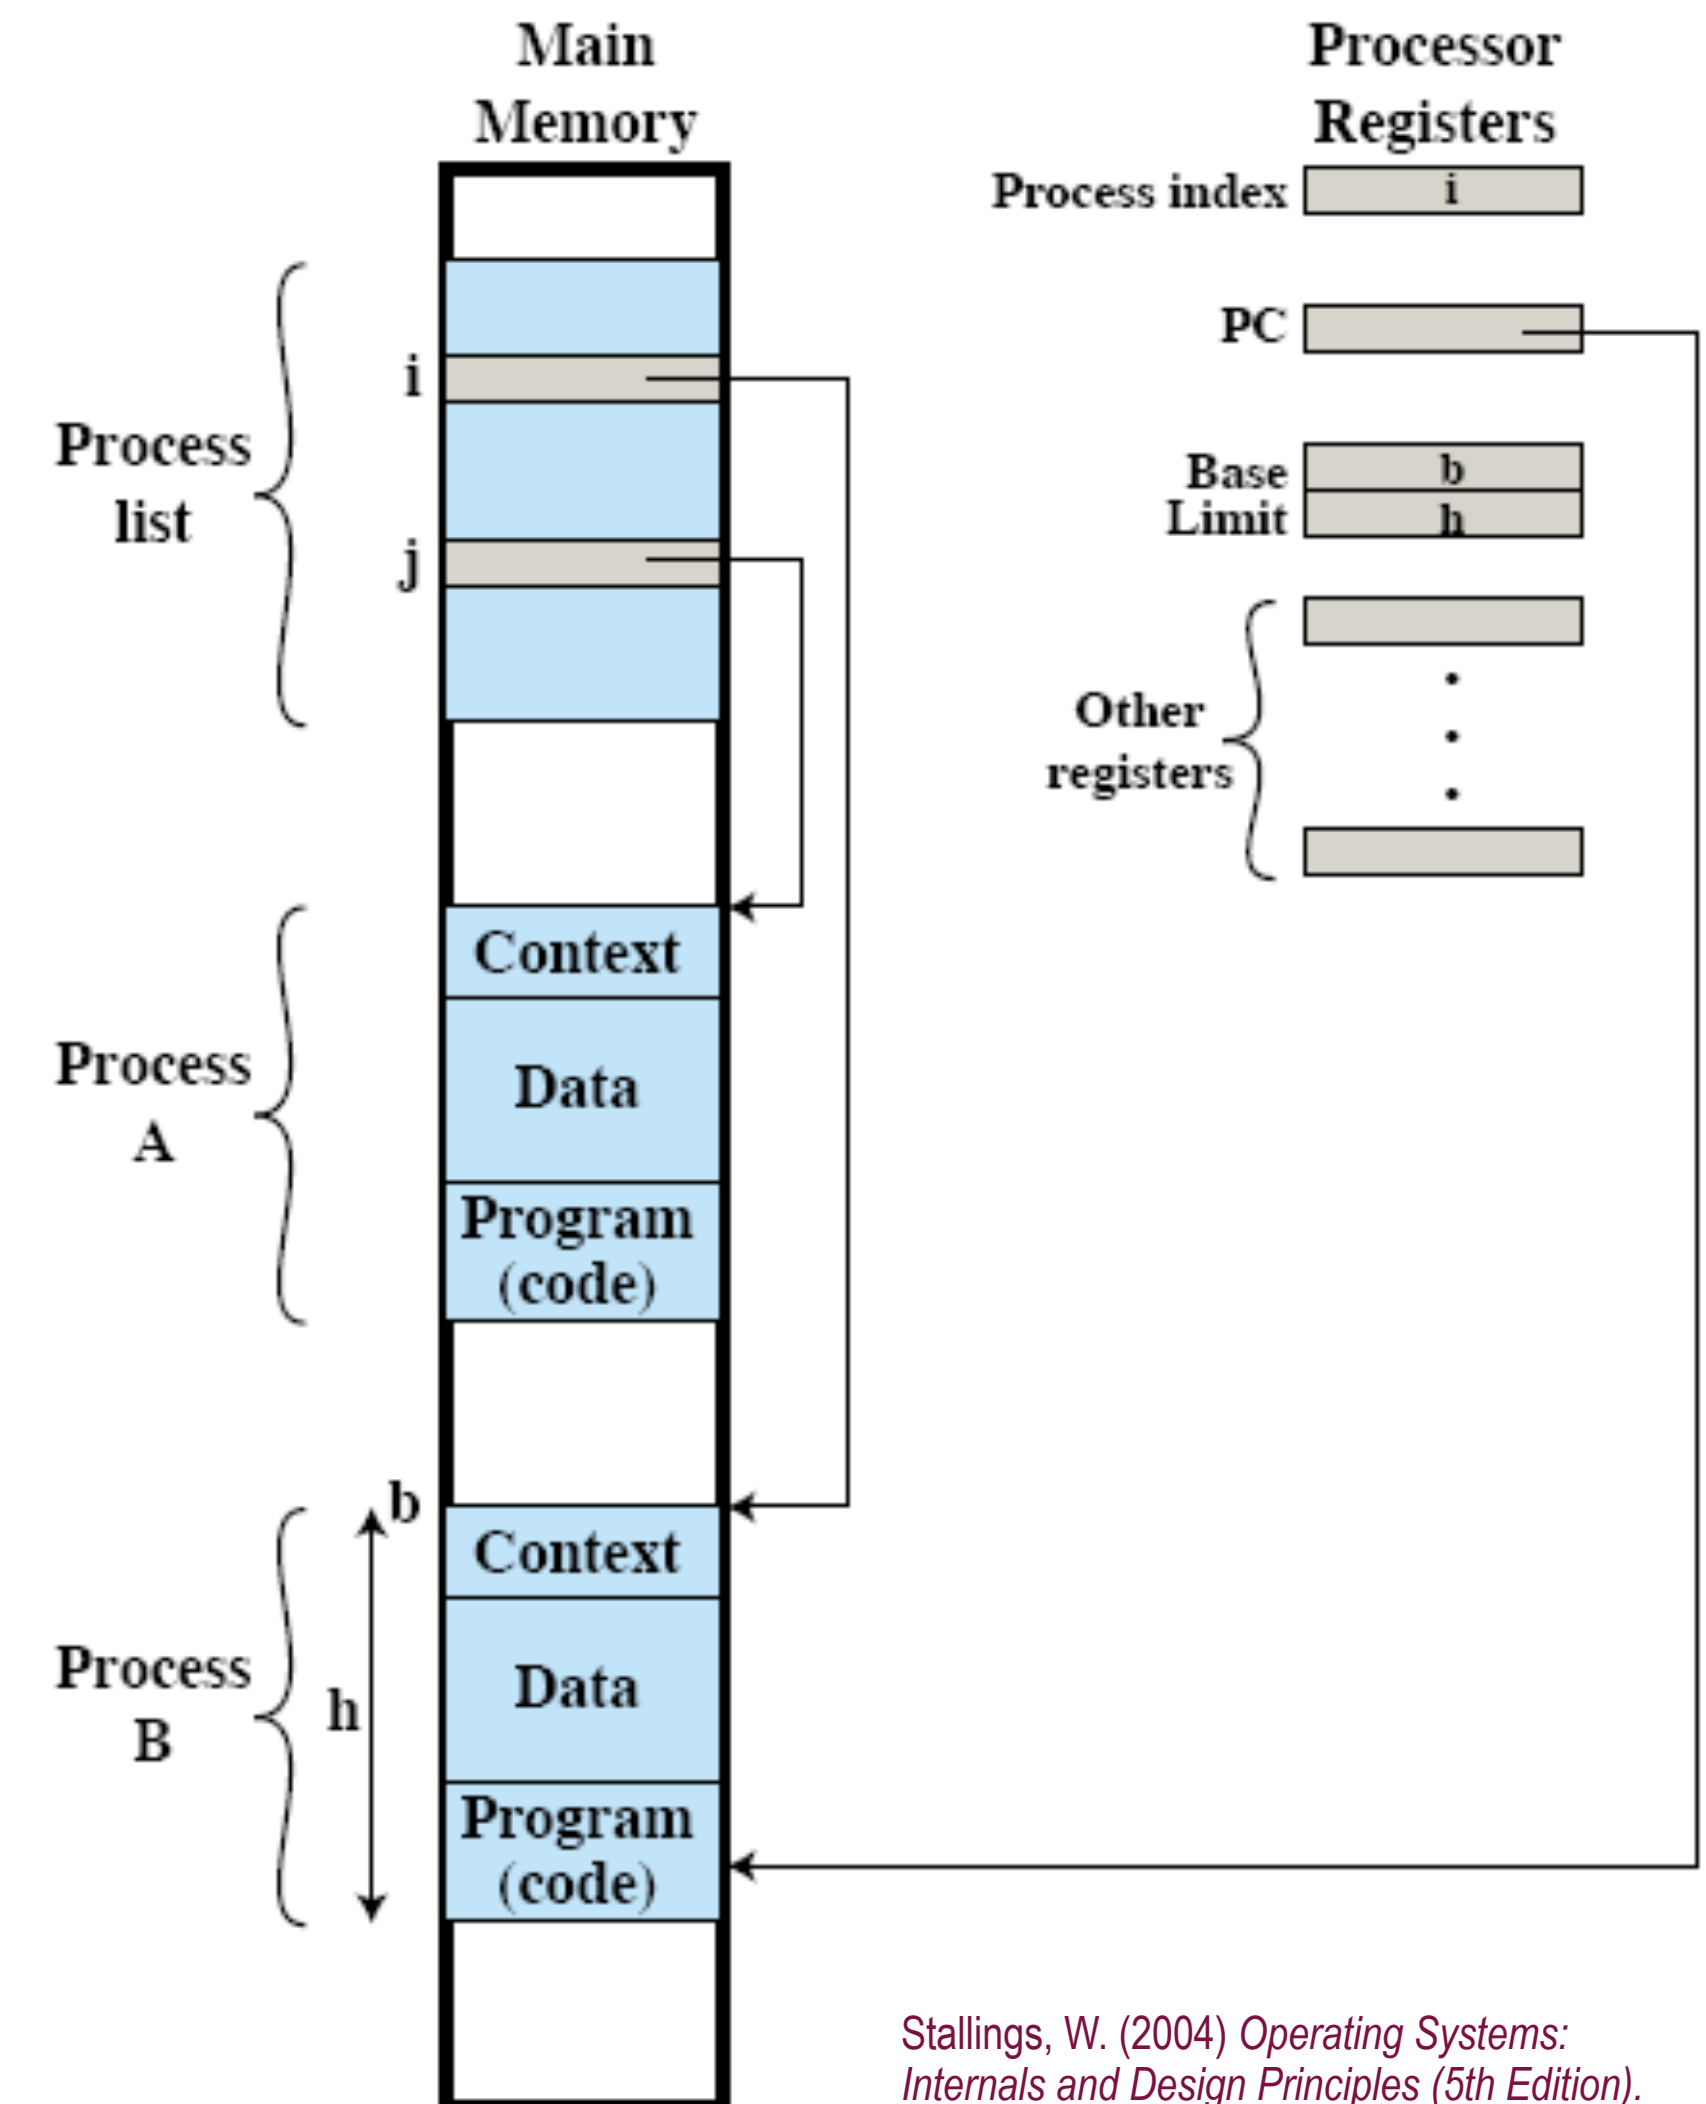
- manage hardware so apps can be machine-agnostic; provide a set of common services



—Thanks to Adam Bates for the nice analogies: <https://courses.engr.illinois.edu/cs423/sp2018/>

The Process


- The notion of a **Process**...
 - arose from multiprogramming, time sharing, real-time transactions;
 - is more general than a “job”; many definitions; e.g., ***an instance of a program running on a computer***;
 - consists of
 - (1) an executable program (**code**),
 - (2) associated **data**,
 - (3) **execution context** (info the OS needs to manage the process)
 - is realized as nothing more than a data structure!
 - A **thread** = cooperative execution within a process; use shared context.
- Process Switching
 - Interrupt = save context (e.g., PC and other registers)
 - ➔ execute interrupt handler
 - ➔ resume processing (same or different process)
 - states (simply put) = **executing** or **awaiting** execution



Stallings, W. (2004) *Operating Systems: Internals and Design Principles* (5th Edition).

Memory Management

- OS responsibilities w.r.t. MM
 - Process Isolation + Automatic Allocation/Mgmt. + Modular + Protection & Access Control + Long-Term Storage
- Virtual Memory
 - address memory **logically**, without regard to the actual, **physical** memory.
 - (recall “paging”) **virtual address** = page # + offset within the page.
 - MM maps between virtual address and real (physical) addresses.
 - pages can be in memory (or not!) — *how is that possible?!*
- File Systems
 - good for persisting information for extended periods of time
 - **files** = named objects





Information Protection & Security

- CIA
 - **C**onfidentiality — protection from unauthorized access (*e.g., snooping*)
 - **I**ntegrity — protection from unauthorized modification (*e.g., tampering*)
 - **A**vailability — protection from interruption (*e.g., denial of service*)
- Also...
 - Authenticity — verification of the origin/identity
 - +...



traditional infosec model...



...increasingly common to include other goals

Scheduling & Resource Management

- Manage Resources!
 - **Manage** system resources (Main Memory, I/O devices, Processors)
 - Keep these resources **utilized** (i.e., schedule processes to utilize them)
 - ...all while being **fair**, **responsive**, and **efficient**
- Operations-Research Problems & Lots of Maths
 - How do we do things good?
 - **Data Structures:** A bunch of queues, and lists, and (mostly) other simple data structures
 - **Algorithms:** A **scheduler** or **dispatcher** to pick which process runs next
 - **round-robin** = everyone gets some time in turn. *Other approaches?*

