

CSCI 460 Operating Systems

Processes (Part II)

Professor Travis Peters

Fall 2019

Some slides & figures adapted from Stallings instructor resources.

*Some slides adapted from Adam Bates's F'18 CS423 course @ UIUC
<https://courses.engr.illinois.edu/cs423/sp2018/schedule.html>*

Goals for Today

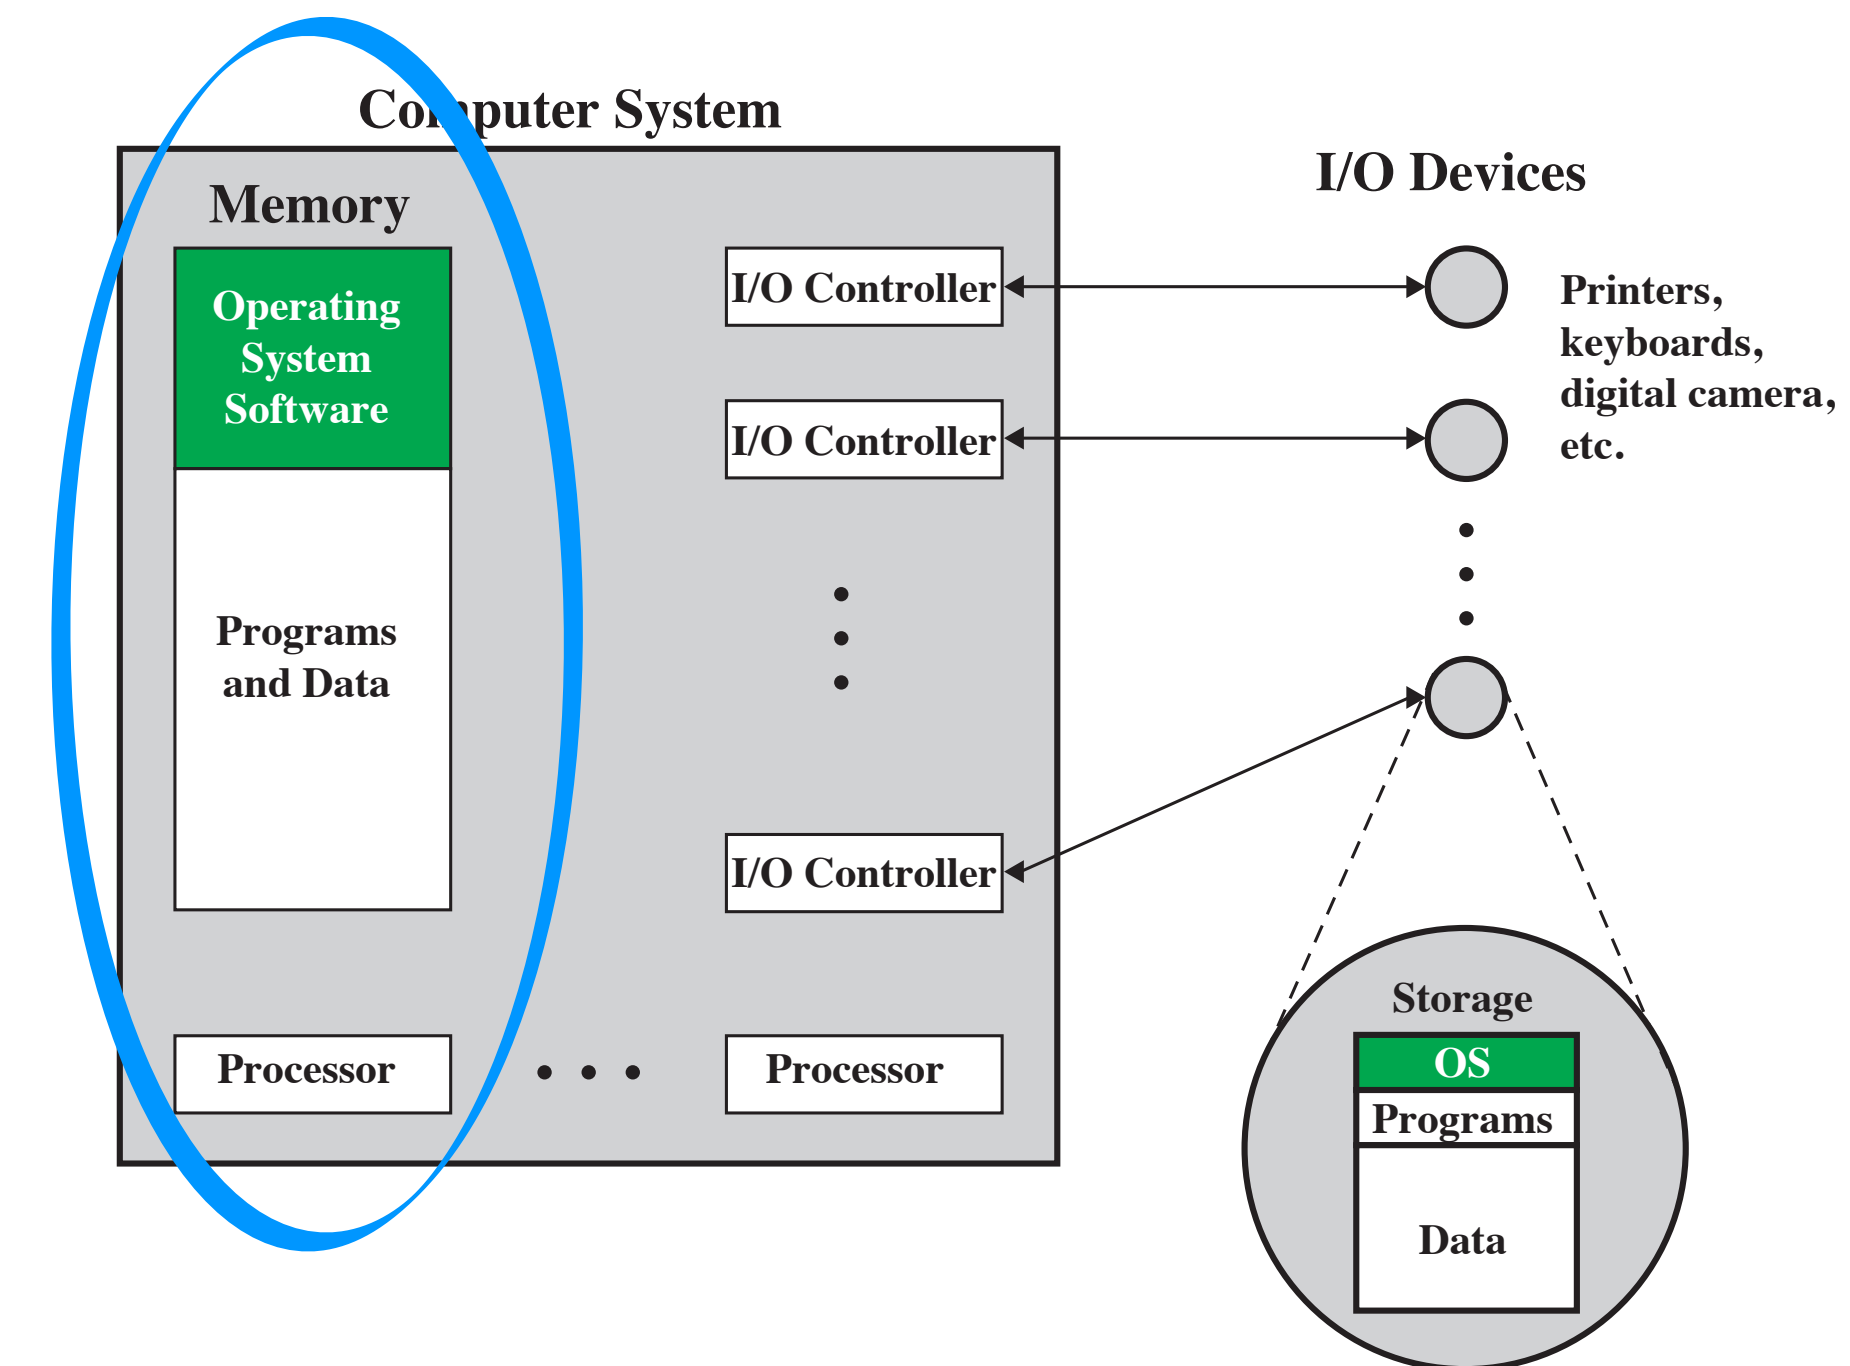
- **Learning Objectives**

- Understand basic concept of process (control info, creation, termination, states, etc.)
- Review some important UNIX syscalls and concepts for sys. programming

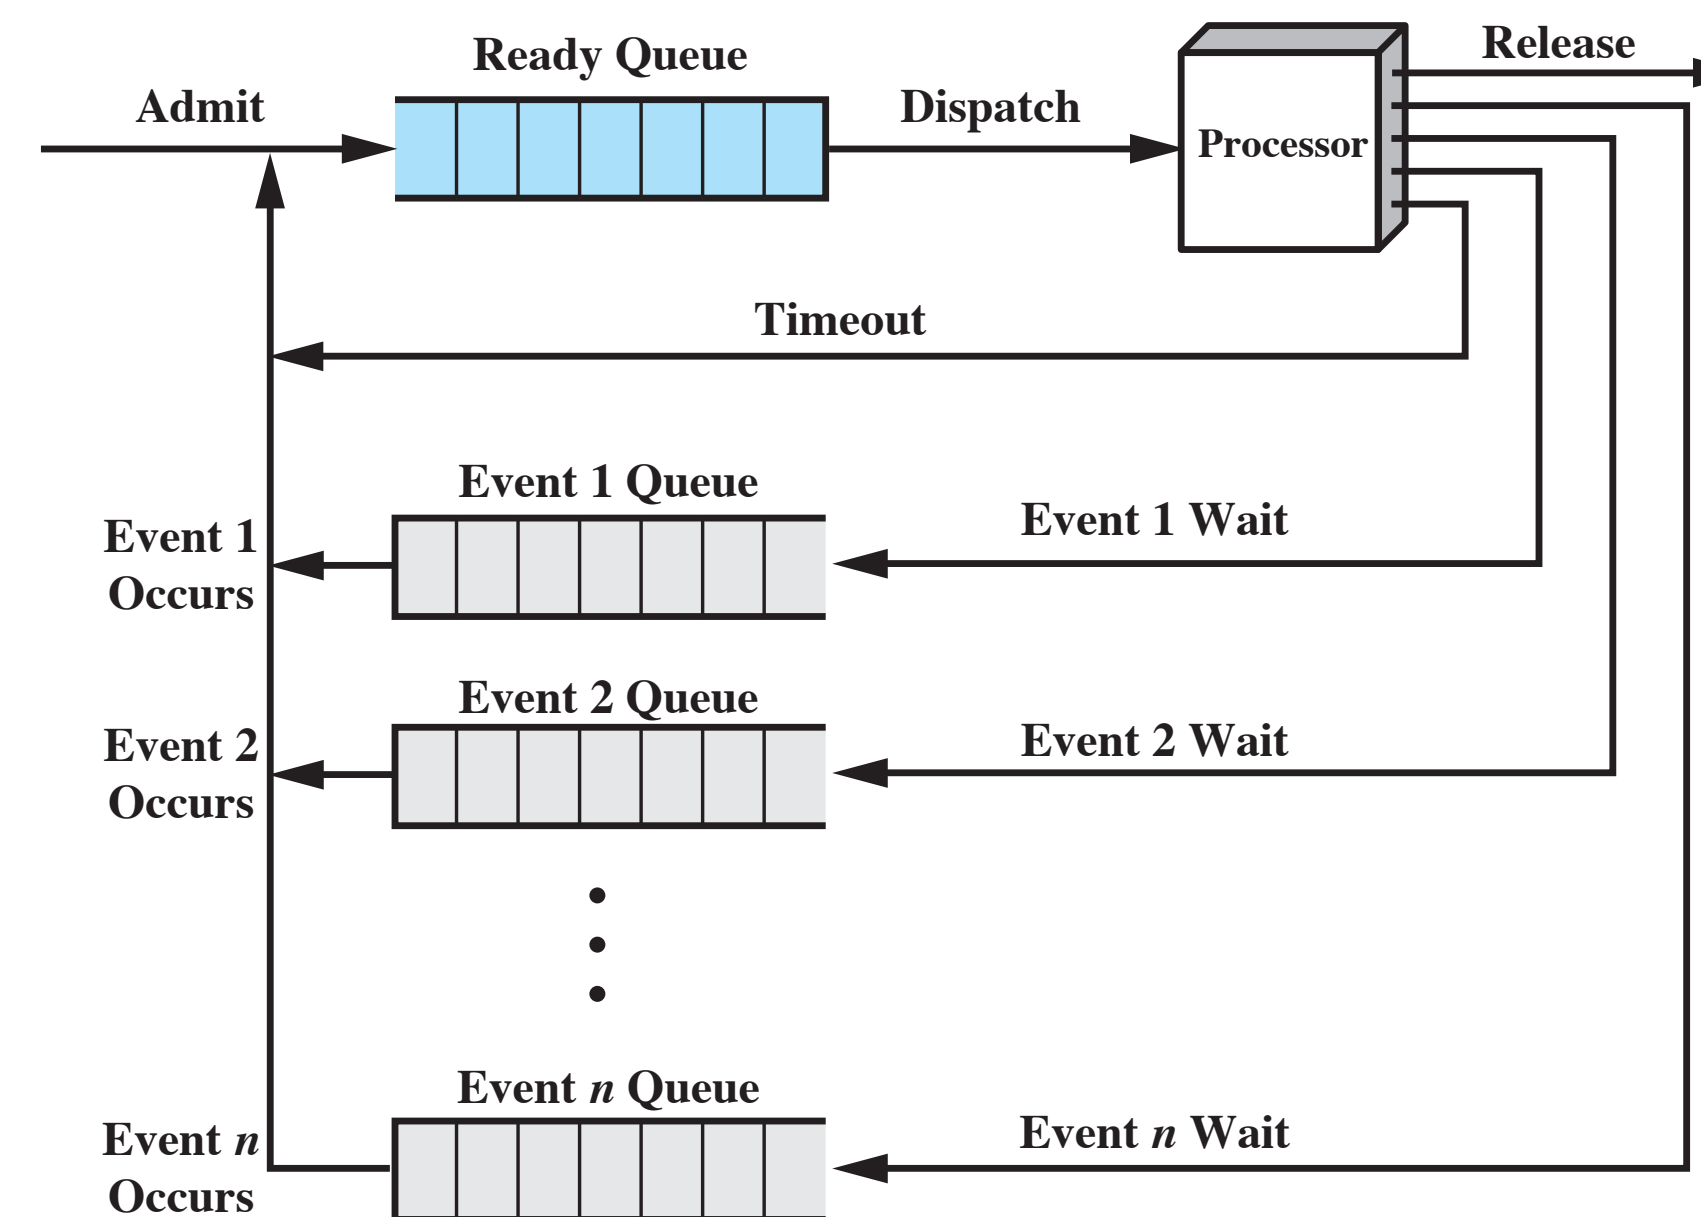
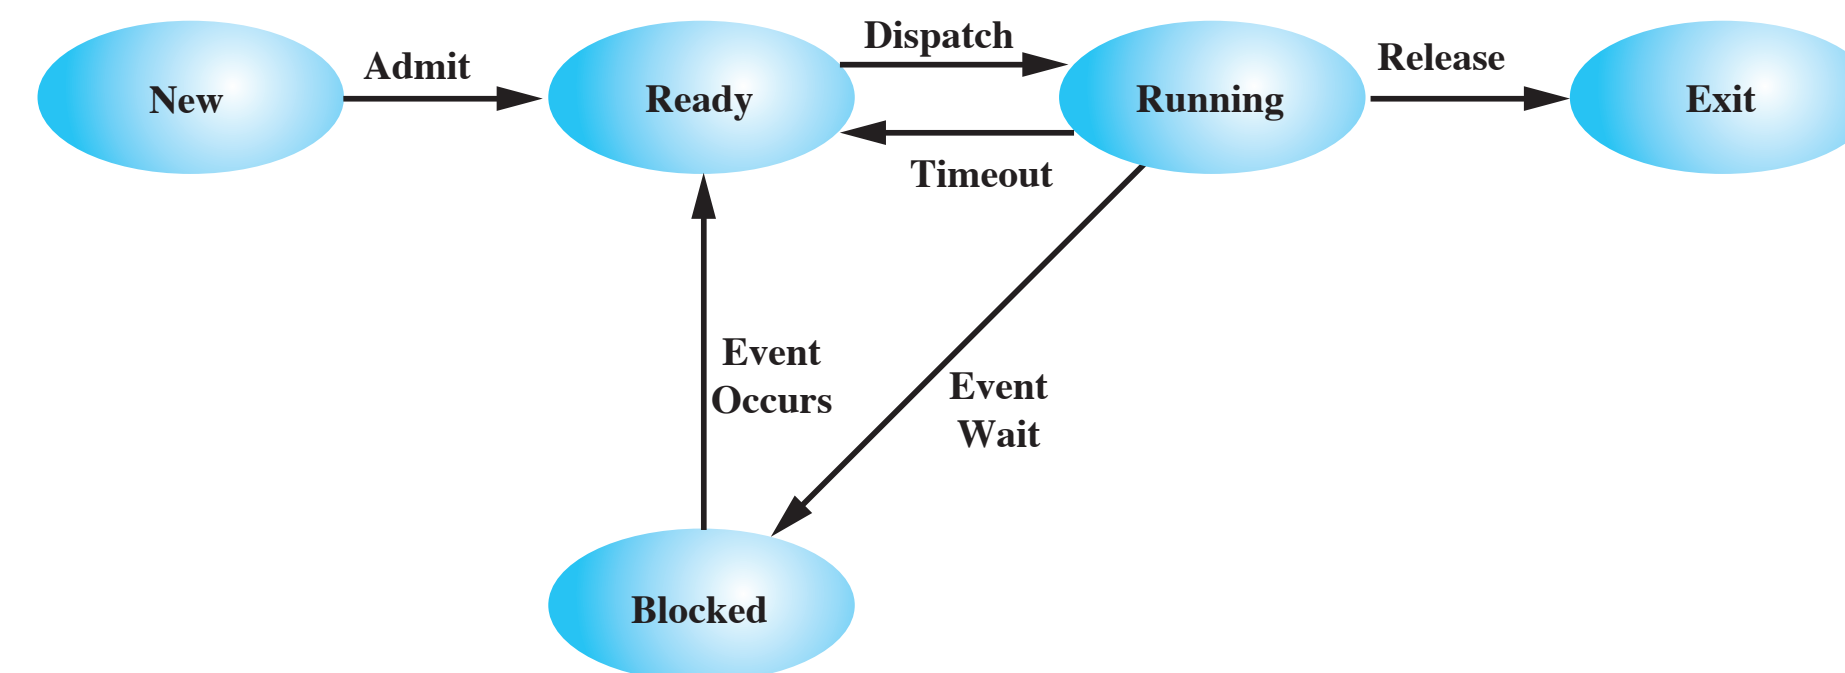
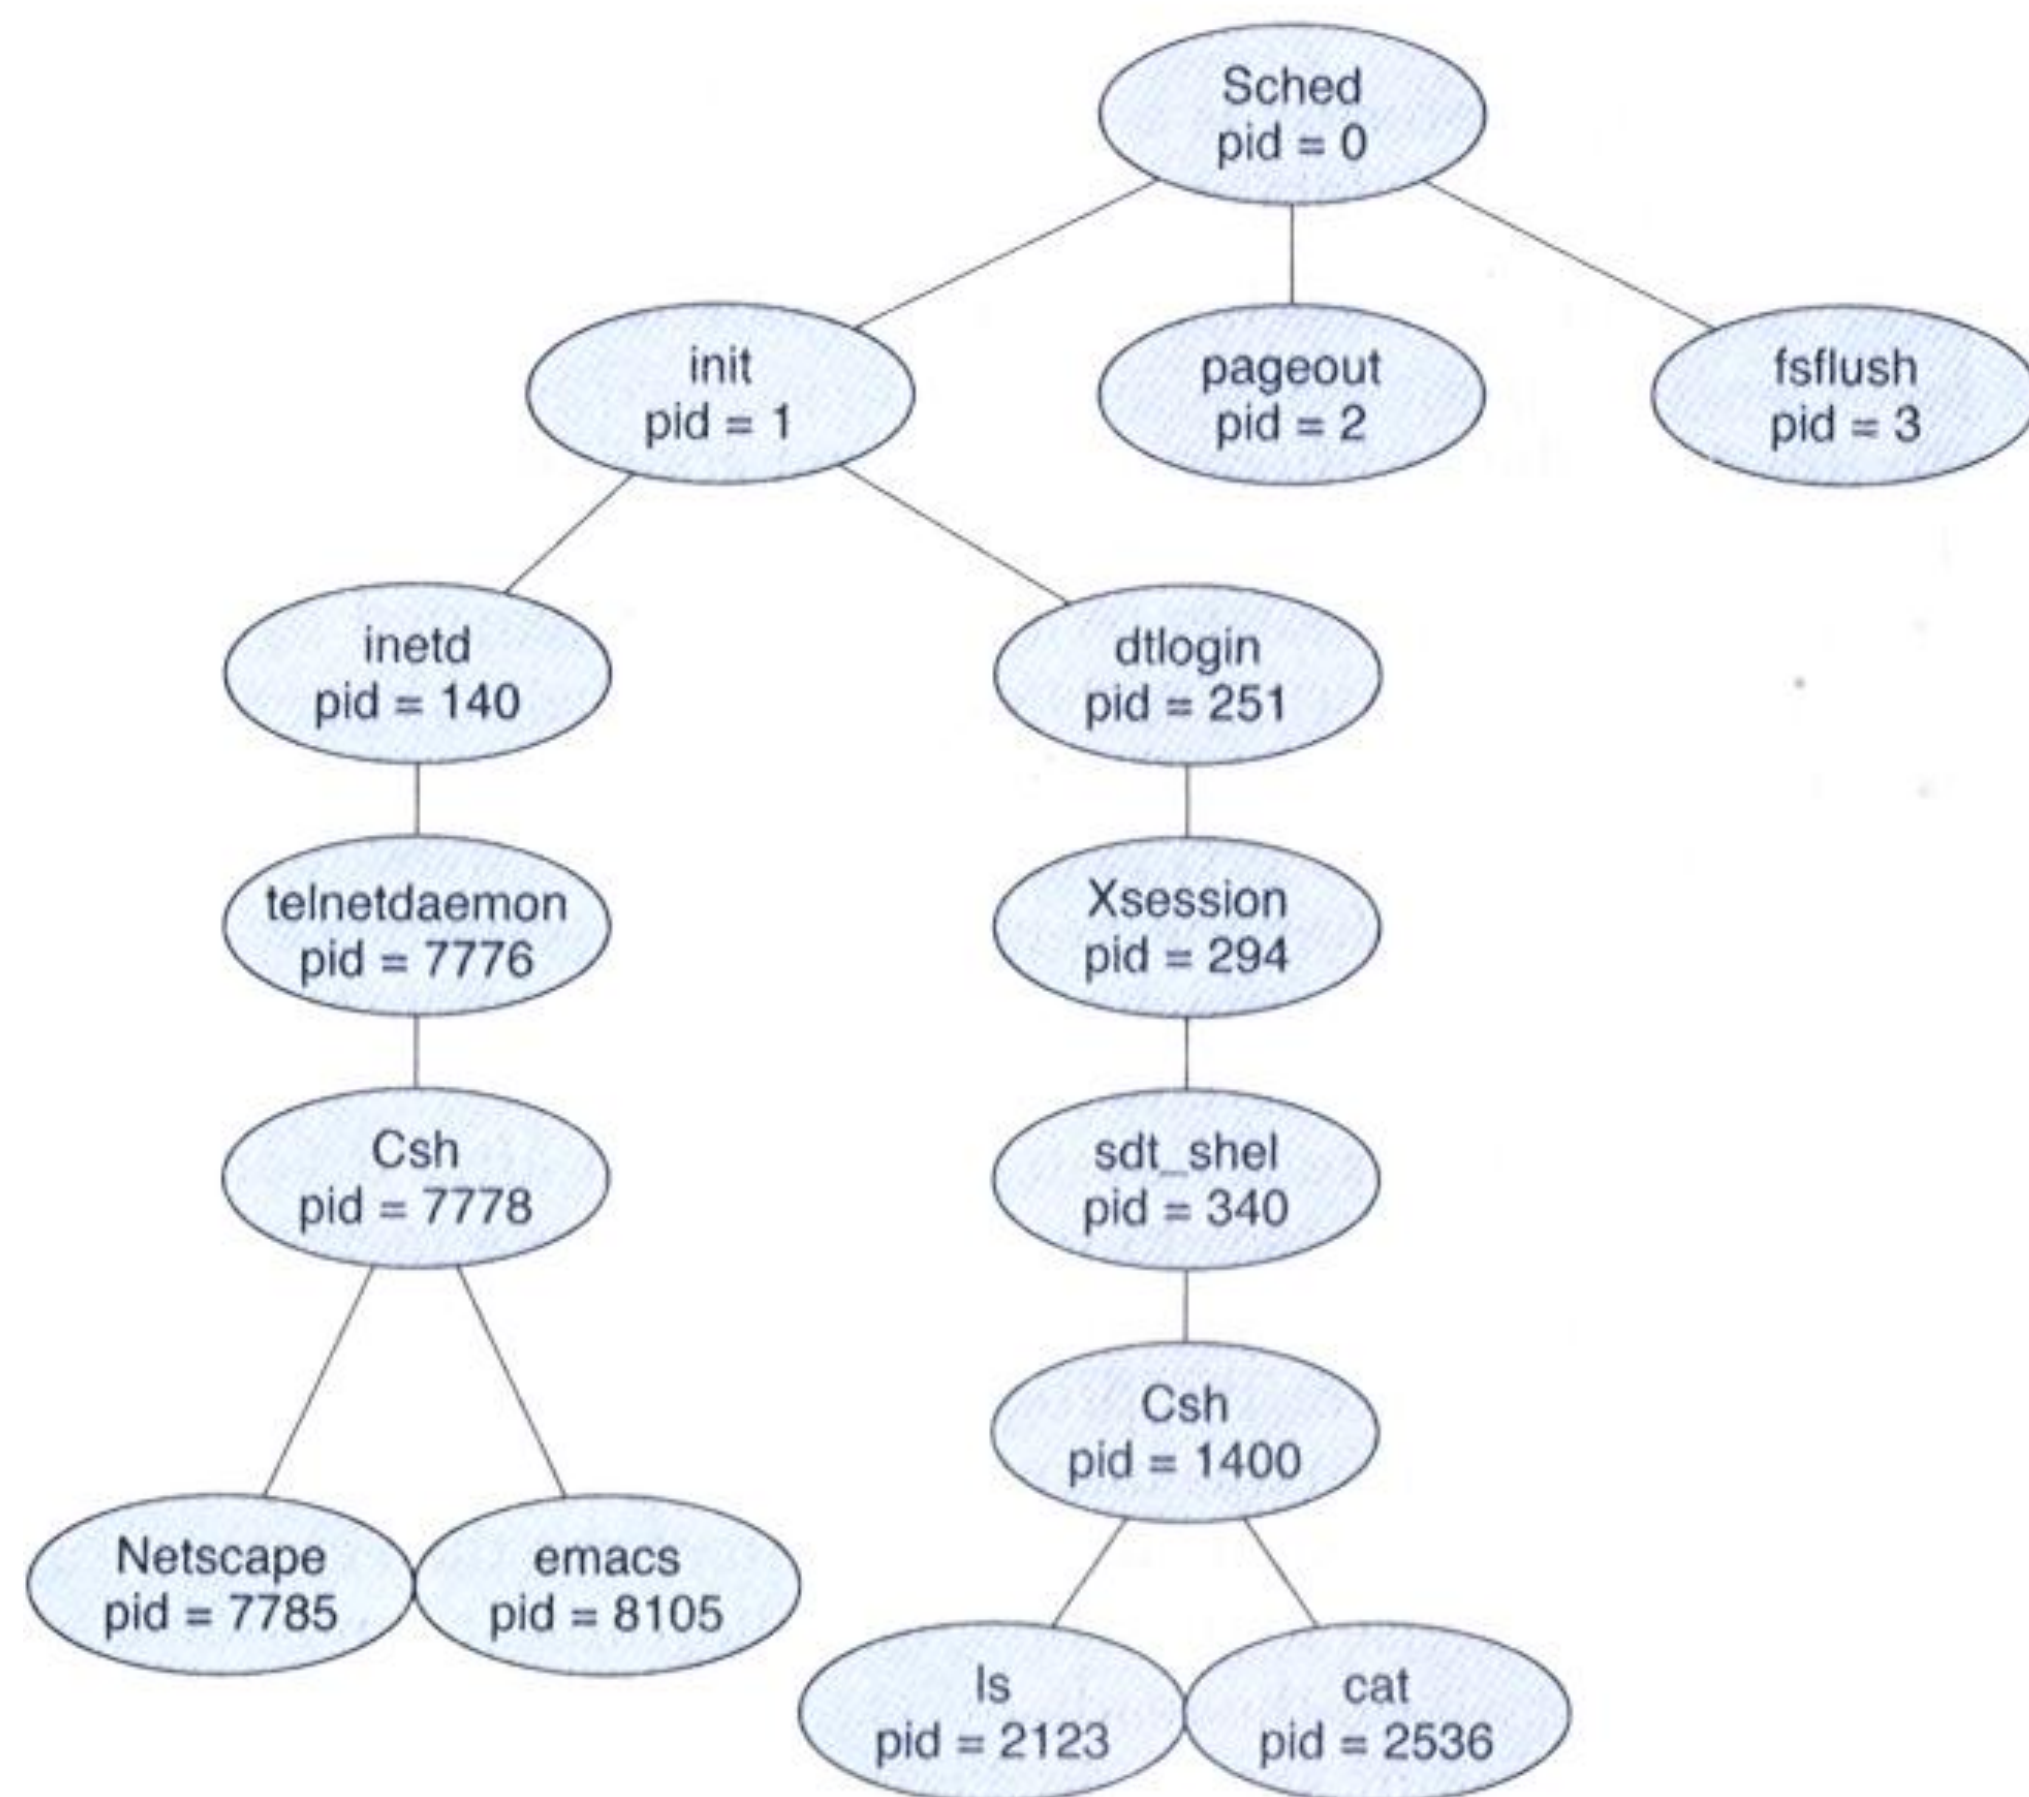
- **Announcements**

- Grades should be fixed...
- Rough schedule posted today
- *Coming Soon...*
 - zyBook for OS (optional resource); link posted soon
 - 1st programming assignment
 - Details about project

<https://www.traviswpeters.com/cs460/>



Last Time... Creating & Managing Processes



Creating a Process

DEMO

*take a look at a process tree
(already-created processes)*

``ps axjf``

Some things to note:

names (e.g., `init`),

relationships (parents, children, grandchildren),

IDs (PID, PPID)

```
vagrant@osbox:~/os$ ps axjf
```

PPID	PID	PGID	SID	TTY	TPGID	STAT	UID	TIME	COMMAND
0	2	0	0	?	-1	S	0	0:00	[kthreadd]
2	3	0	0	?	-1	S	0	0:00	\ [ksoftirqd/0]
2	4	0	0	?	-1	S	0	0:00	\ [kworker/0:0]
2	5	0	0	?	-1	S<	0	0:00	\ [kworker/0:0H]
2	7	0	0	?	-1	S	0	0:01	\ [rcu_sched]


0	1	1	1	?	-1	Ss	0	0:00	/sbin/init
1	524	524	524	?	-1	Ss	0	0:00	dhclient -1 -v -pf /run/dhclient.eth0.pid -lf /va
1	621	621	621	?	-1	Ss	0	0:00	rpcbind
1	663	663	663	?	-1	Ss	107	0:00	rpc.statd -L
1	741	741	741	?	-1	Ss	102	0:02	dbus-daemon --system --fork
1	806	806	806	?	-1	Ss	0	0:00	rpc.idmapd
1	826	826	826	?	-1	Ss	0	0:00	/lib/systemd/systemd-logind
1	832	832	832	?	-1	Ssl	101	0:00	rsyslogd
1	937	937	937	tty4	937	Ss+	0	0:00	/sbin/getty -8 38400 tty4
1	940	940	940	tty5	940	Ss+	0	0:00	/sbin/getty -8 38400 tty5
1	944	944	944	tty2	944	Ss+	0	0:00	/sbin/getty -8 38400 tty2
1	945	945	945	tty3	945	Ss+	0	0:00	/sbin/getty -8 38400 tty3
1	947	947	947	tty6	947	Ss+	0	0:00	/sbin/getty -8 38400 tty6
1	987	987	987	?	-1	Ss	0	0:00	/usr/sbin/sshd -D
987	8560	8560	8560	?	-1	Ss	0	0:00	\ sshd: vagrant [priv]
8560	8637	8560	8560	?	-1	S	1000	0:00	\ sshd: vagrant@pts/0
8637	8638	8638	8638	pts/0	14242	Ss	1000	0:00	\ -bash
8638	14242	14242	8638	pts/0	14242	R+	1000	0:00	\ ps axjf
1	989	989	989	?	-1	Ss	0	0:00	acpid -c /etc/acpi/events -s /var/run/acpid.socke
1	990	990	990	?	-1	Ss	0	0:00	cron
1	991	991	991	?	-1	Ss	1	0:00	atd
1	1183	1183	1183	?	-1	Ssl	0	0:01	/usr/bin/ruby /usr/bin/puppet agent
1	1199	1196	1196	?	-1	Sl	0	0:12	/usr/sbin/VBoxService --pidfile /var/run/vboxadd-
1	1322	1319	1319	?	-1	Sl	0	0:00	ruby /usr/bin/chef-client -d -P /var/run/chef/cli
1	1349	1349	1349	tty1	1349	Ss+	0	0:00	/sbin/getty -8 38400 tty1
1	2176	2175	2175	?	-1	S	0	0:00	upstart-file-bridge --daemon
1	2179	2178	2178	?	-1	S	0	0:00	upstart-socket-bridge --daemon
1	3800	3799	3799	?	-1	S	0	0:00	upstart-udev-bridge --daemon
1	3803	3803	3803	?	-1	Ss	0	0:00	/lib/systemd/systemd-udevd --daemon

Creating a Process

- **Q:** But *how* to create a process? What UNIX call creates a process?

Creating a Process - `fork()`

- **Q:** But *how* to create a process? What UNIX call creates a process?
- `fork()` duplicates a process so that instead of one process, you get two!
 - *P1 and P2 continue in parallel from the statement that follows the `fork()`*
- **Q:** How can you tell P1 and P2 apart?
...the return value of `fork()`!
 - `rval == 0` //to child
 - `rval == child_pid` //to parent
 - `rval == -1` //fork() failed
- **Q:** Will ***child*** see changes to global variable made by the ***parent***?
No! On `fork()`, child gets new PC, stack, heap, globals, PID!

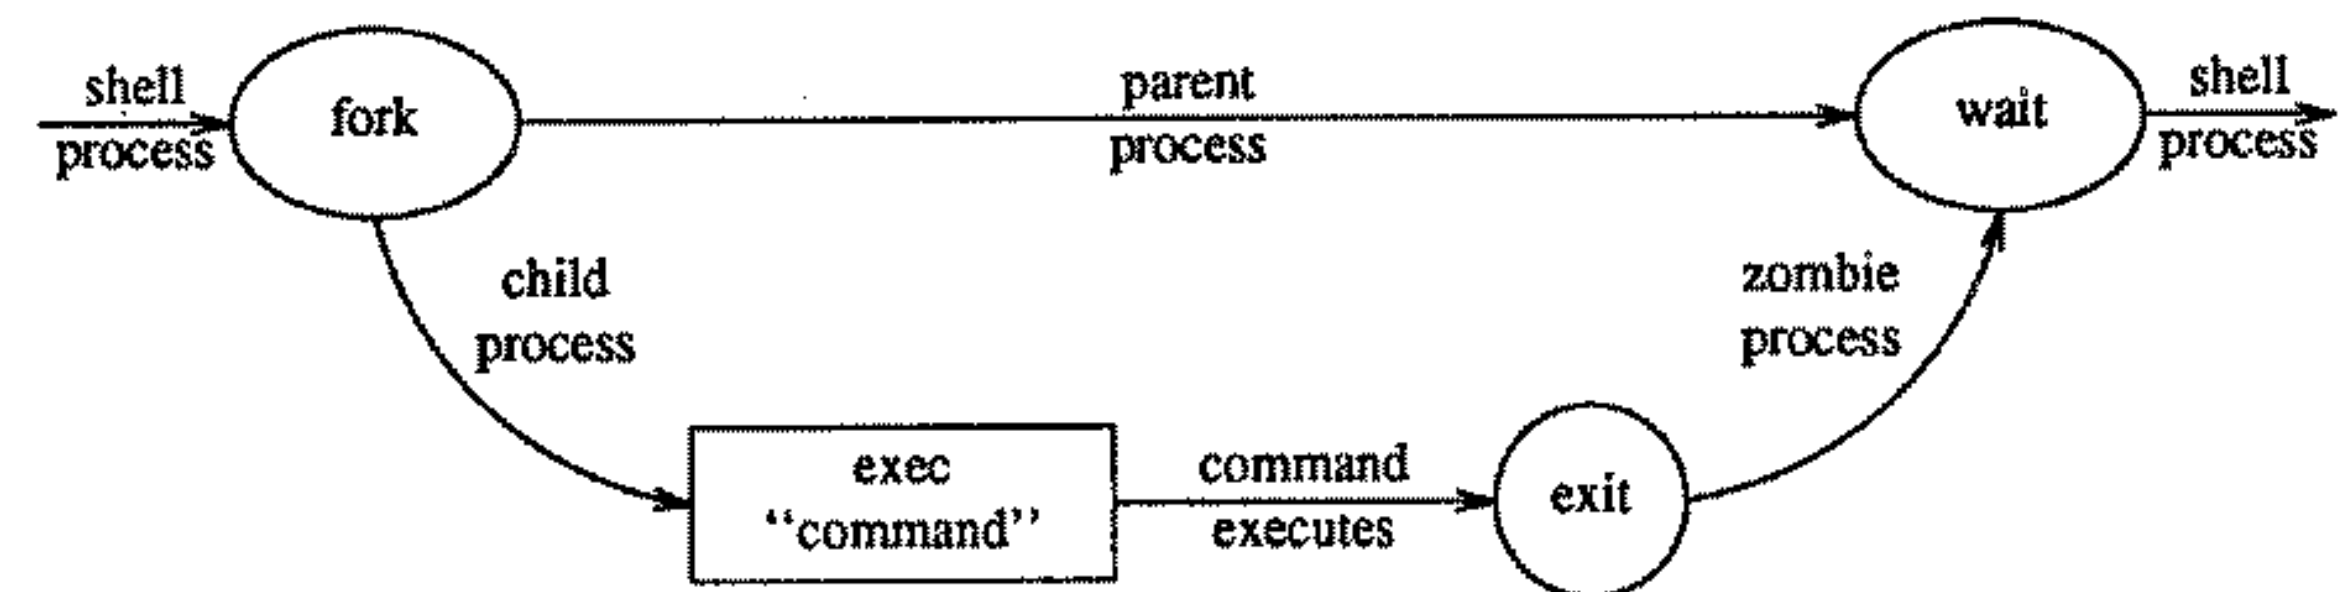


Creating a (Different) Process

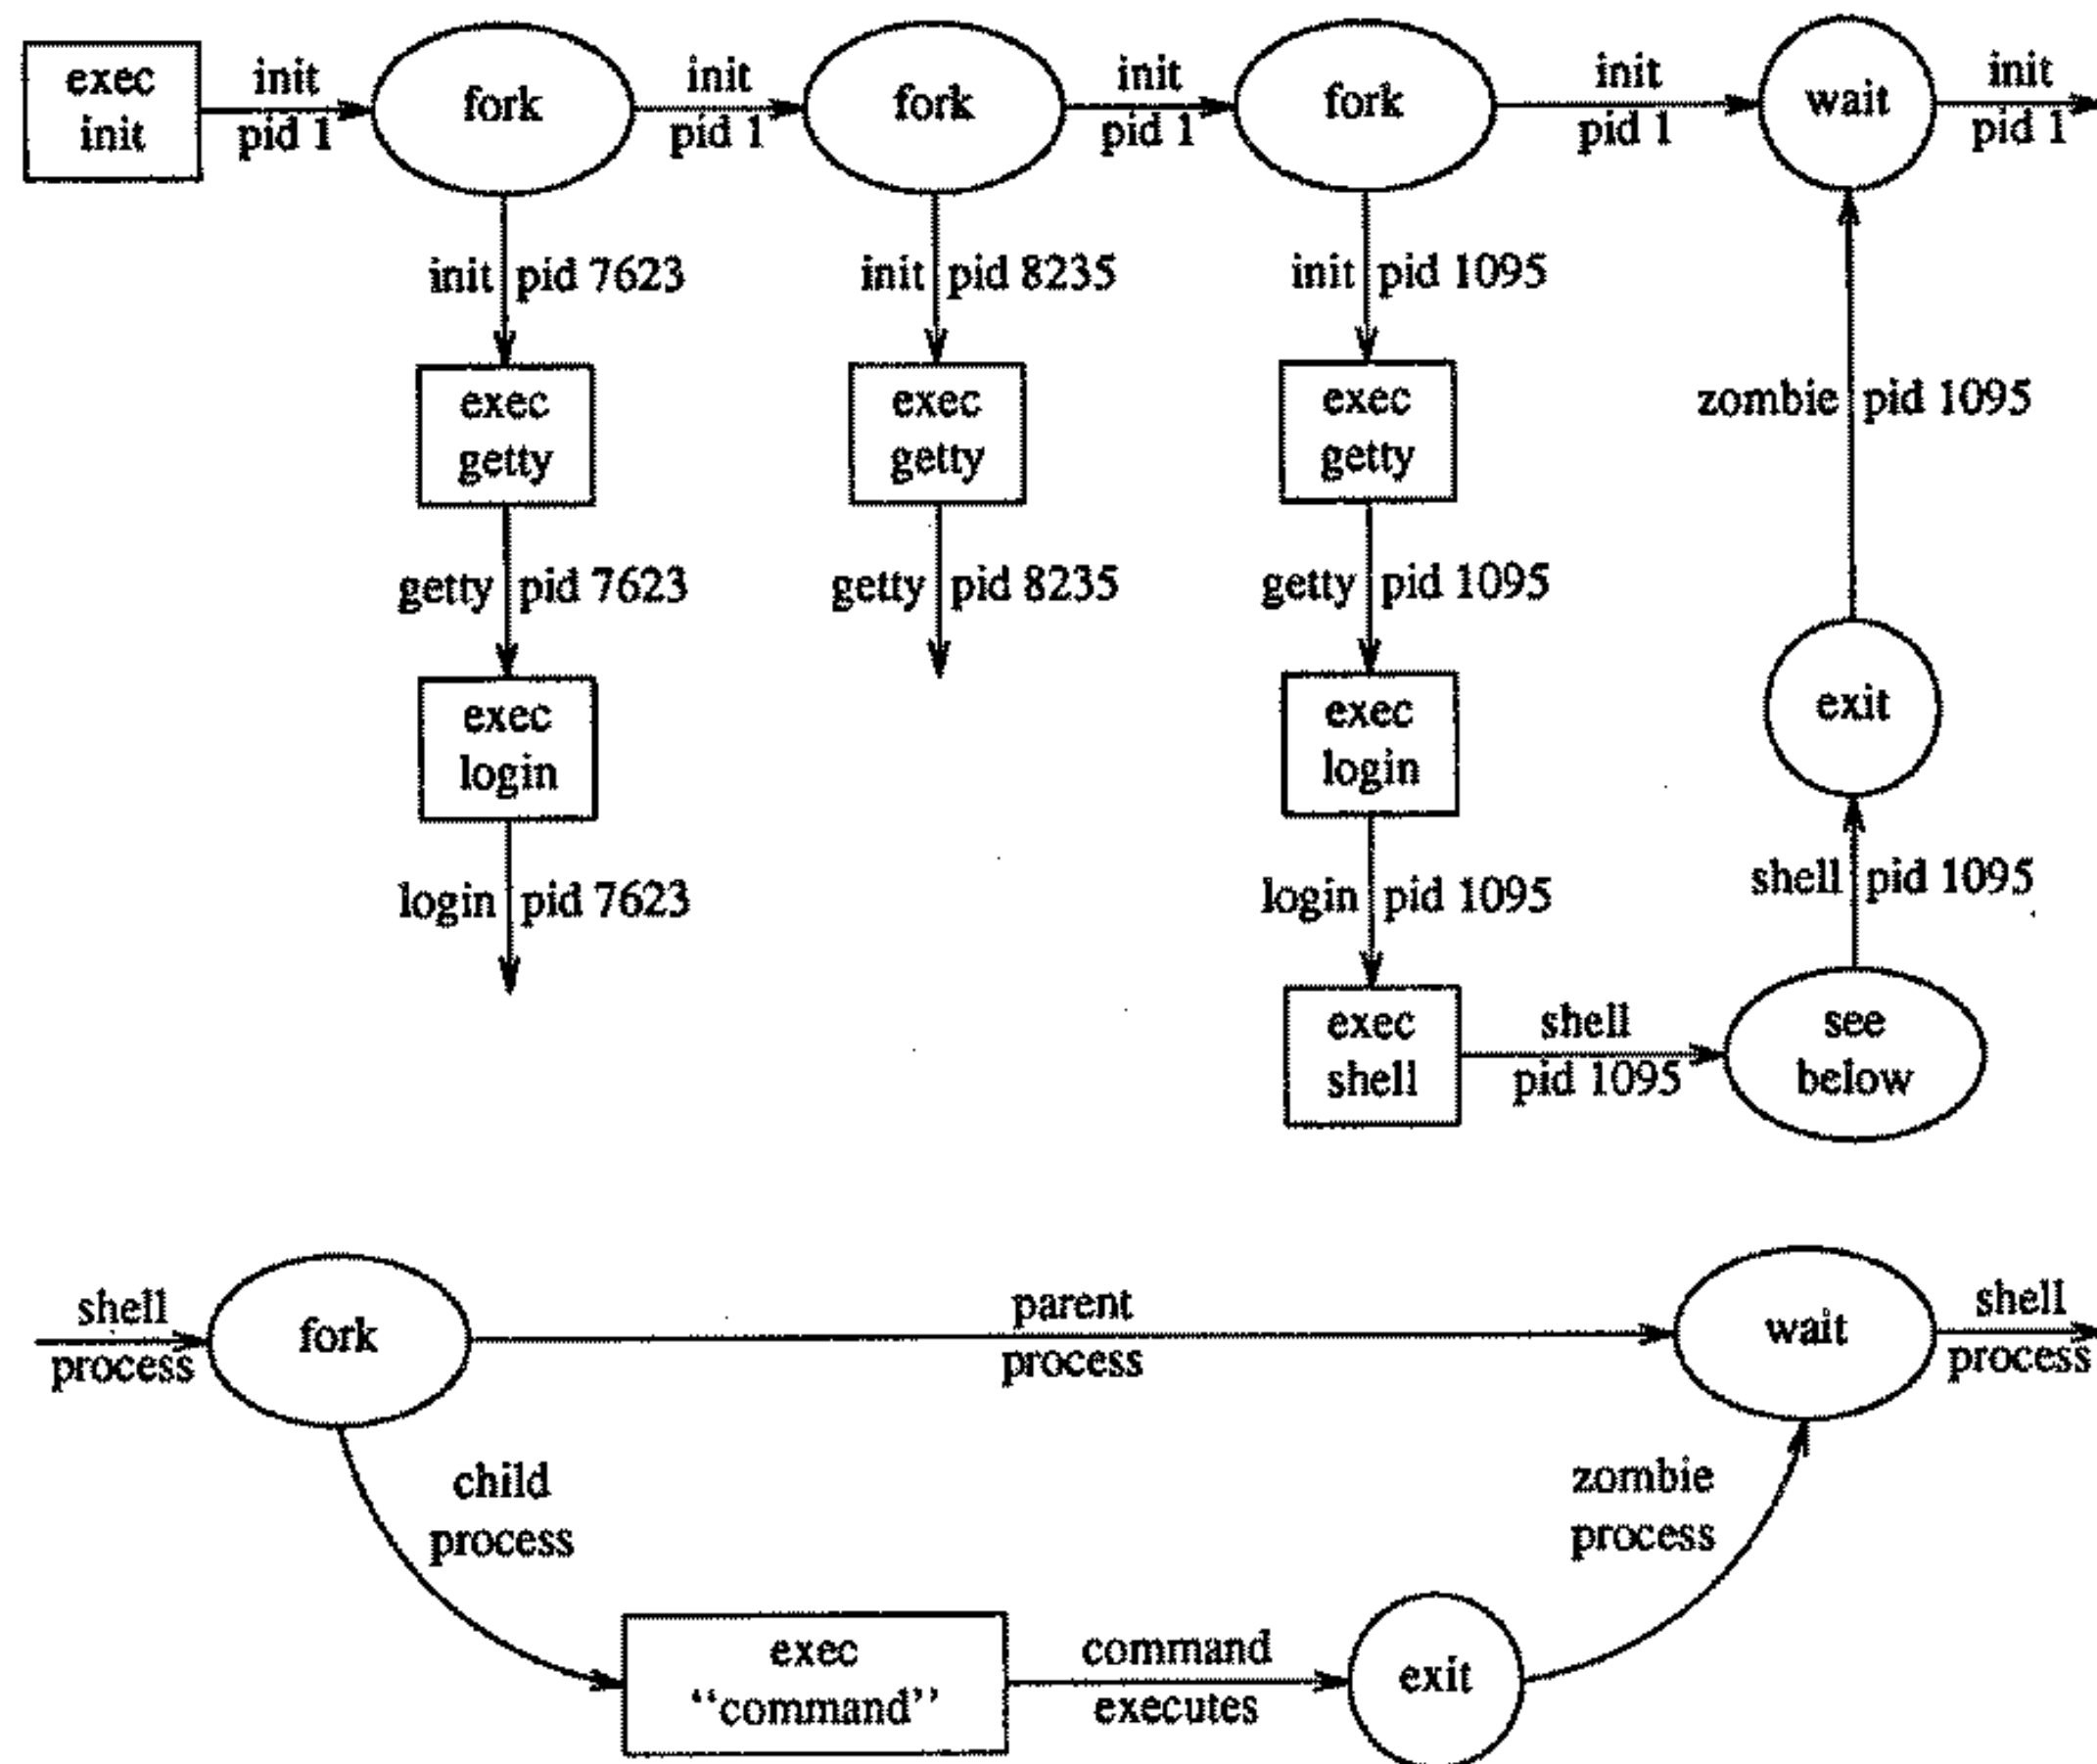
- **Q:** What if we want the child process to execute different code than the parent process?

Creating a (Different) Process - `exec()`

- **Q:** What if we want the child process to execute different code than the parent process?
- `exec()` ...
 - ...allows child to execute code that is different from the parent's code
 - ...has a family of functions (`execl()`, `execv()`, `execlp()`, `execvp()`) that provide a facility for overlaying the process image of the calling process with a new image
 - ...returns -1 and sets `errno` if unsuccessful



Example: fork() and exec()



```

/* Start a new process to do the job. */
cpid = fork();
// printf("process id is %d\n", cpid);

if(cpid < 0) {
    perror("fork");
    free(main_ptr); // clean-up
    return;
}

/* Check for who we are! */
if(cpid == 0) {
    /* We are the child! */
    execvp(main_ptr[0], main_ptr);
    perror("exec");
    free(main_ptr); // clean-up
    exit(127);
}

/* Have the parent wait for child to complete */
if(wait(&status) < 0)
    perror("wait");
// printf("wait result for process id %d is %d\n", cpid, status);

```

—fork() and exec() code snippet that Travis wrote back when he was an undergrad...

—John S. Quarterman, Abraham Silberschatz, and James L. Peterson, "4.2BSD and 4.3BSD as Examples of the UNIX System", Computing Surveys, Volume 17, Number 4, (December 1985), pages 379-418; translated to Japanese, Computer Science (BIT), Volume 18, Number 3, (1986), pages 175-213.


Terminating a Process

- There must (*at least should...*) be a means for a process to indicate its completion
- A batch job should include a HALT instruction (or something similar)
- An interactive process will terminate when a user, e.g., logs off, quits an app
- Again, *there is some discussion in the text on when/why processes terminate...*
 - *...normal completion*
 - *...timeout*
 - *...killed due to error/violation*
 - *...parent terminates*
 - *...explicitly requested (e.g., parent's request)*
 - *etc.*

Terminating a Process

- **Q:** What is the mechanism in UNIX to explicitly terminate a process?

HANDLING NON-RESPONDING & FROZEN APPLICATIONS



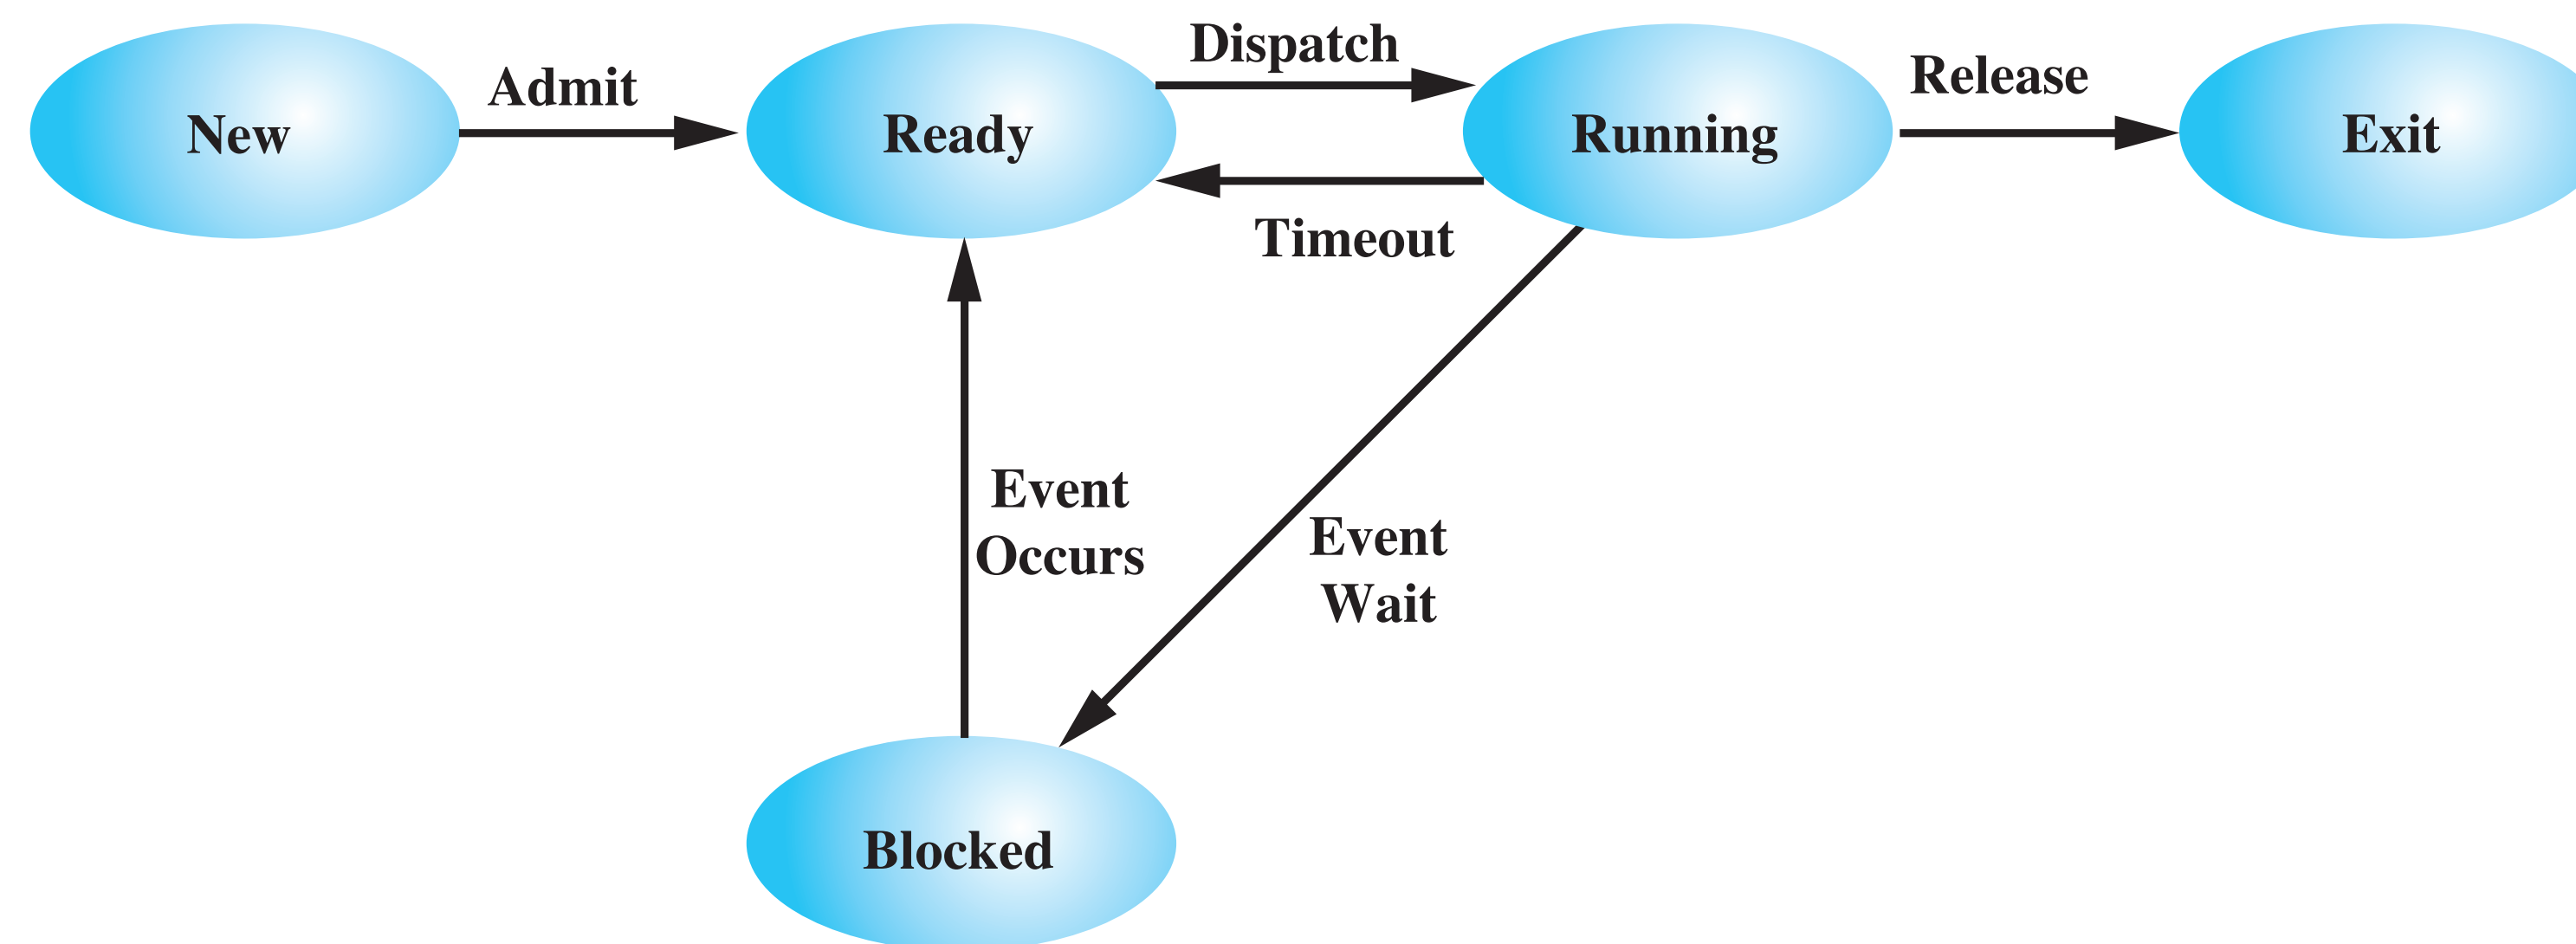
Terminating a Process - `kill()`

- **Q:** What is the mechanism in UNIX to explicitly terminate a process?
- `kill()` ...
 - ...enables terminating a process by specifying the **PID** and a **signal**
 - ...default at commandline is to send signal 15 (SIGTERM) — “*please terminate...*”
- Read up on **signals** and different types of signals...
 - ...signals are basically just software interrupts that can be sent to a program to indicate that an important event has occurred.
 - **Exercise:** run the following command in a terminal: `kill -l`

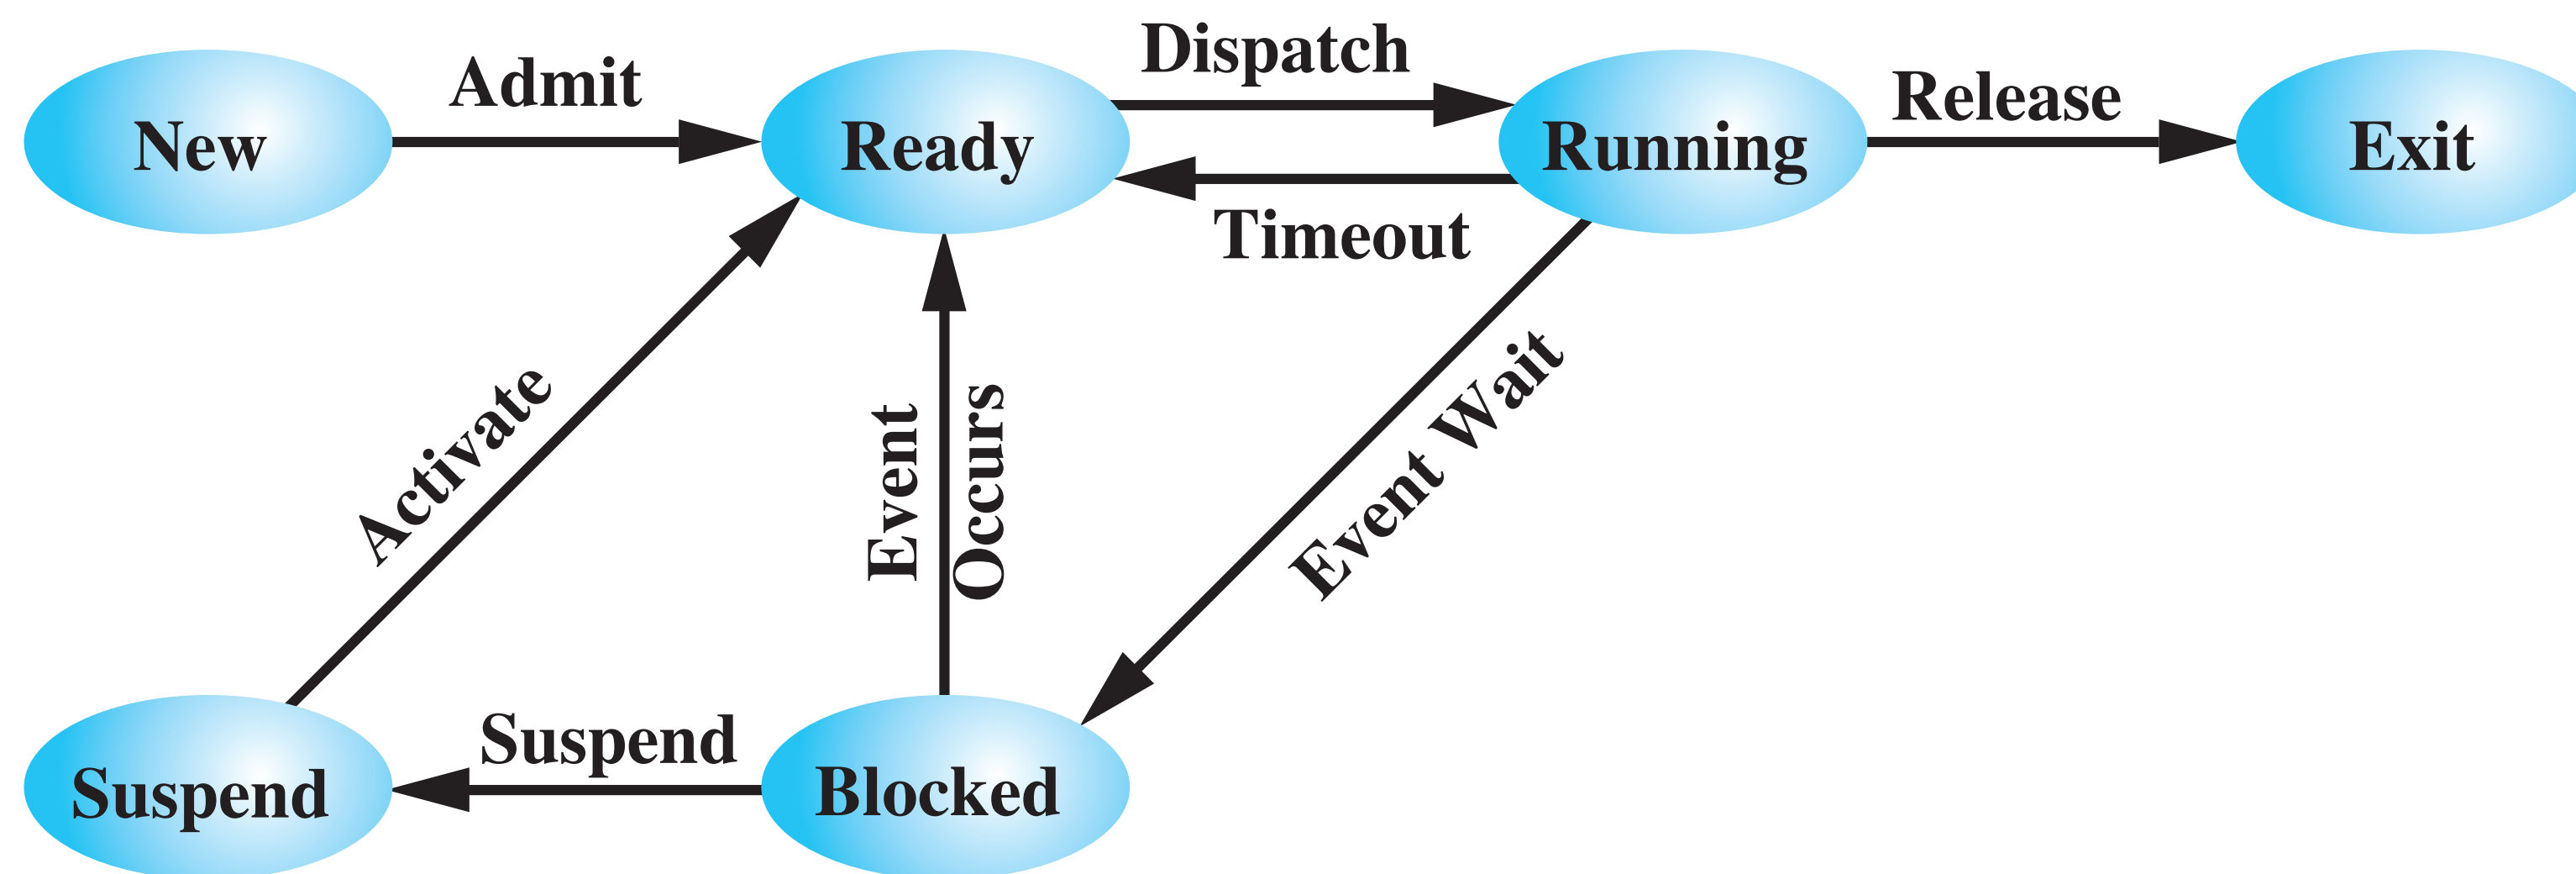
Suspended Processes & Swapping

- Swapping
 - Move parts of the process from main memory to disk to free up resources for other processes
 - OS will choose to swap blocked process(es) out to disk into a ***suspended state***
 - Swapping is an I/O operation; has the potential to make the problem worse...
 - on-system I/O is pretty fast though (relative to, e.g., network I/O)
 - Swapping *usually* improves performance
- **Q:** Why swap?
OS needs to release resources, debugging, periodic task, etc.
- **Q:** Why not make memory bigger?!
Cost; not more processes, usually bigger processes

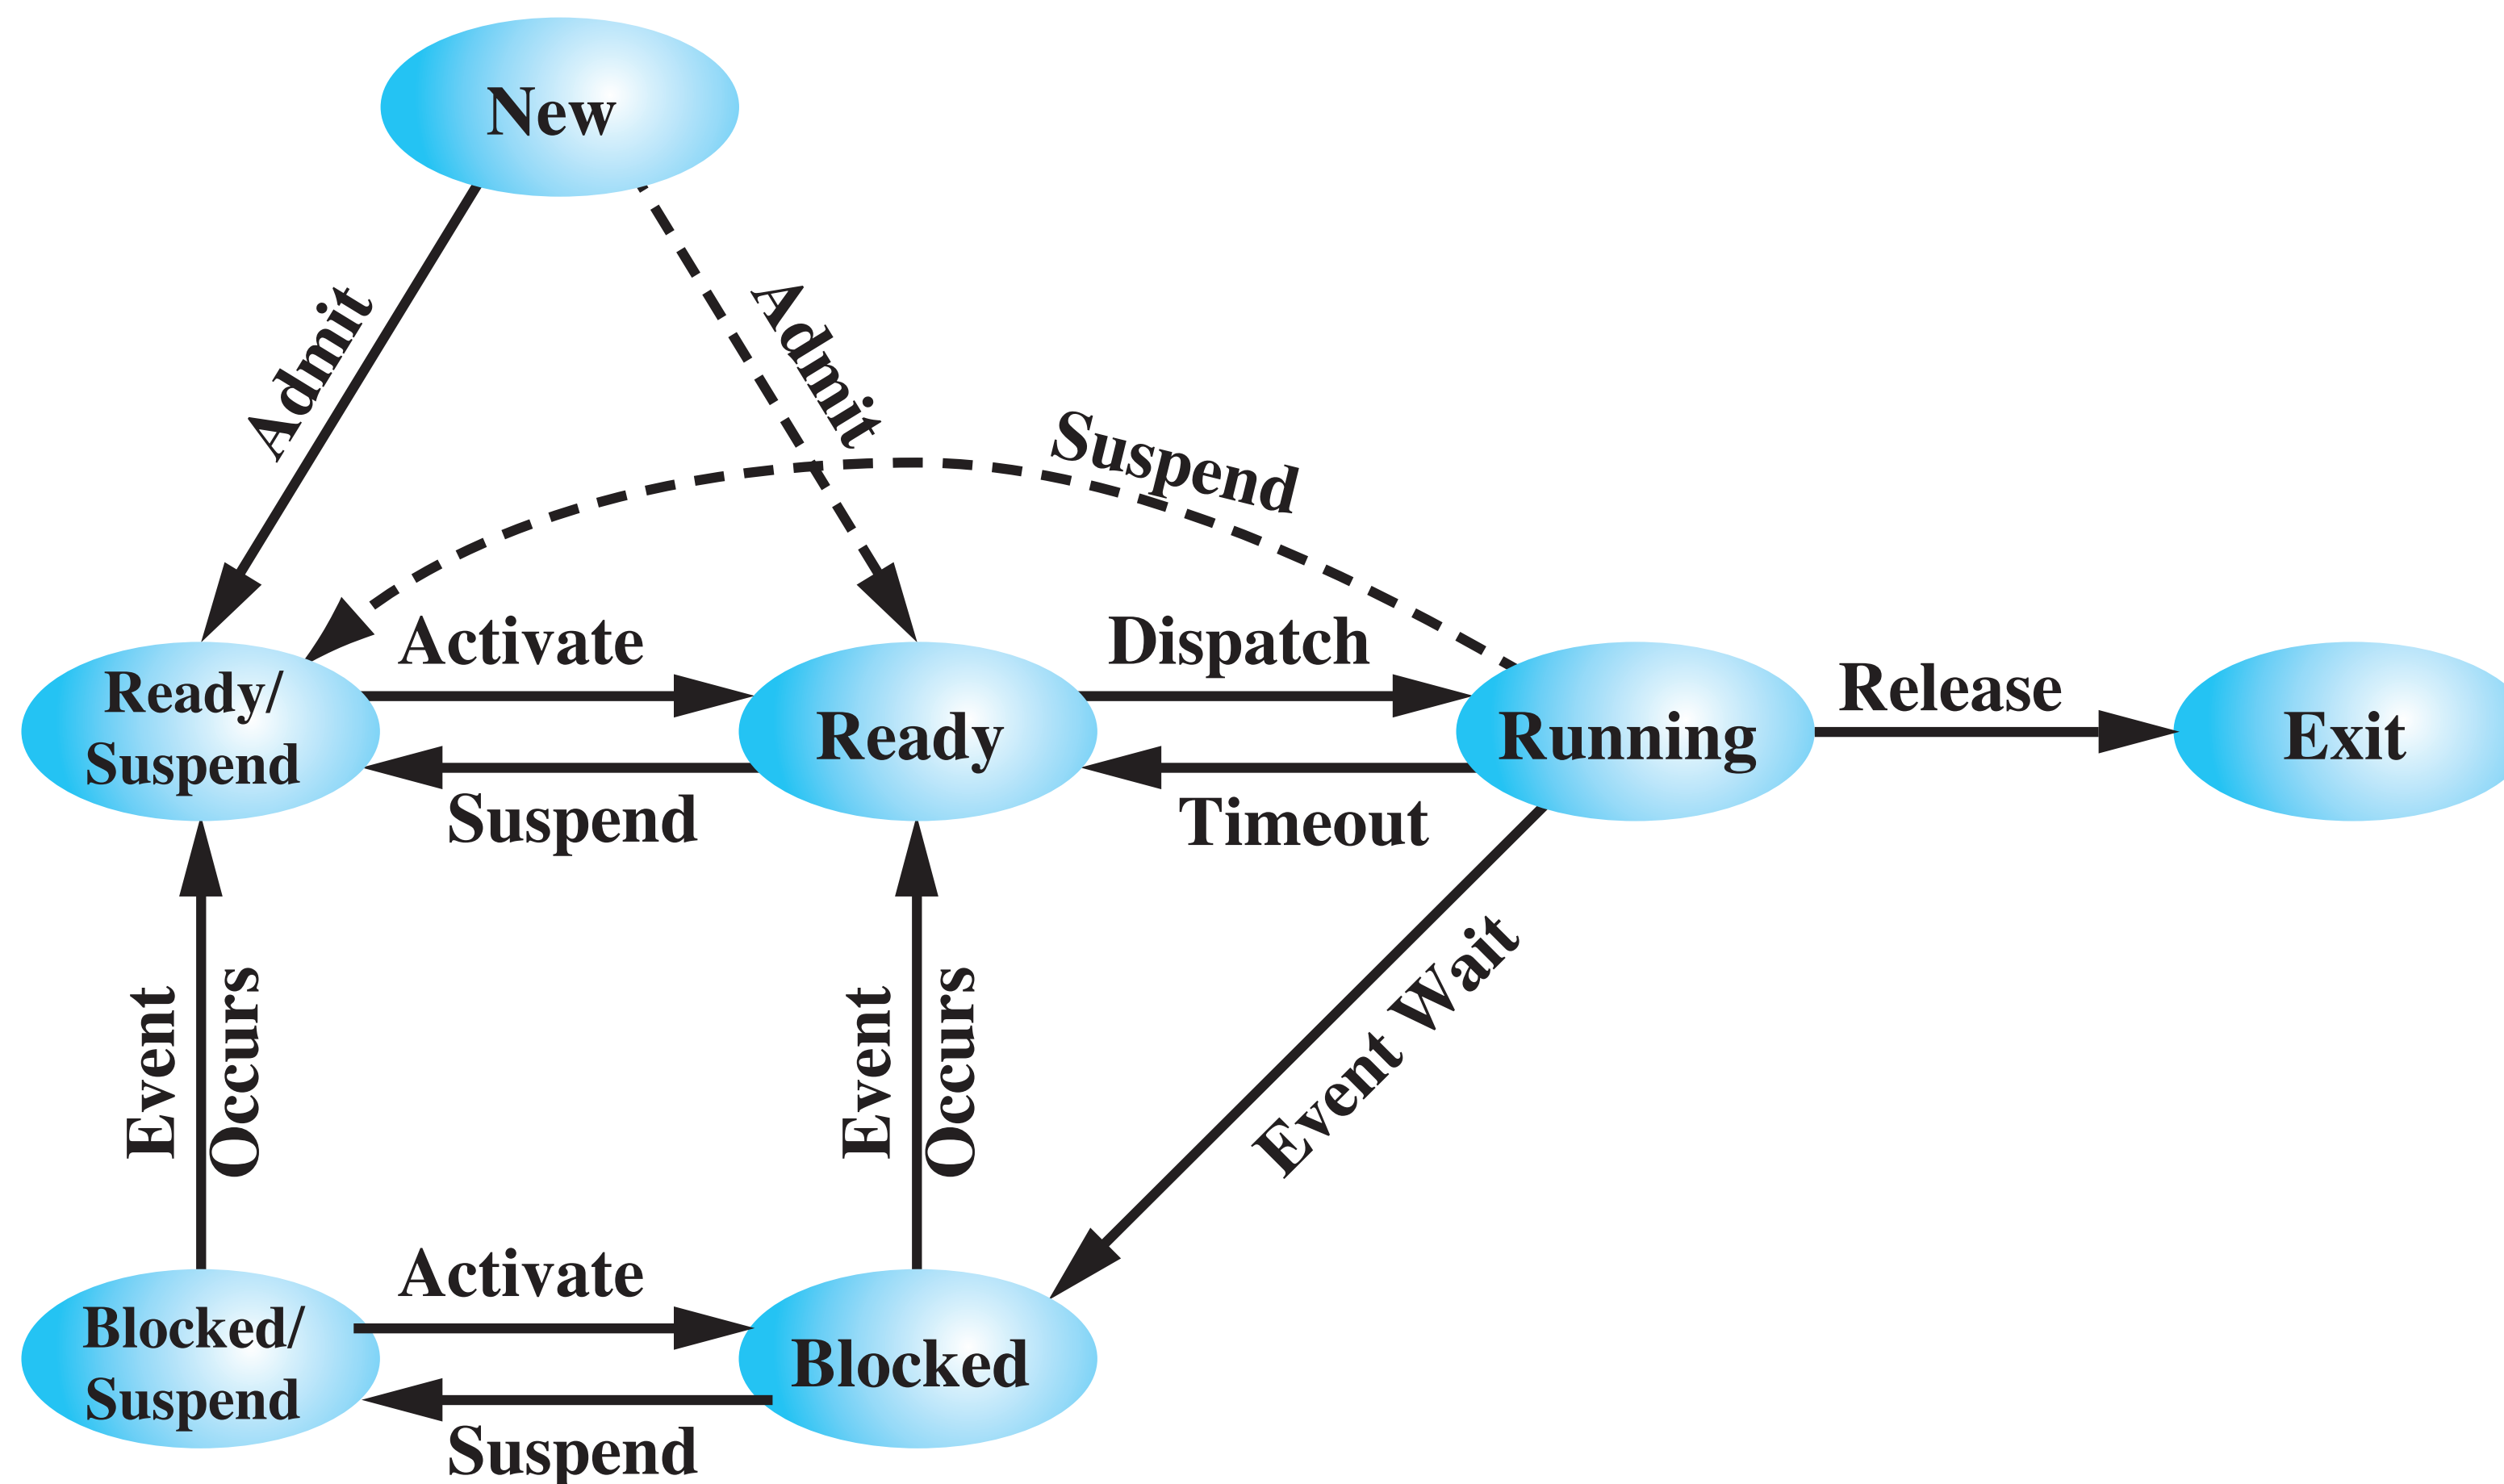
Suspended Processes & Swapping (**Before**)



Suspended Processes & Swapping (**After**)



Suspended Processes & Swapping (**After +Ready/Blocked**)



Summary: Process States

- It is **important for the OS to maintain information** about each process and its state **to manage system resources effectively**.
- Many, *many* **trade-offs** to consider...
 - More states (and more data structures for managing processes) are needed as we want to **reduce unnecessary operations**, such as searching through queues to identify ready processes, higher-priority processes, etc.
 - But more data structures means **more memory is used** to maintain this state information, and **more overhead operations** to move processes between different data structures when they are blocked, suspended, ready, etc.

Suggestion:

*know the differences between various process states,
and understand the reasons for (and pros/cons of) each.*

Fun: The `fork()` Bomb

- **WARNING:** Run at your own risk. I'm running on a VM...

```
#include <sys/types.h>
#include <unistd.h>

int main()
{
    while(1) {
        fork();
    }
    return 0;
}
```

—https://en.wikipedia.org/wiki/Fork_bomb

```
forkbomb() { forkbomb | forkbomb & } ; forkbomb
```