

# CSCI 460 Operating Systems

## Processes (Part III)

---

Professor Travis Peters

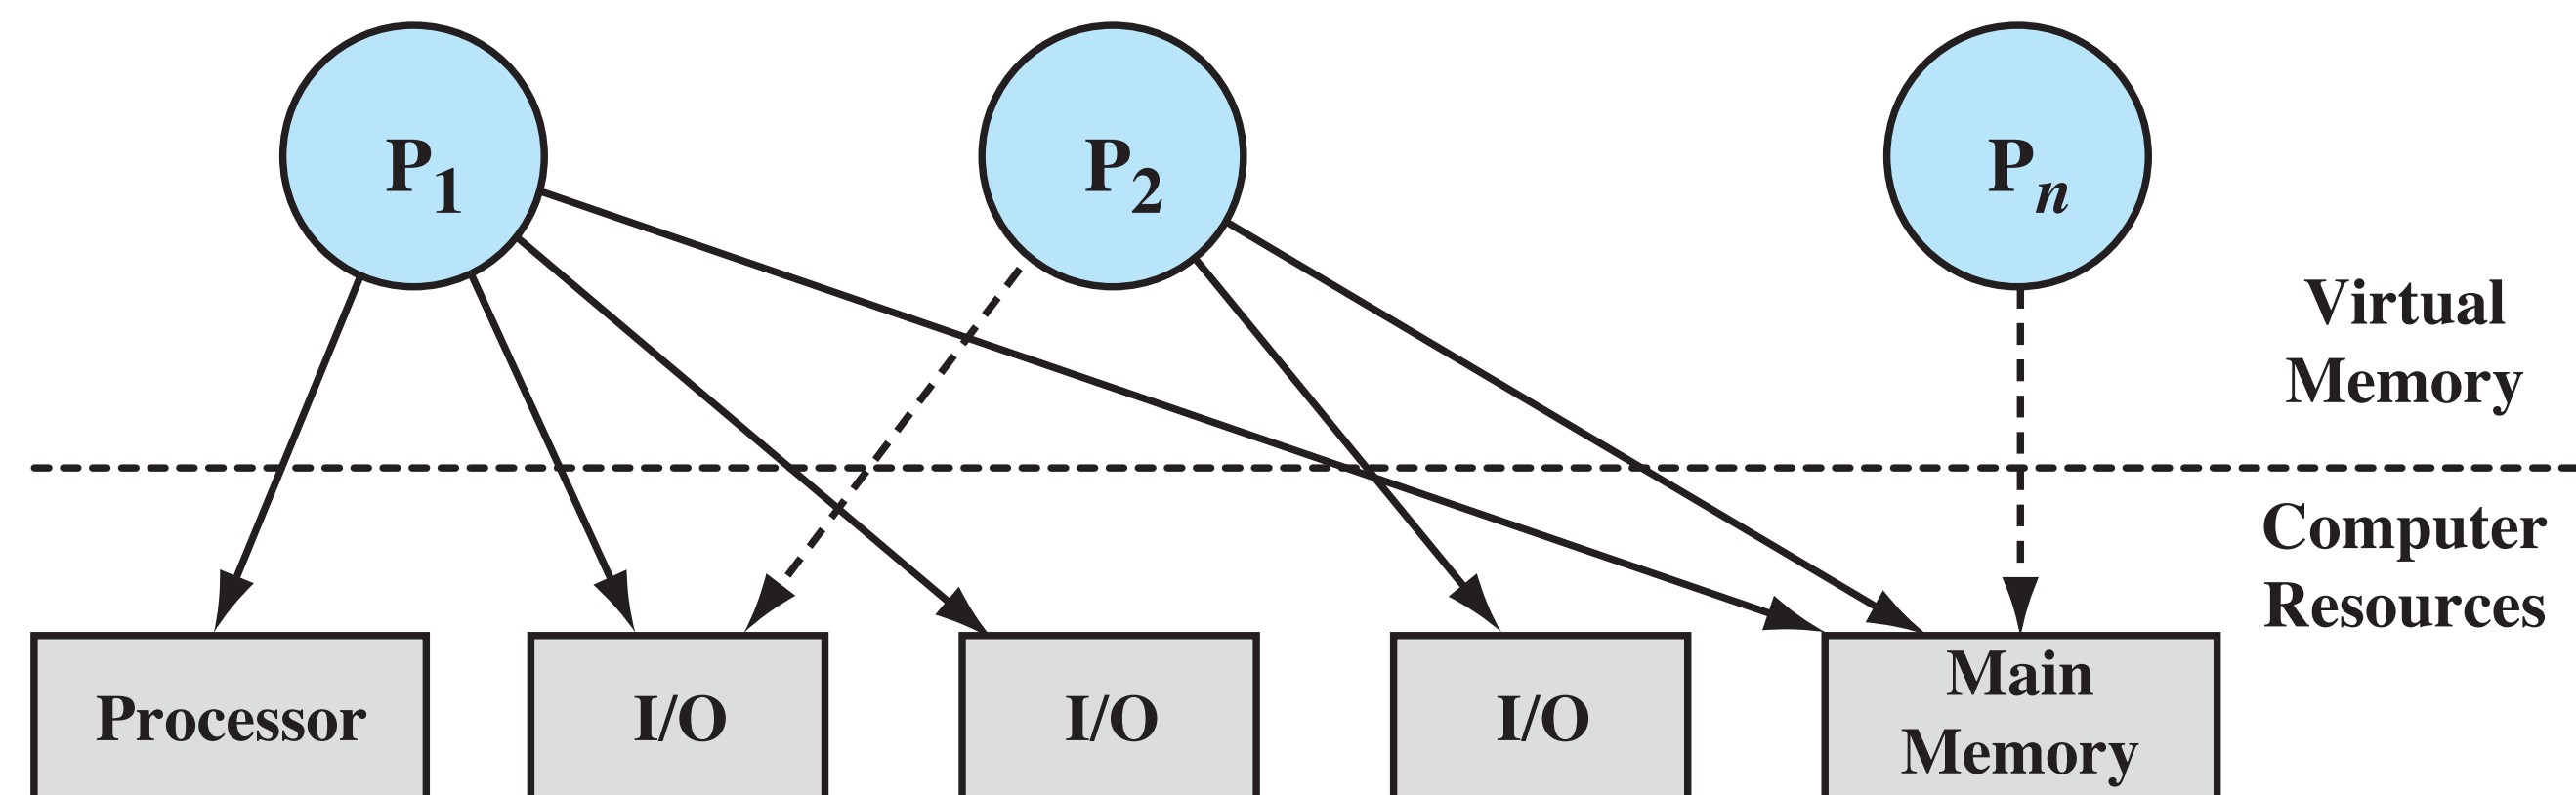
Fall 2019

*Some slides & figures adapted from Stallings instructor resources.*

*Some slides adapted from Adam Bates's F'18 CS423 course @ UIUC  
<https://courses.engr.illinois.edu/cs423/sp2018/schedule.html>*

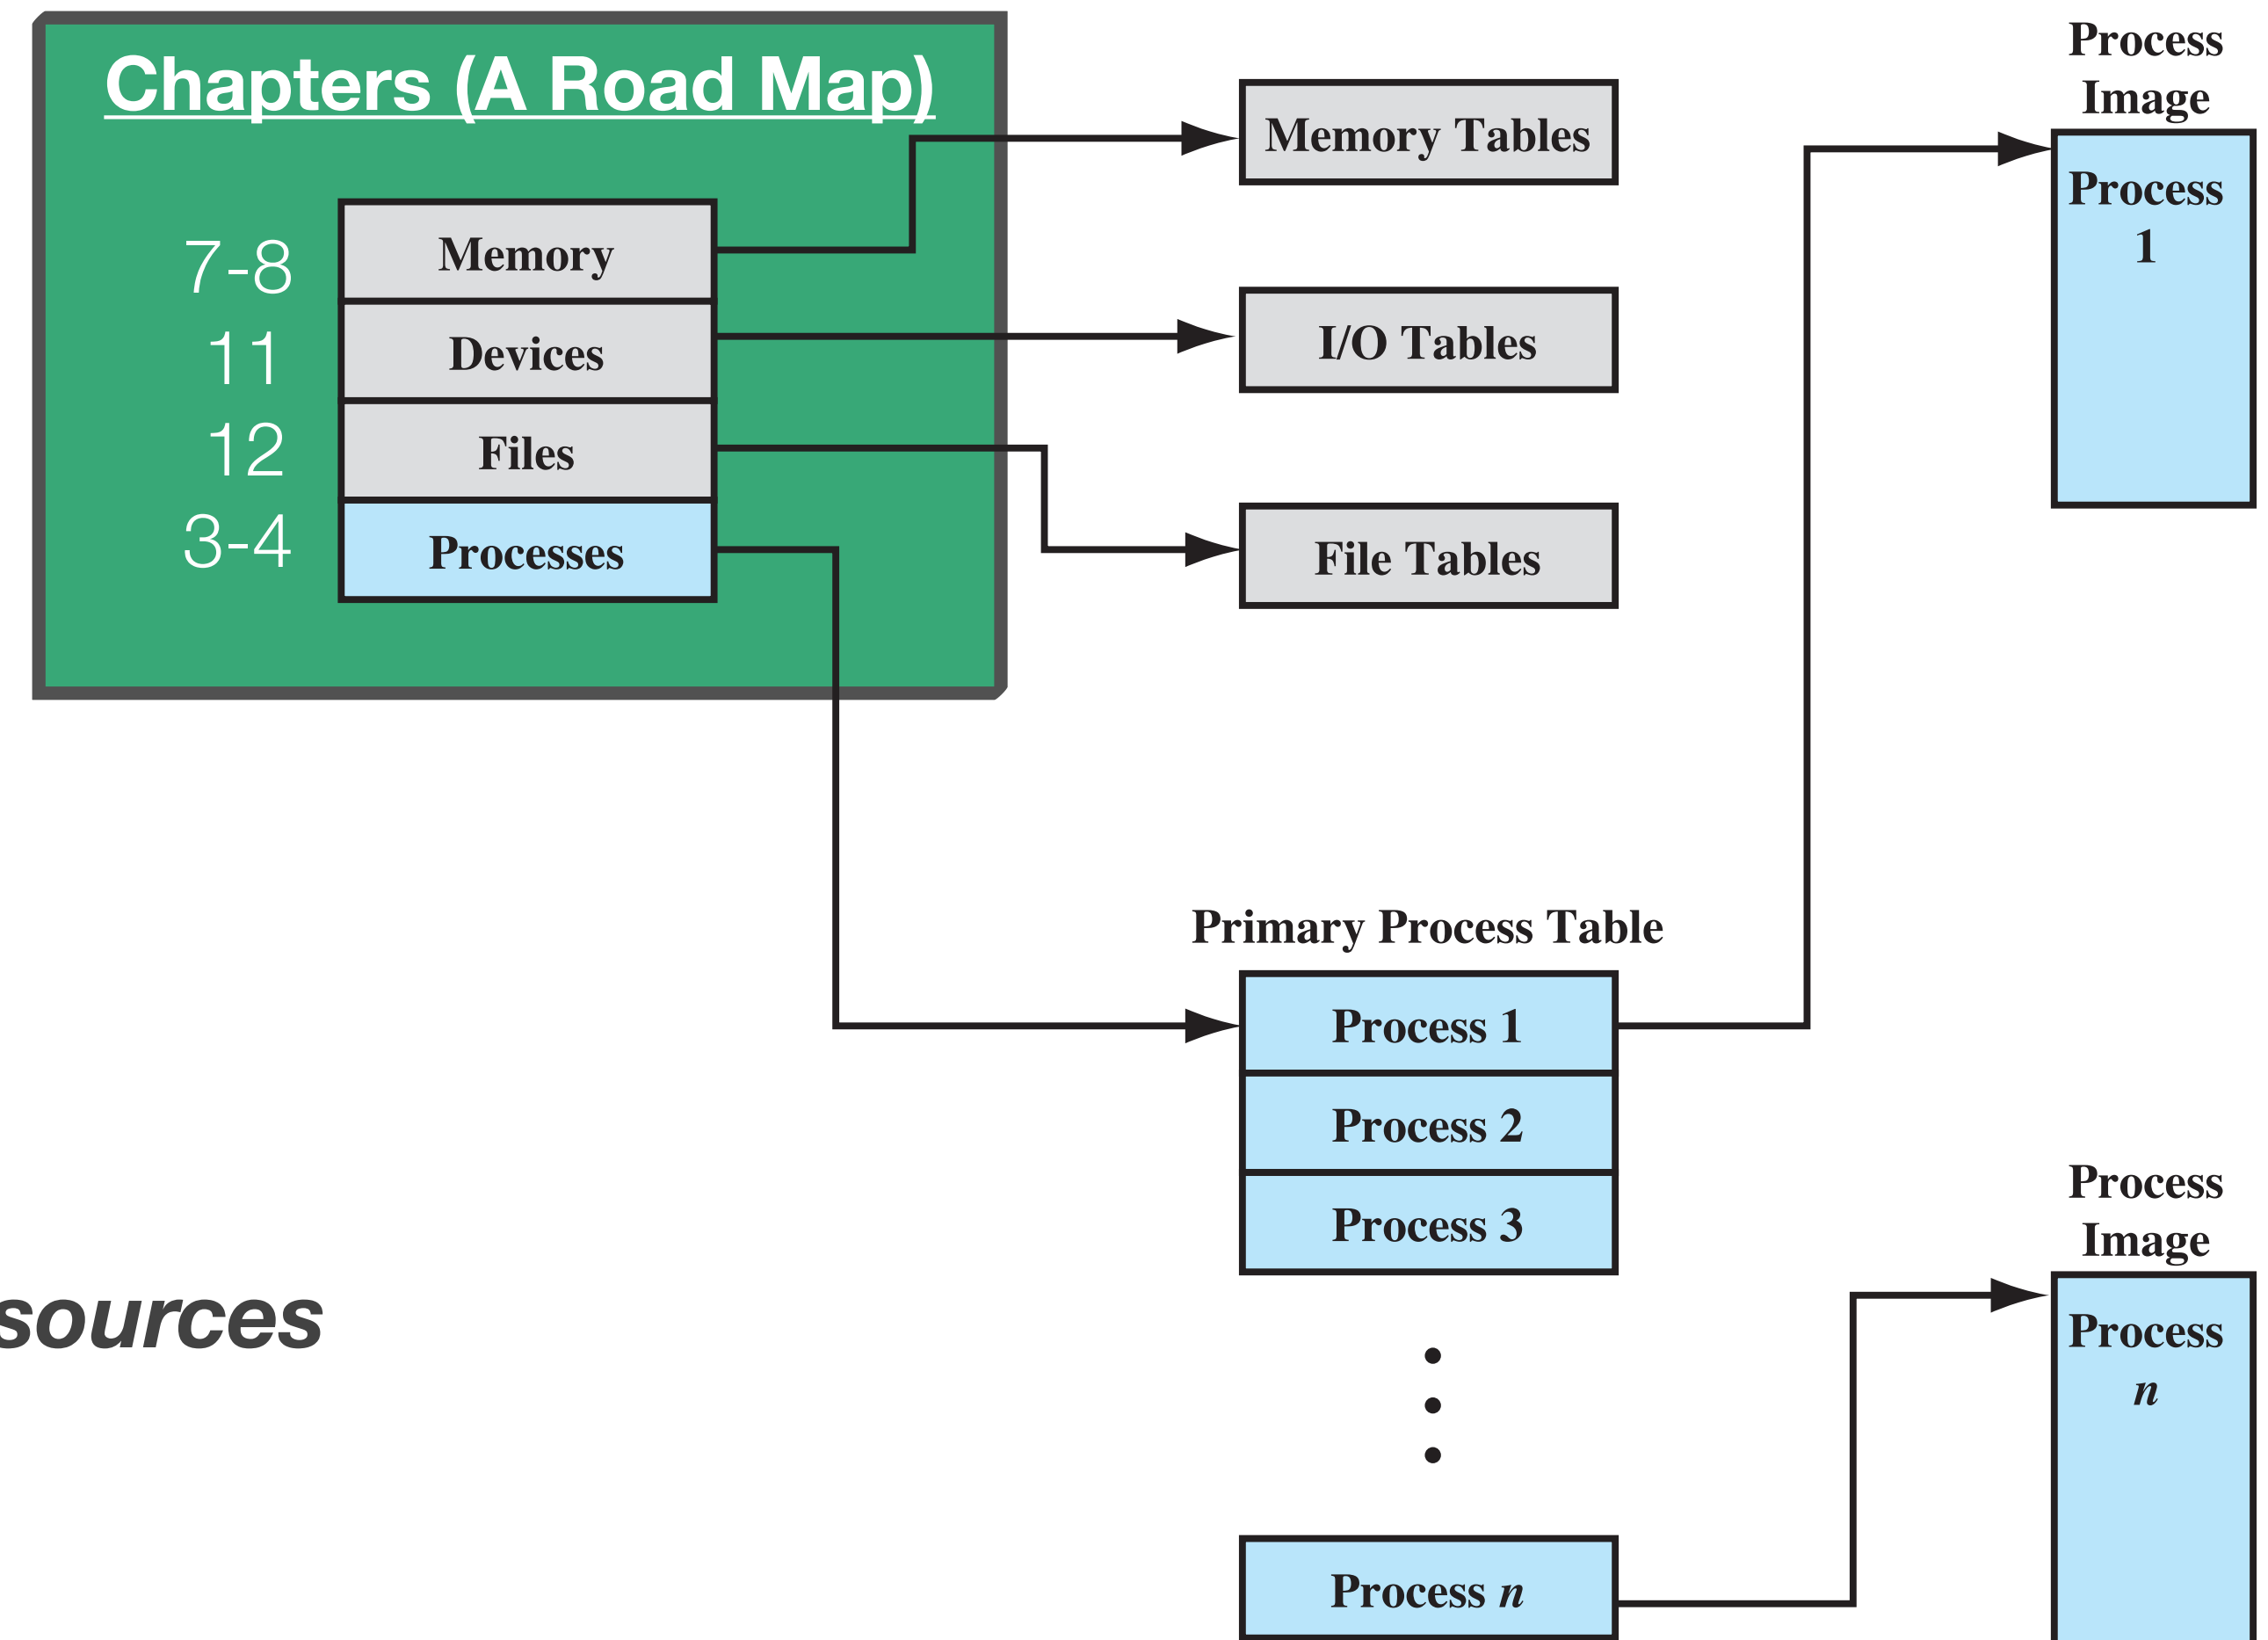
# The OS as the Manager of Processes & Resources

- **Basically:** The OS manages the use of system resources by processes.
- **Q:** But ***what information*** does the OS need to control processes and manage system resources for them?



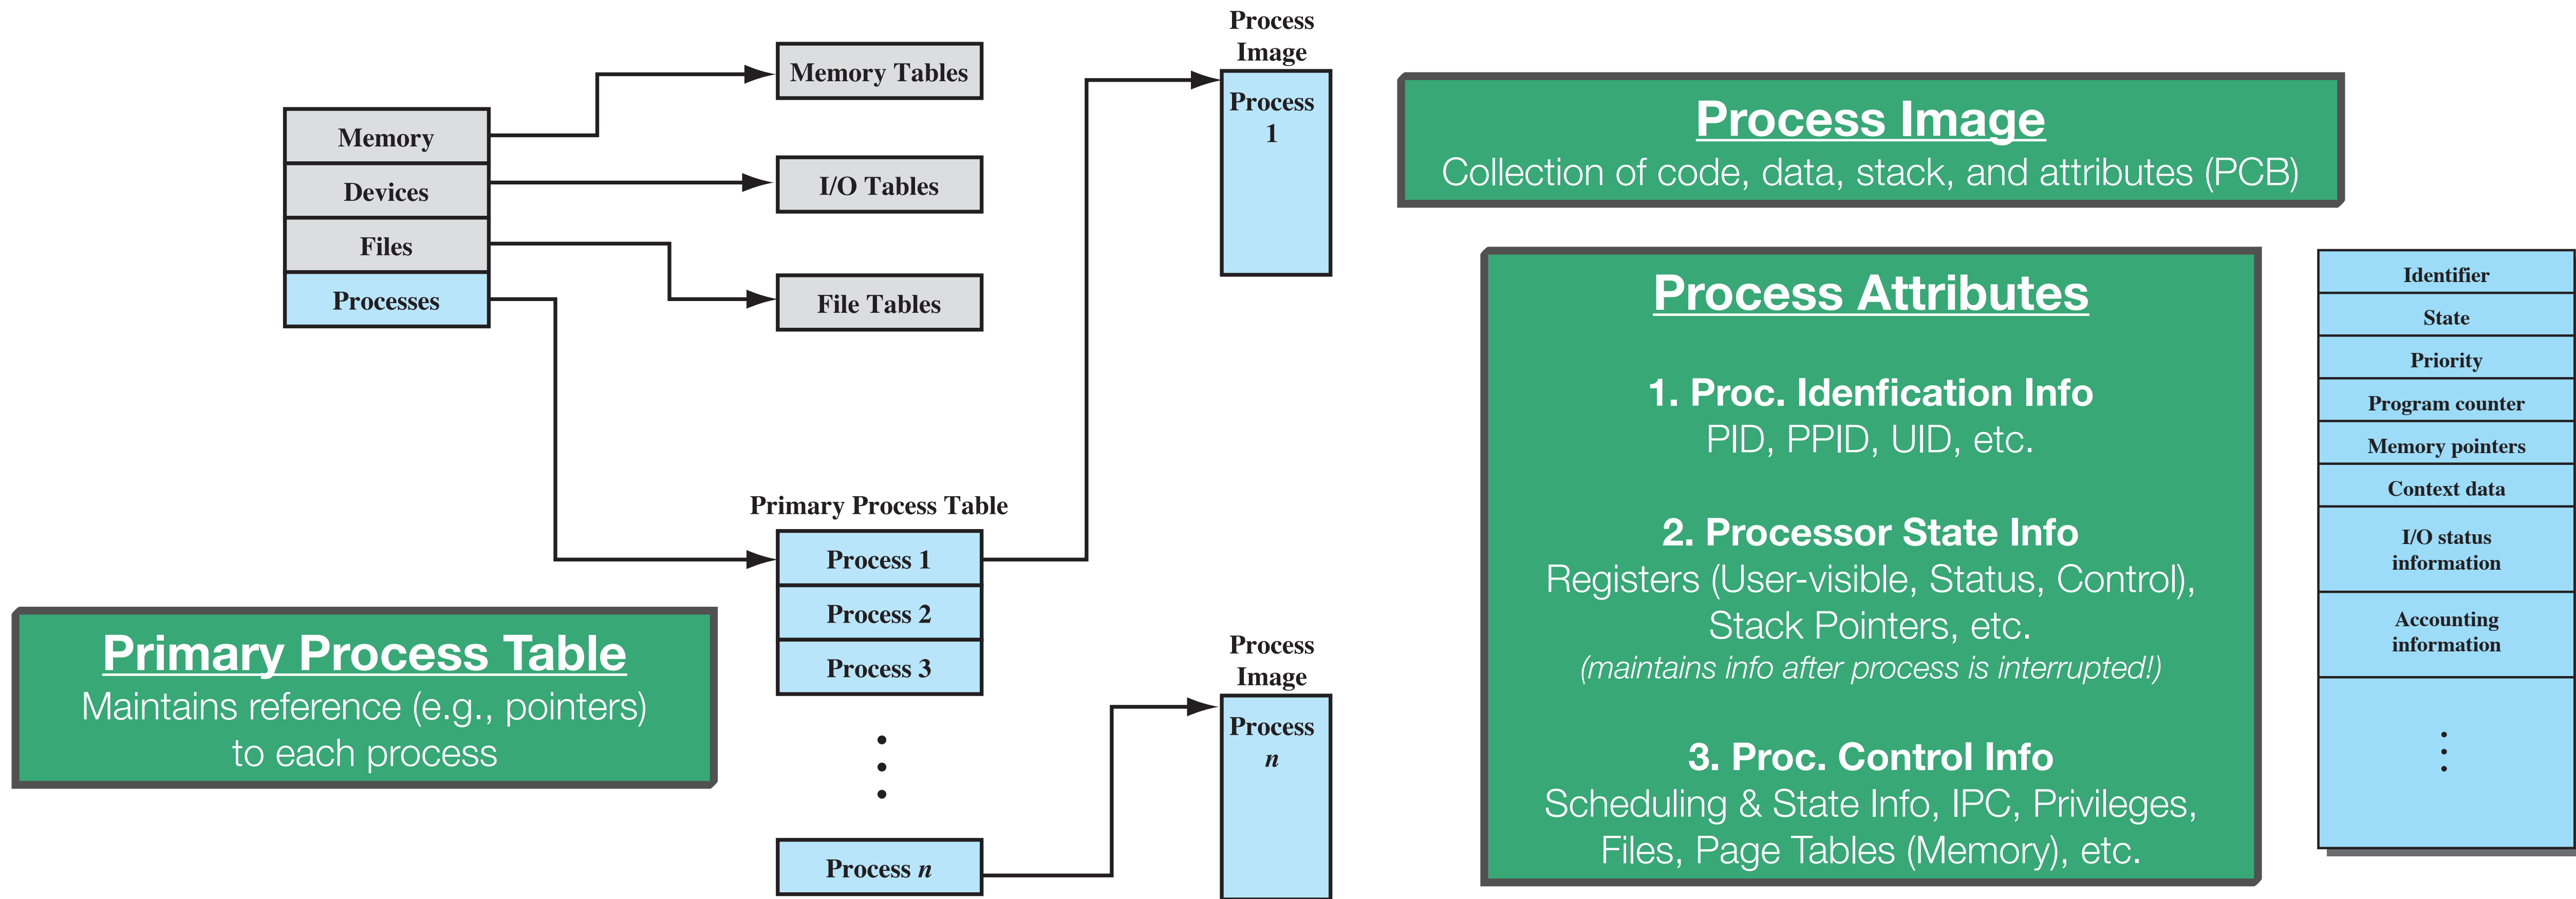
# The OS as the Manager of Processes & Resources

- OS Control Structures (*tables...*)
- Specifics vary, but generally we have four types of tables →
- **NOTE: Tables are Interconnected**
  - The tables all cross-reference each other  
*e.g., P1 has access to files in the file table*
- **NOTE: Need Initial Configuration**
  - The OS must have access to relevant ***information about the system and its resources*** in order to create these tables.
    - *How much main memory exists?*
    - *What I/O devices exist? How to access them?*
    - *etc.*



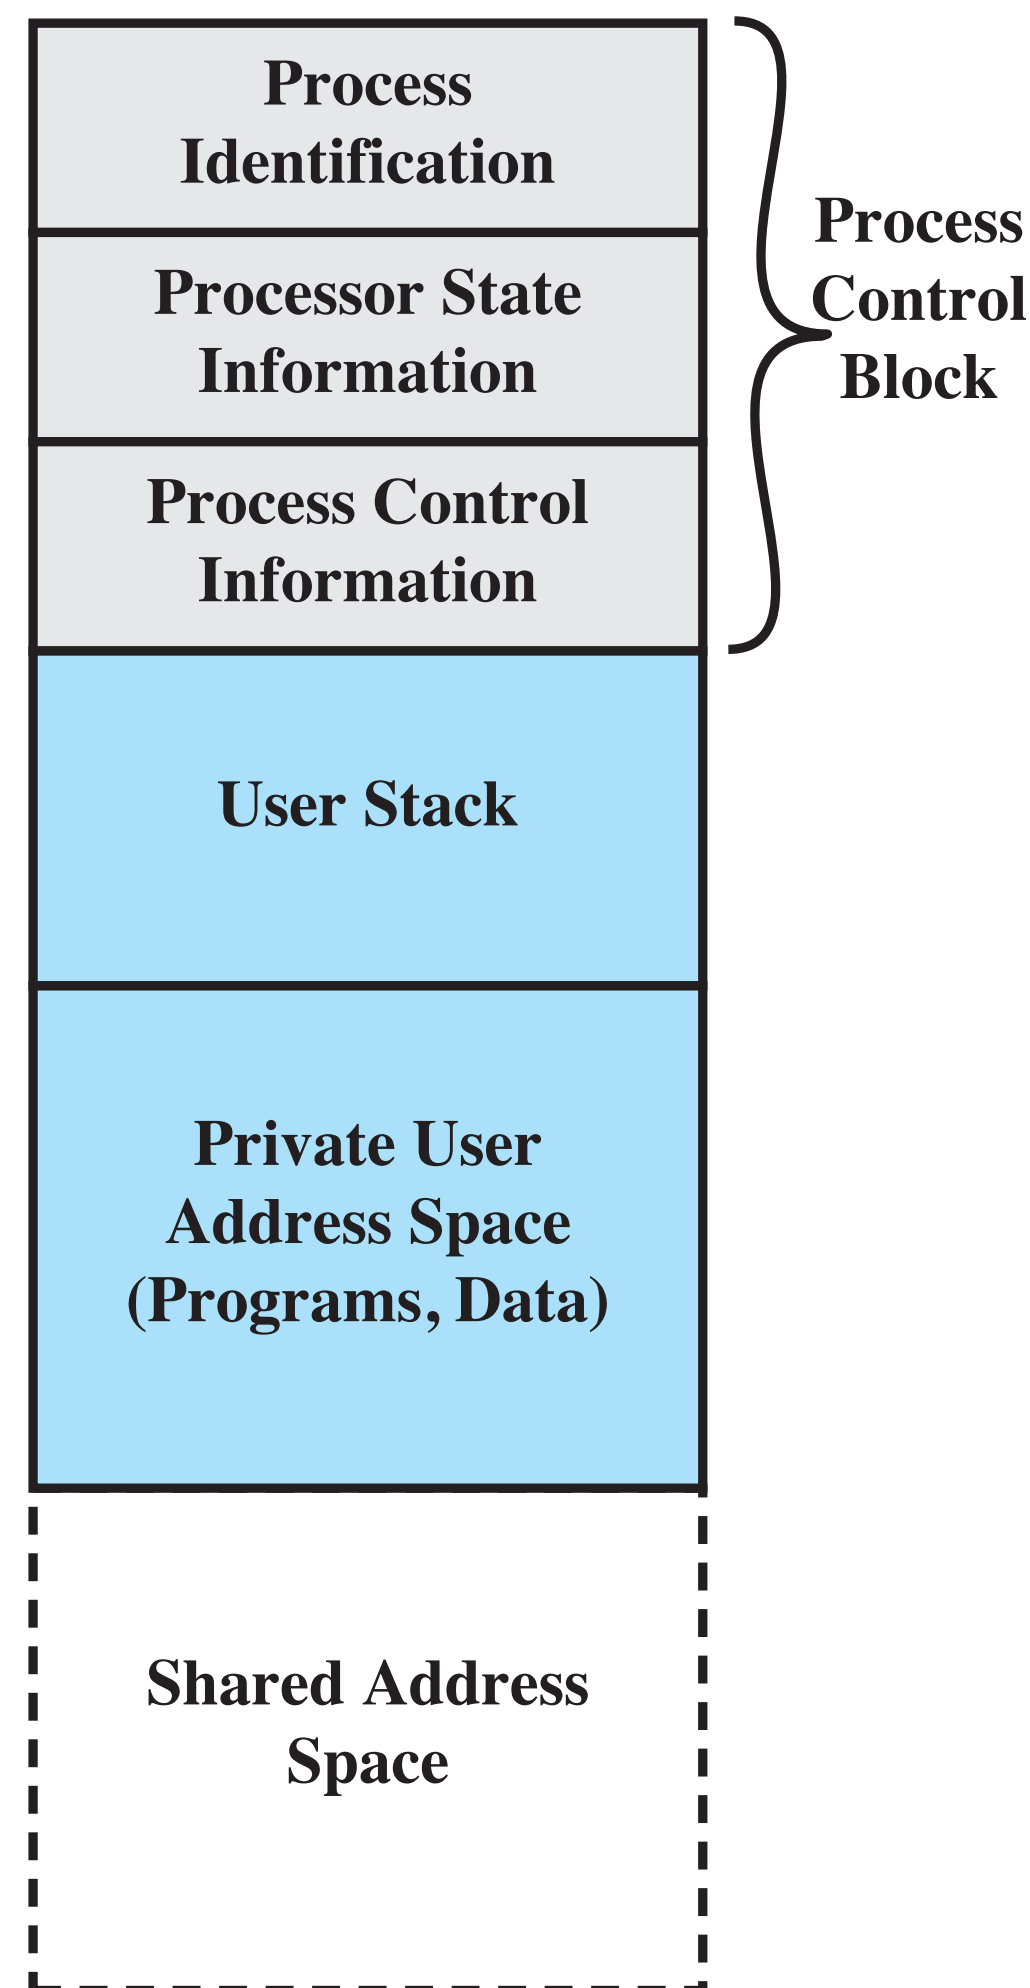
# The OS as the Manager of Processes & Resources

To manage and control a process the OS must know:



—W. Stallings (2018) Operating Systems: Internals and Design Principles (9th Edition).

# Processes in Memory



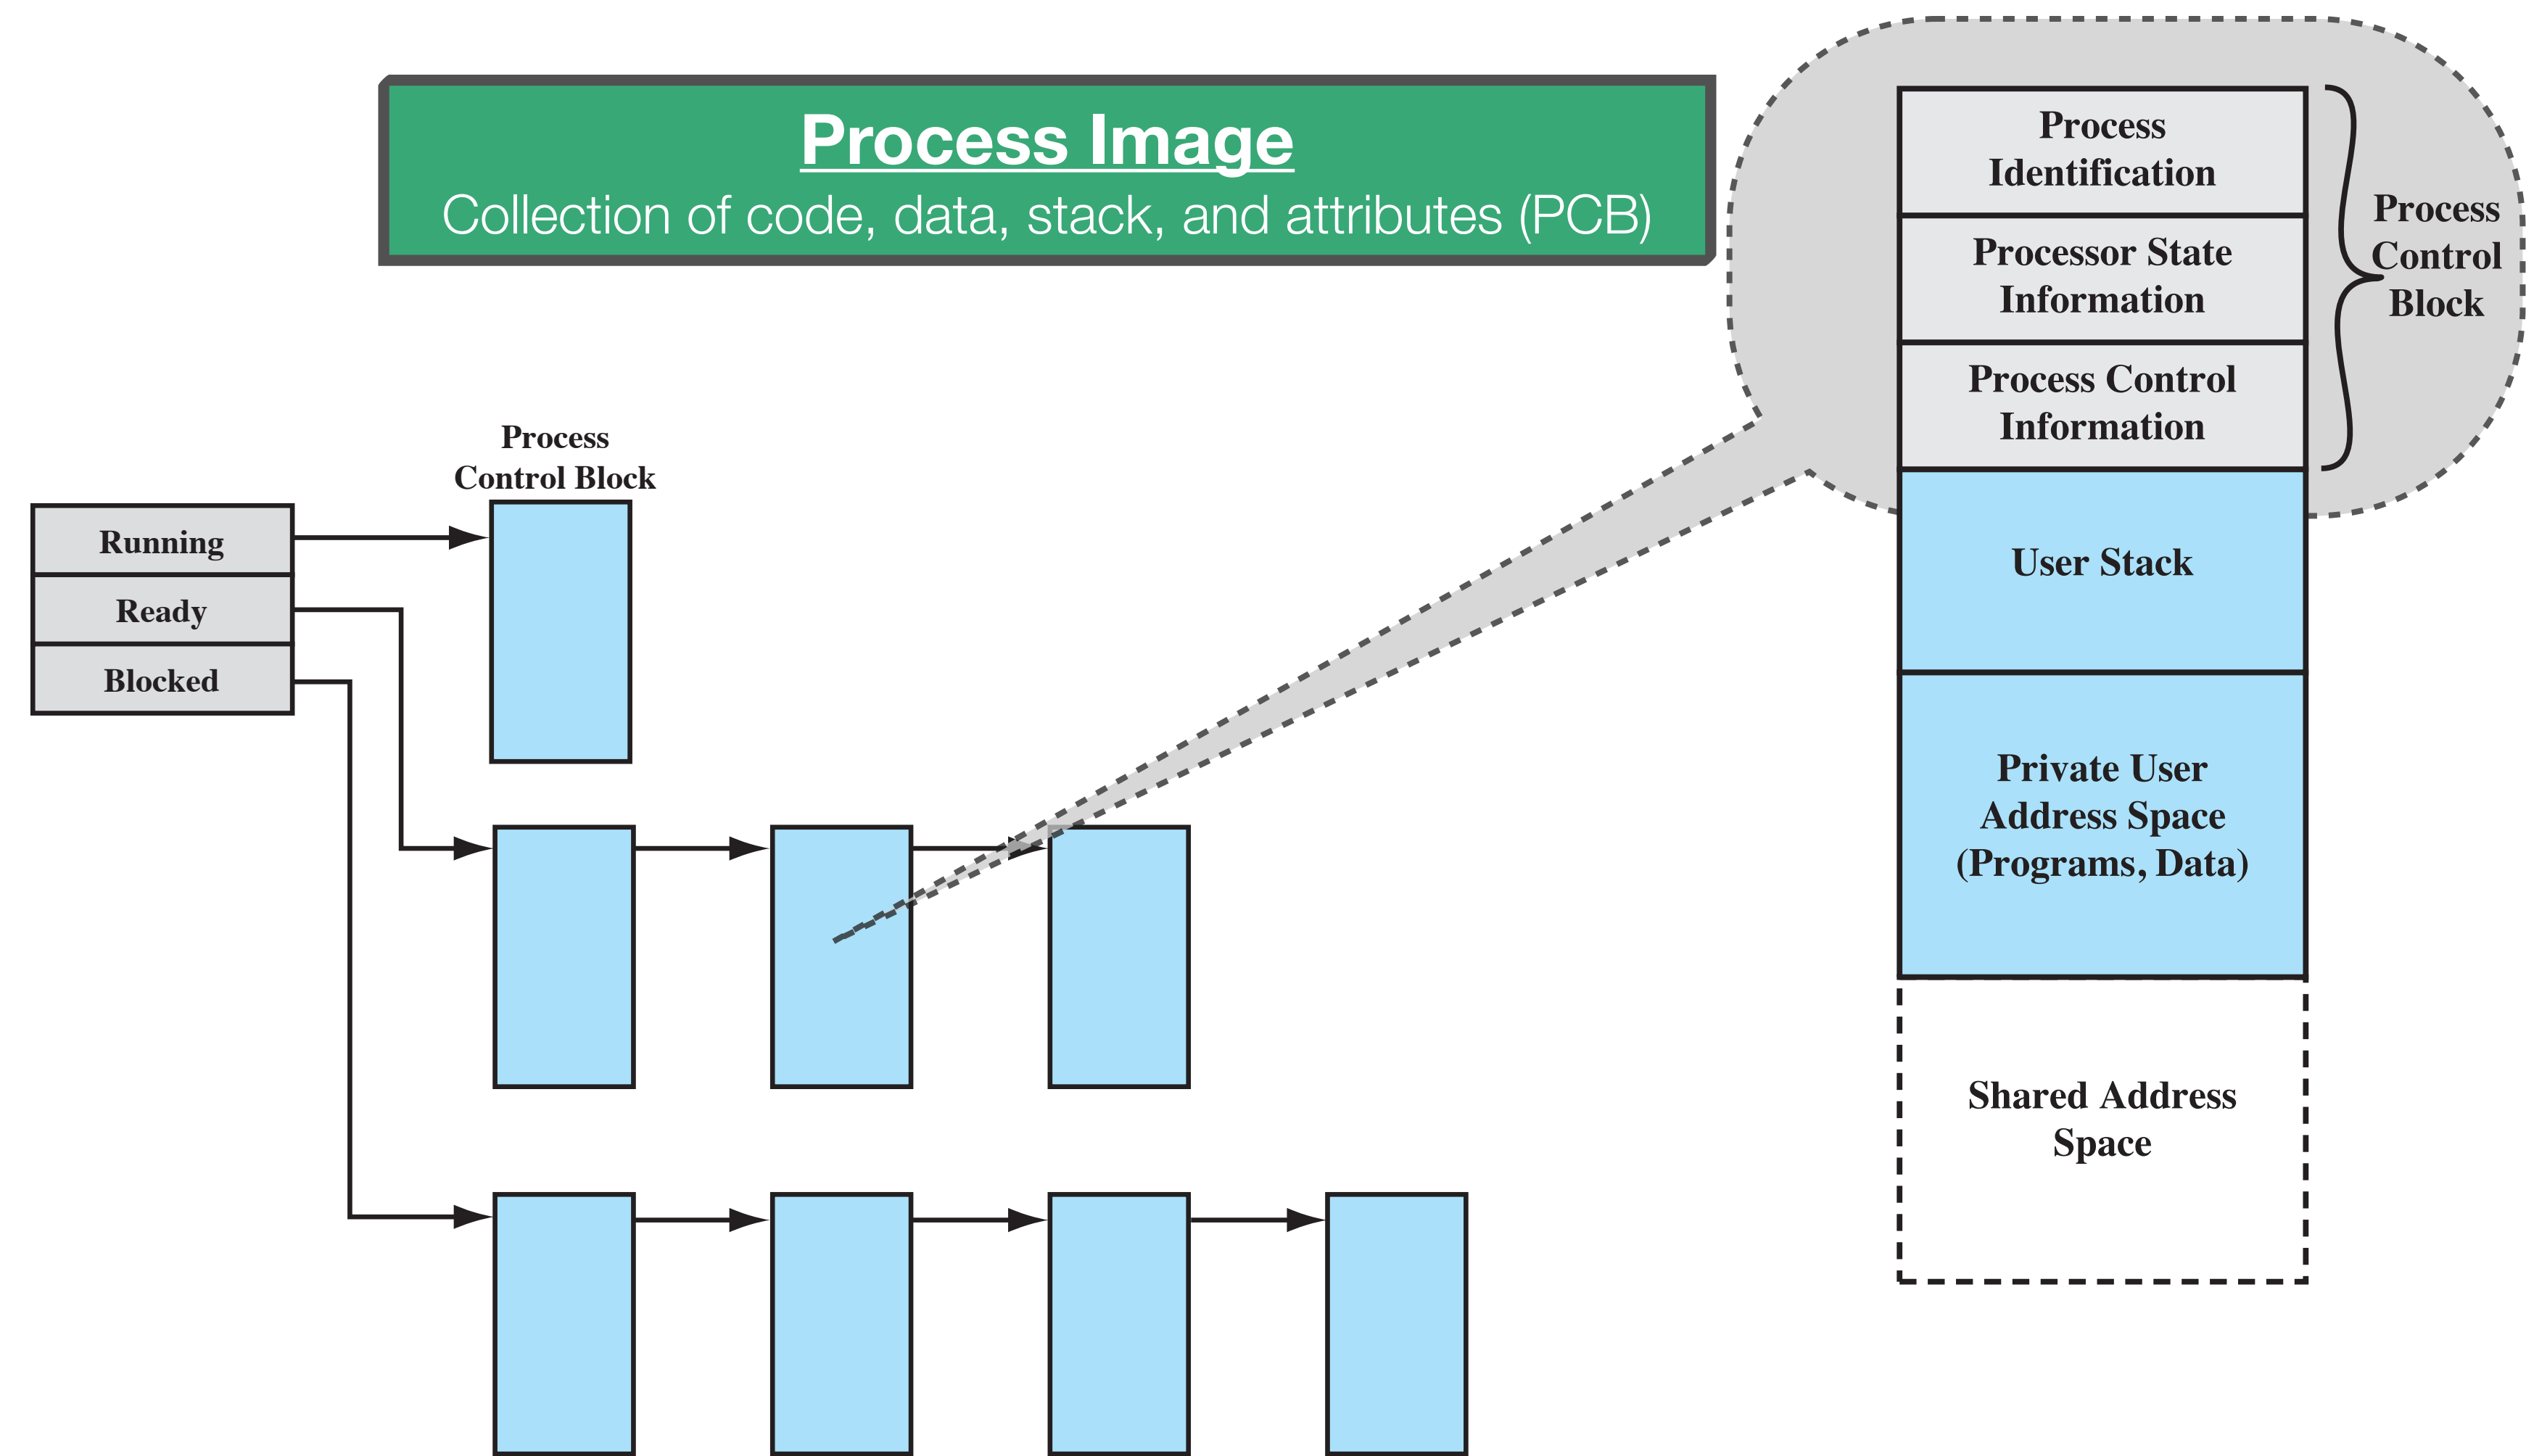
## Process Image

Collection of code, data, stack, and attributes (PCB)

*More specifically...*

- memory
- open streams/files
- devices, including abstract ones like windows
- links to condition handlers (signals)
- processor registers (single thread)
- process identification
- process state - including waiting information
- priority
- owner
- which processor
- links to other processes (parent, children)
- process group
- resource limits/usage
- access rights
- process result - may be waited for by another process
- ...

# Processes in Memory & Process List Structures





# Summary: The Role of the PCB

---

- PCB is the most important data structure in a modern OS
  - contains ALL of the information about a process that is needed by the OS
  - Blocks are read and/or modified by virtually *every* module in the OS
  - Defines the state of the OS
- **Q:** What could go wrong?!

Difficulty is ***not access***, but ***protection***

- A bug in a single OS routine could damage PCBs, which could destroy the system's ability to manage the affected processes
- A design change in the structure or semantics of the PCB could affect many modules in the OS
- Must tightly control access to the PCB...  
→ *all PCB changes happen through a designated handler routine*

# Privilege & Modes of Execution

---

- Processors support at least two modes of execution:  
***privileged*** (*kernel mode / system mode*) VS. ***non-privileged*** (*user mode*)
  - Sometimes referred to as ***rings***.
  - kernel mode = complete control over/access to processor, registers, memory
- **Q:** Why do we need different modes of execution?
- *Things* running in non-privileged mode should not...
  - be able to **access arbitrary memory** (e.g., other processes, I/O devices)
  - be able to **execute certain instructions** (e.g., manipulating control-flow registers)



# Privilege & Modes of Execution

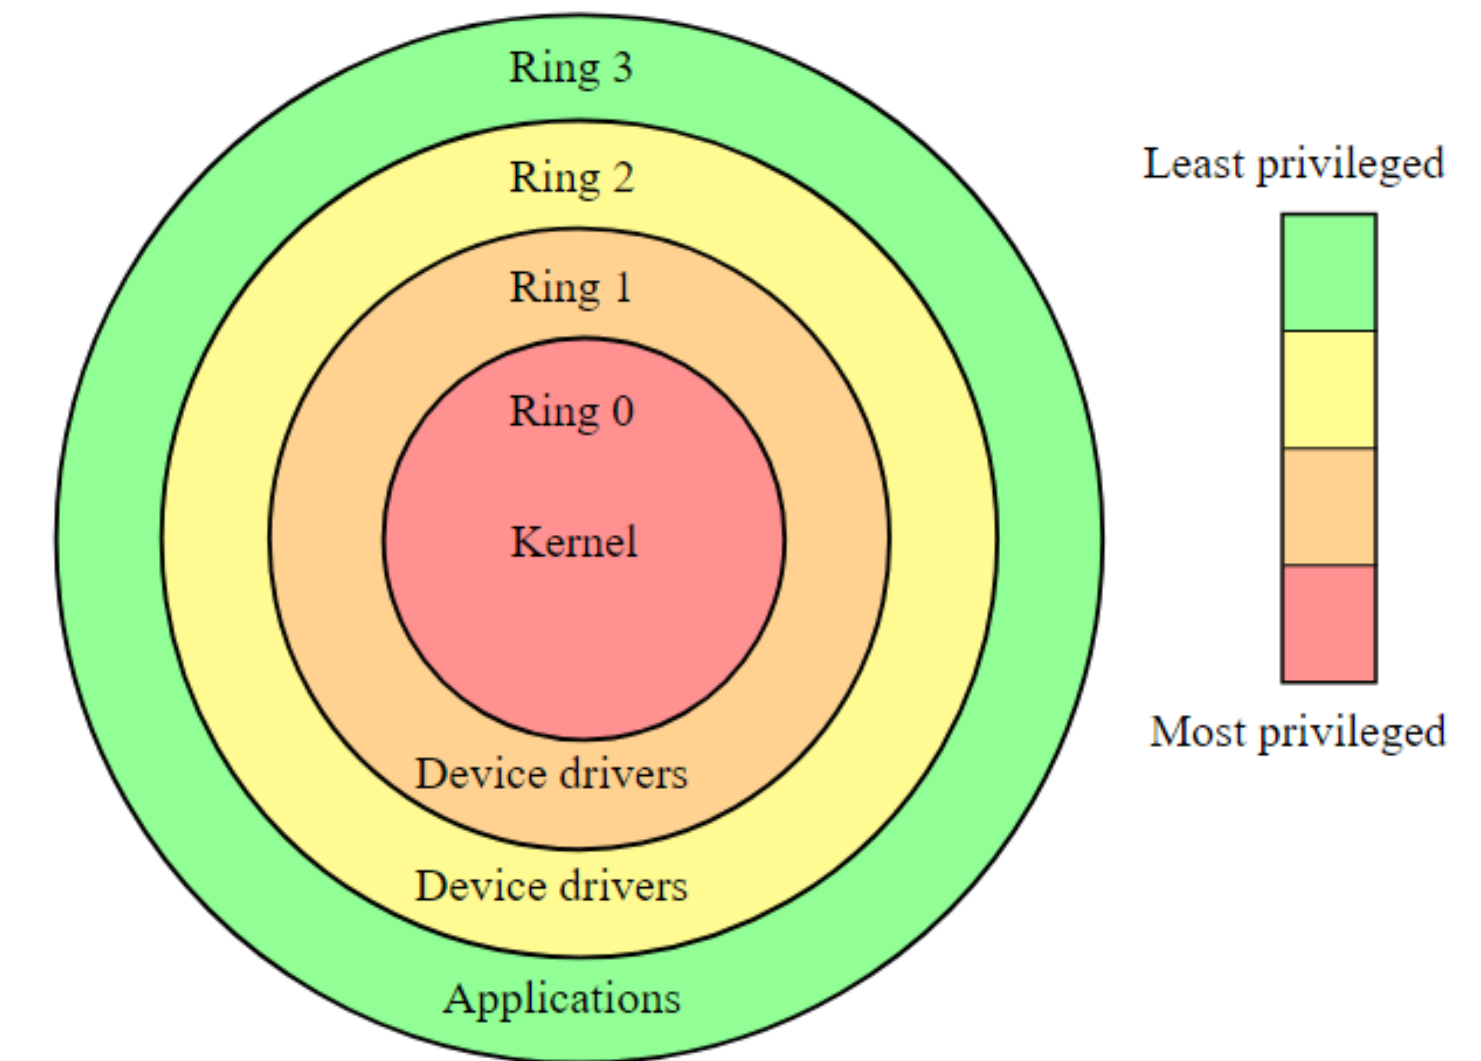
- **Q:** How does the processor know which mode to execute code in?

**Usually a single bit can be used  $\rightarrow$  0/1 == privileged/non-privileged**

*Example:*

- *Intel uses 2-bit flag (CPL = current privilege level)*
  - *00 = most privileged = kernel*
  - *01 = ?*
  - *10 = ?*
  - *11 = least privileged = user mode*
- **Q:** How is the mode changed?

**Short Answer: Certain instructions clear/set bits**



# Process Switching

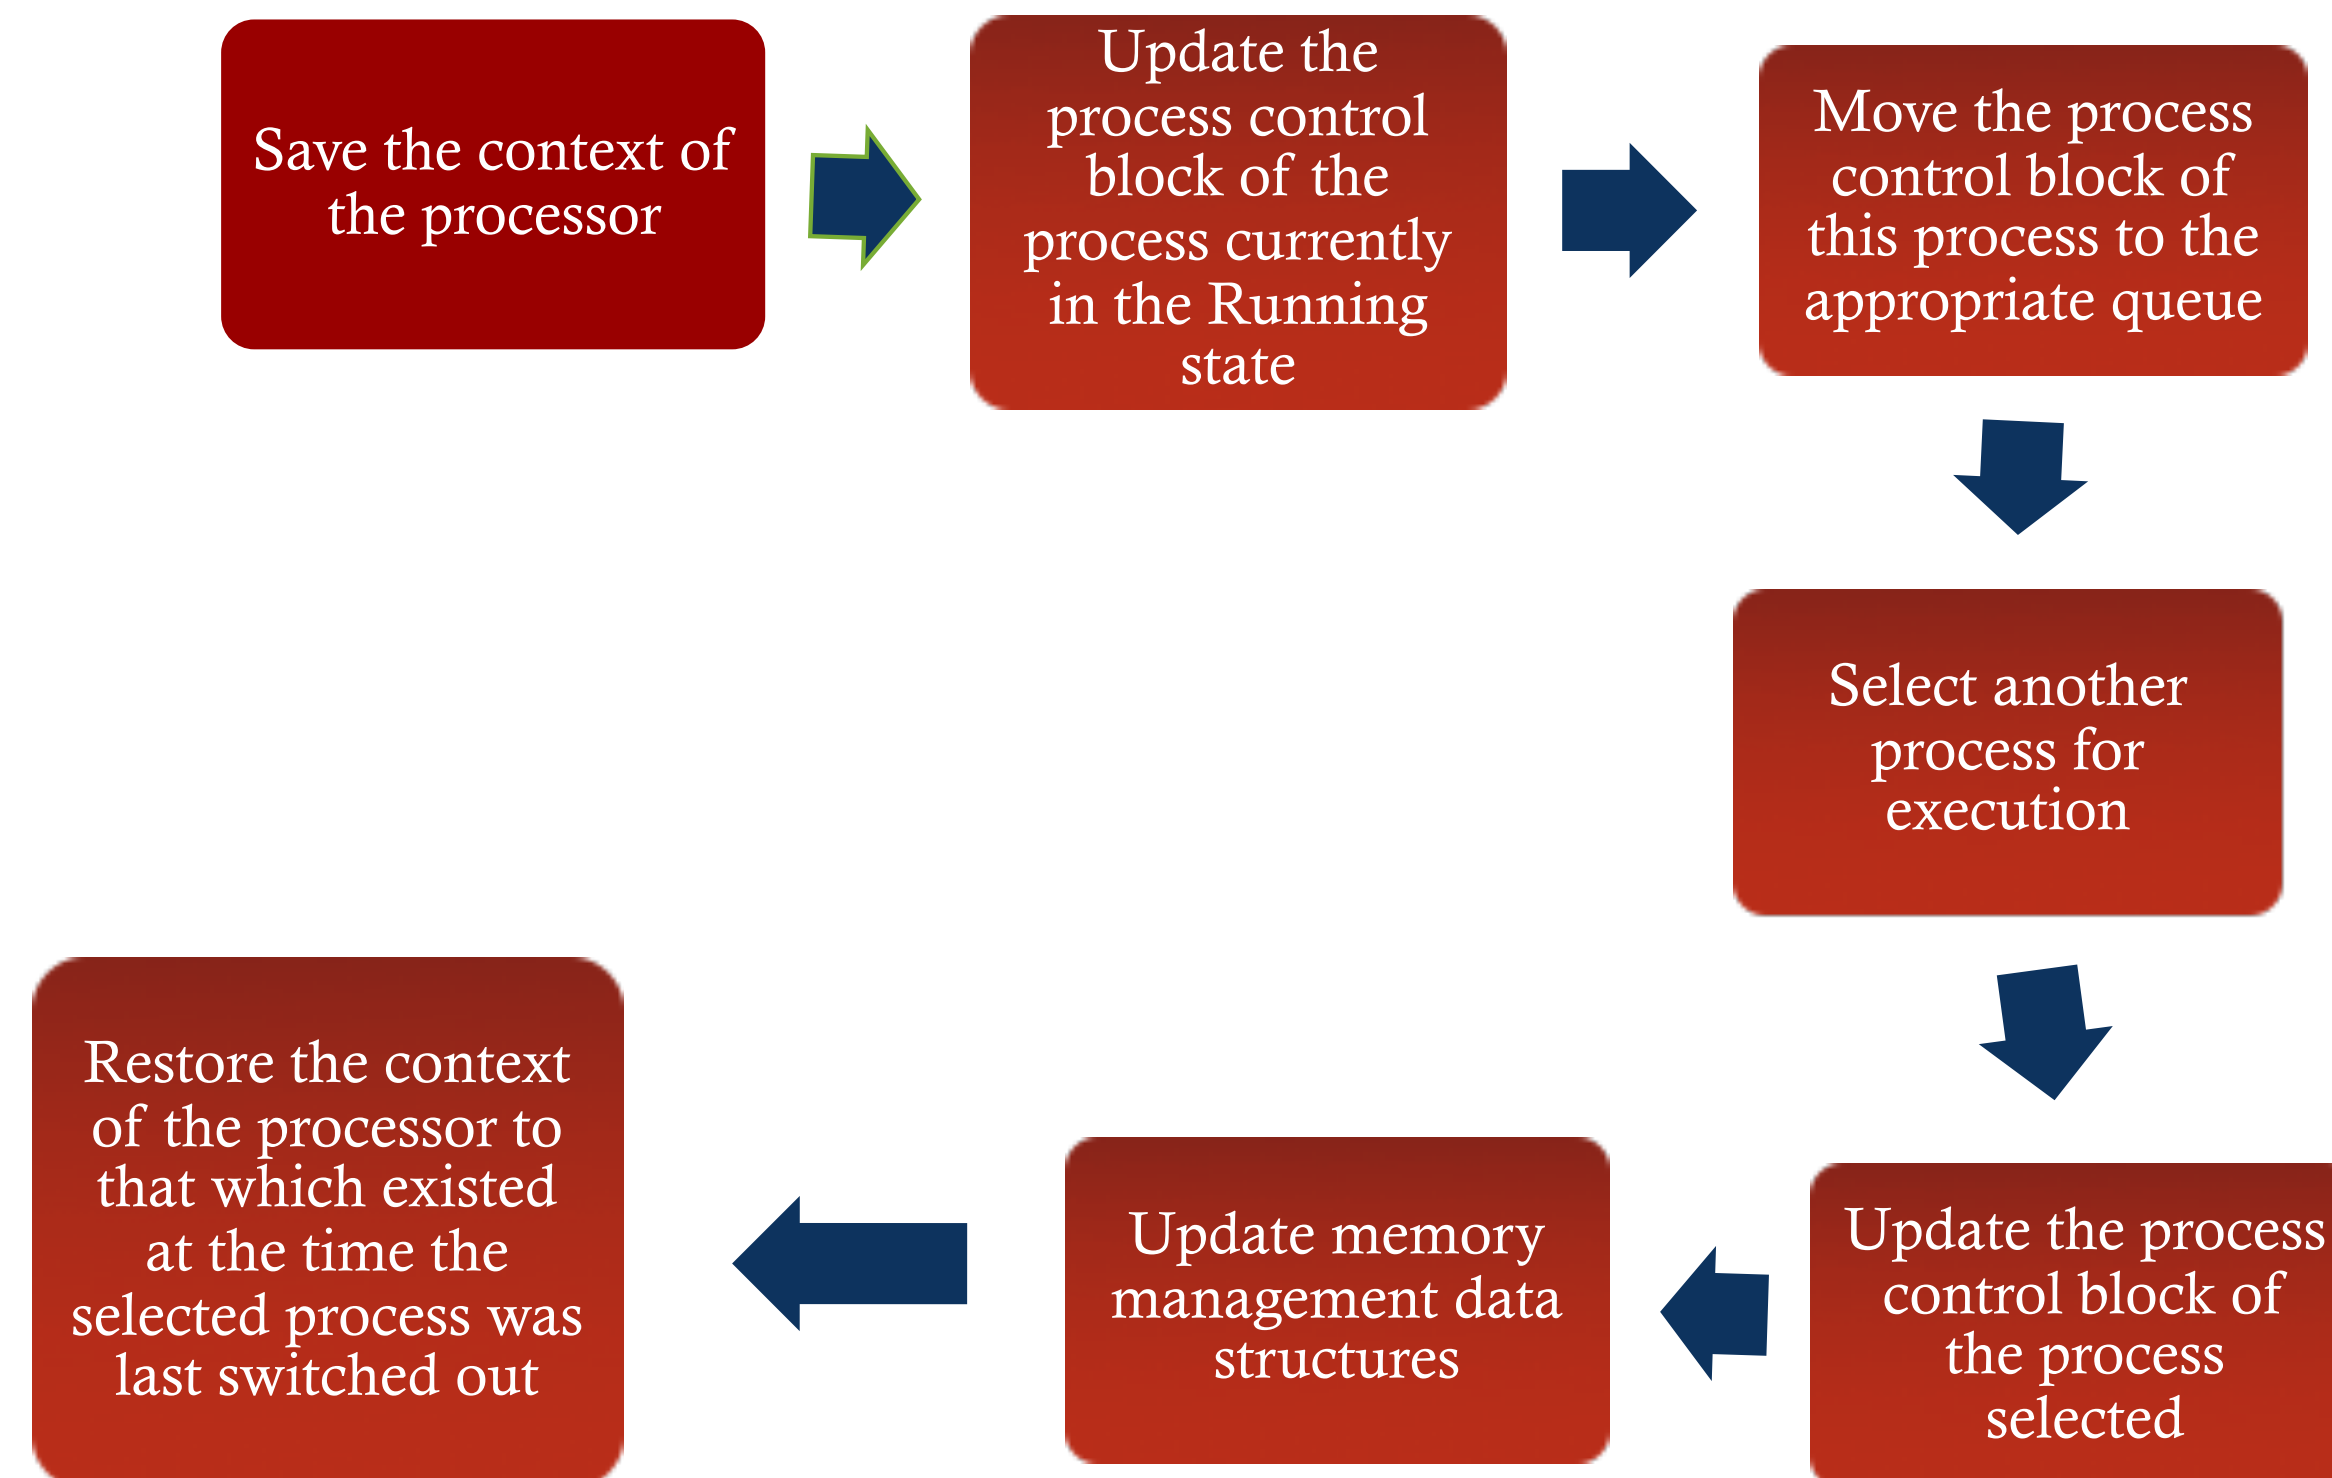
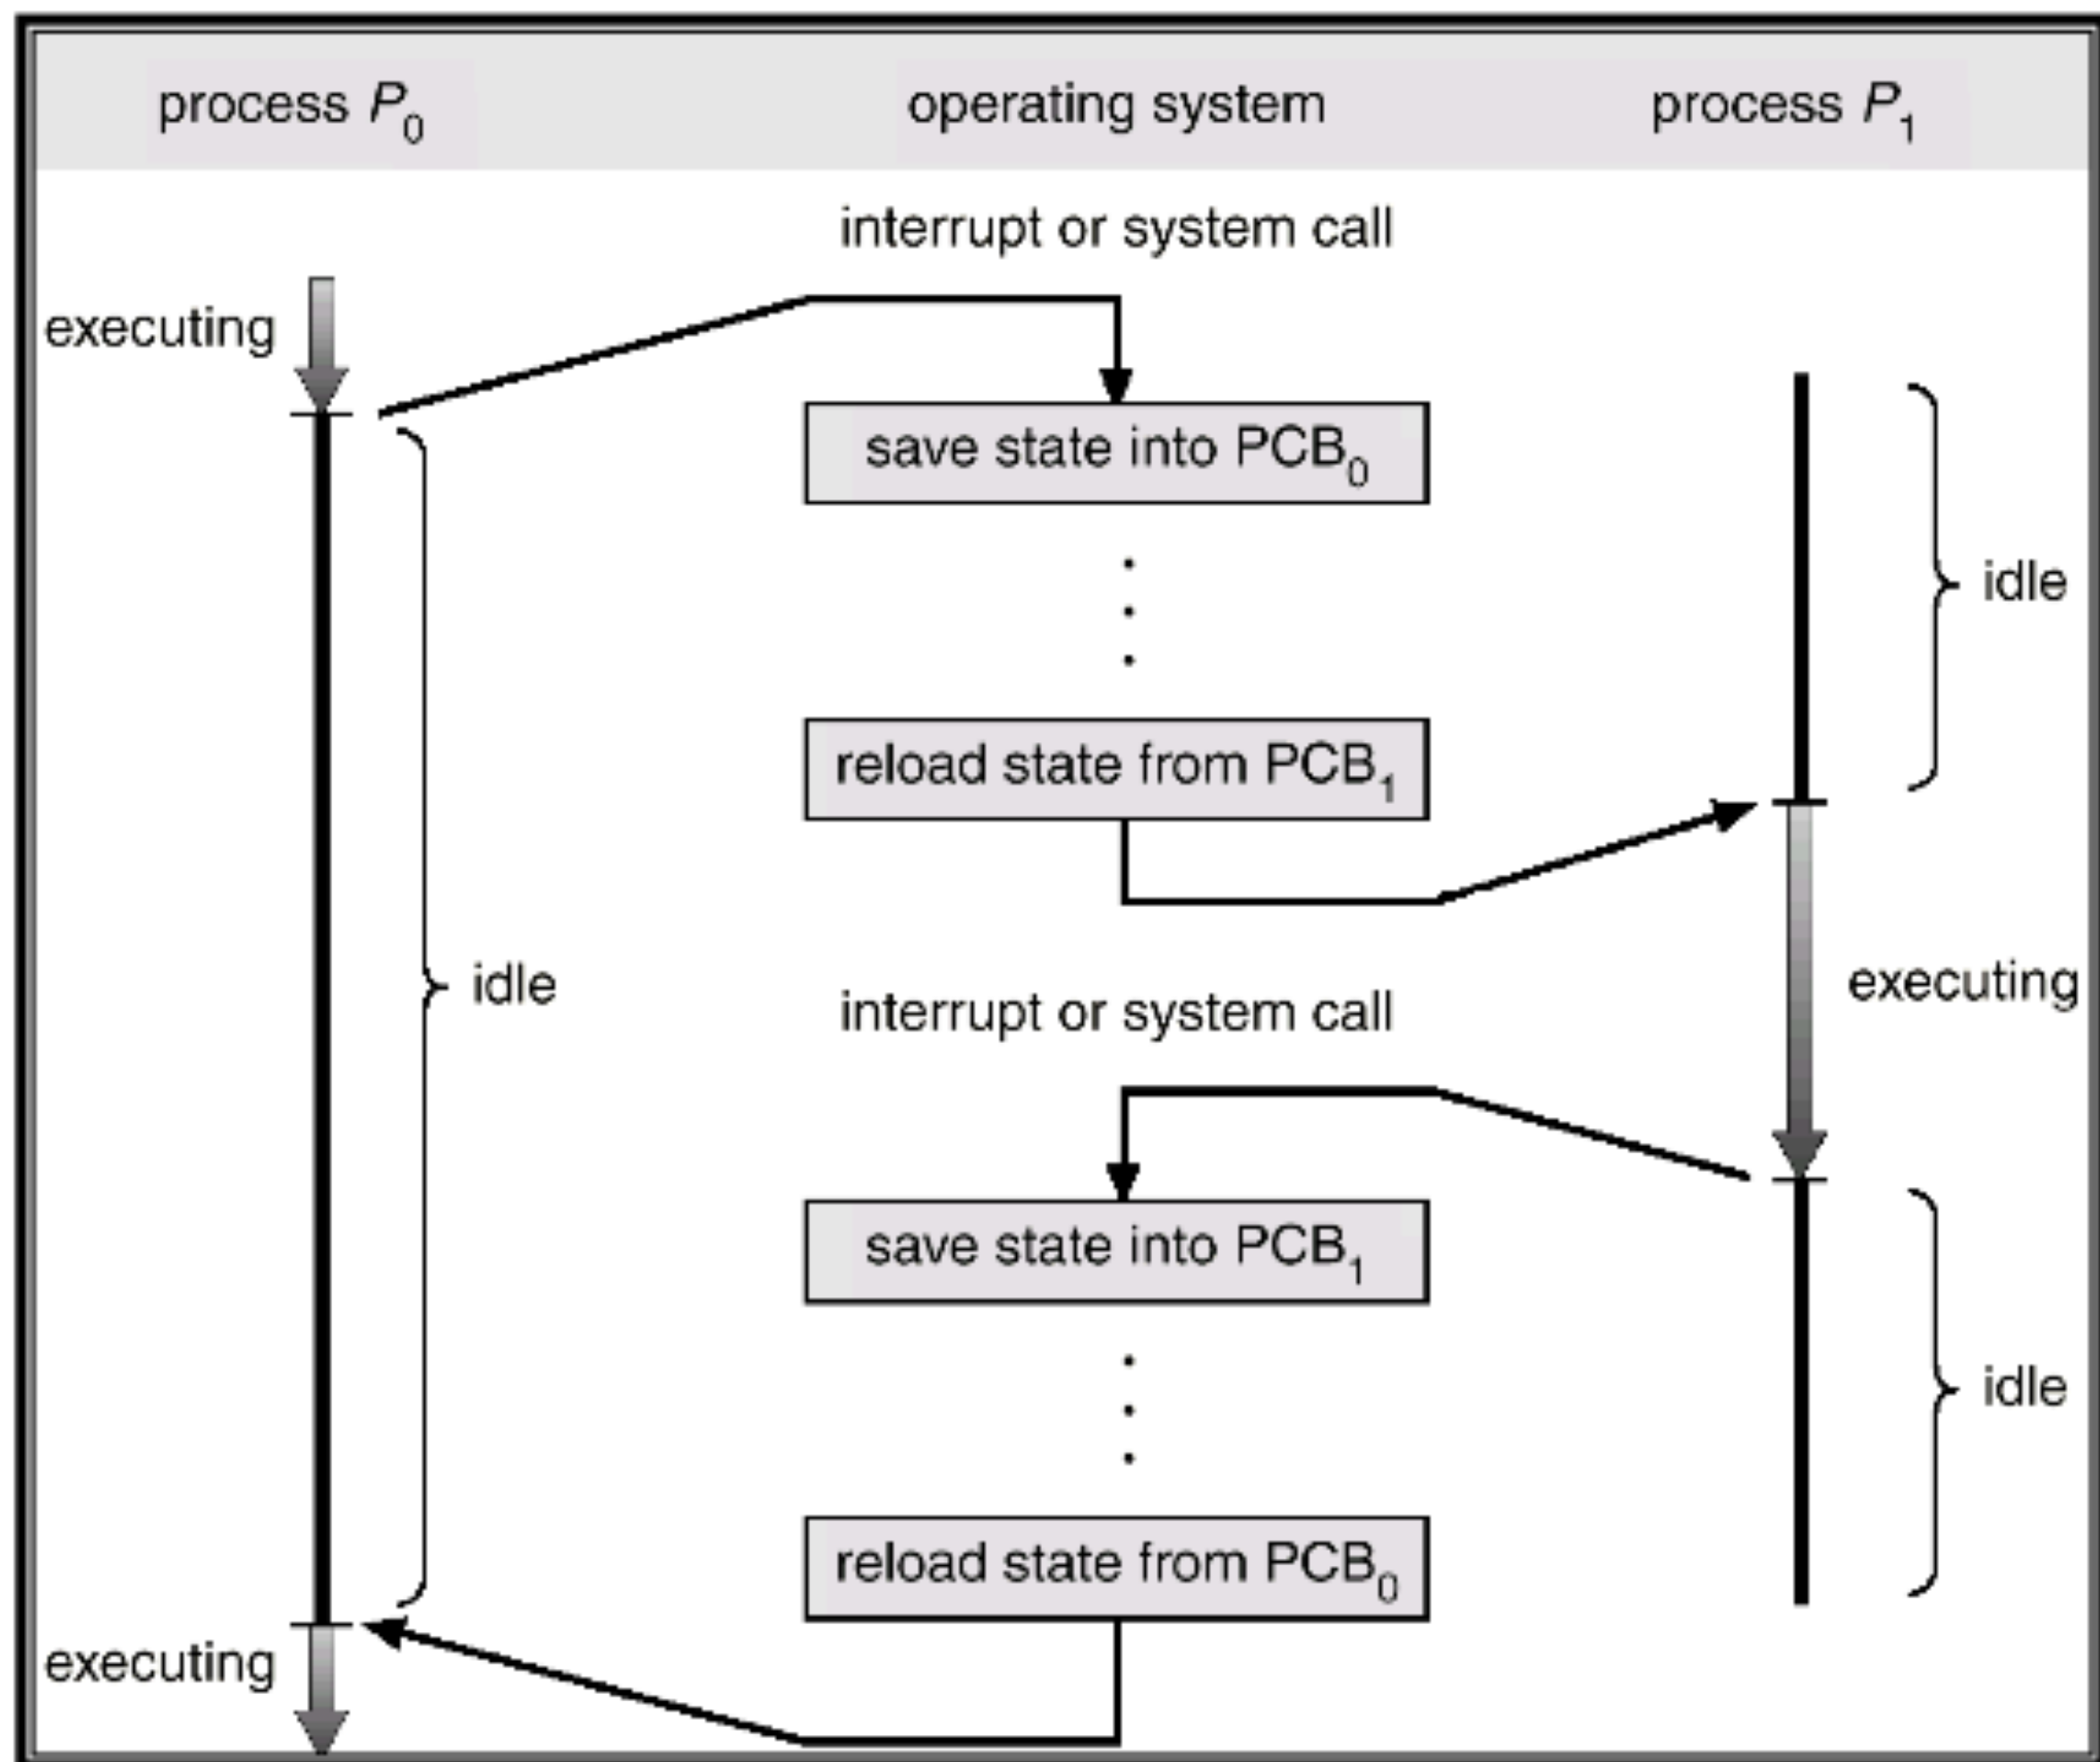
- What events **trigger a process switch**?

Mechanism	Cause	Use
Interrupt	<b>External to</b> the execution of the <b>current instruction</b>	Reaction to an asynchronous external event
Trap	<b>Associated with</b> the execution of the <b>current instruction</b>	Handling of an error or an exception condition
Supervisor call	Explicit request	Call to an operating system function

- What is the difference between **mode switching** and **process switching**?
  - mode switch = change of privilege; process isn't necessarily changed
  - process switch = swap out all process-related info (e.g., registers, page tables)
  - (Recall the fetch, execute, check interrupt cycles)*
  - (Also, see next slide)*
- What is **the role of the OS**? (i.e., **HOW** is a process switched?)
  - saving context, updating PCB(s), updating resource tables, restoring context of next process, etc.


# Process Switching

*A high-level “sketch” of a full process switch a.k.a. “Context Switch”*




# OS Execution


*So... is the kernel executed as a process?*



(a) Separate kernel



(b) OS functions execute within user processes

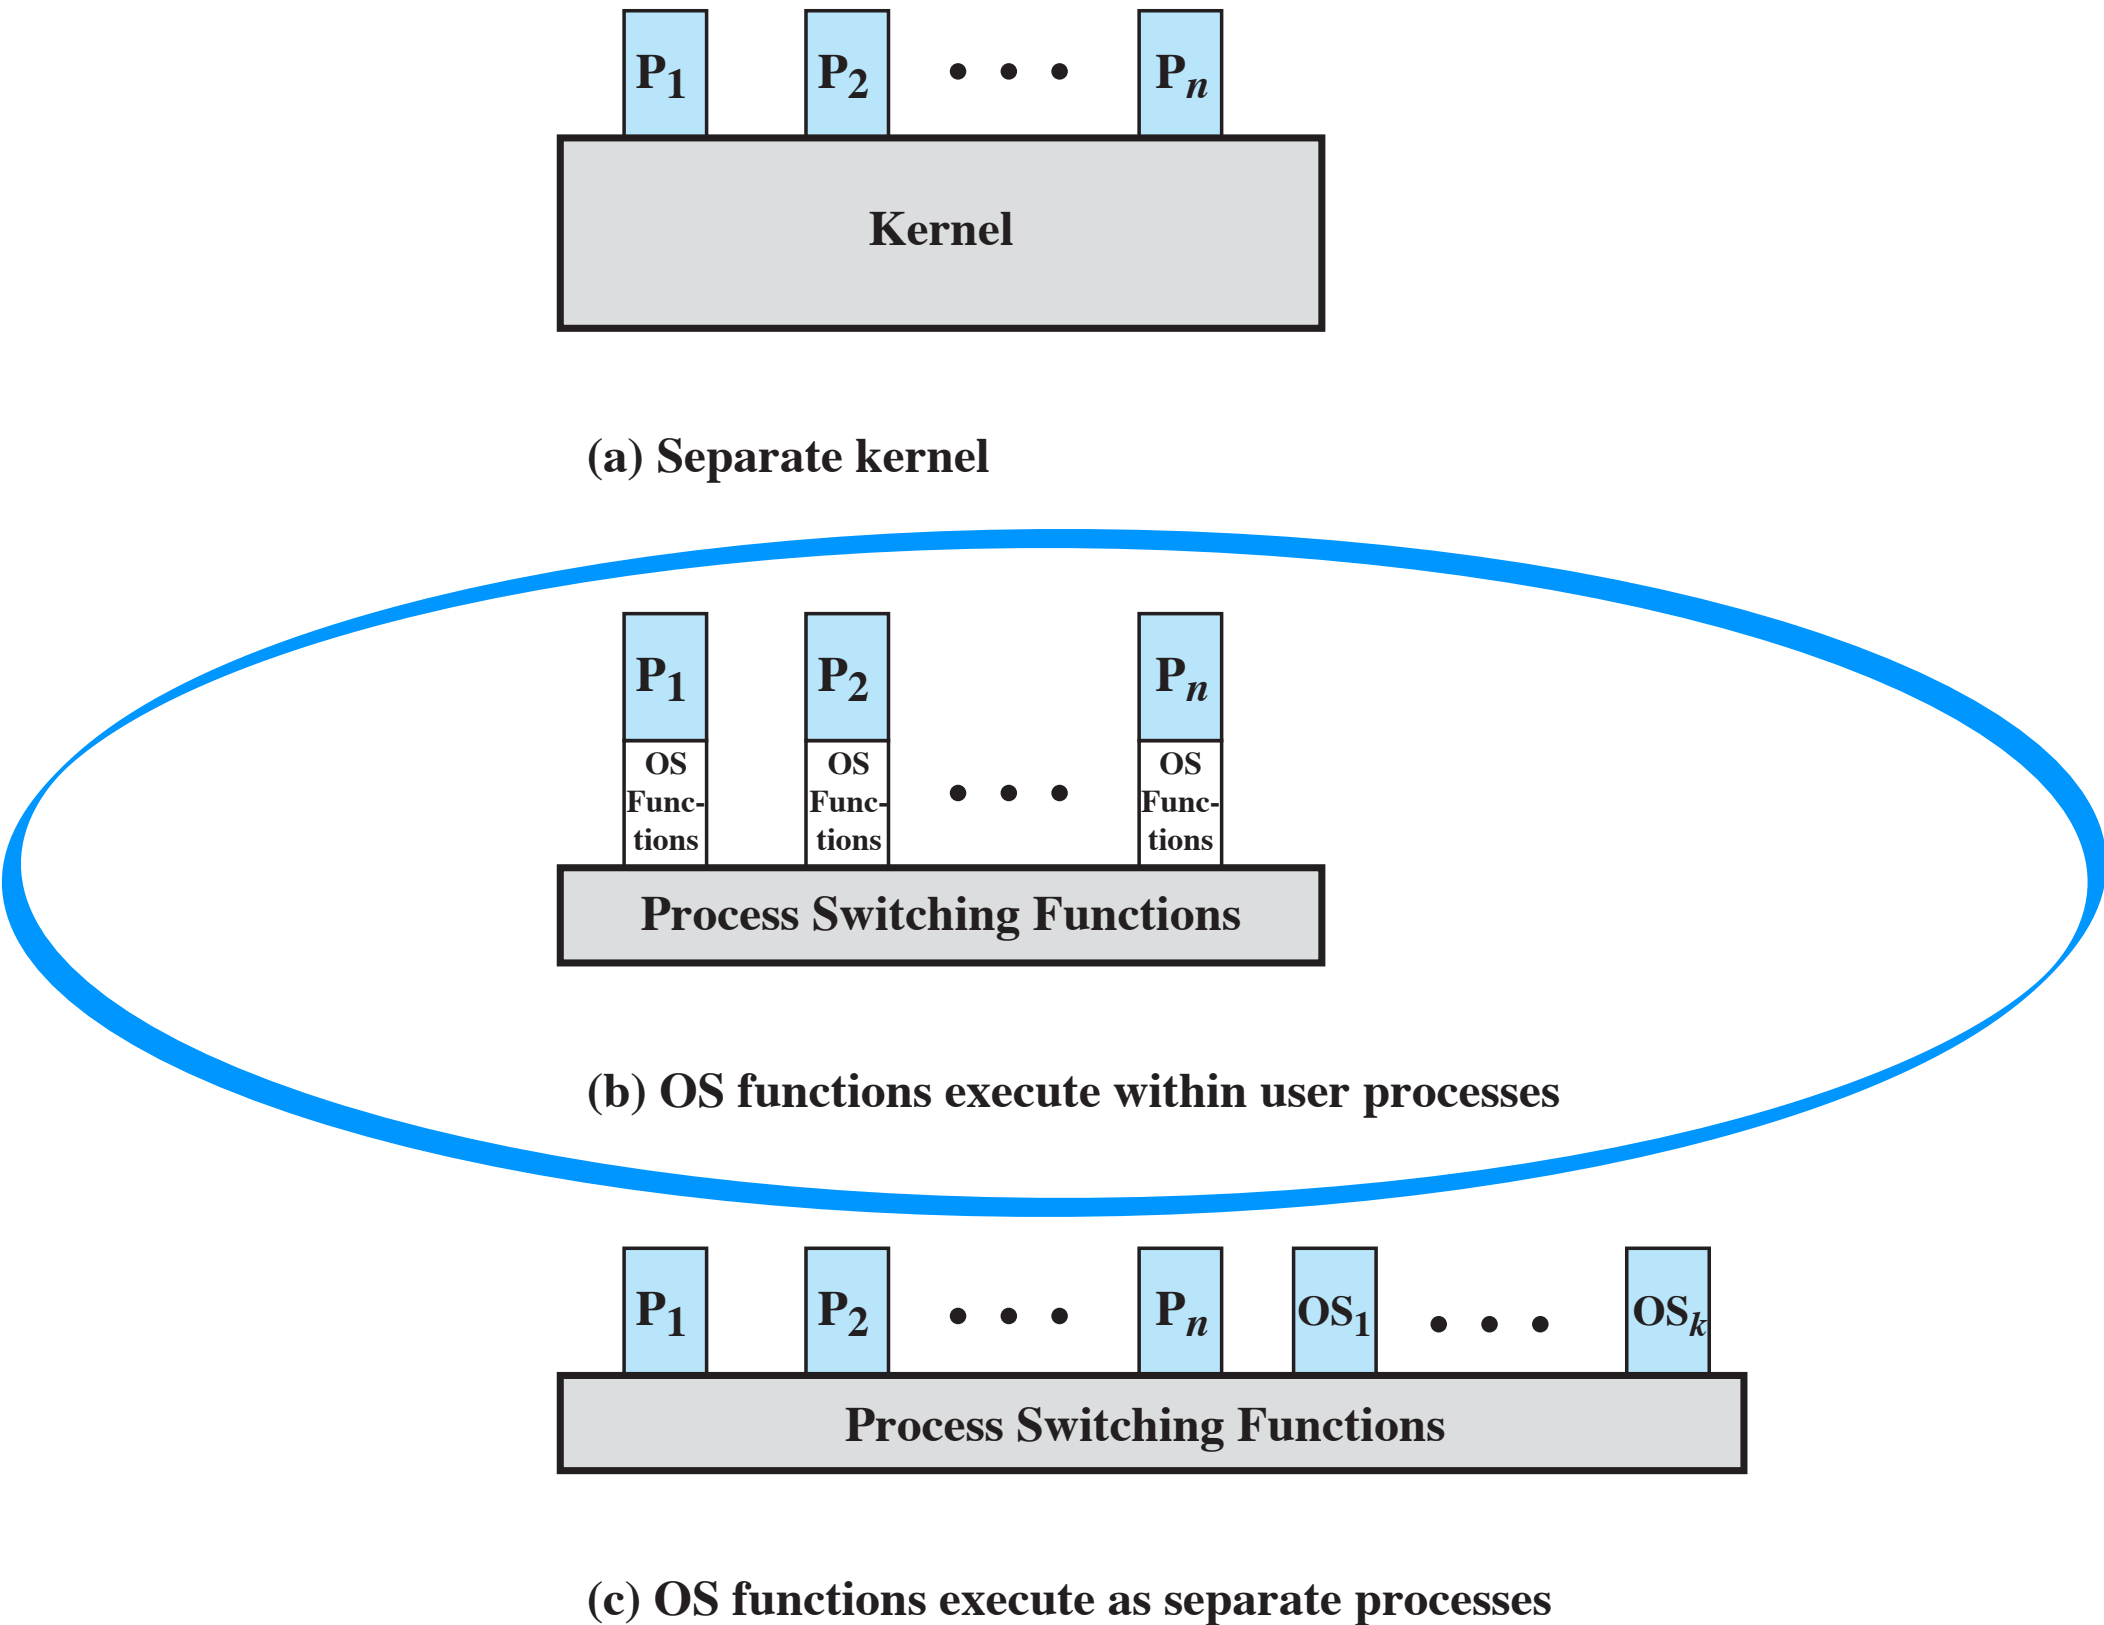


(c) OS functions execute as separate processes

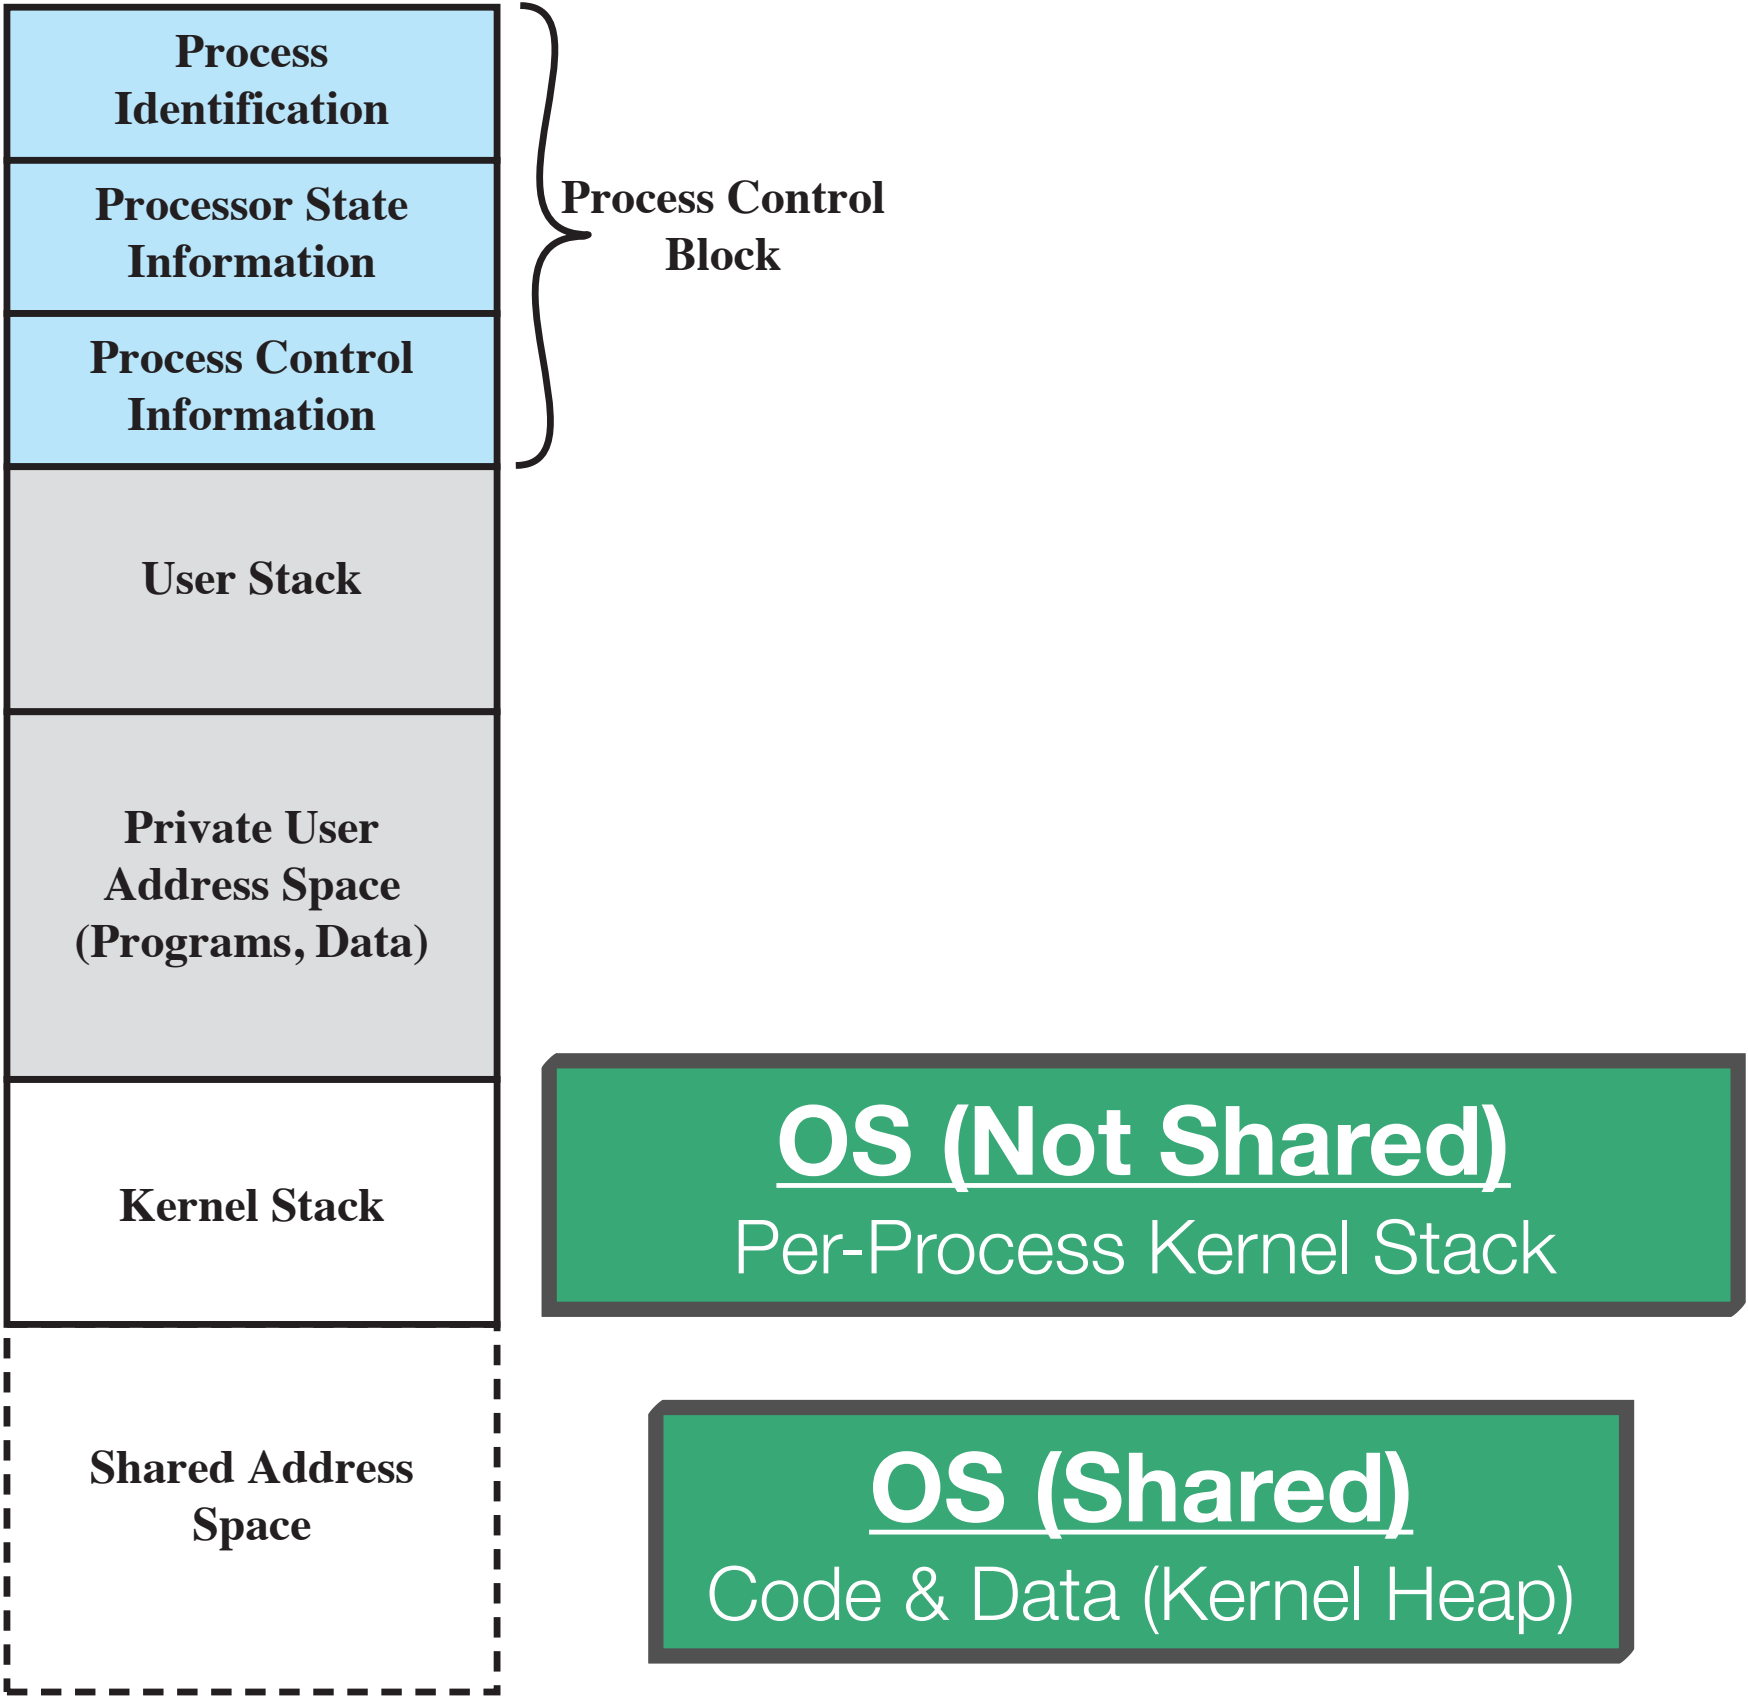
**Figure 3.15 Relationship Between Operating System and User Processes**

# OS Execution

*So... is the kernel executed as a process? Sort of...*



**Figure 3.15 Relationship Between Operating System and User Processes**

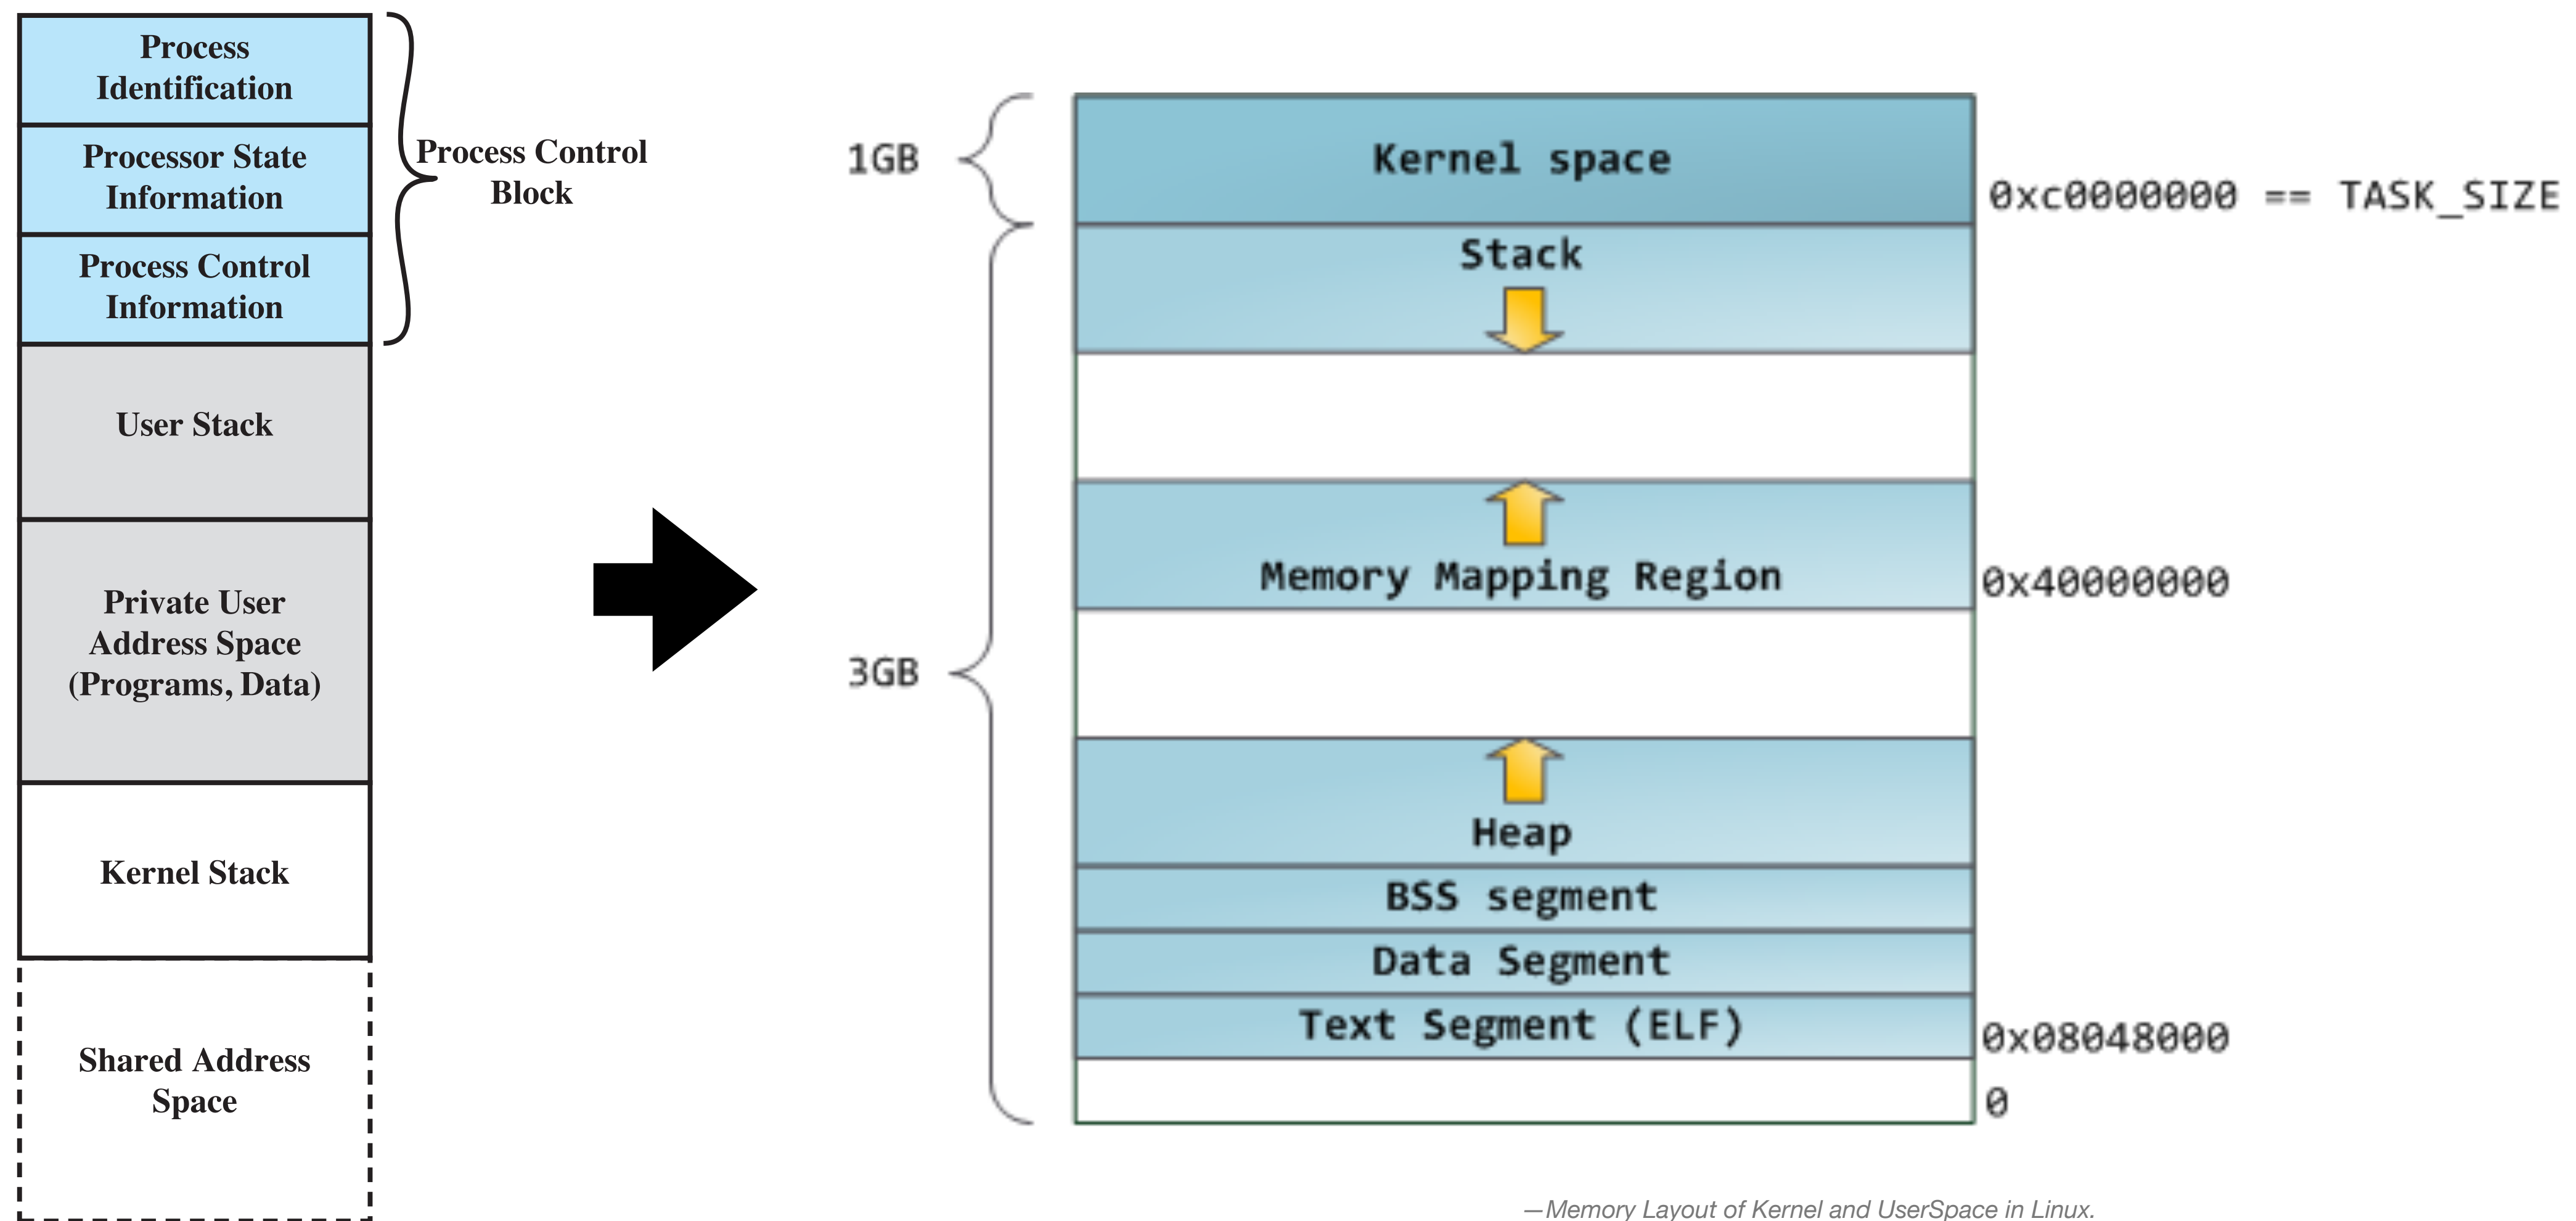


**Figure 3.16 Process Image: Operating System Executes Within User Space**



# OS Execution

*From abstract representation to modern OSs (e.g., Linux)*



—Memory Layout of Kernel and UserSpace in Linux.  
<https://learnlinuxconcepts.blogspot.com/2014/03/memory-layout-of-userspace-c-program.html>



# UNIX Example @ End of Chapter (Process Management in a Real OS)

<b>User Running</b>	Executing in user mode.
<b>Kernel Running</b>	Executing in kernel mode.
<b>Ready to Run, in Memory</b>	Ready to run as soon as the kernel schedules it.
<b>Asleep in Memory</b>	Unable to execute until an event occurs; process is in main memory (a blocked state).
<b>Ready to Run, Swapped</b>	Process is ready to run, but the swapper must swap the process into main memory before the kernel can schedule it to execute.
<b>Sleeping, Swapped</b>	The process is awaiting an event and has been swapped to secondary storage (a blocked state).
<b>Preempted</b>	Process is returning from kernel to user mode, but the kernel preempts it and does a process switch to schedule another process.
<b>Created</b>	Process is newly created and not yet ready to run.
<b>Zombie</b>	Process no longer exists, but it leaves a record for its parent process to collect.

