


Concurrency (Part II): Mutual Exclusion, **Synchronization**, Deadlock, and Starvation

Professor Travis Peters
CSCI 460 Operating Systems
Fall 2019

Some slides & figures adapted from Stallings instructor resources.

*Some slides adapted from Adam Bates's F'18 CS423 course @ UIUC
<https://courses.engr.illinois.edu/cs423/sp2018/schedule.html>*

SIMPLY EXPLAINED



—<http://www.datamation.com/news/tech-comics-quantum-physics-2.html>

Goals for Today

Learning Objectives

- Dive a bit deeper into core topics in concurrency
- Discuss common mechanisms for achieving mutual exclusion & synchronization

Announcements

• Schedule Updates

- Exam #1 — delayed...
 - ...career fair...
 - topics to include everything up through concurrency (at least... possibly scheduling as well...)
- Exam #2 — topics will include:
 - (Scheduling)
 - Memory Management & Virtual Memory
 - File Systems & I/O
 - Selected OS Security Topics
- No Exam #3!
- **Homework 2** (Chapters 3-4) DUE 10/11 —*trust me, you don't need that long!*
- **Homework 3** (Chapters 5-6) DUE 10/18
- **Programming Assignment 1** (Concurrency + C programming/threads) — should be posted by Friday

Recap: Concurrency & Solutions for Mutual Exclusion

To achieve correct & meaningful solutions to concurrency problems, **mutual exclusion** is a must!

Software Support

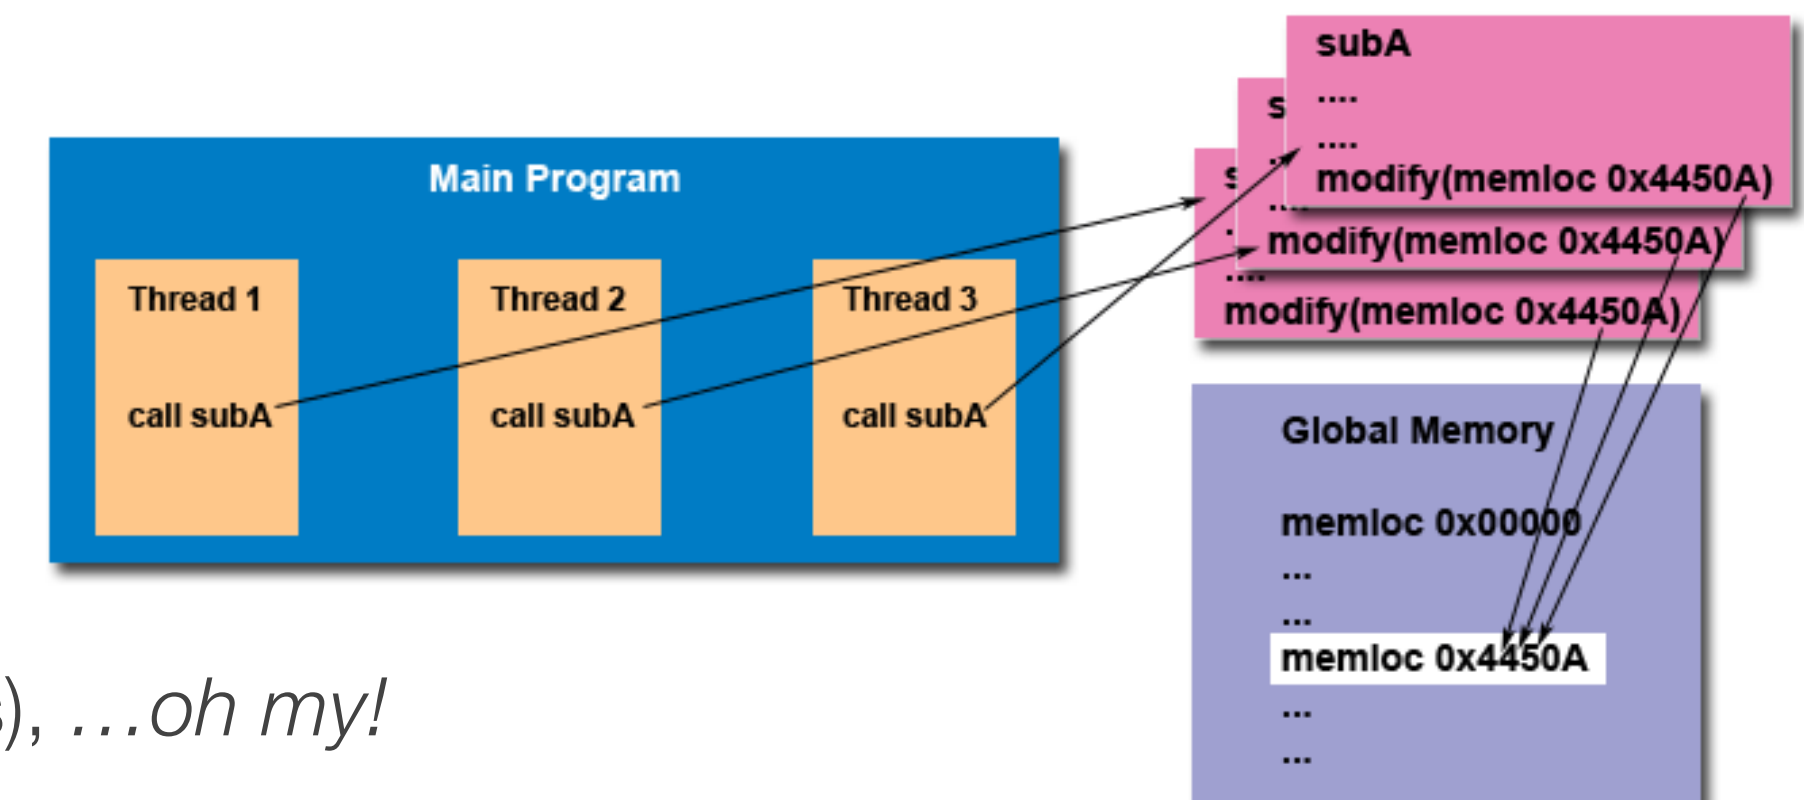
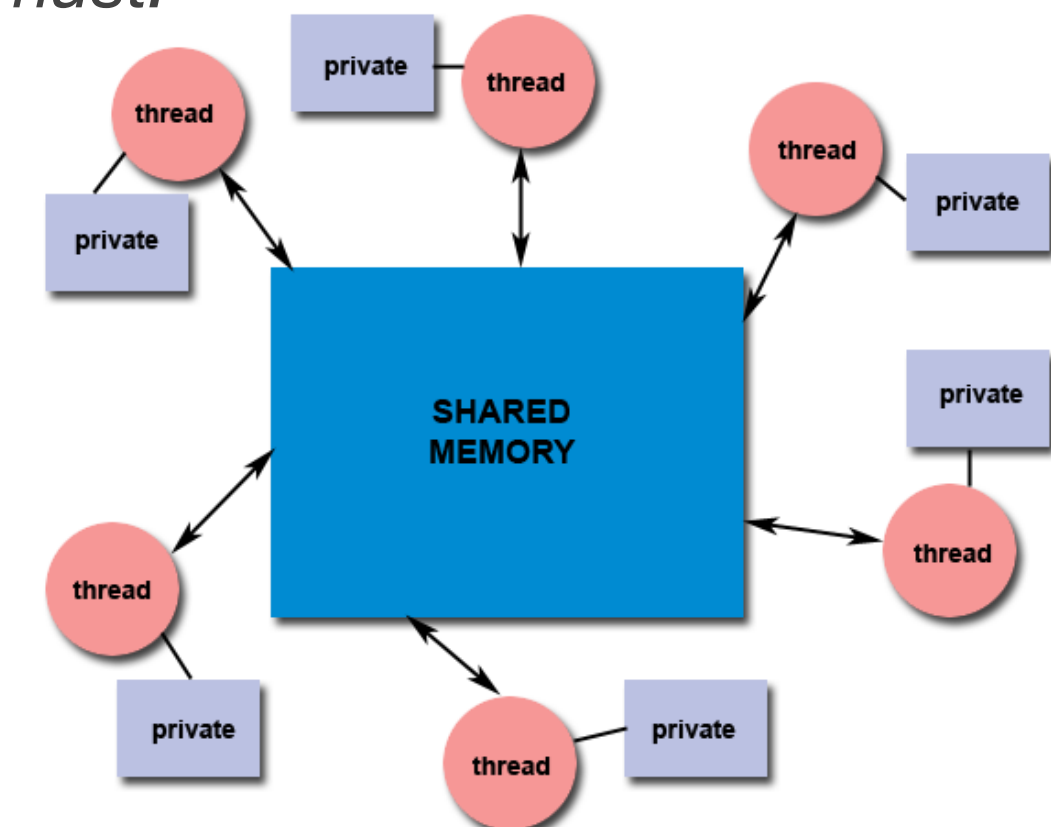
- Assume elementary mutual exclusion at the memory access level; serialized by “memory arbiter”
- Decker’s Algorithm, Peterson’s Algorithm

Hardware Support

- Interrupt Disabling
 - **Disadvantages:** inhibits processor’s ability to interleave processes; doesn't work across processors.
- Special Machine Instructions
 - **Compare&Swap:** compare values => if values are the same, swap!
 - **Exchange (XCHG):** exchanges the contents of a register w/ that of a memory location
 - **Advantages:** simple & easy to implement; can be used on multi-processor machines
 - **Disadvantages:** possibly expensive busy-waiting; starvation & deadlock are still possible

Programming Language Mechanisms

- Examples using pthreads
- Semaphores, Monitors, Condition Variables, Message Passing, Mutexes (Locks), ...oh my!



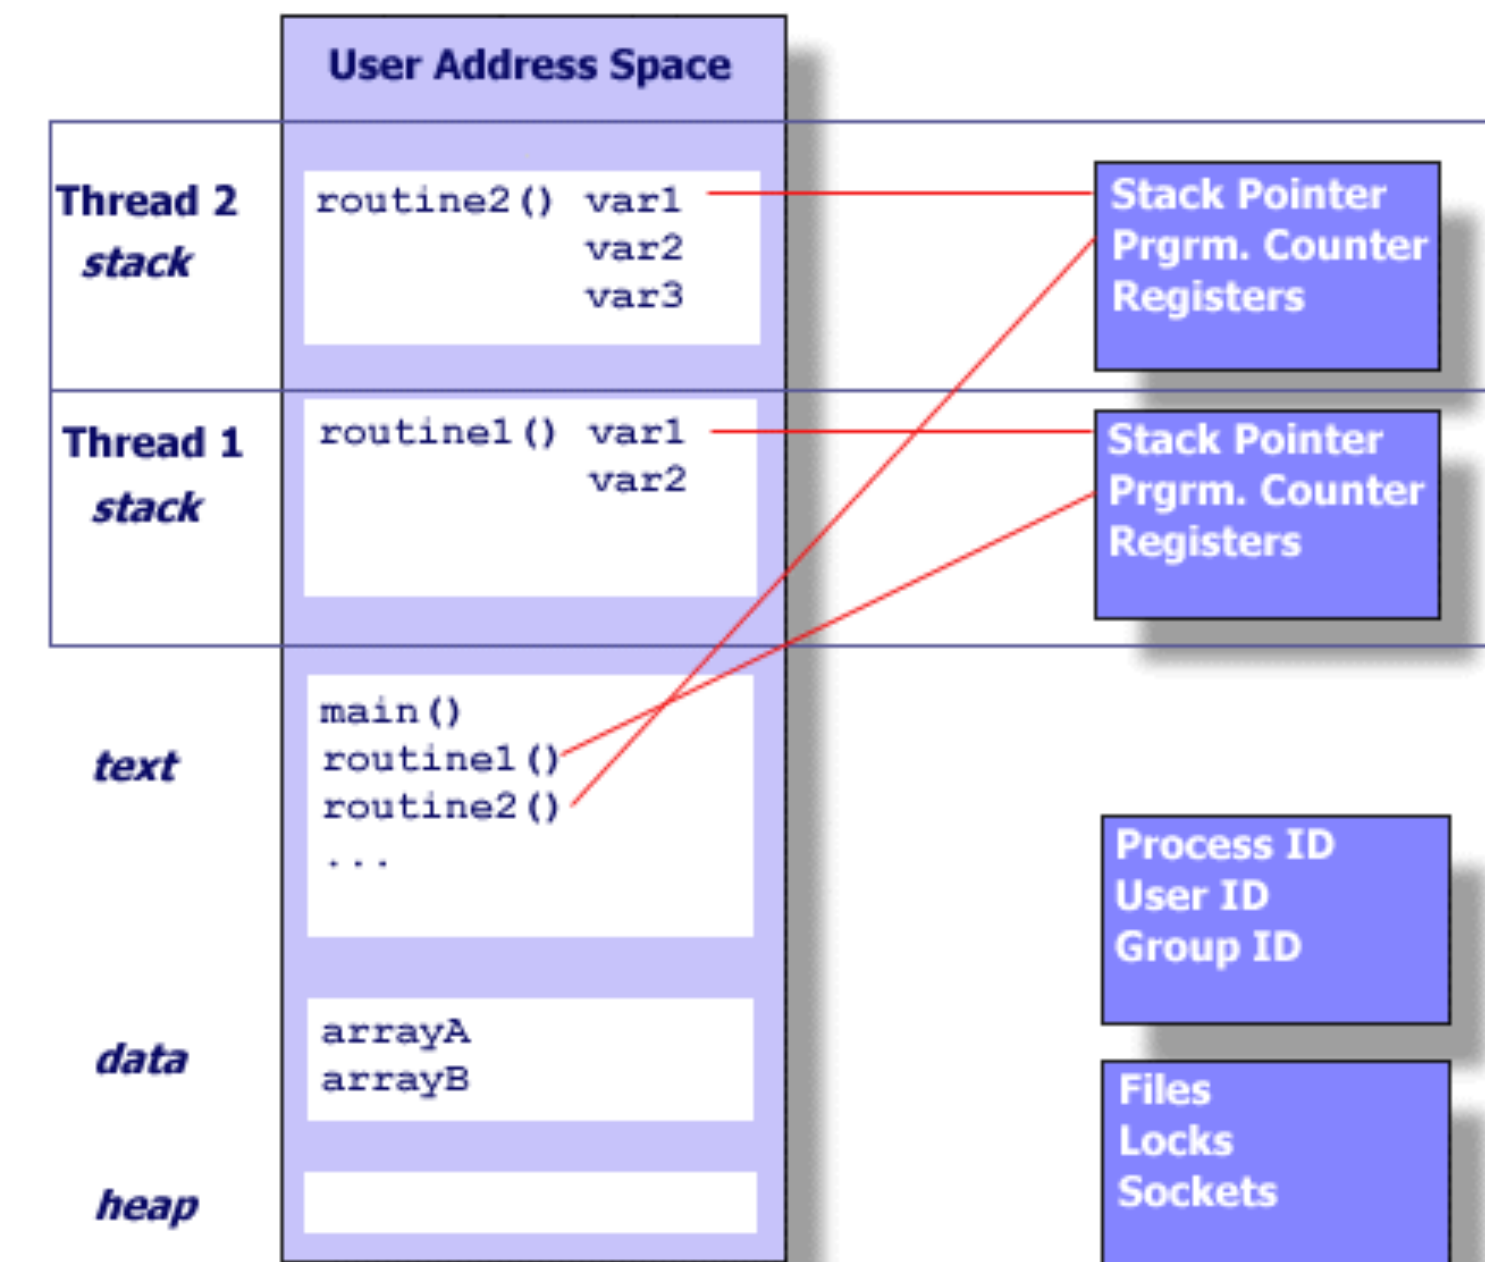
— <https://computing.llnl.gov/tutorials/pthreads/>

➔ **Famous Problems:** Producer/Consumer Problem, Dining Philosophers Problem

Programming w/ pthreads

A good write-up on pthreads: <https://computing.llnl.gov/tutorials/pthreads/>

- **pthreads** (POSIX Threads)
 - defines a set of C programming language types, functions and constants.
 - is implemented with a `pthread.h` header and a thread library.
 - includes mutexes, condition variables, etc.
- E.g., A typical sequence in the use of a **mutex** is as follows:
 - ➔ Create and initialize a mutex variable
 - ➔ Create and start several threads
 - ➔ Several threads attempt to lock the mutex (only 1 succeeds and "owns" the lock)
 - ➔ The owner thread performs some set of actions
 - ➔ The owner unlocks the mutex
 - ➔ Another thread acquires the mutex and repeats the process
 - ➔ Finally, the mutex is destroyed



— <https://computing.llnl.gov/tutorials/pthreads/>

NOTE: When several threads compete for a mutex, the losers block at that call (there does exist a non-blocking call: "trylock")

NOTE: It is the programmer's responsibility to make sure **every thread** that needs to use a mutex does so.
For example, if 10 threads are updating the same data, but only one uses a mutex, the data can still be corrupted!

Programming w/ pthreads

A good write-up on pthreads: <https://computing.llnl.gov/tutorials/pthreads/>

Most of class featured in-class demo:

- <Code Walk-Through - **threads01.c** & **threads02.c**>
- What happens if I **vary the number of threads**?
 - 1 vs. 2 vs. 5 vs. 10 vs.
- What happens if the **target number is small**?
 - 100 vs. 1000000000000
- What happens if we **don't wait** for the threads to complete?
- What happens if I compile & run on **different machines** (e.g., native machine/OS vs. my local Linux VM)?
- **Why** are these things happening?!
 - NOTE: make sure you (re-)build the executable wherever you are going to run it....
 - `objdump -d t1`
- With proper (simple) synchronization, why does the program appear to execute **slower**?

