

Scheduling (Part III)

Professor Travis Peters
CSCI 460 Operating Systems
Fall 2019

Some slides & figures adapted from Stallings instructor resources.

*Some slides adapted from Adam Bates's F'18 CS423 course @ UIUC
<https://courses.engr.illinois.edu/cs423/sp2018/schedule.html>*

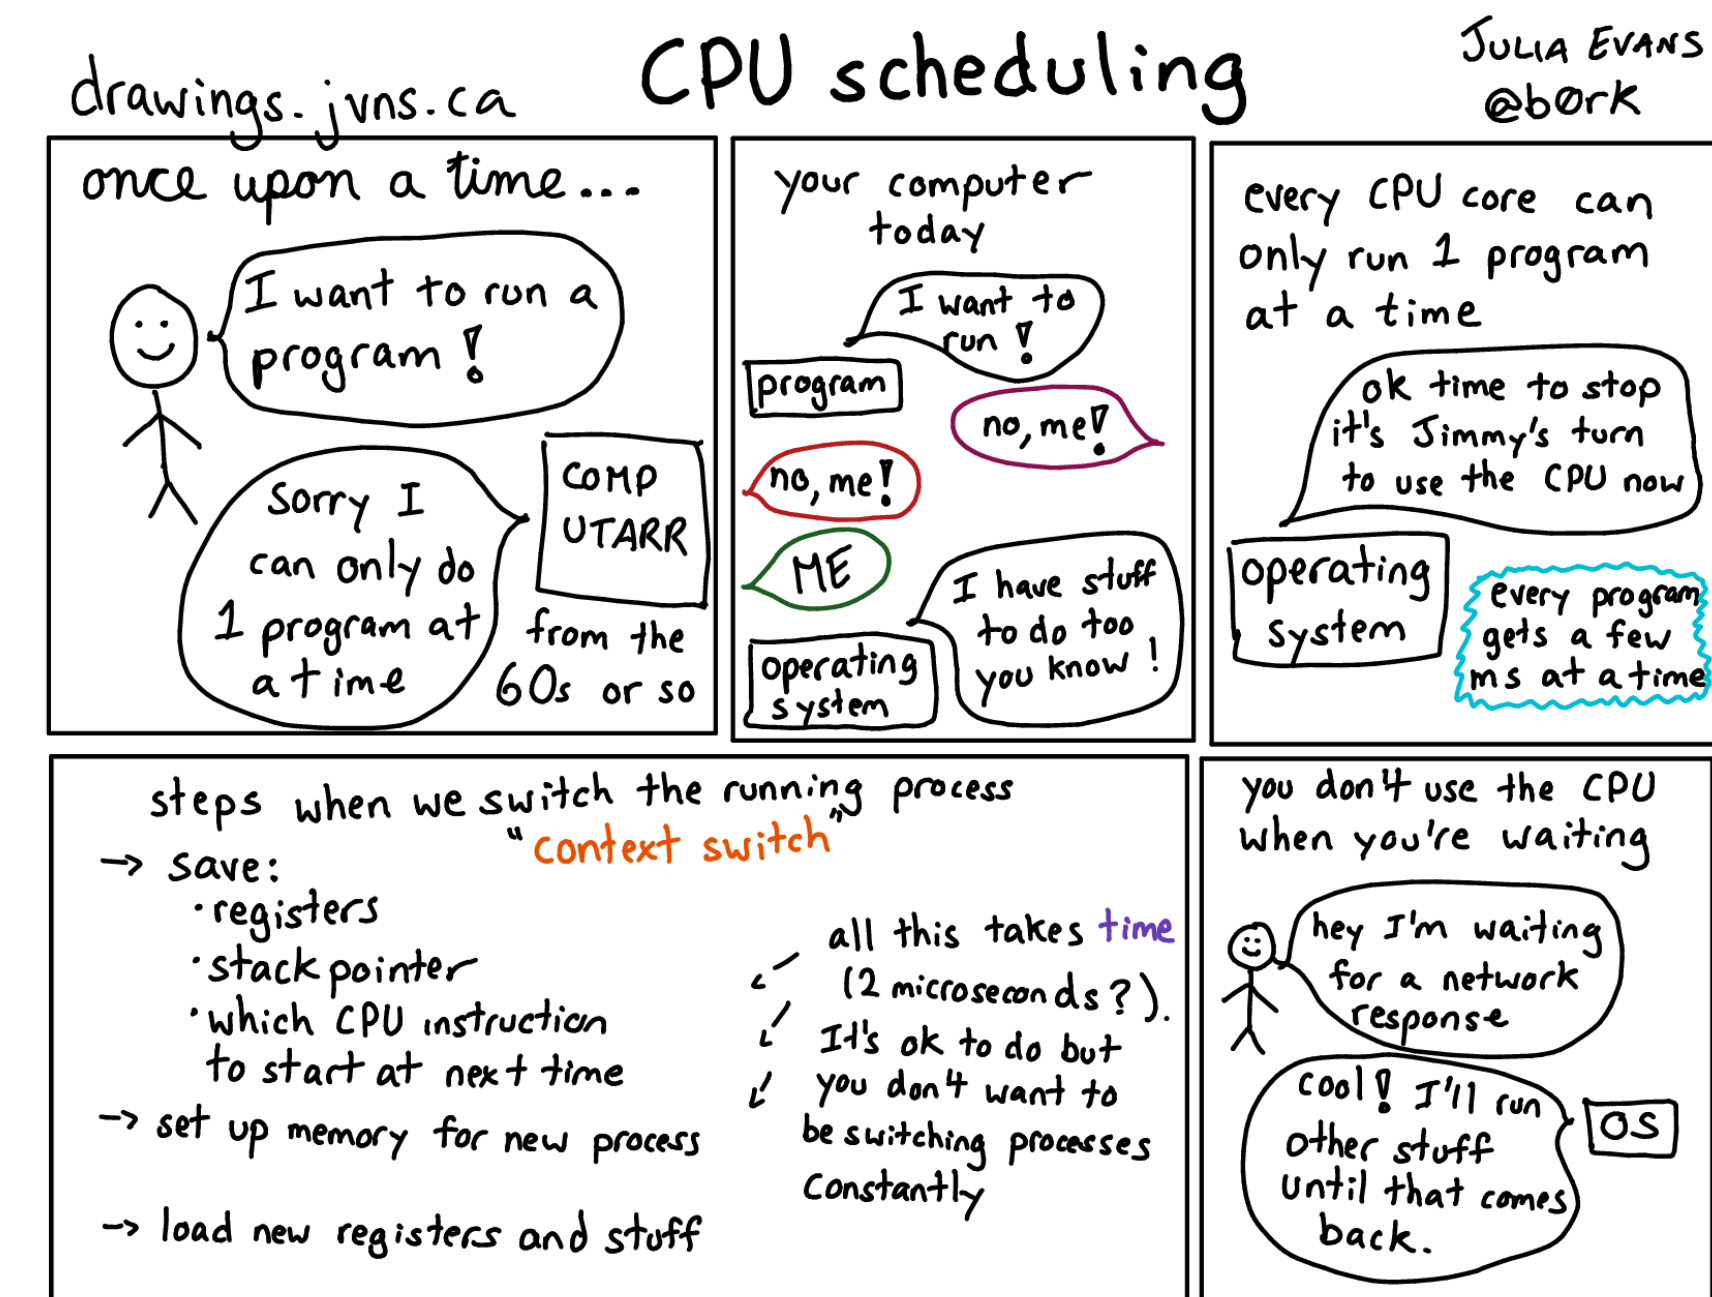
Goals for Today

Learning Objectives

- Discuss the key design issues in multiprocessor scheduling + some of the key approaches to scheduling
- Understand the requirements imposed by real-time scheduling
- Look at some of the scheduling methods used in Linux & UNIX

Announcements

- Programming Assignment was Due Wednesday @10pm!
→ Please upload as a **zipped folder**... upload issue was fixed ;)
- Project deadlines posted
- **REMINDER:** You have 1 free late pass... (see the syllabus)
- **NOTE:** If using GitHub (**good!**), but make sure code isn't public (**BAD!**)



<https://drawings.jvns.ca/scheduling/>

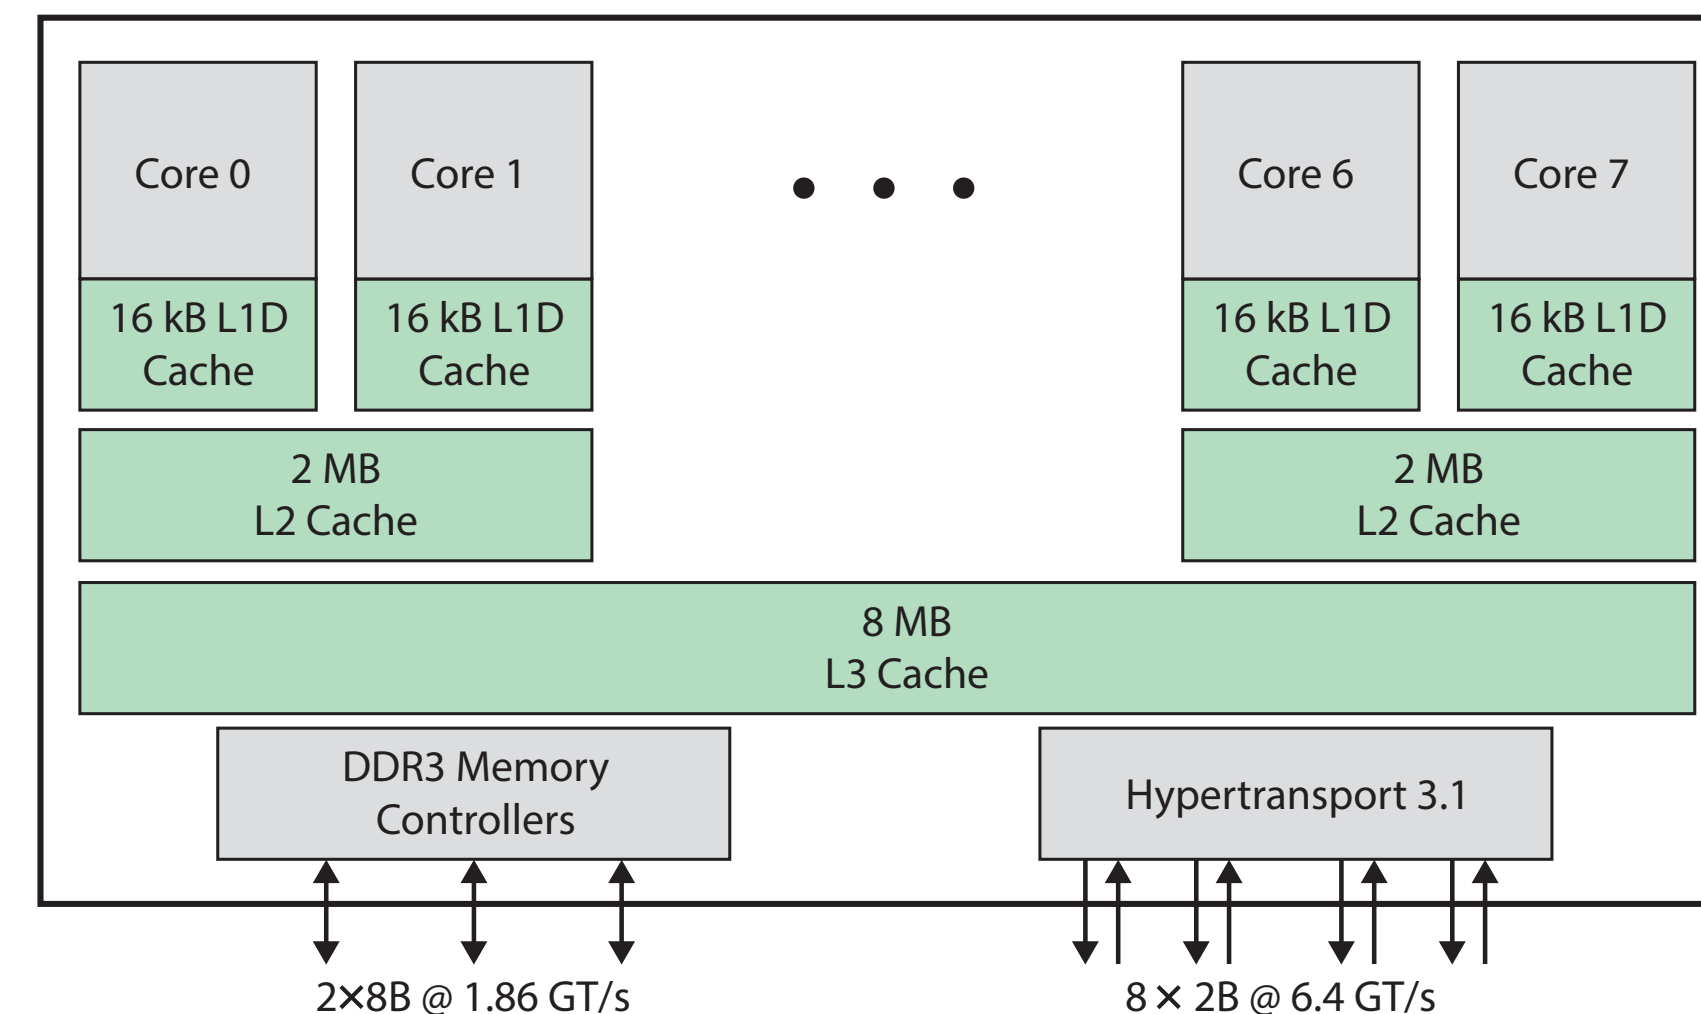
Multiprocessor and Multicore Scheduling (Summary)

- See the text for discussion of various considerations for
 - Synchronization
 - Parallelism
 - etc.
- At the end of the day, its all about assigning processes**es** to processors**ors**, and being cognizant of the trade-offs
 - Static Assignment — processes are assigned to a processor-specific queue upon initialization; assignments don't change
 - Dynamic Load Balancing — keep workload distributed (as equally as possible) across all processors
 - Master/Slave (*simple*) vs. Peer (*more complex*) vs. Hybrids

Multiprocessor and Multicore Scheduling (Summary)

Punchline(s):

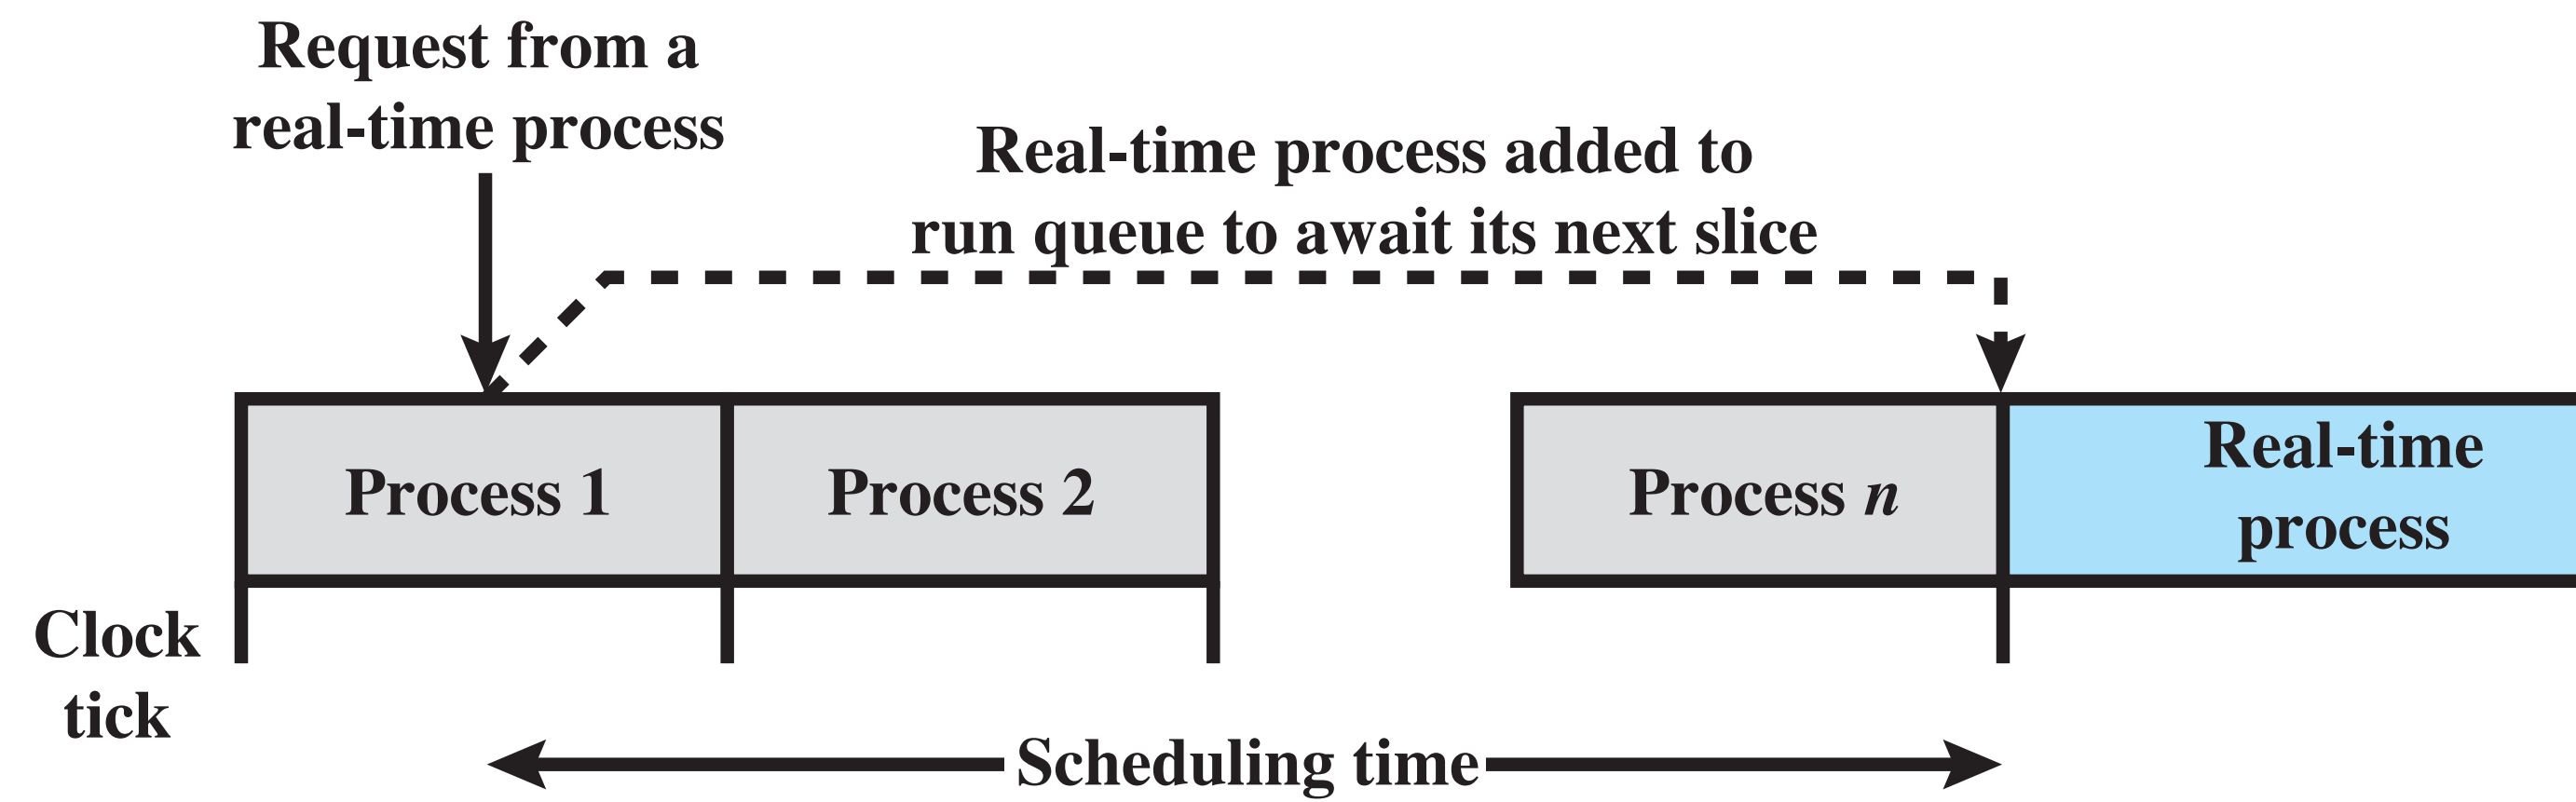
- more processors → less emphasis on the efficiency of your scheduling algorithm.
- **Load Sharing** (*processors pull from global queue of ready-to-run processes*) is probably the most common approach to scheduling on multiprocessor systems.
 - Recall Concurrency—MP scheduling comes w/ all the same advantages and disadvantages... ;)
 - To do this, you **must enforce** mutual exclusion
- Try to do things s.t. cost of process switching is avoided as much as possible
- Try to limit the number of threads in an application to no more than the number of processors available on the system
- Most (e.g., Linux, Windows) hold multiprocessor scheduling \approx multicore scheduling



Real-Time Processes, Scheduling, and OSs

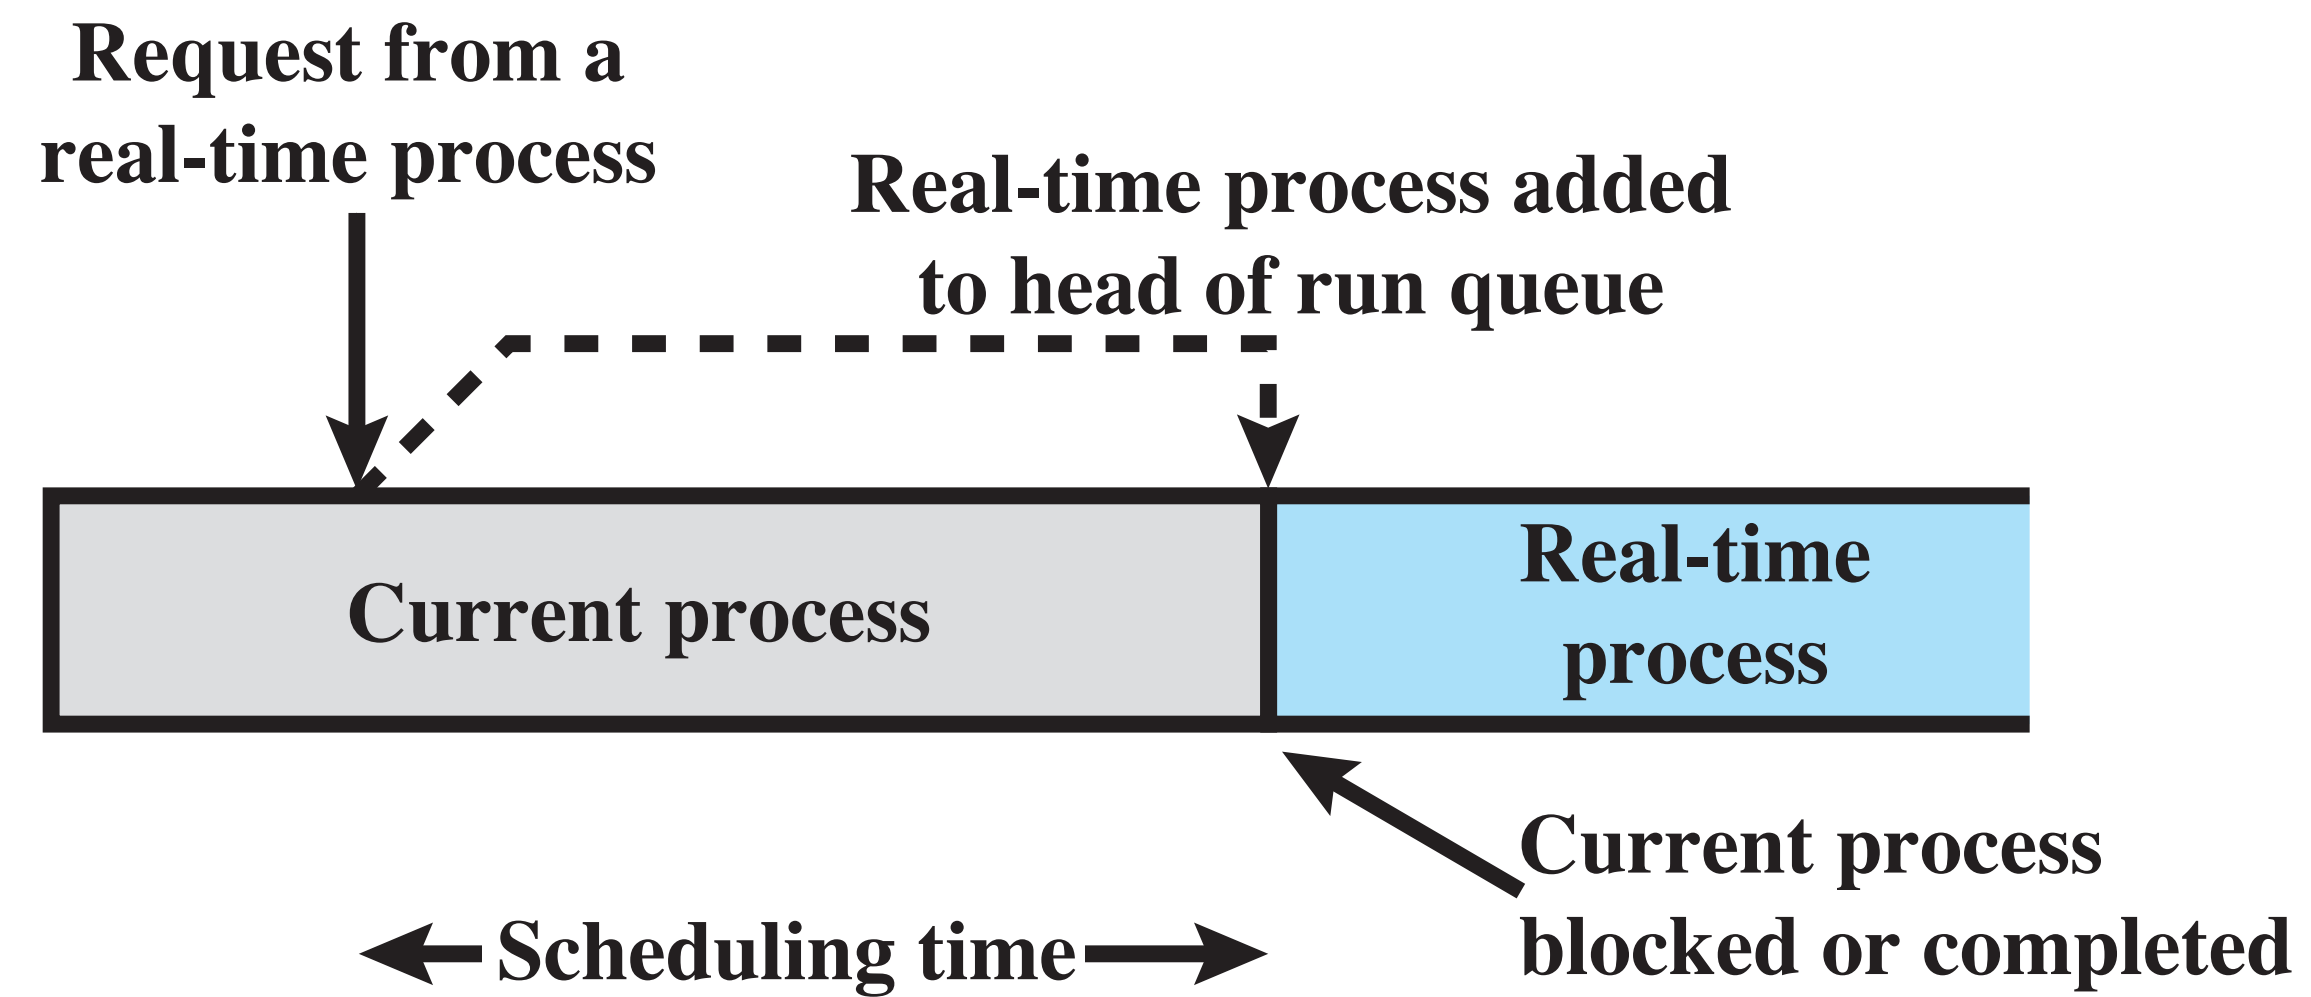
- A **real-time process** must produce nearly instantaneous output based on new inputs.
 - Each arriving input item is subject to a deadline.
 - **Ex:** Streaming of audio or video, control of robots.
 - **hard real-time tasks** vs. **soft real-time tasks**
- A **period** is a **time interval** (typically in *ms* or μs) within which each input item must be processed.
 - The end of each period is the **implicit deadline** for processing the current item.
- Characteristics of RTOSs
 - Deterministic, Responsiveness, User Control, Reliability, Fail-Soft Operation
 - Most contemporary RTOSs don't deal directly with "deadlines."
Instead, they try to be as responsive as possible

Real-Time Scheduling



(a) Round-robin Preemptive Scheduler

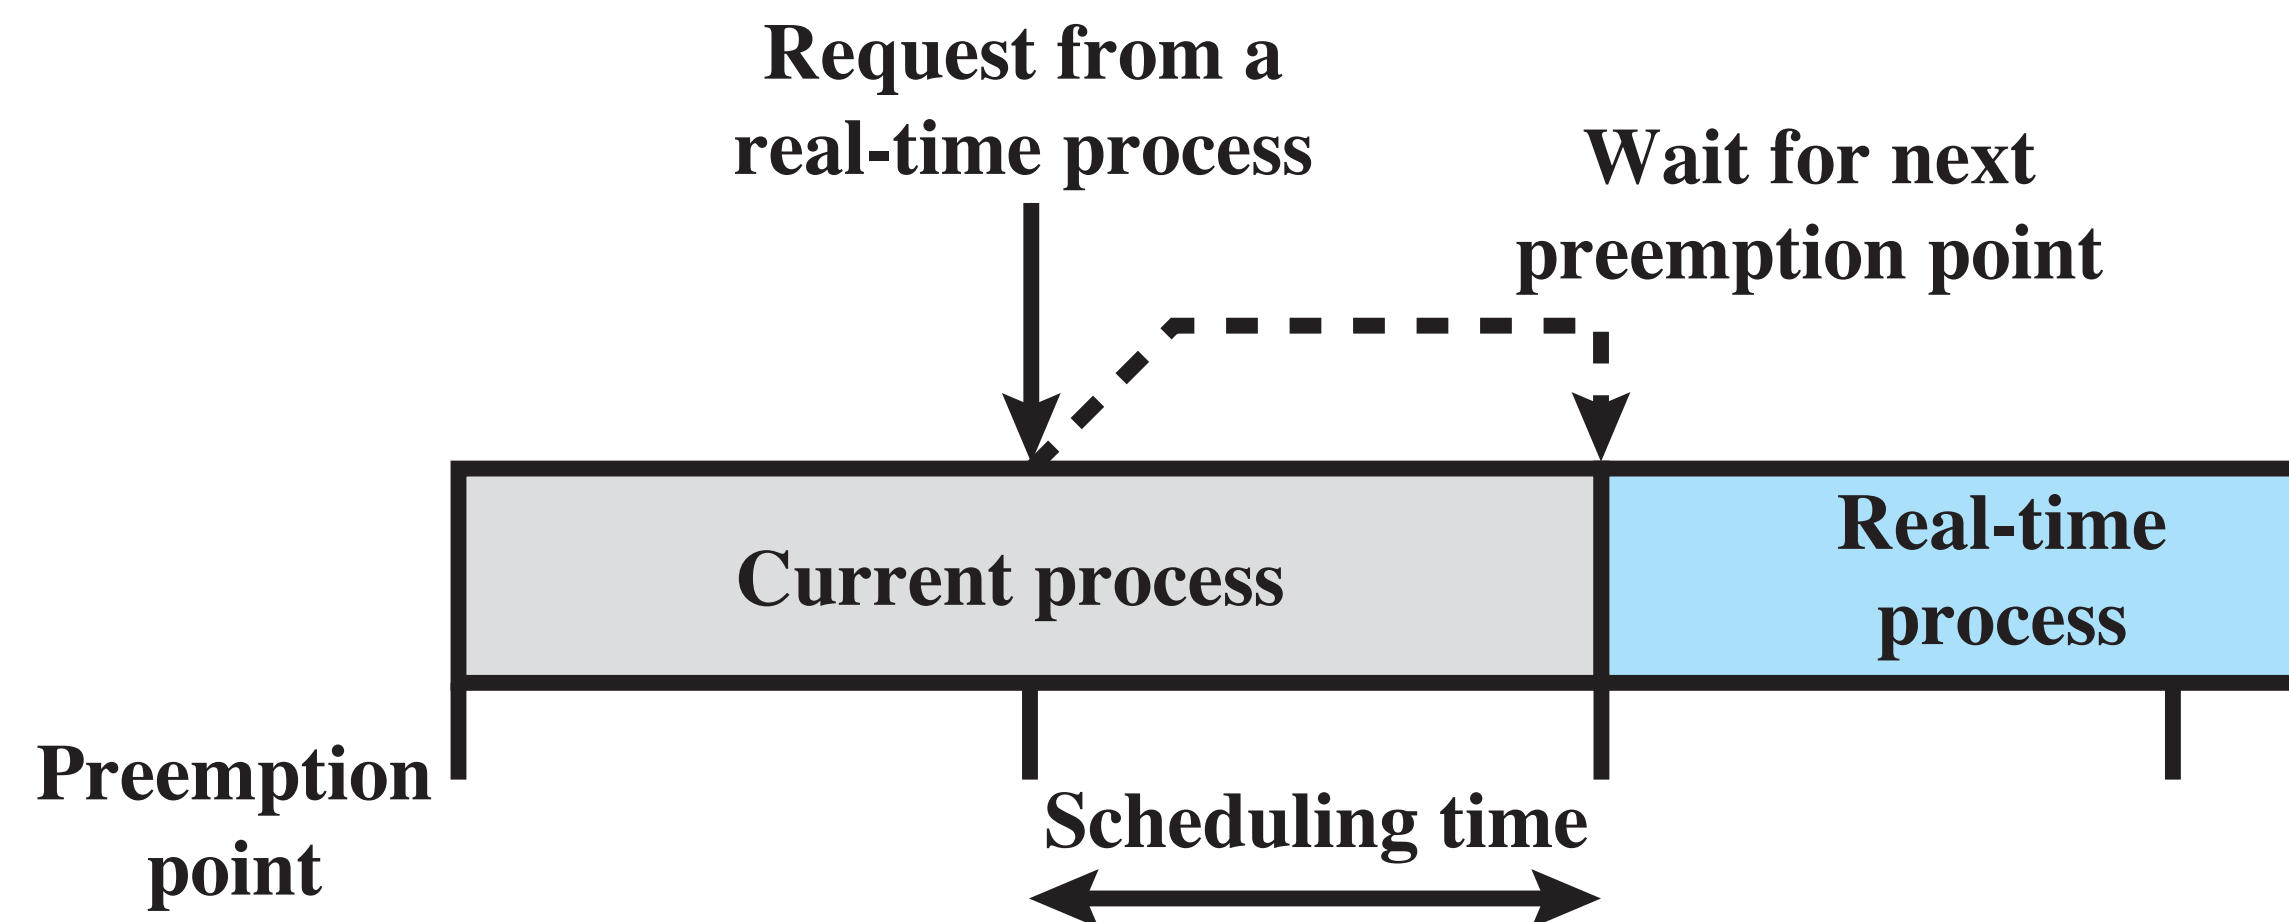
Real-Time Scheduling



(b) Priority-Driven Nonpreemptive Scheduler

Figure 10.4 Scheduling of Real-Time Process

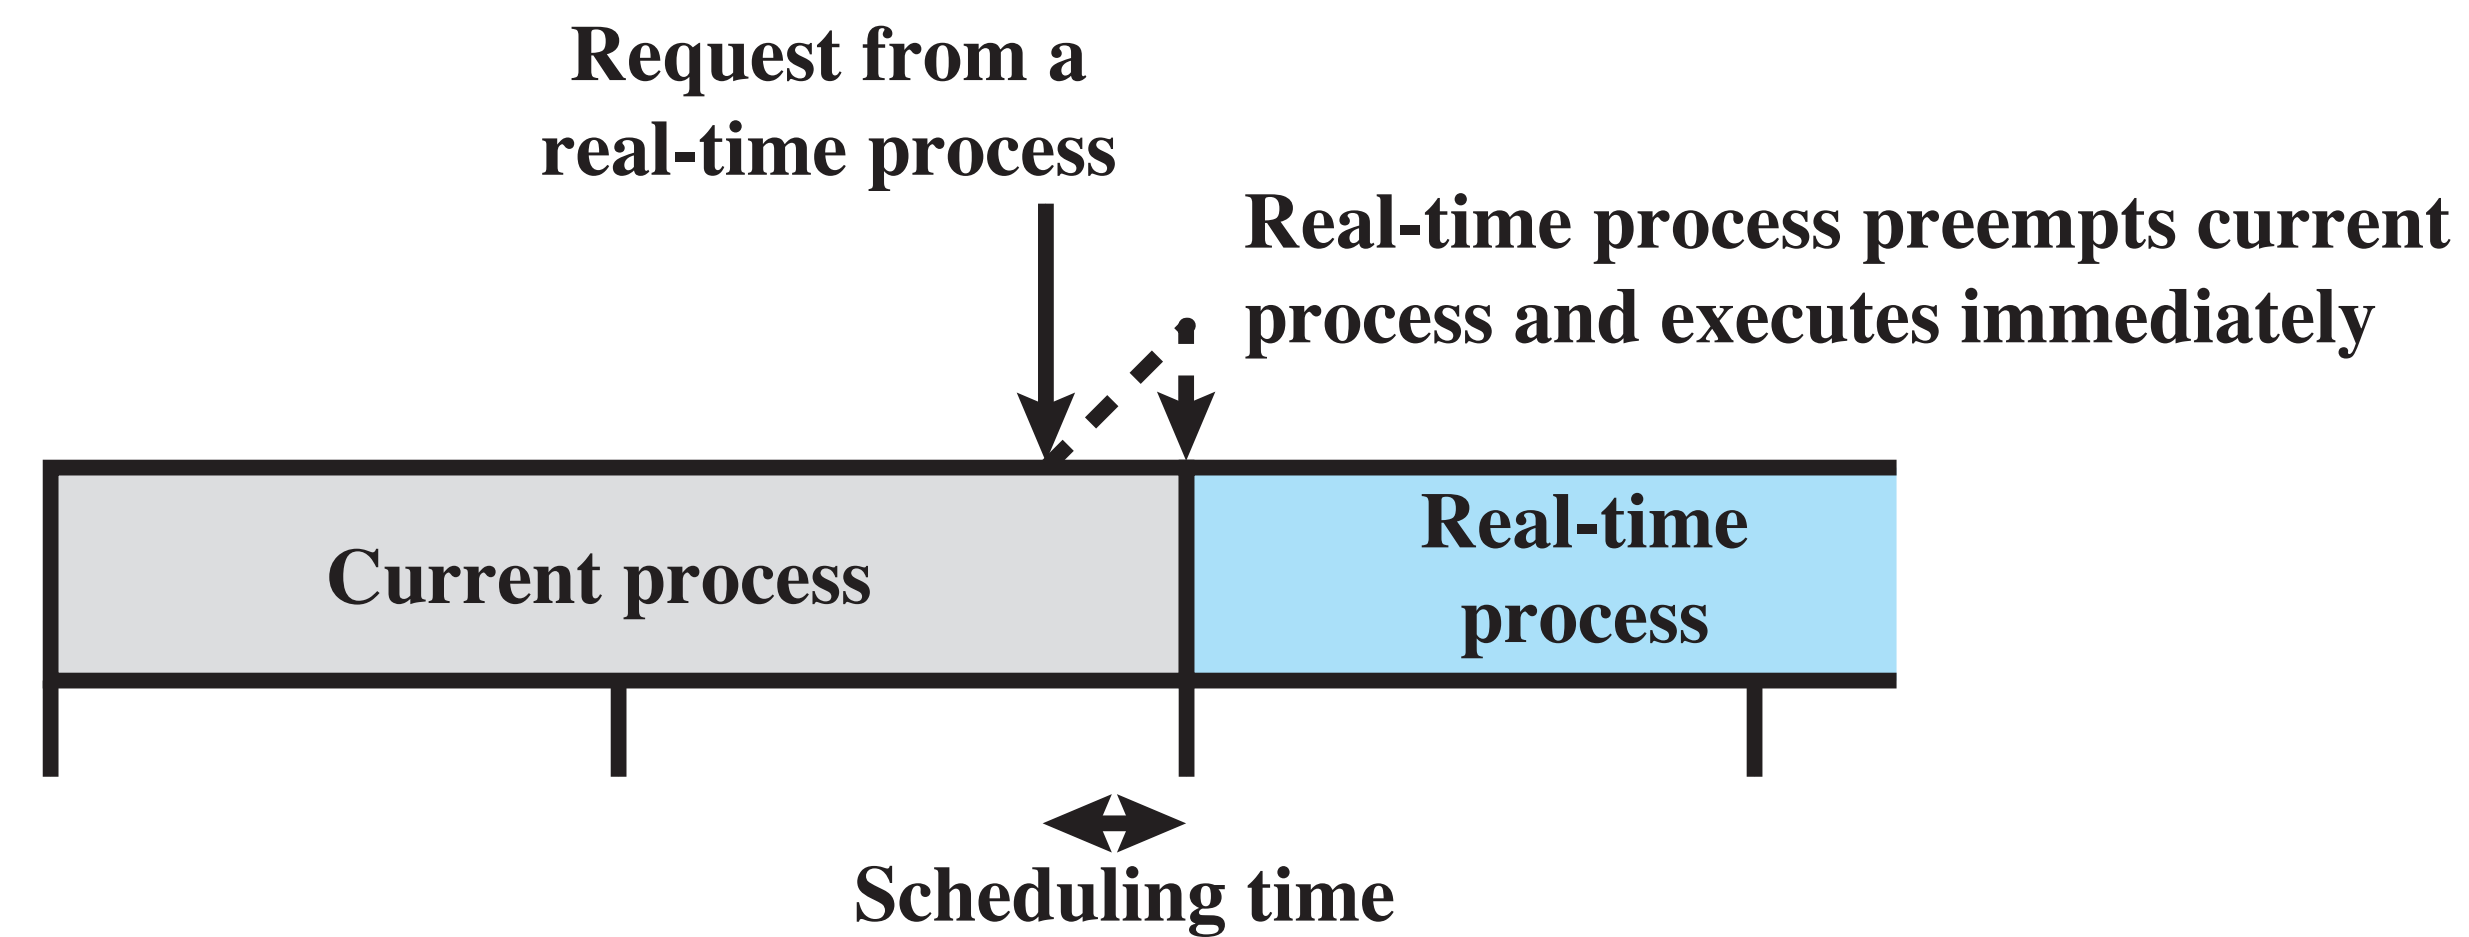
Real-Time Scheduling



(c) Priority-Driven Preemptive Scheduler on Preemption Points

Figure 10.4 Scheduling of Real-Time Process

Real-Time Scheduling



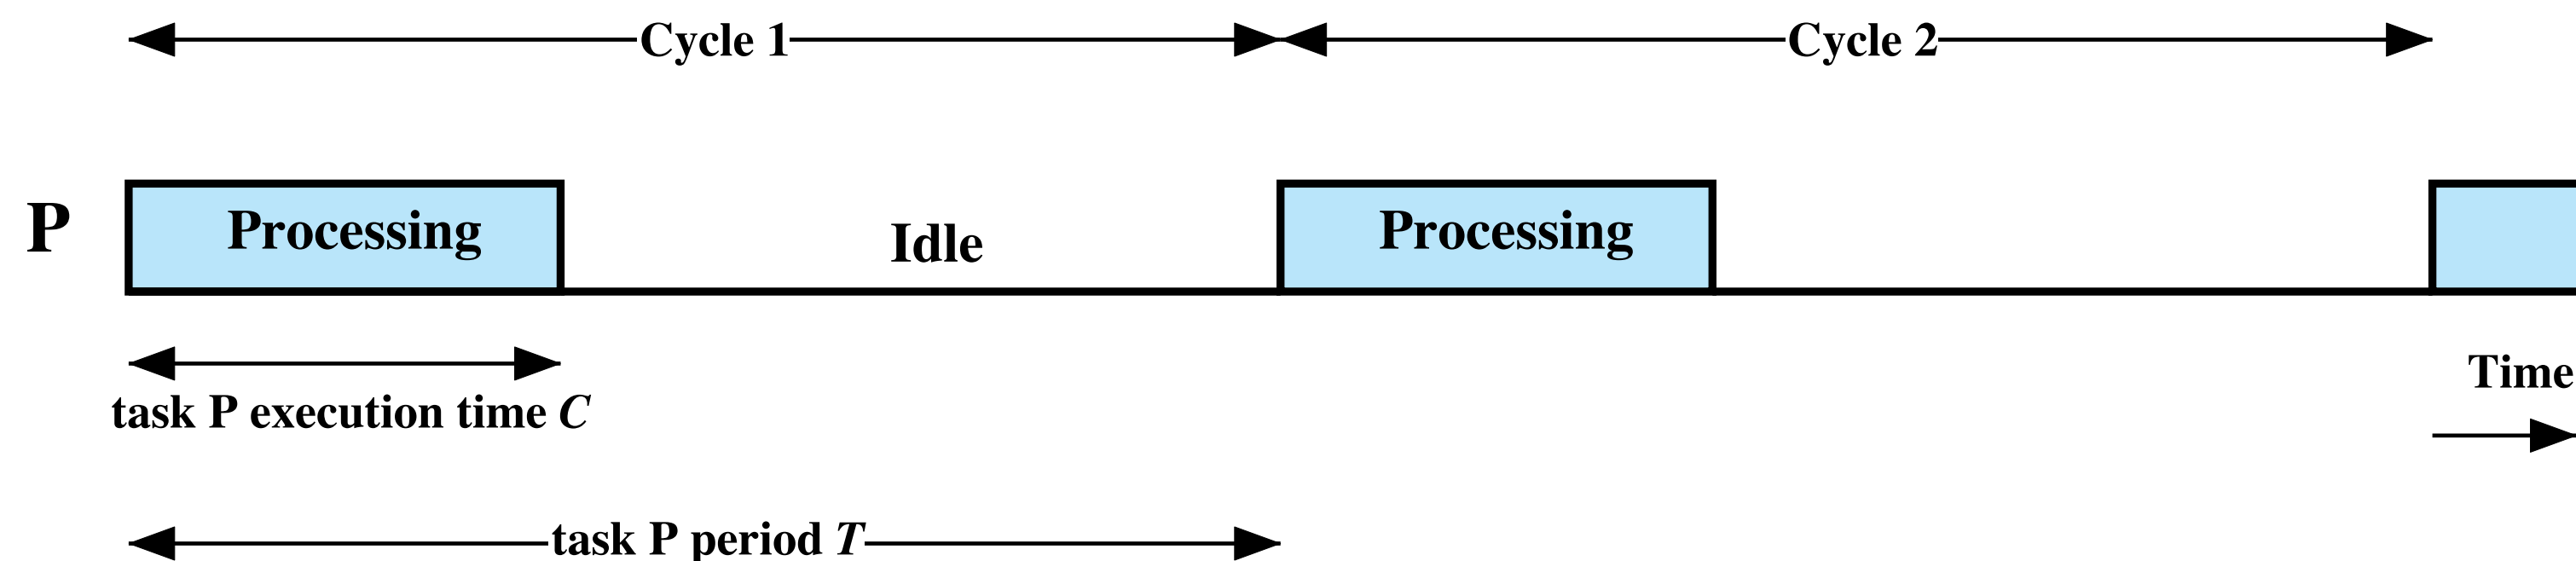
(d) Immediate Preemptive Scheduler

Figure 10.4 Scheduling of Real-Time Process

The *Rate Monolithic (RM)* Scheduling Algorithm


- Schedules processes according to the period.
- The shorter the period, the higher the priority.
- RM is preemptive.

(But only higher priority processes can preempt lower priority processes)

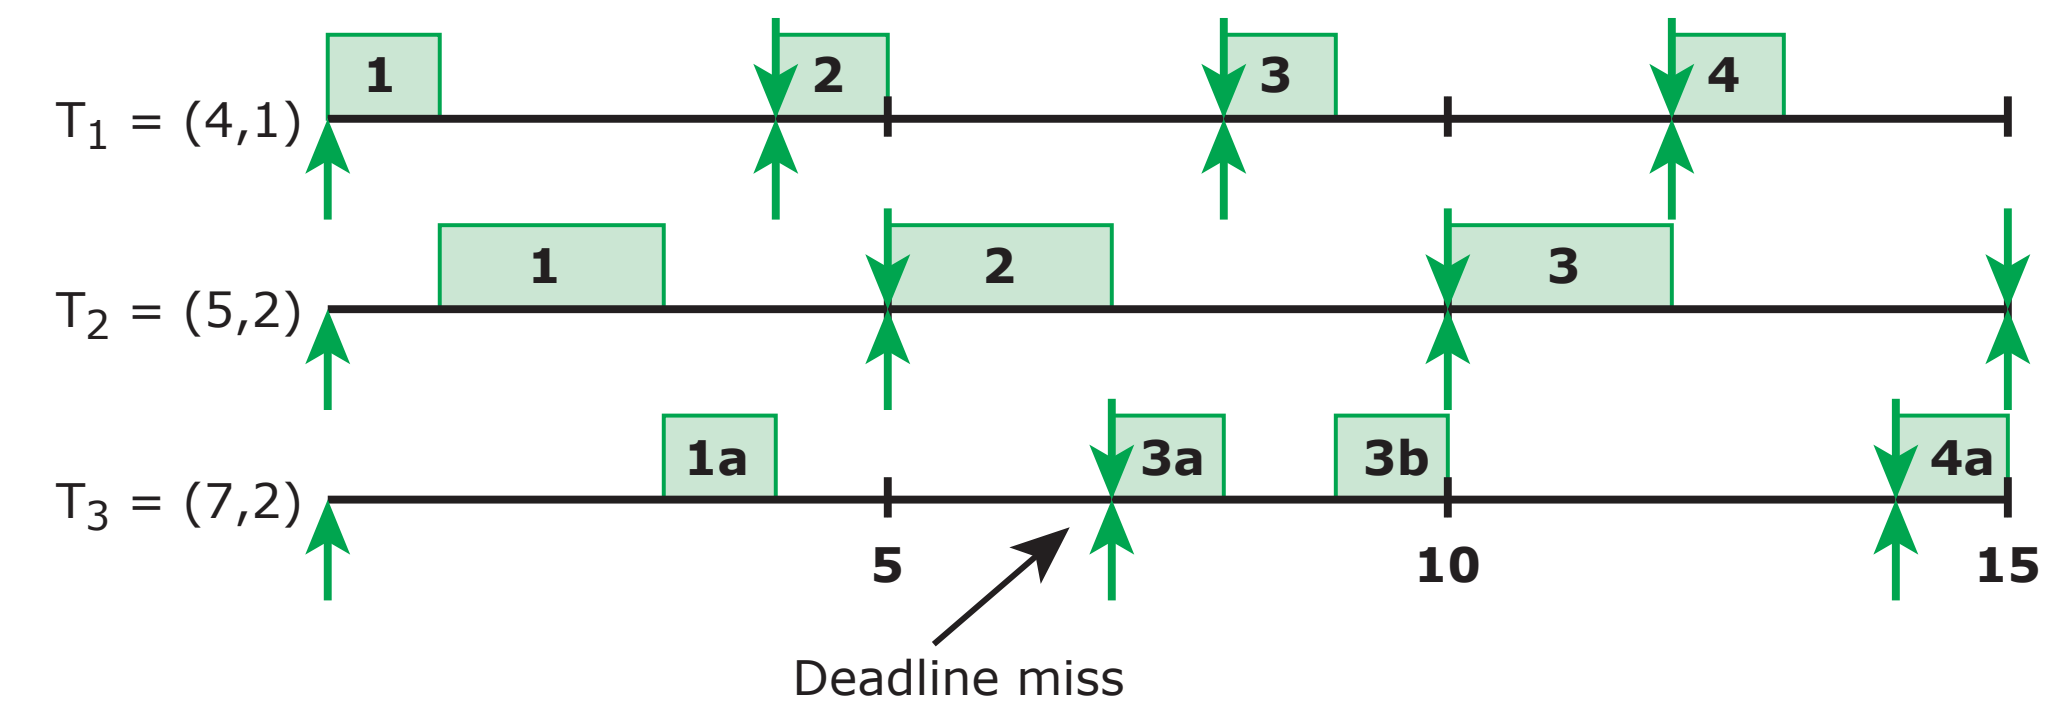


The *Rate Monolithic (RM)* Scheduling Algorithm

- Schedules processes according to the period.
- The shorter the period, the higher the priority.
- RM is preemptive.
(But only higher priority processes can preempt lower priority processes)



(a) Arrival times and deadlines for task $T_i = (P_i, C_i)$;
 P_i = period, C_i = processing time



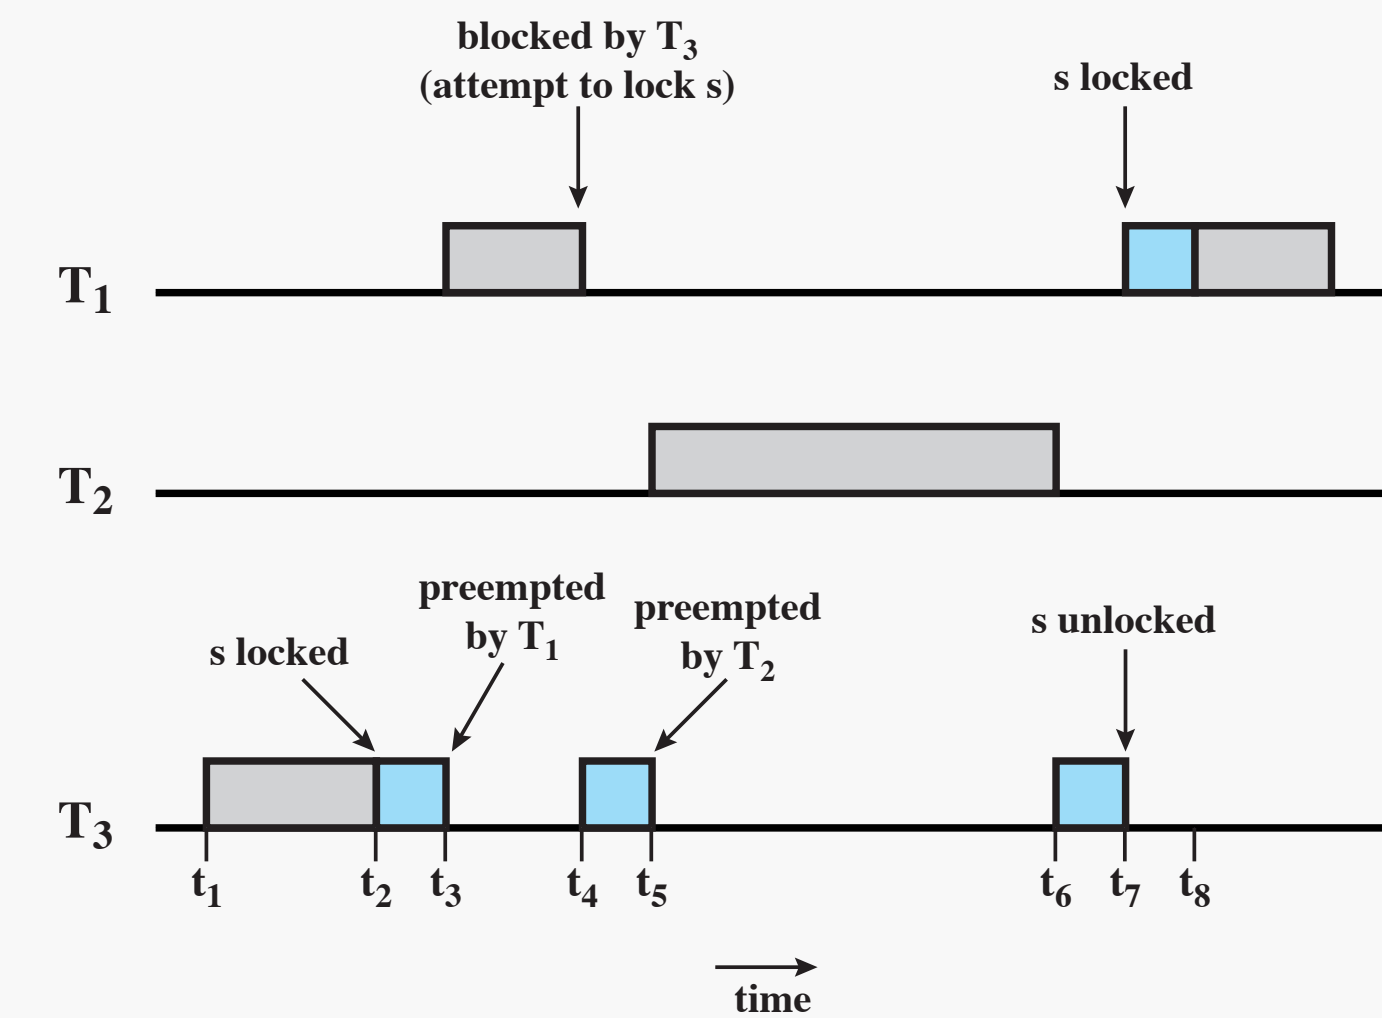
(b) Scheduling results

Figure 10.8 Rate Monotonic Scheduling Example

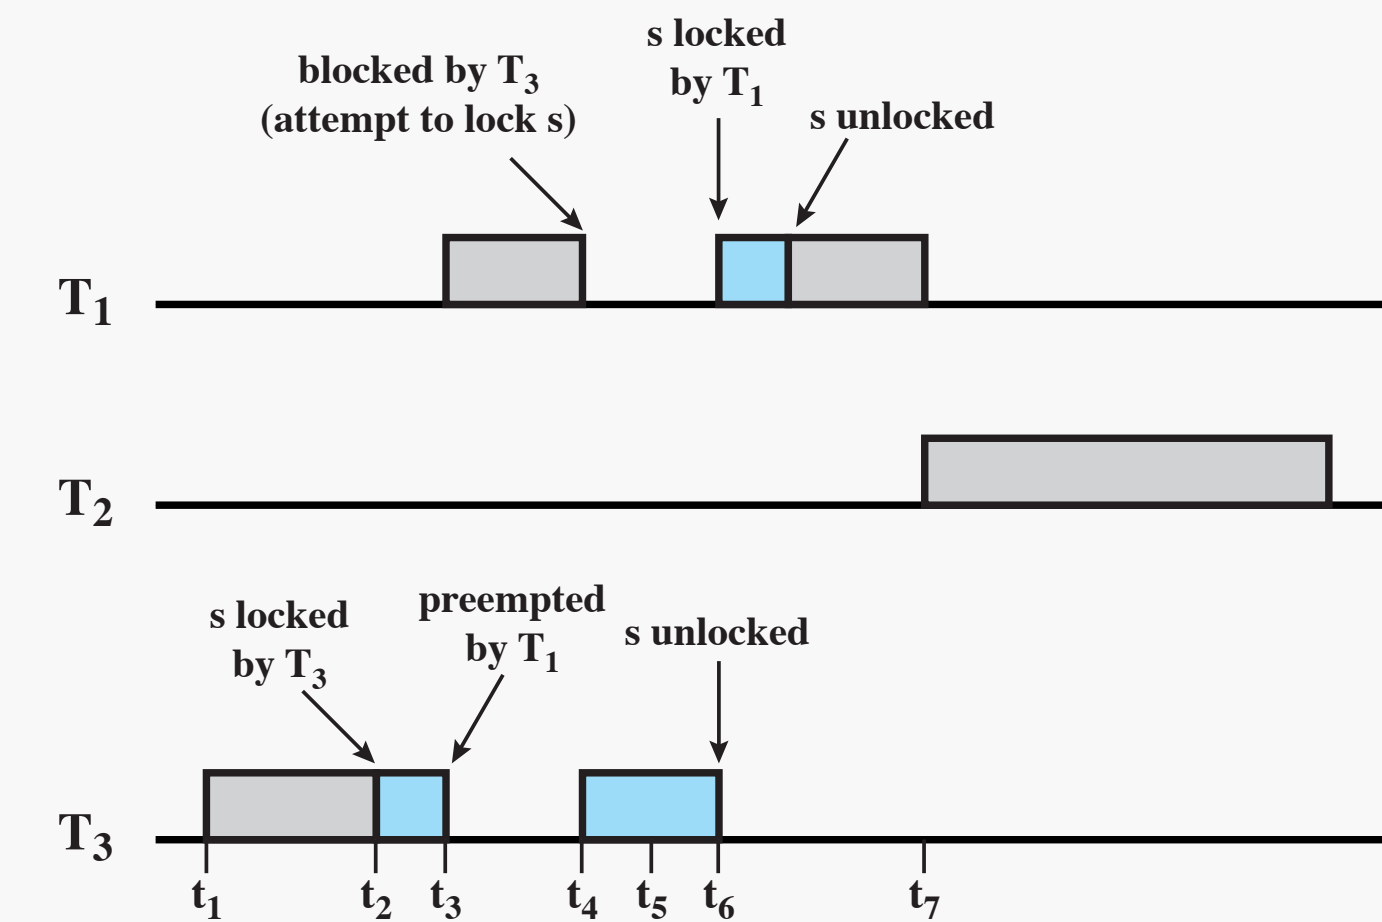
Priorities & Scheduling

- **Priority Inversion** — when a high priority task has to wait on a lower priority task
- **Priority Inheritance** — lower priority task inherits the priority of any higher-priority task pending on a shared resource
- **Priority Ceiling** — associate priorities with resources

Path Finder Fail...



(a) Unbounded priority inversion



(b) Use of priority inheritance

■ normal execution ■ execution in critical section

Figure 10.9 Priority Inversion

Linux Scheduling

- Scheduling Classes

- **SCHED_FIFO** (limited preemption) & **SCHED_RR** (time-sliced) — (0-99)
- **SCHED_NORMAL** (100-139)
 - lower value == higher priority
 - non-RT thread only executes if no RT threads are ready to execute

- Linux 2.6+ → **O(1) Scheduler**

- Named so because time to select & run a process takes constant time.
- Complex and not good to run in the kernel....

- Linux 2.6.23+ → **Completely Fair Scheduler (CFS)**

- Maintain **virtual runtime** value for each task (normalized amt. of time spent executing so far)
- **Sleeper Fairness** (i.e., ensure that processes that wait on I/O get fair access to processor)
- Use Red-Black Tree to order runnable tasks (efficient inserts/deletes, and searches)

