

Memory (Part IV): **Virtual Memory**

Professor Travis Peters
CSCI 460 Operating Systems
Fall 2019

Some slides & figures adapted from Stallings instructor resources.

Some slides adapted from Adam Bates's F'18 CS423 course @ UIUC
<https://courses.engr.illinois.edu/cs423/sp2018/schedule.html>

Today

Announcements

- Project Proposal ***Due Friday (11/01)!***
- HW5 posted — ***Also Due Friday (11/01)!***
- PA2 posted — ***Due NEXT Friday (11/08)...*** get started ASAP... ;)

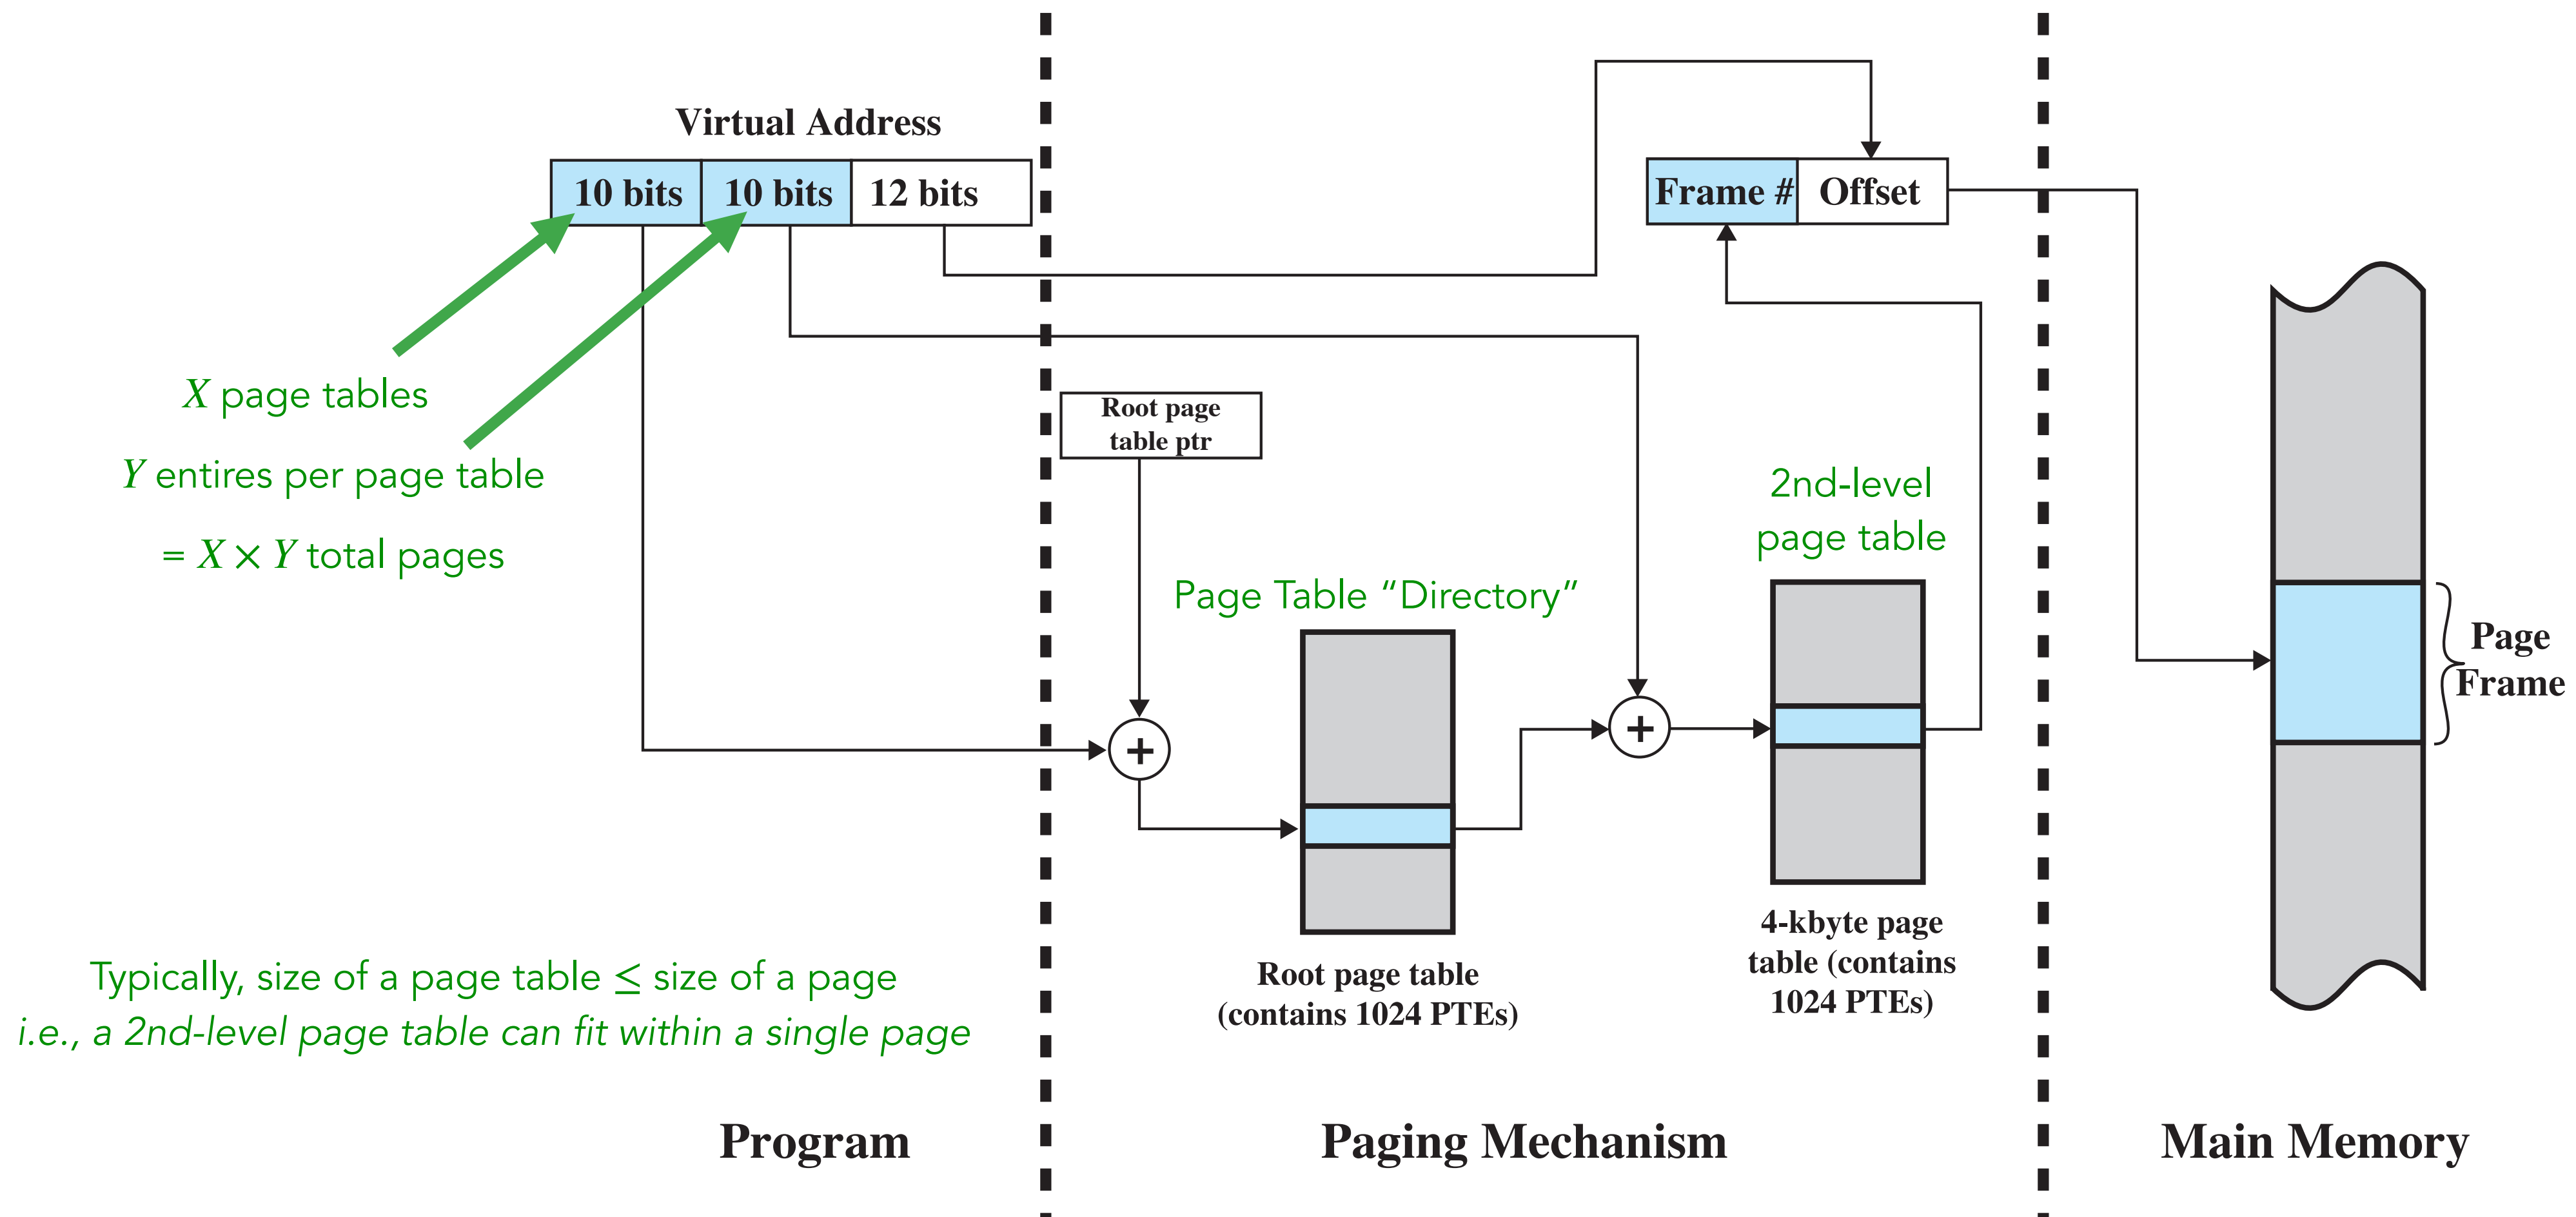
Goals & Learning Objectives

- Wrapping up virtual memory
- Discuss some issues closer to real-world implementations of virtual memory

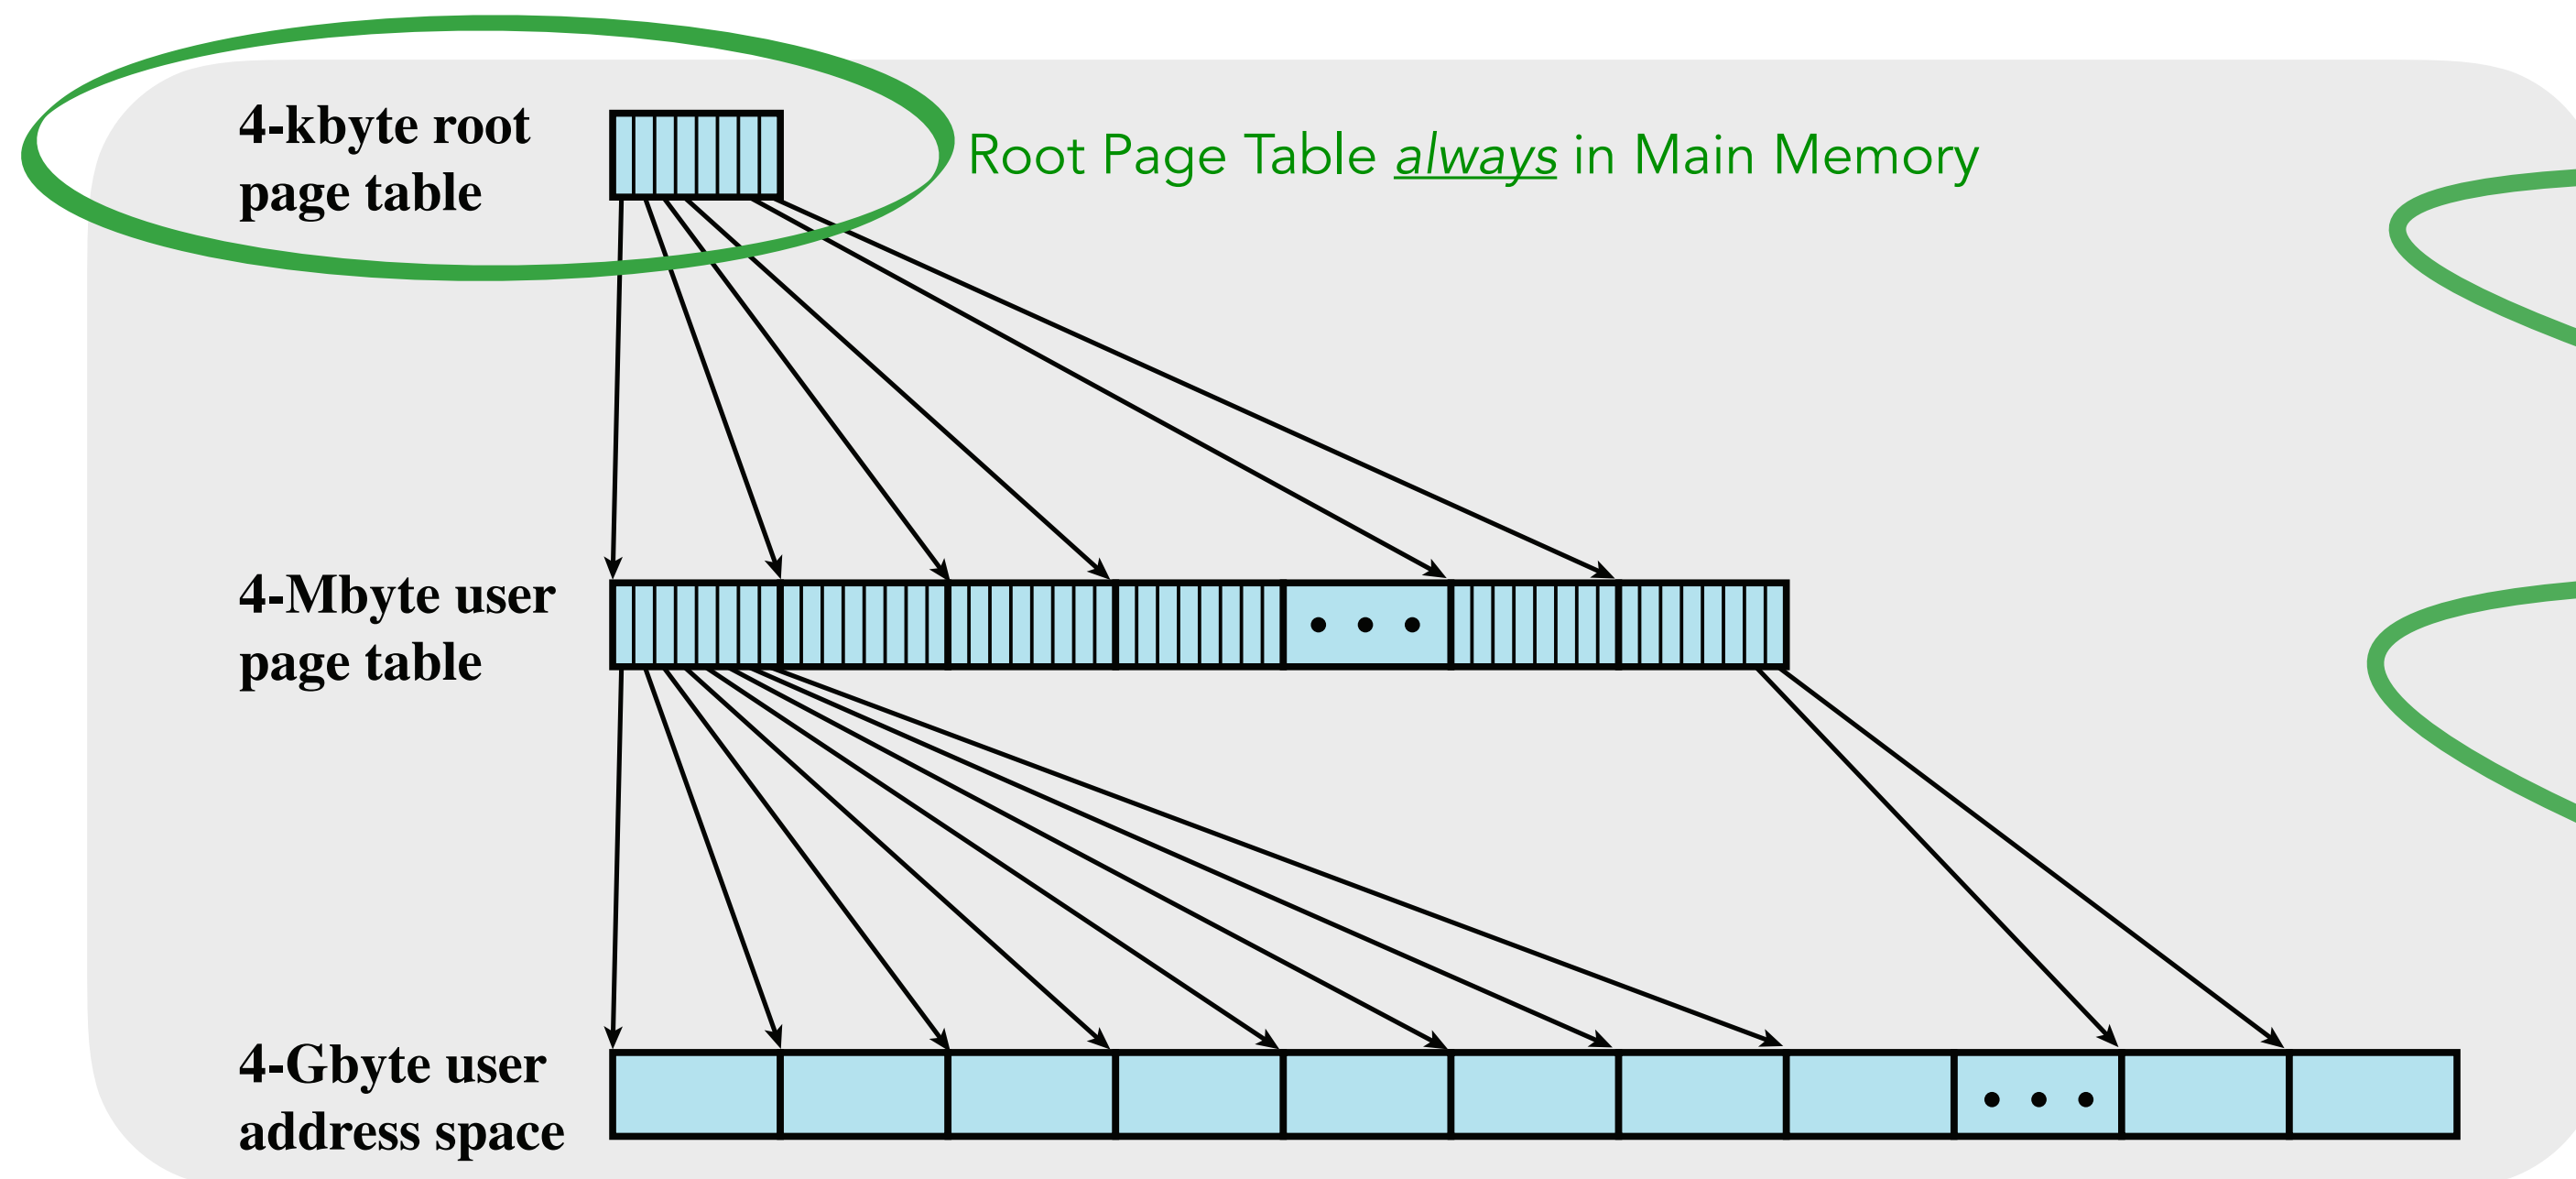
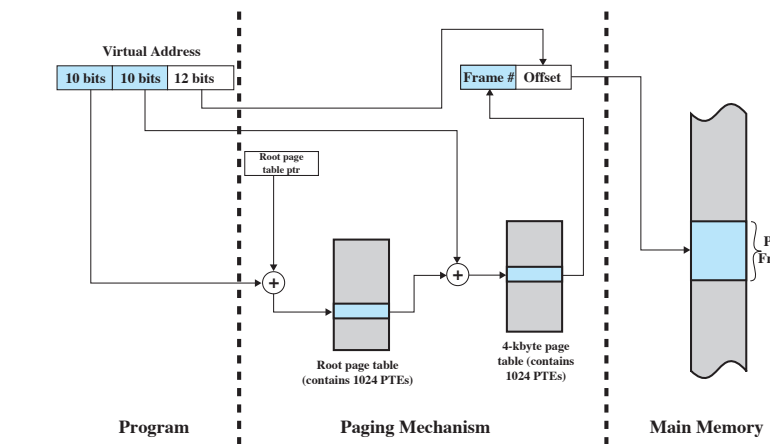
Last Time in CS 460...

Address Translation in 2-Level Paging System

Page tables themselves could be paged in/out...
 but at least *part* of a page table must be in MM for a process to run... → **2-Level Paging System**



Address Translation in 2-Level Paging System (cont.)



Each **4 MB page table** is mapped by root page table → 1024 page tables mapped by 4-byte PTE → root page table size = 4 kB

Each **4 kB page** is mapped by 4-byte PTE → user page table requires 4 MB (2^{22}) ... or according to this scheme, $2^{22} \div 2^{12} = 2^{10} = 1\text{K}$ pages

Example: 2-level scheme,

- 4 GB (2^{32}) address space,
- byte-level addressing,
- 4 kB (2^{12}) page size
→ 2^{20} total pages

Inverted Page Table Structure

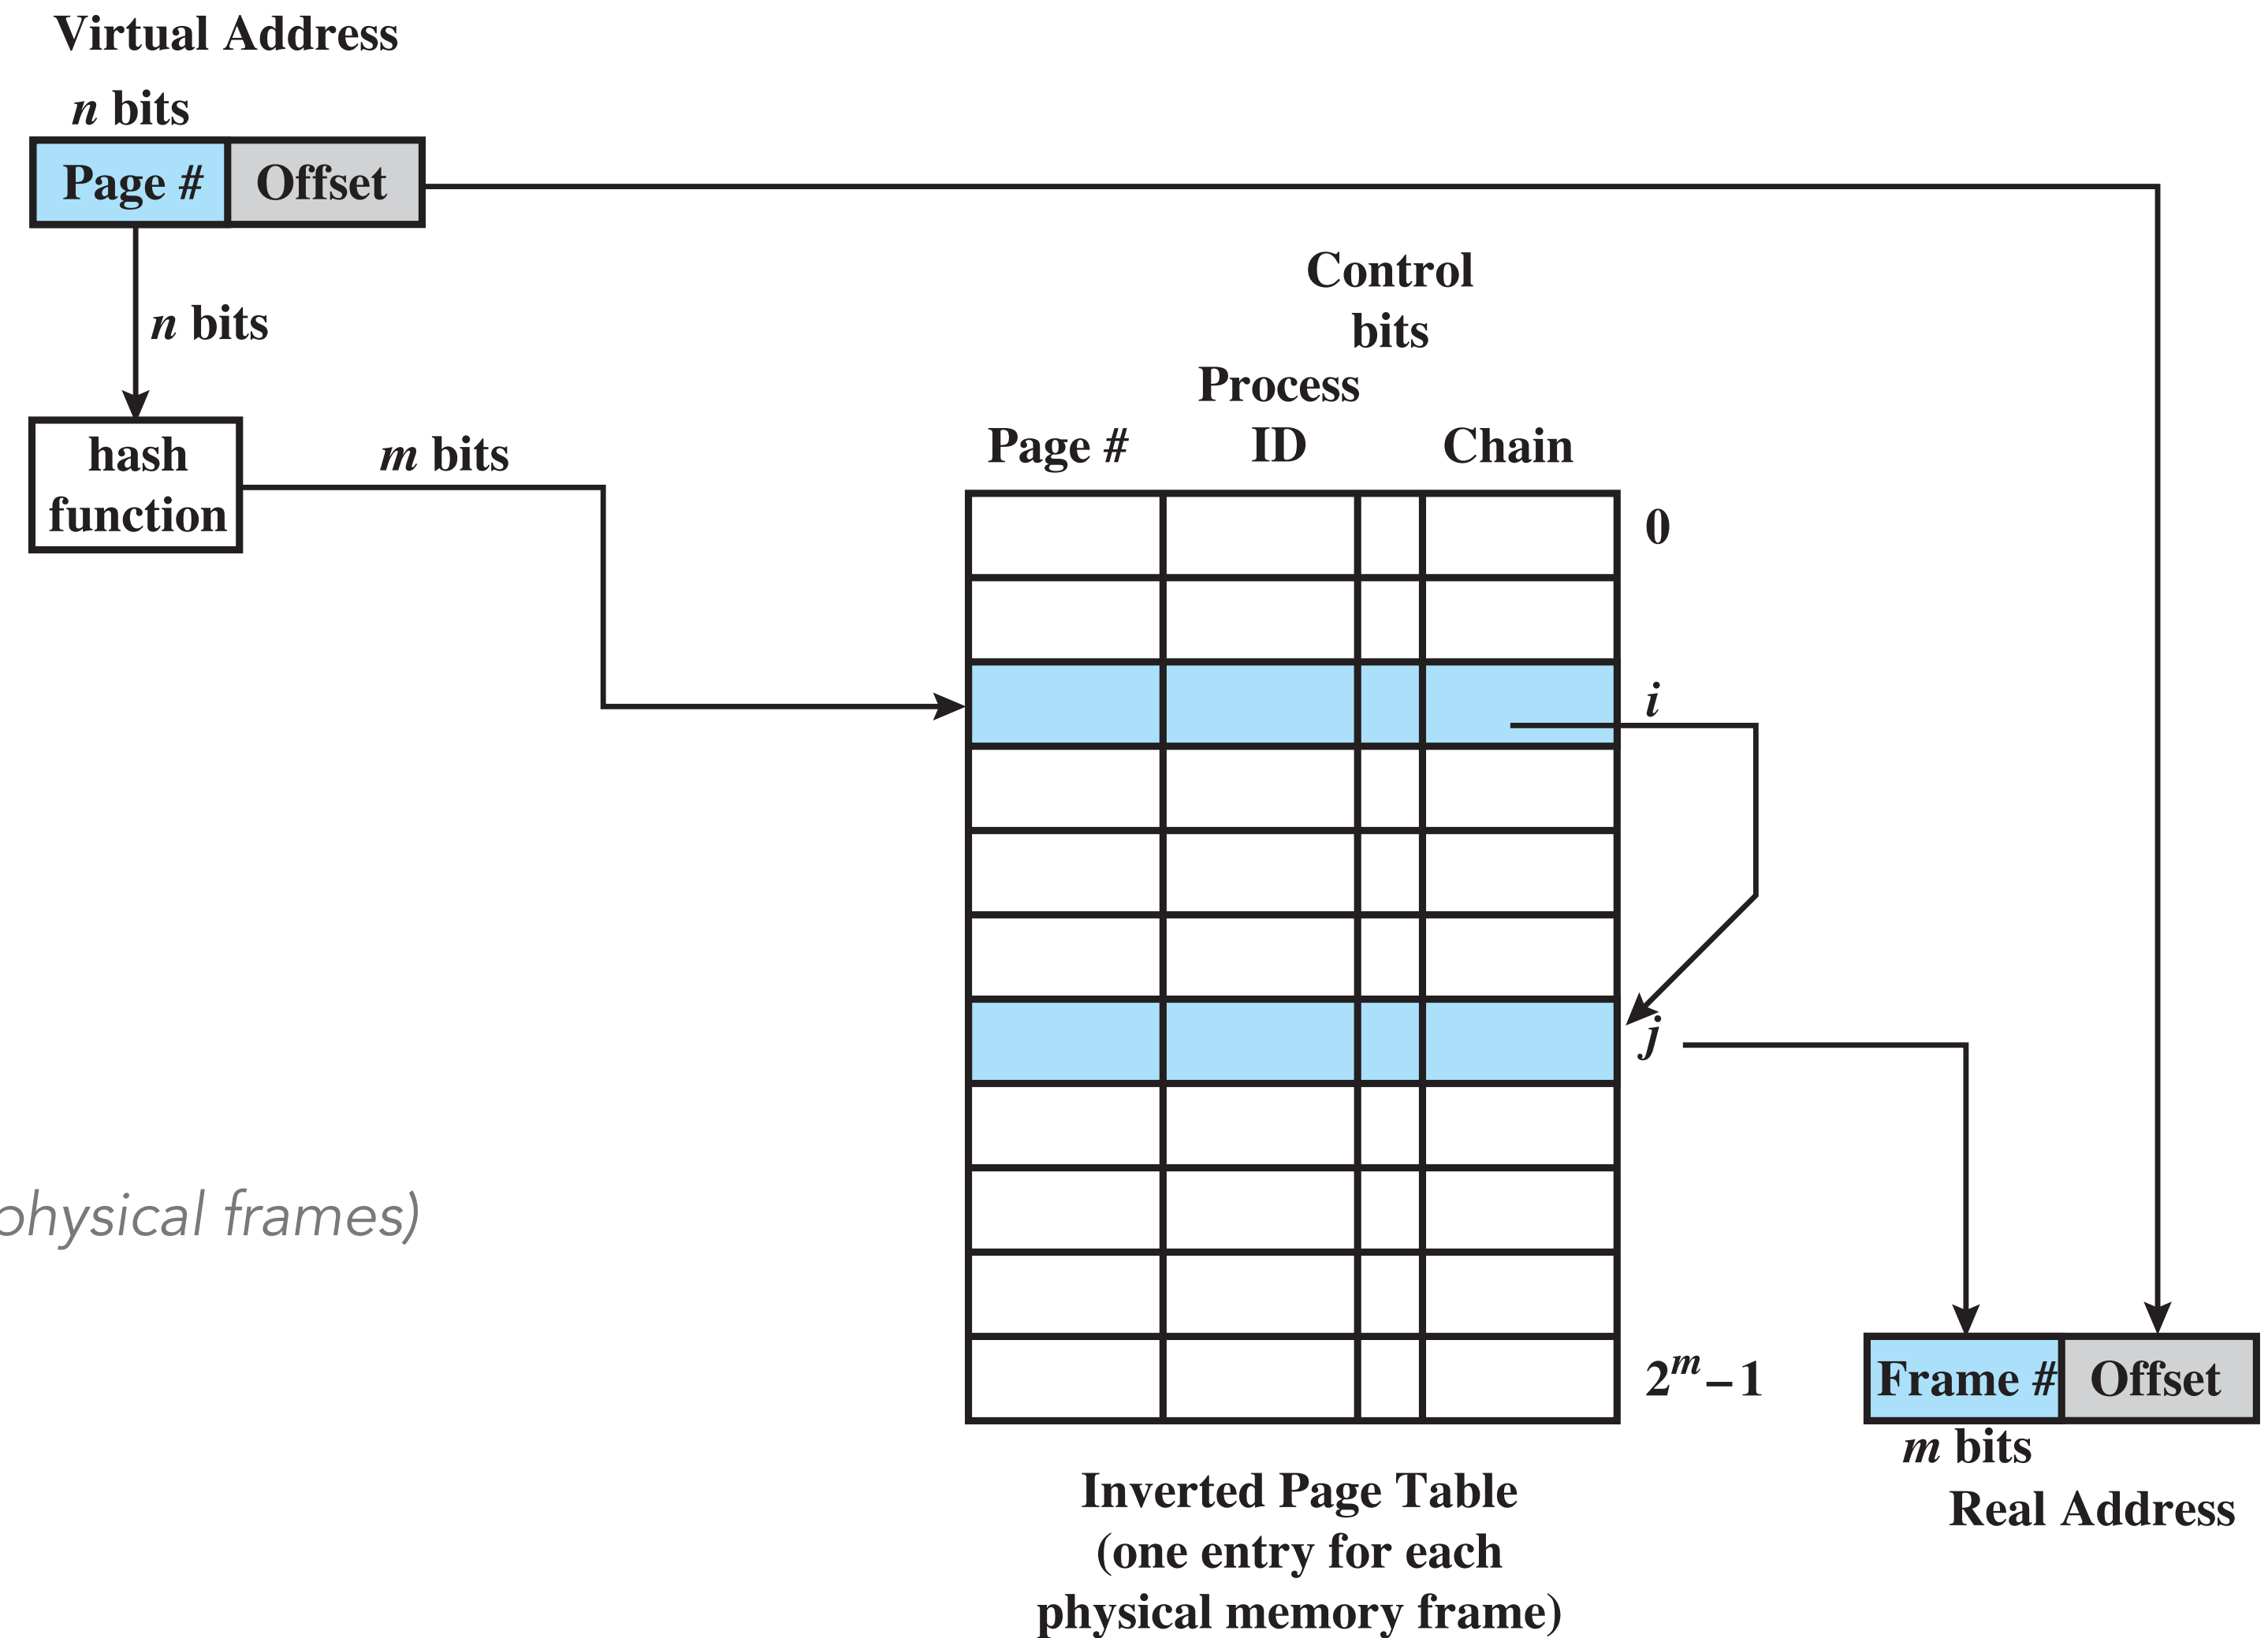
Problem

size of 1+ level page tables is that they are proportional to the size of the virtual address space

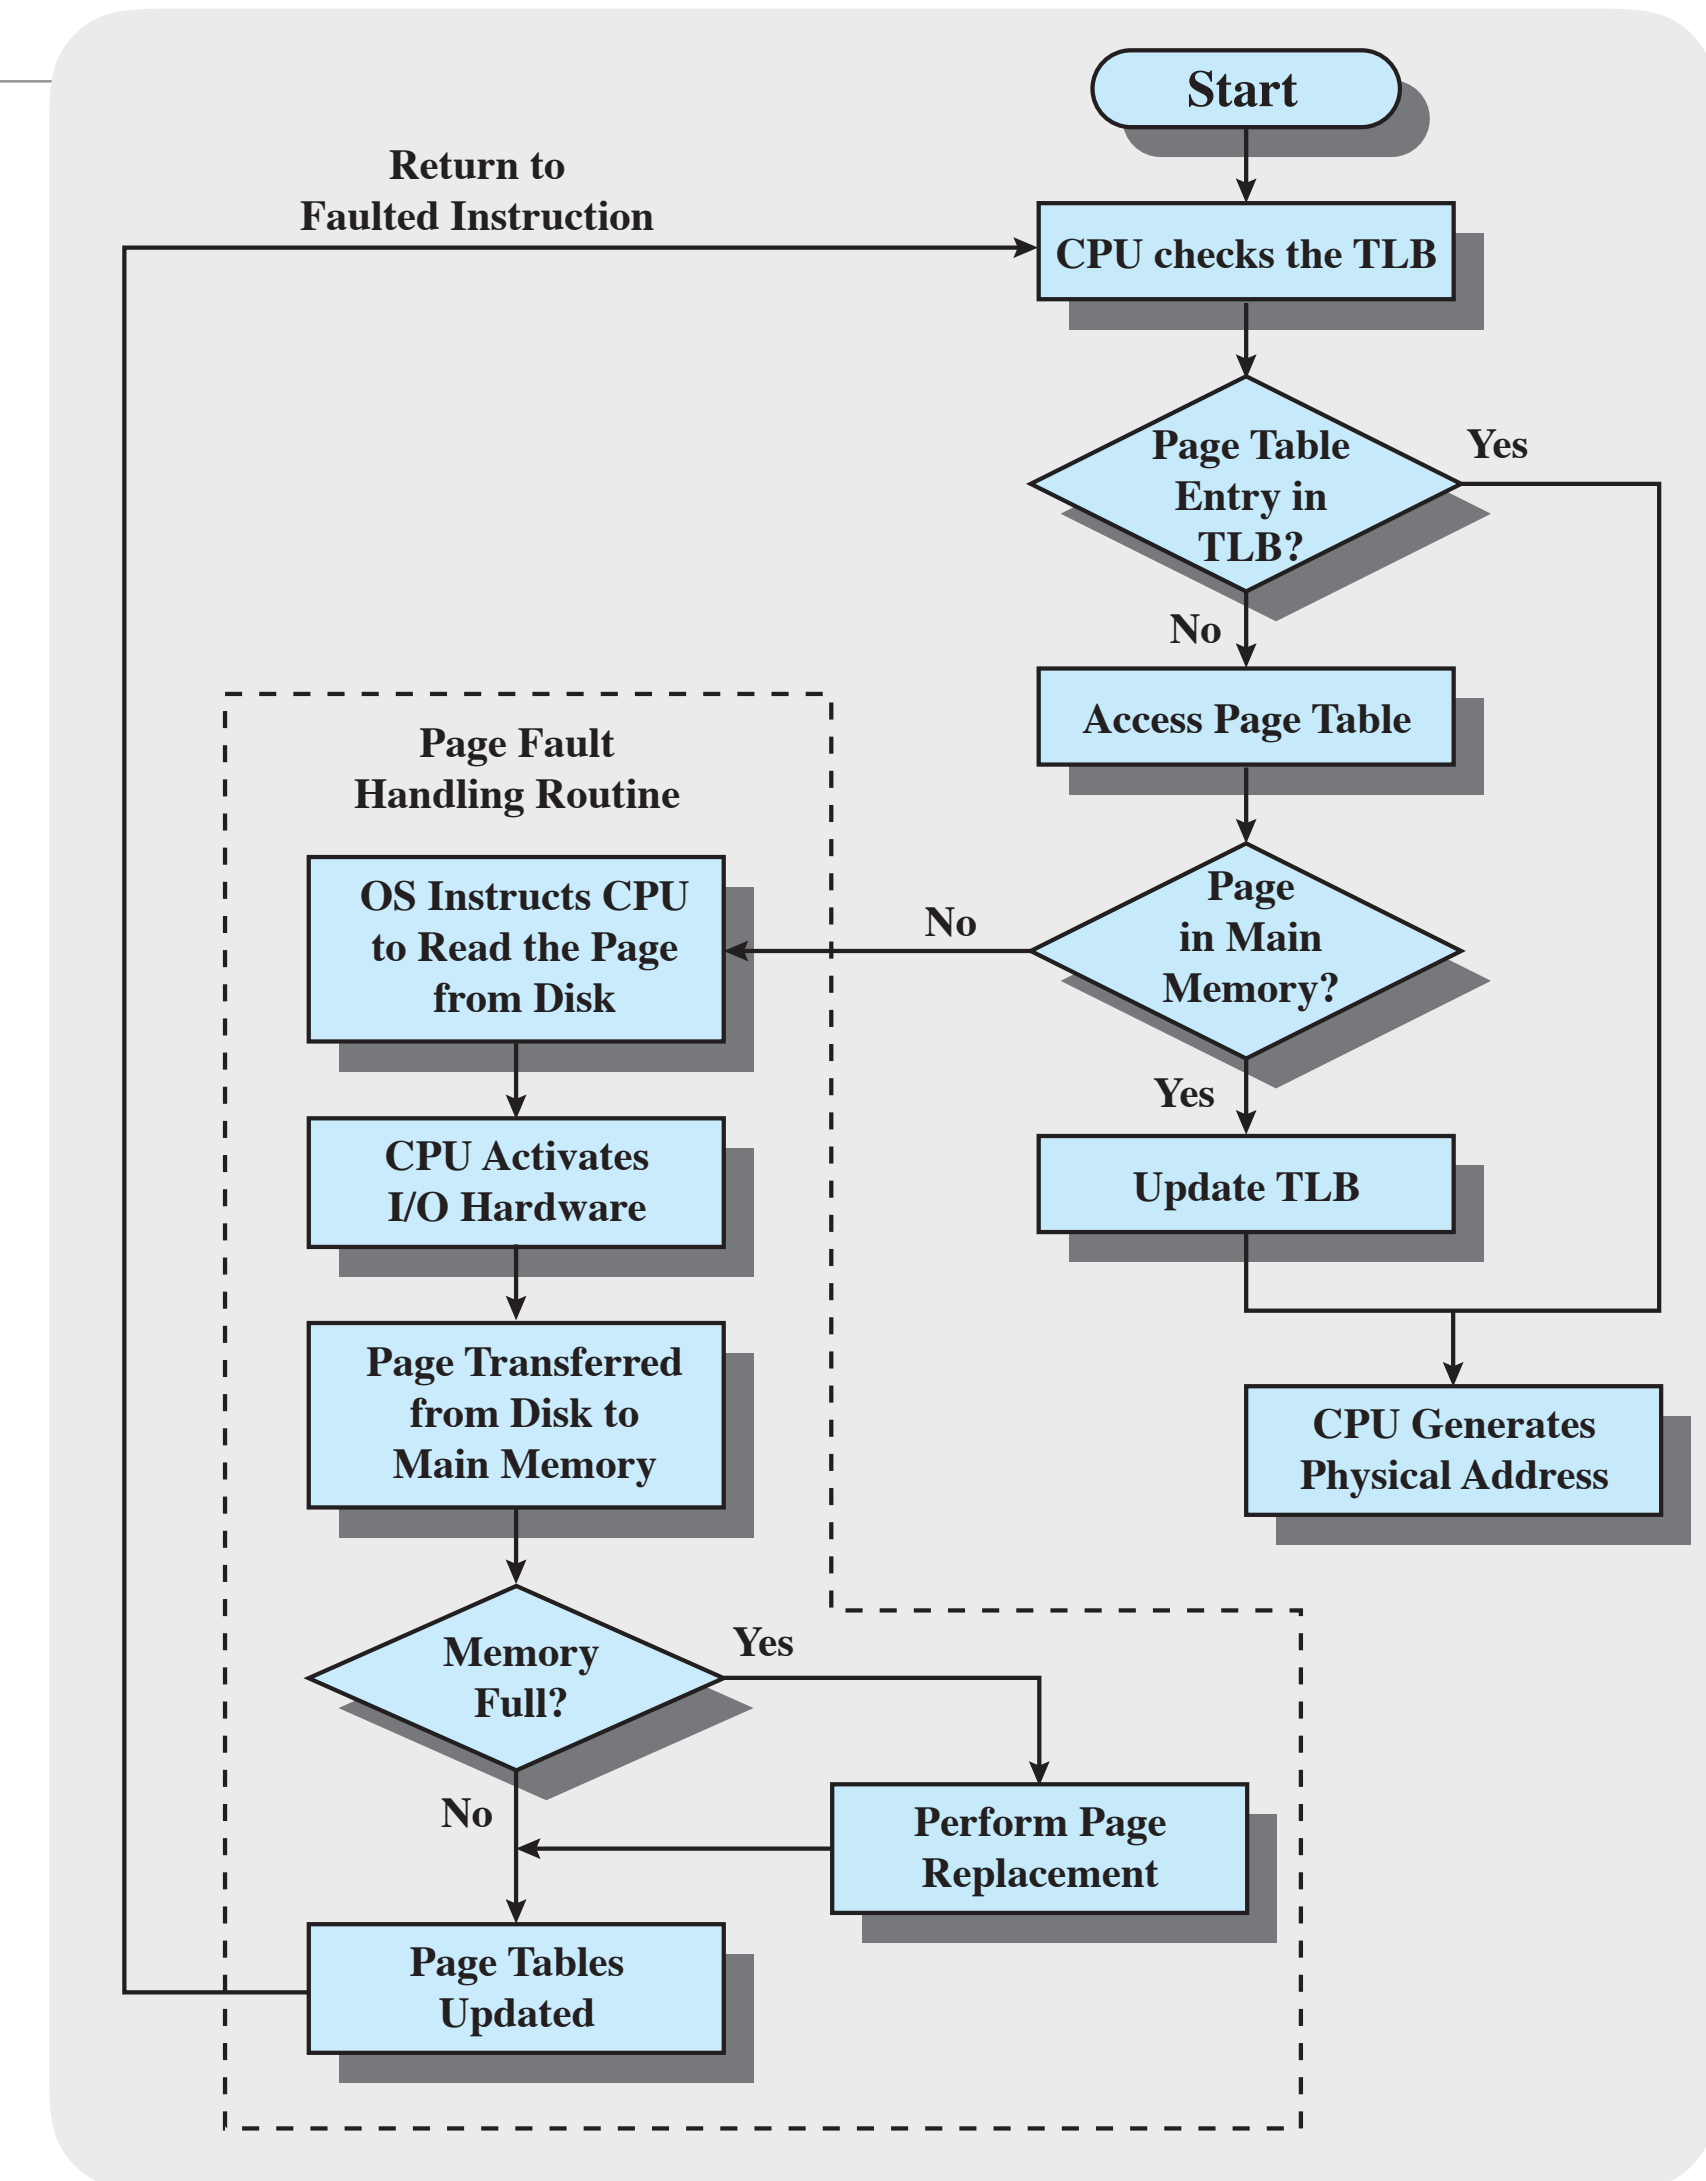
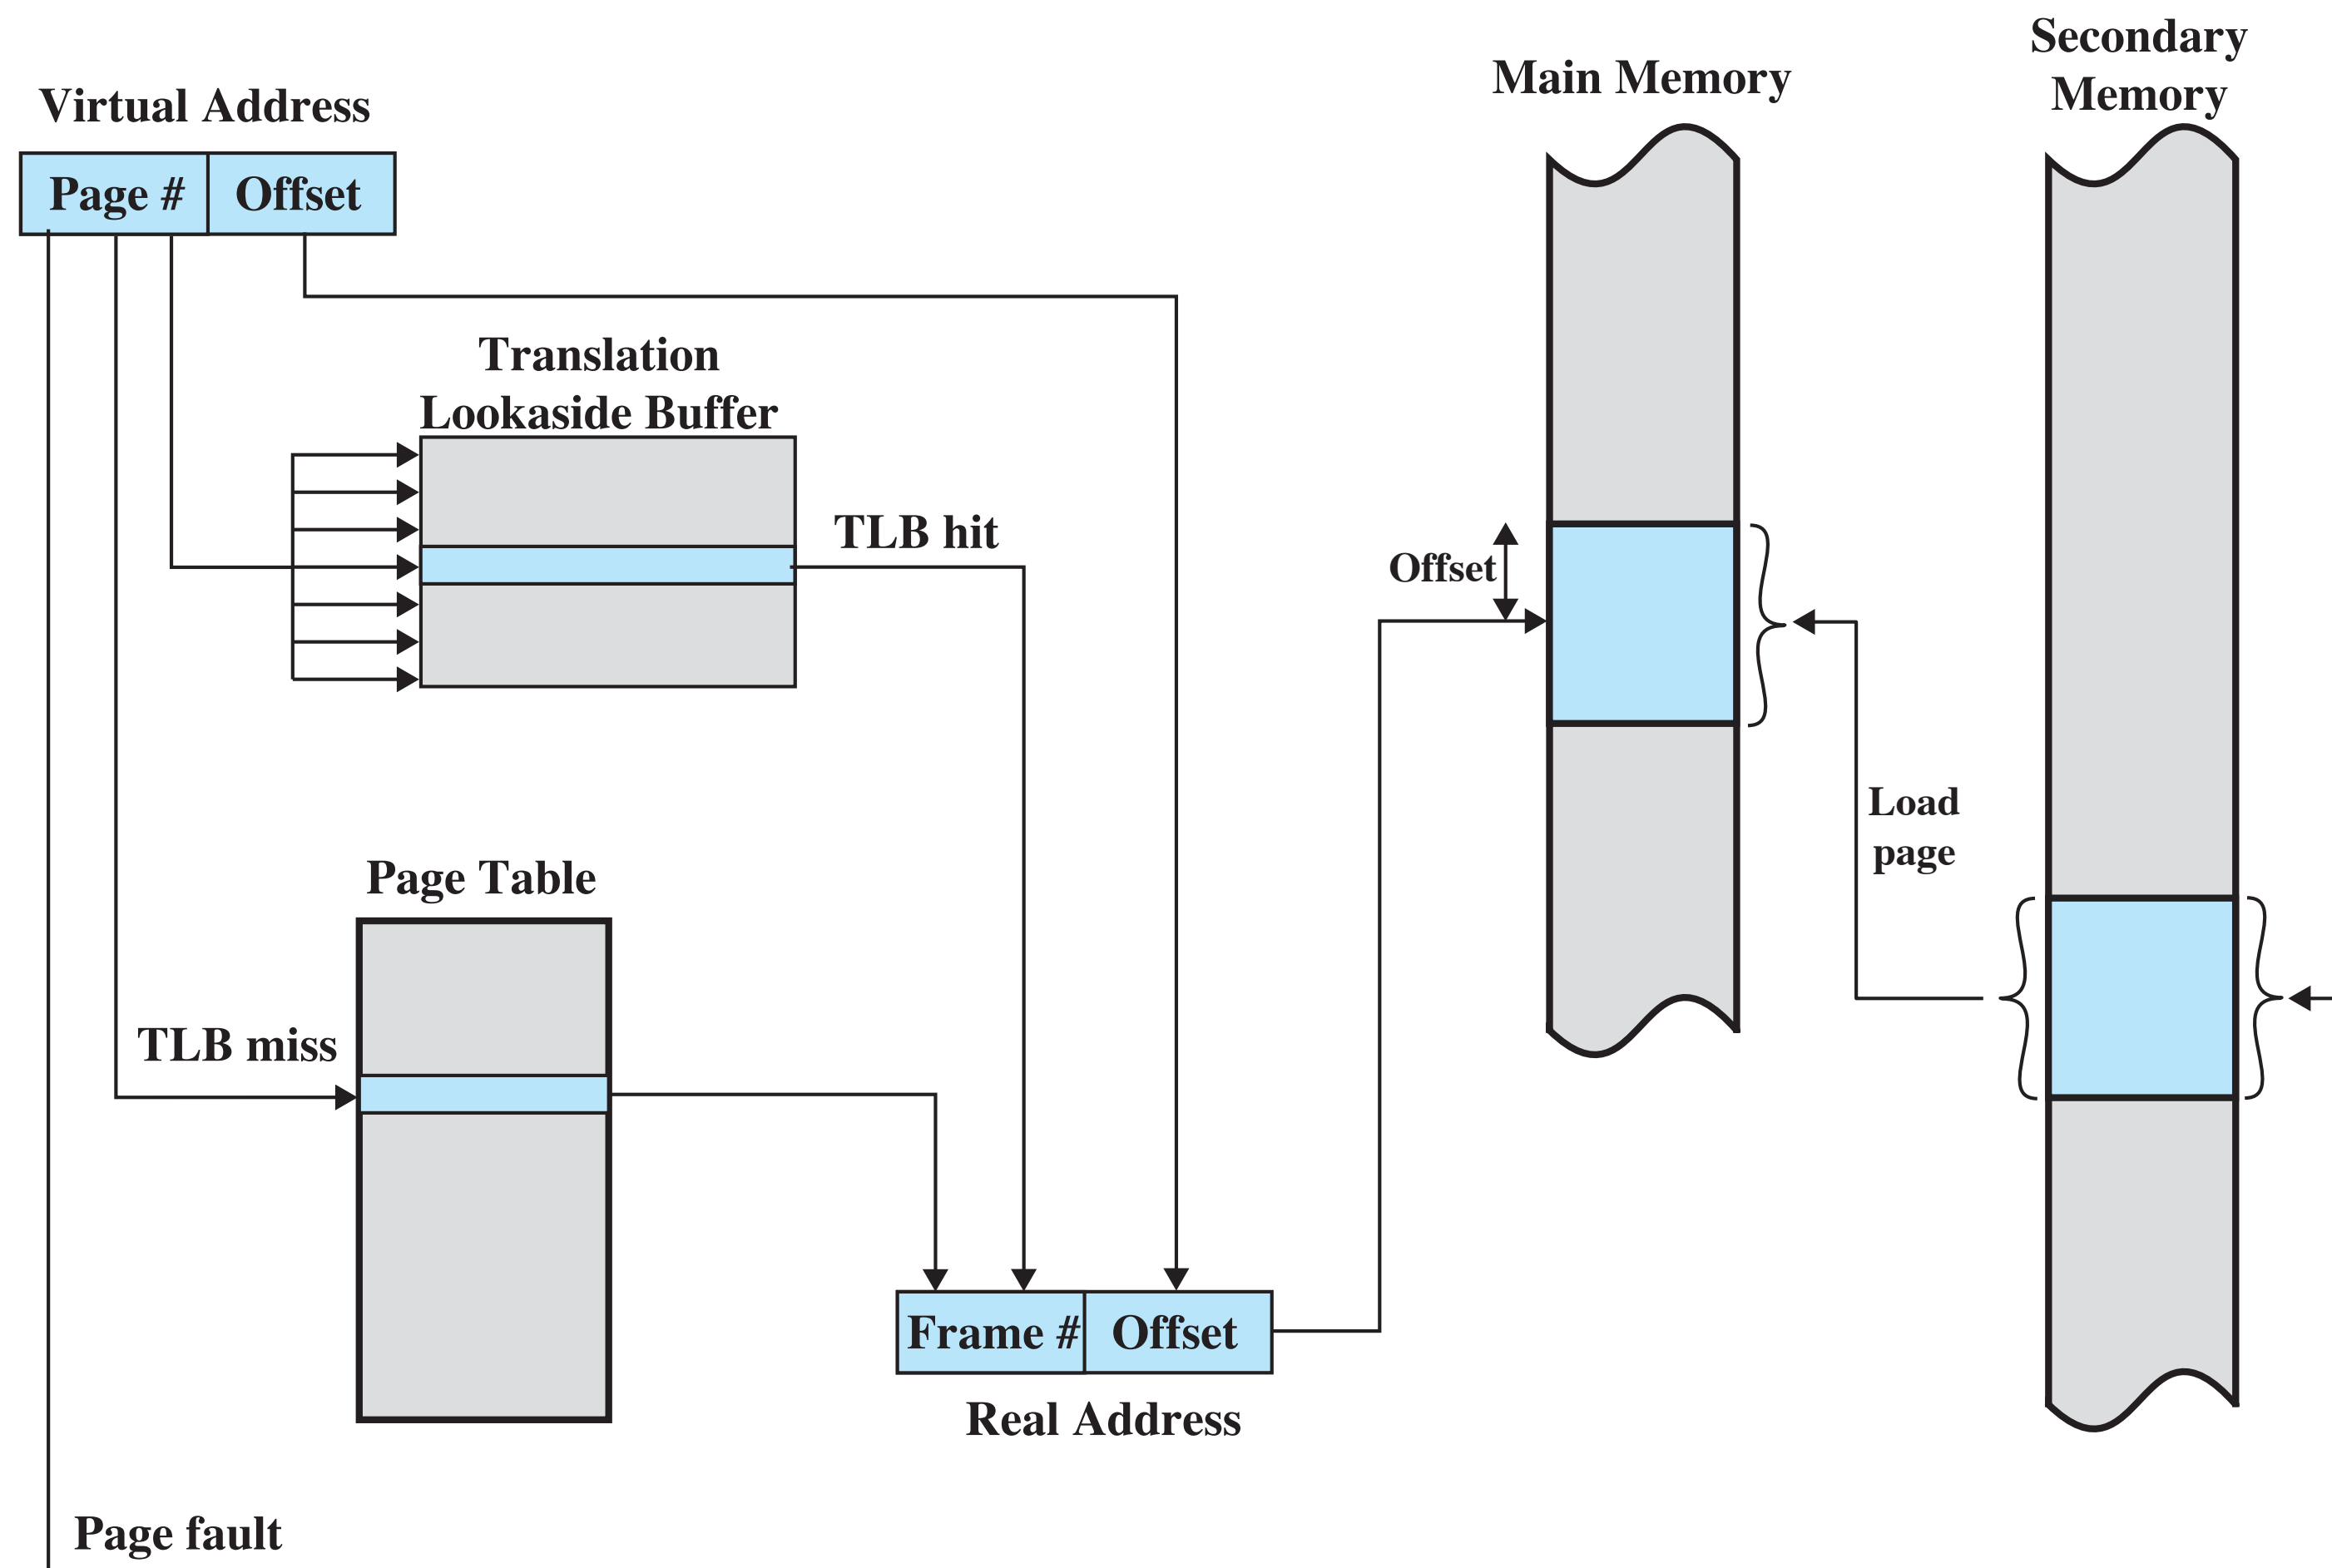
Idea!

Invert page table (index by physical frames)

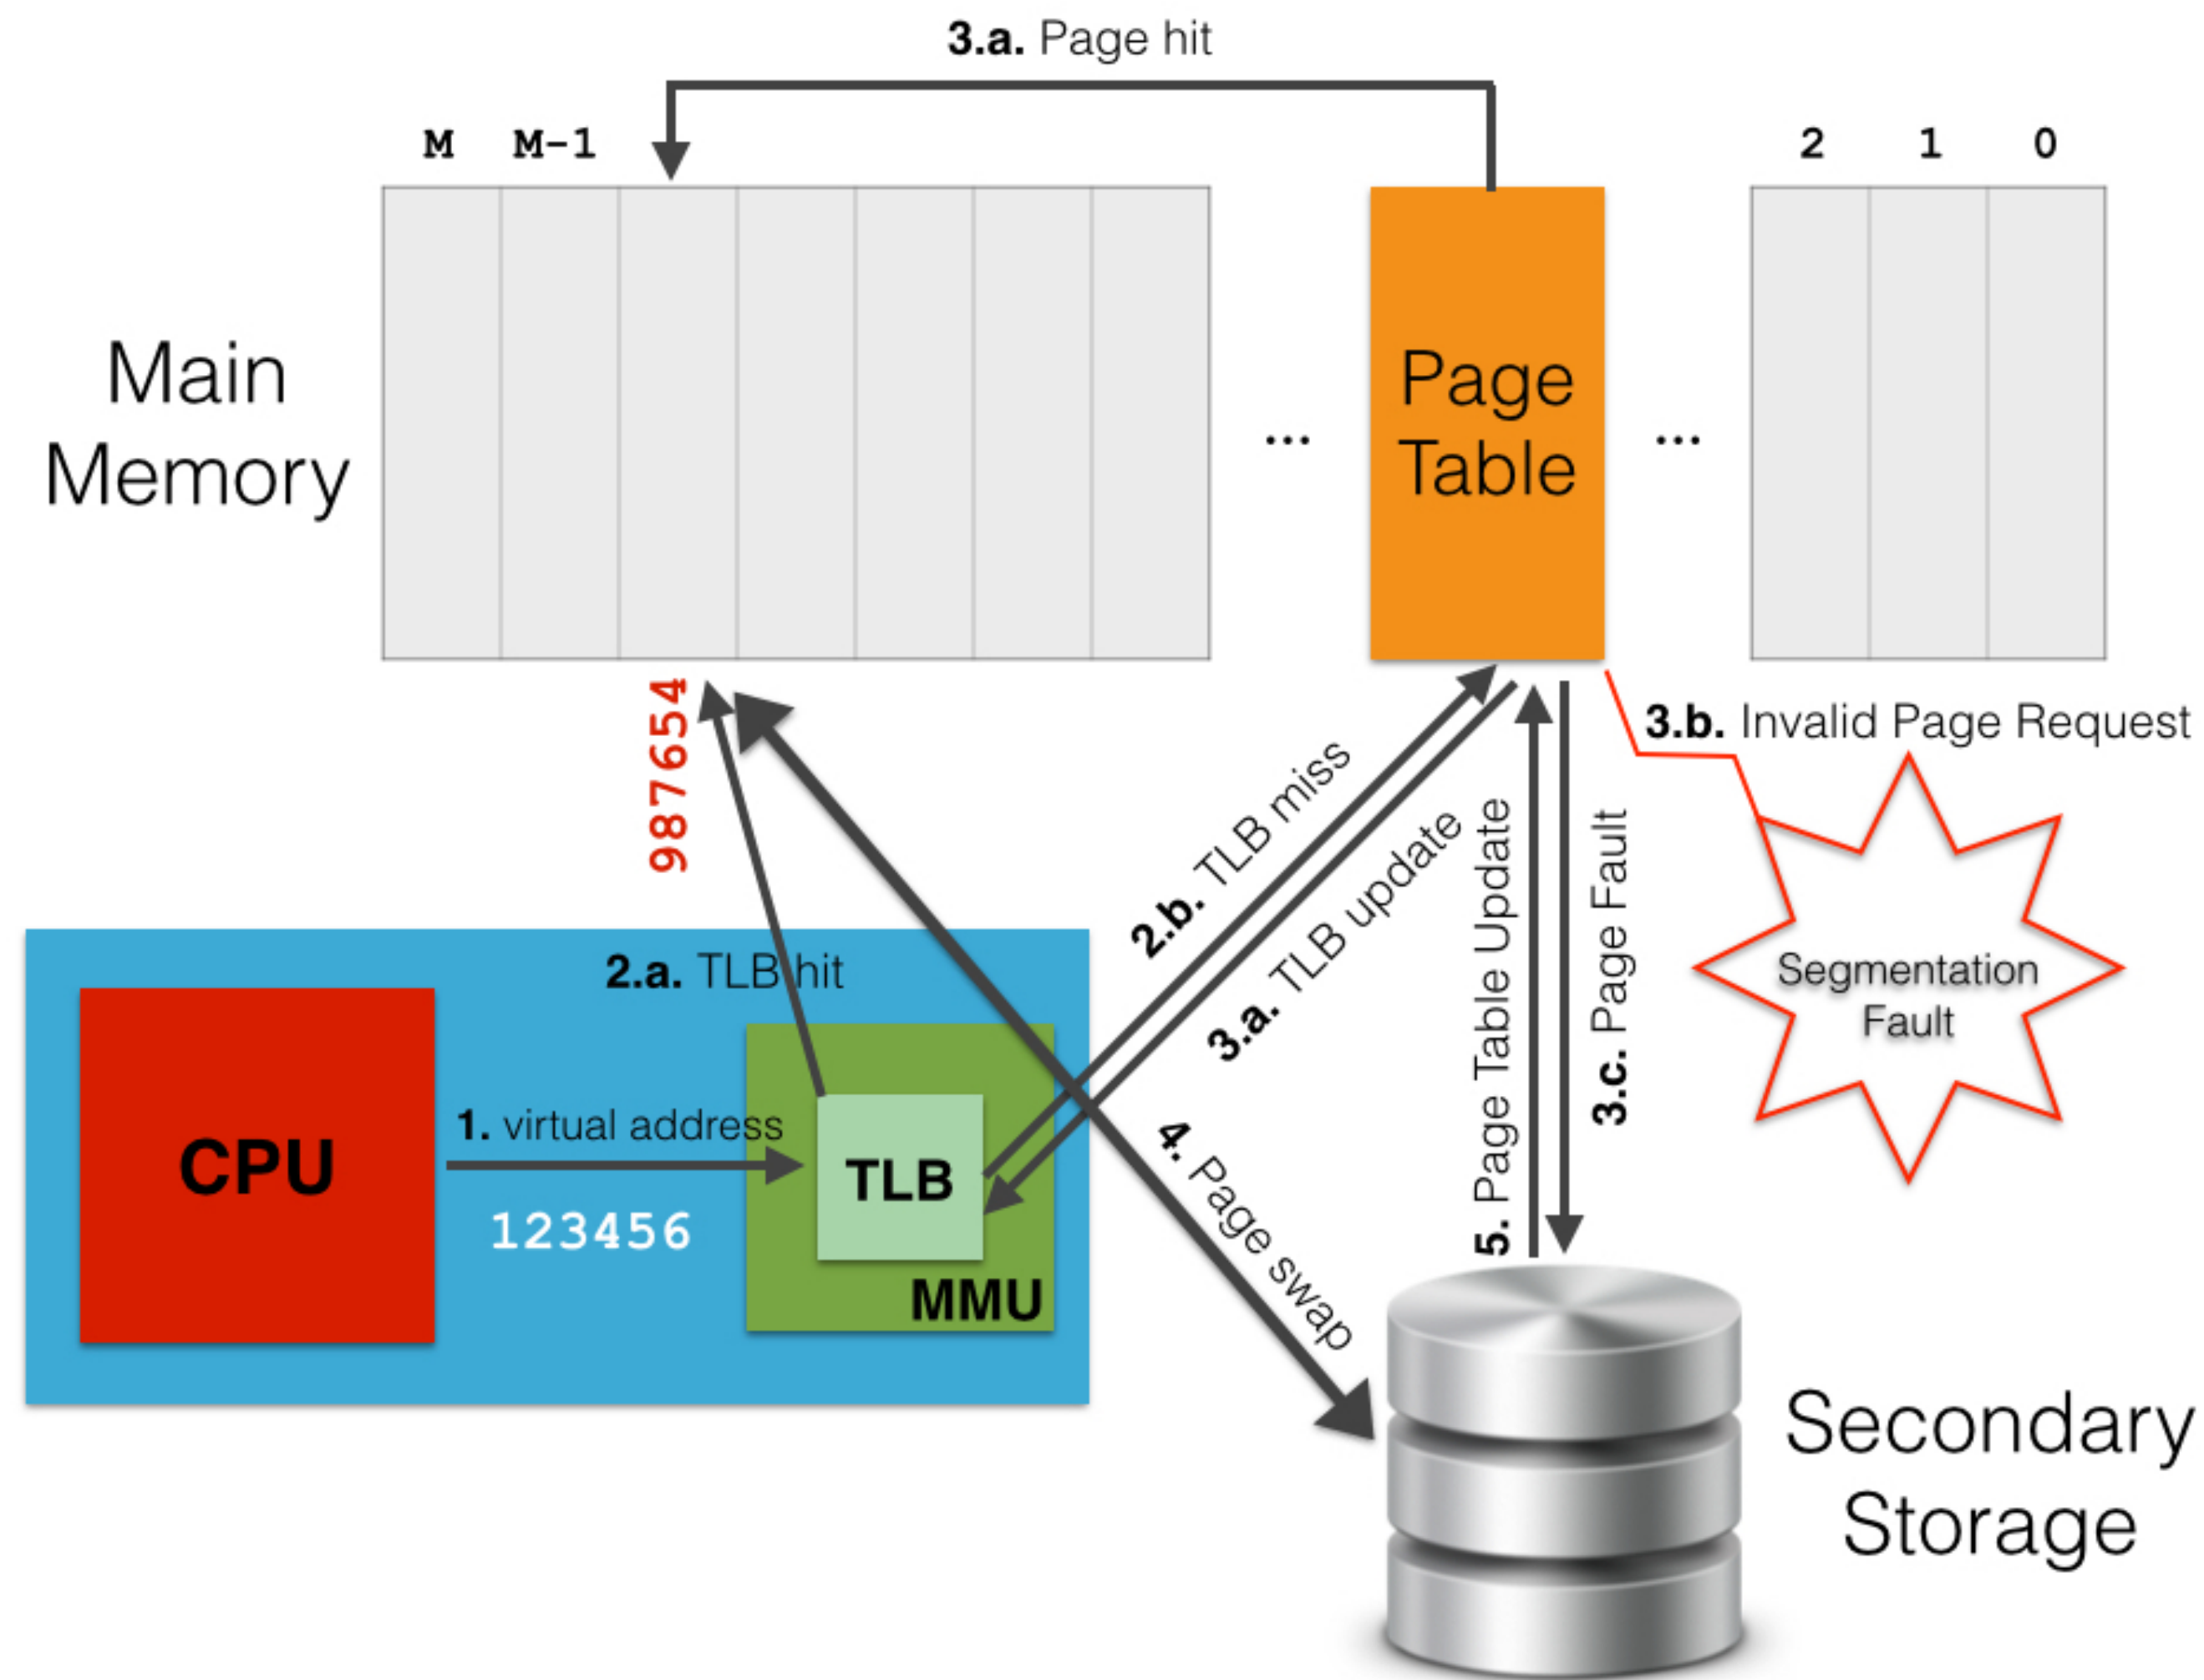
- $n \gg m$
 - hash page # into smaller page table
 - use chaining for collisions
- Regardless of virtual address space size, the size of an inverted page table remains the same (i.e., inverted page table size is proportional to # of physical frames)



Using a Translation Lookaside Buffer (TLB)

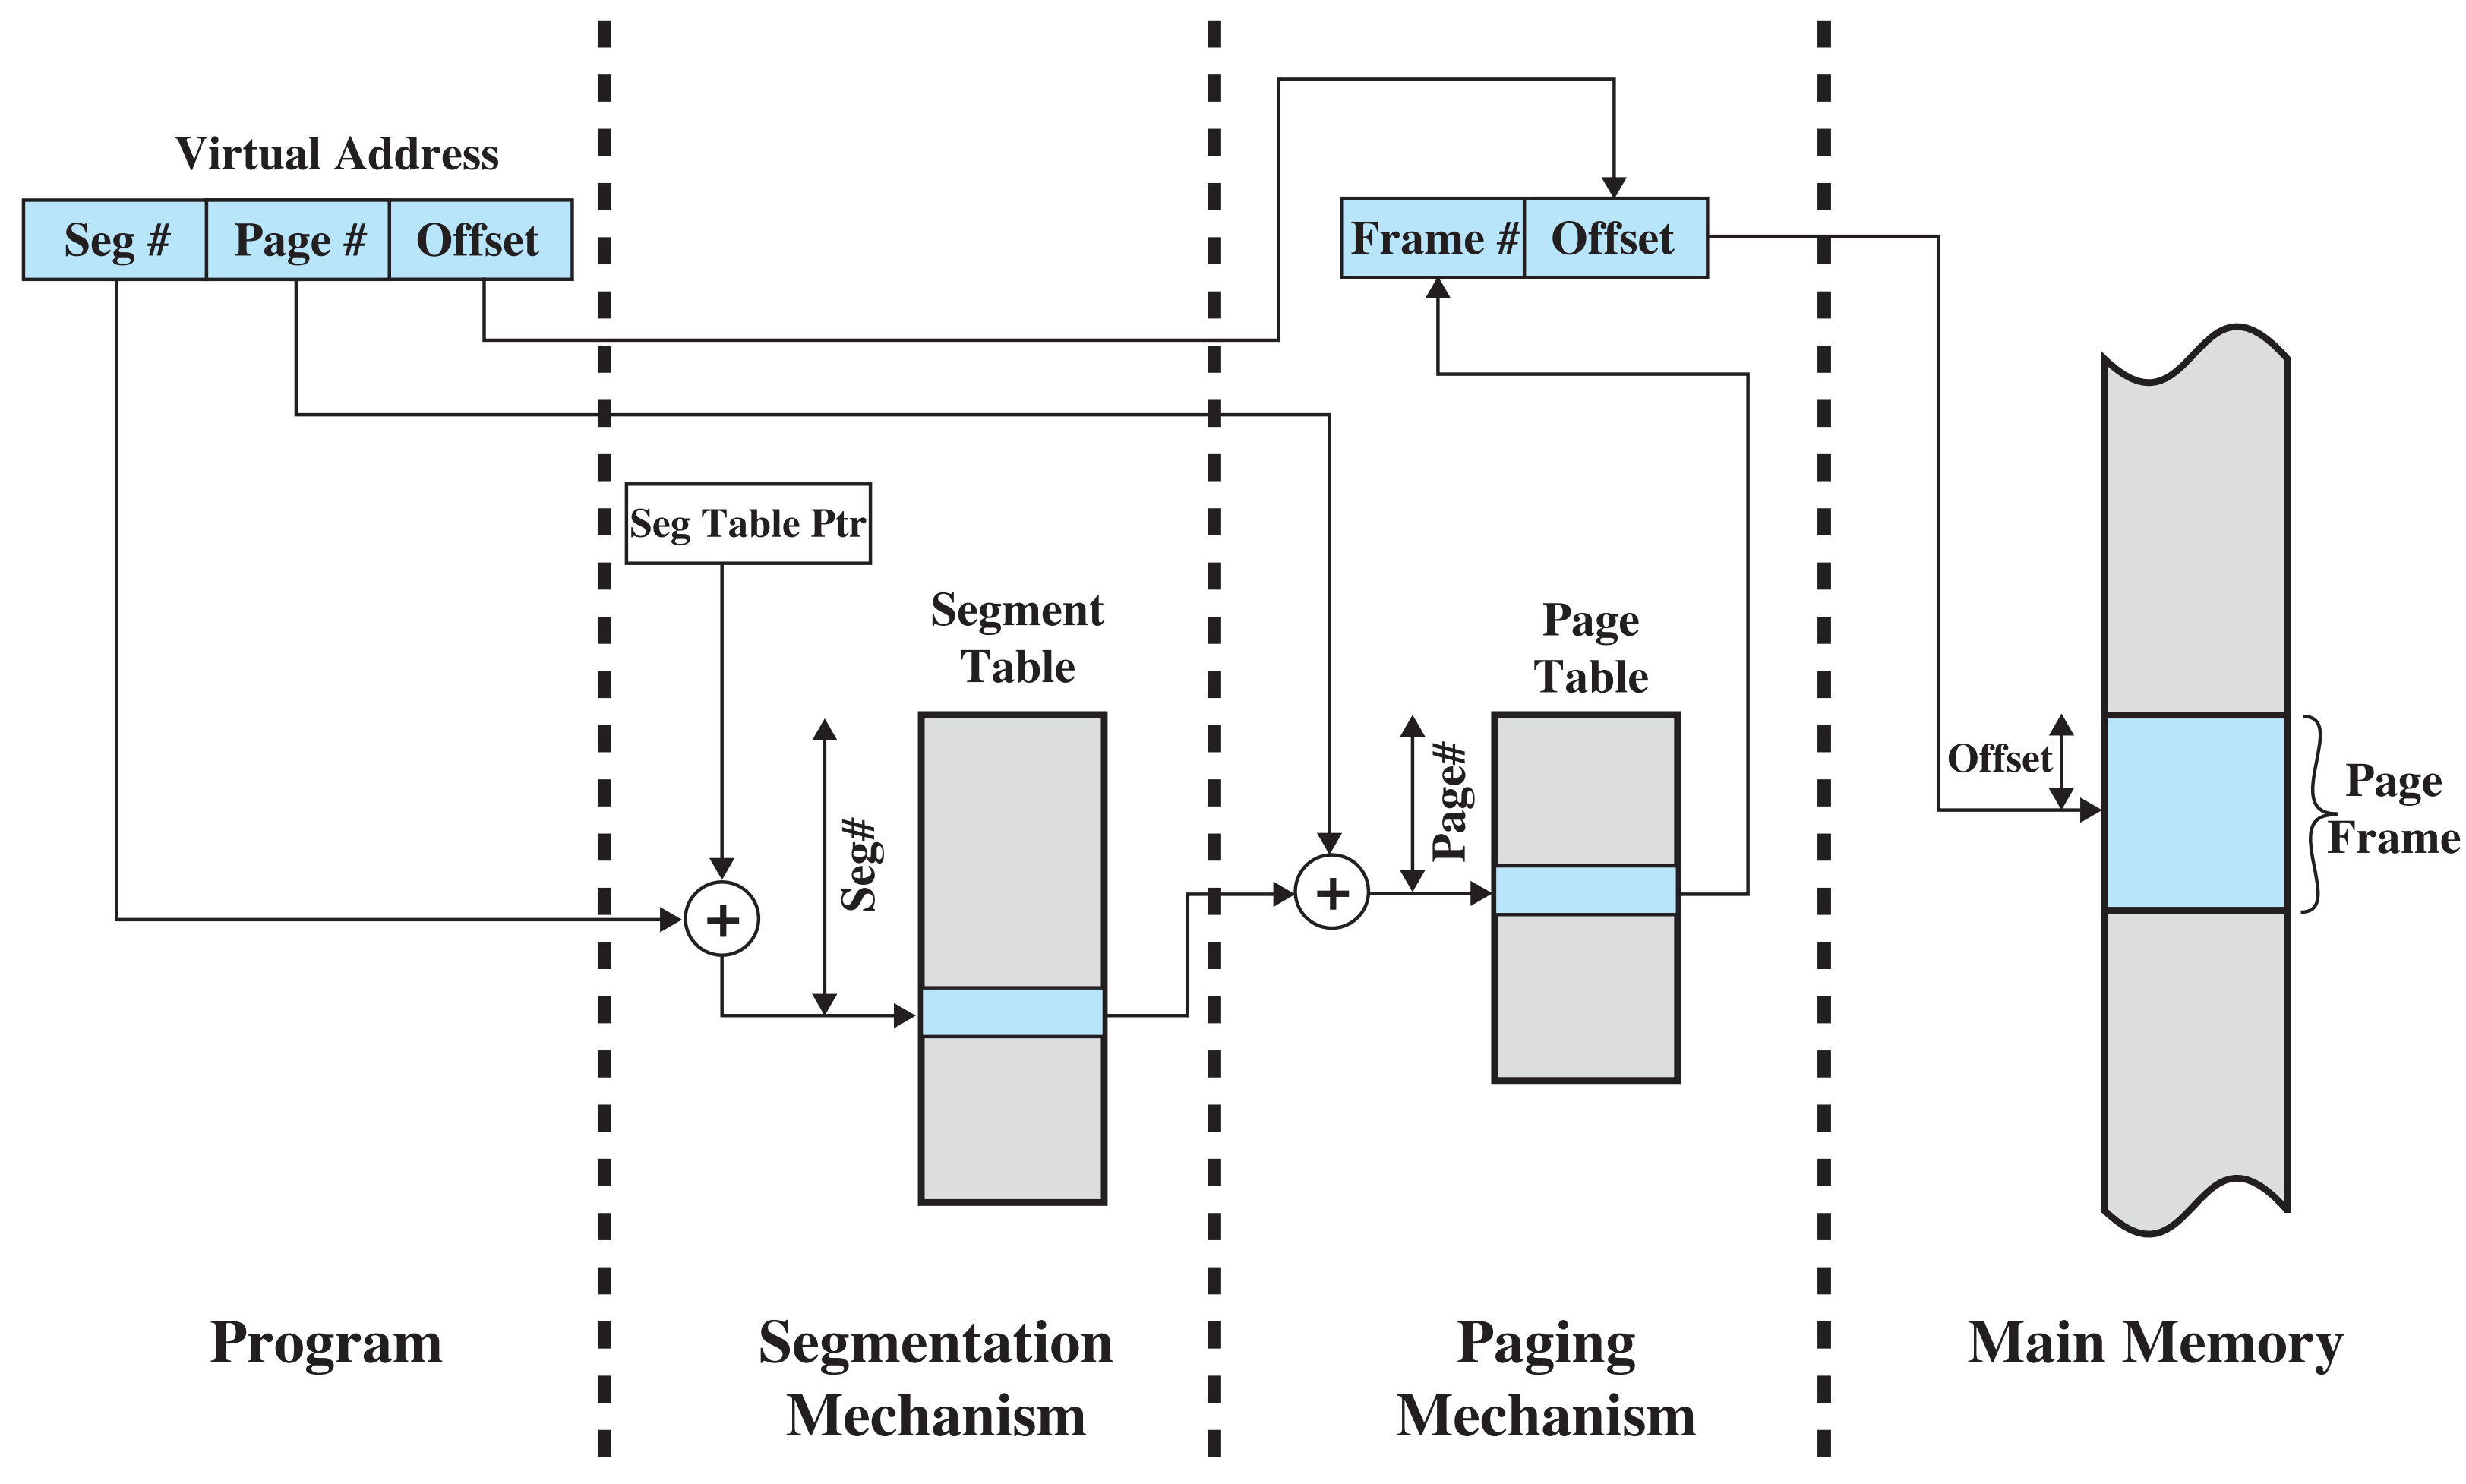
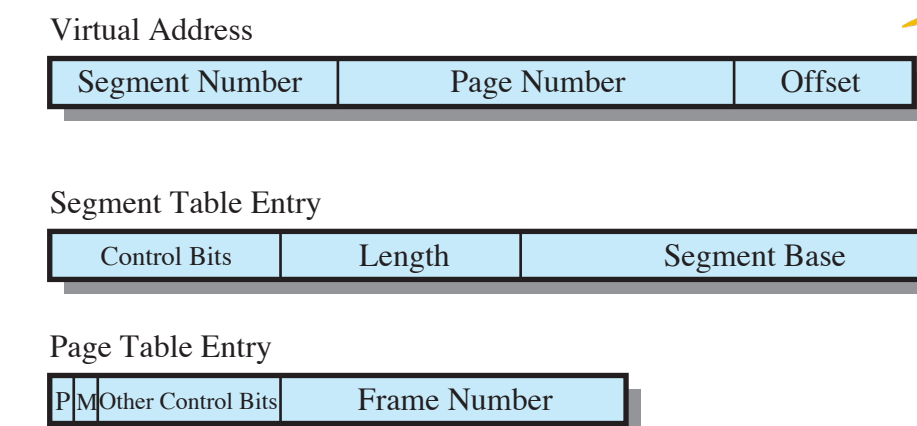


Using a Translation Lookaside Buffer (TLB) — *another look*



<https://gabrieletolomei.files.wordpress.com/2013/10/mmu.jpg>

Address Translation in a Segmentation/Paging System



Memory Management Formats

Virtual Address



Page Table Entry



Paging

Virtual Address



Segment Table Entry



Segmentation

If Page is Present (P)
→ *use frame number*

If Page Modified (M)
→ *write out*

Virtual Address



Segment Table Entry



Page Table Entry

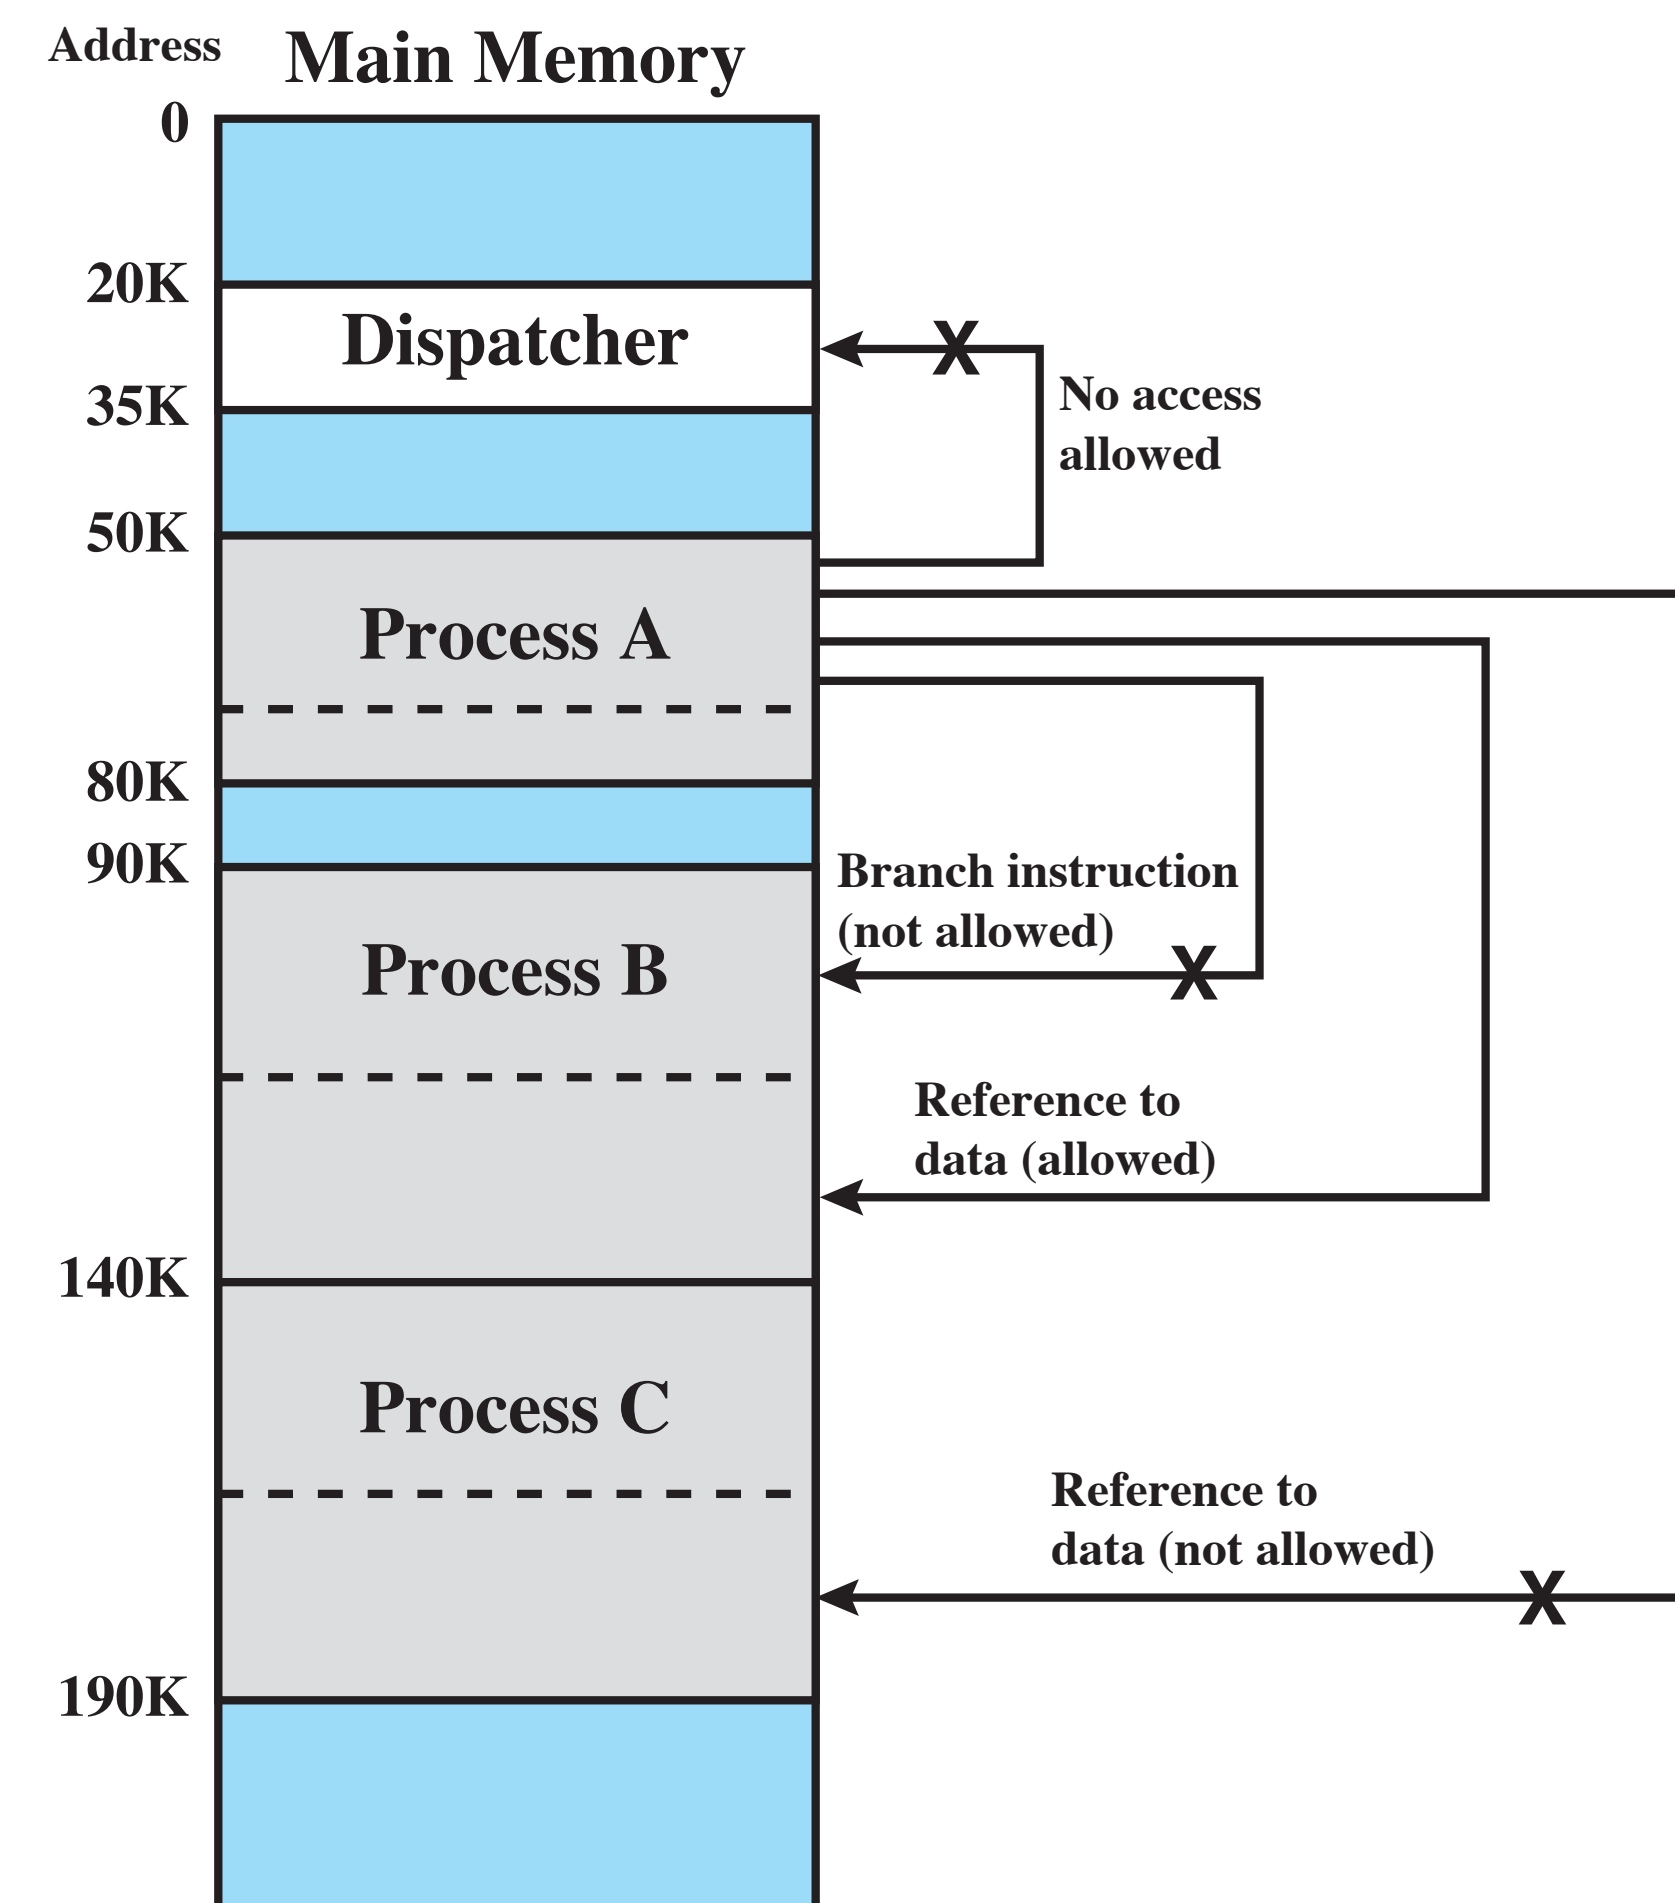


Paging and Segmentation

P= present bit
M = Modified bit

Protection Between Segments

- Configurable settings for segments
 - readable
 - writable
 - executable
- This is possible on a per-page level, but it is less intuitive/kind of awkward since paging structure is not visible to programmers.



Designing an OS with Virtual Memory

Goal: Reduce Frequency of Page Faults

Designing an OS with Virtual Memory

Goal: Reduce Frequency of Page Faults

Fetch Policy — *when* should a page be brought into main memory?

- **Demand Paging**

Bring pages in when a reference is made (as needed / on demand)

- **Prepaging**

Bring in a number of pages at once (e.g., other nearby pages on disk)

Designing an OS with Virtual Memory

Goal: Reduce Frequency of Page Faults

Placement Policy — *where should a page be put in main memory?*

(Re: Chapter 7)

- **Segmentation**

→ *Best-Fit, First-Fit, Etc.*

- **Paging**

→ *Not an issue (page-frame mappings are equally efficient)*

Designing an OS with Virtual Memory

Goal: Reduce Frequency of Page Faults

Replacement Policy — *which page^{*} should be replaced when a new page must be brought into main memory?^{**}*

- **Random Page Replacement**
Choose a page randomly
 - **Optimal (OPT)**
Replace the page that will not be used the most time in the future
 - **LRU**
Replace the page that has not been used for the longest time
 - **FIFO**
Replace the page that has been in primary memory the longest; use a circular buffer
 - **Clock** (approximation of LRU)
Basically, FIFO w/ a “use” bit; cycle through pages, replace pages not used recently
 - **LFO**
Replace the page that is used least often
 - **2nd Chance**
Choose a victim to evict; “mark” it; spare it until later; if encountered again, then page out
- } Benchmarks

} Nearly optimal
Complex

} Simple
Often Incorrect

^{*}Want to replace an **unlocked** page that is least likely to be referenced in the near future.

^{**}Most interesting if we assume that all frames in main memory are occupied — then a page fault occurs.

Example: Page-Replacement Algorithms

Order of page references: 2 3 2 1 5 2 4 5 3 2 5 2

Page address stream

2 3 2 1 5 2 4 5 3 2 5 2

OPT

2	2	2	2	2	2	4	4	4	2	2	2
	3	3	3	3	3	3	3	3	3	3	3
			1	5	5	5	5	5	5	5	5
				F		F			F		

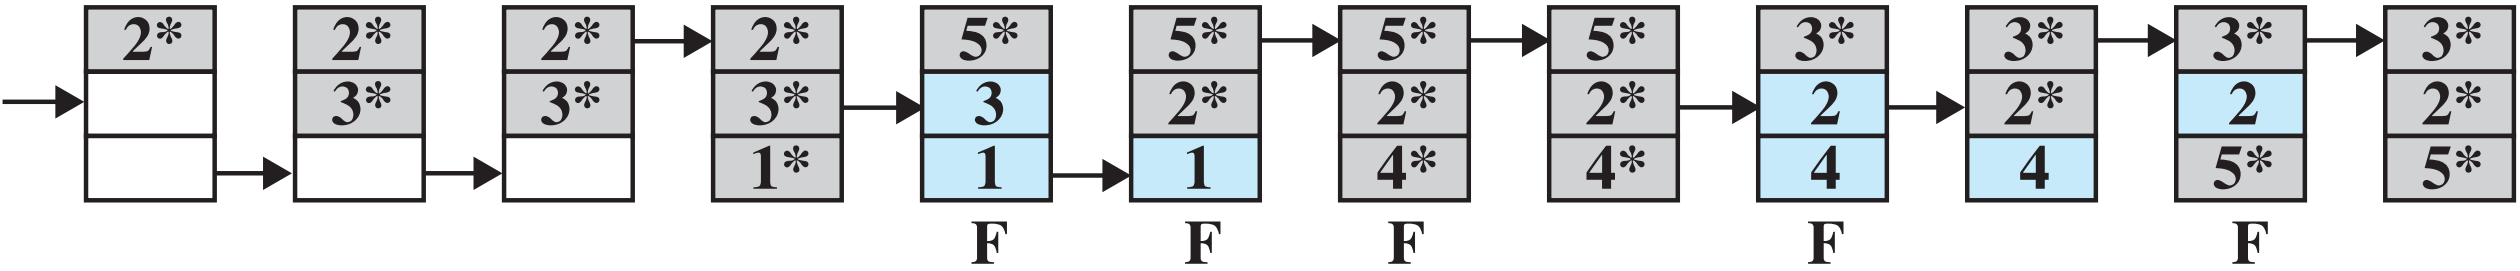
LRU

2	2	2	2	2	2	2	2	3	3	3	3
	3	3	3	5	5	5	5	5	5	5	5
			1	1	1	4	4	4	2	2	2
				F		F		F	F		

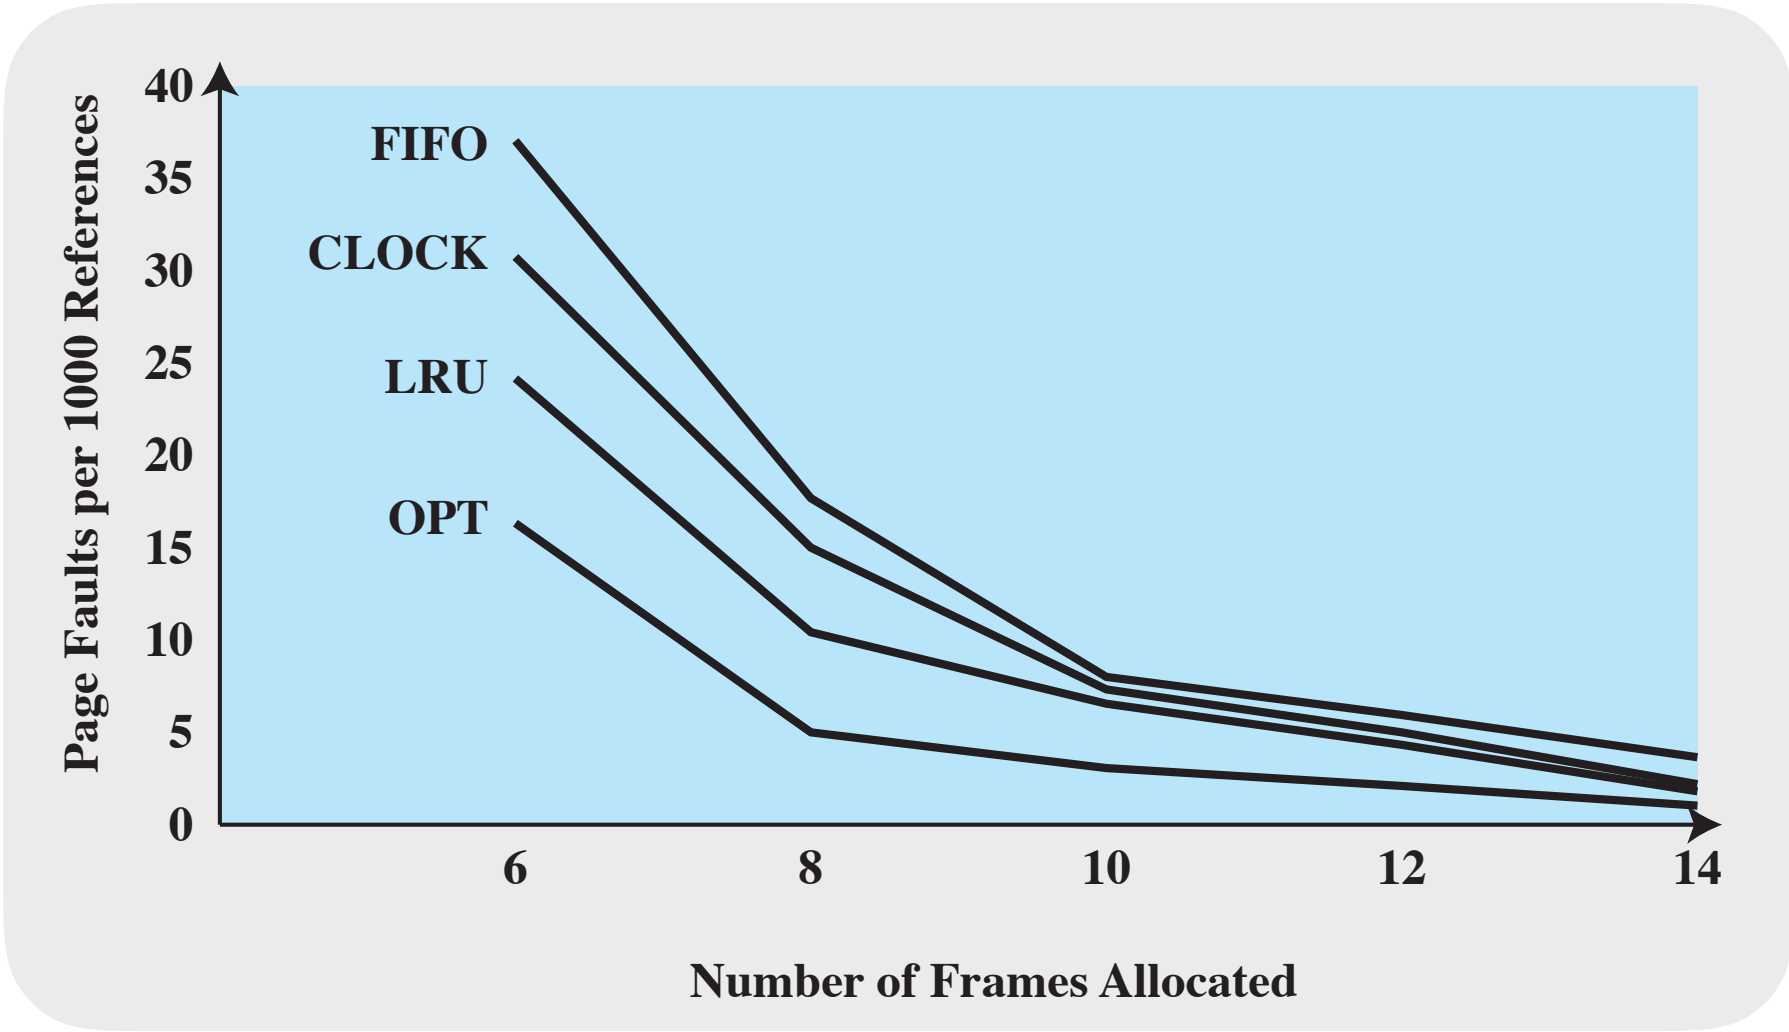
FIFO

2	2	2	2	5	5	5	5	3	3	3	3
	3	3	3	3	2	2	2	2	2	5	5
			1	1	1	4	4	4	4	4	2
				F	F	F		F		F	F

CLOCK



F = page fault occurring after the frame allocation is initially filled



Designing an OS with Virtual Memory

Goal: Reduce Frequency of Page Faults

Resident Set Management — *how many pages should be in main memory?*

- Resident Set Size (Fixed vs. Variable)
- Replacement Scope (Local [current process] vs. Global [everything])

Cleaning Policy — *when should a page be written out from main memory to secondary memory? (opposite of “Fetch Policy”)*

Load Control — *how many processes should be in main memory? (i.e., the level/degree of multiprogramming)*