

# ***I/O, Files, & Storage Devices (Part I)***

---

Professor Travis Peters  
CSCI 460 Operating Systems  
Fall 2019

*Some slides & figures adapted from Stallings instructor resources.*

*Some slides adapted from Adam Bates's F'18 CS423 course @ UIUC*  
*<https://courses.engr.illinois.edu/cs423/sp2018/schedule.html>*

*Some content adapted from the Disk Scheduling Algorithms tutorial*  
*<http://www.cs.iit.edu/~cs561/cs450/disksched/disksched.html>*

# Today

---

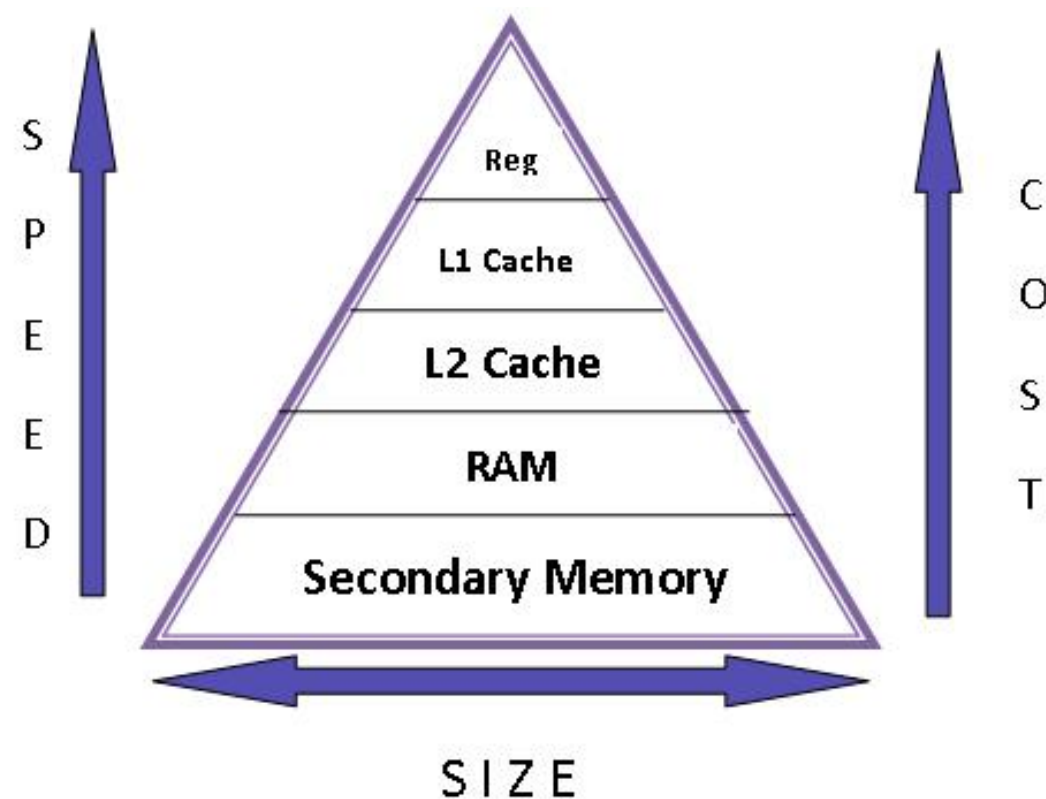
## Announcements

- Project Proposals Due **FRIDAY!**

## Goals & Learning Objectives

- Understand key concepts behind I/O devices and functions
- Understand some of the key issues in the design of OS support for I/O
- Understand basics of secondary storage (emphasis on disks and disk scheduling)
- Understand basics behind files and file systems

# Memory Hierarchy: Respective Sizes & Access Times



LEVEL	ACCESS TIME	TYPICAL SIZE
Registers	"instantaneous"	under 1KB
Level 1 Cache	1-3 ns	64KB per core
Level 2 Cache	3-10 ns	256KB per core
Level 3 Cache	10-20 ns	2-20 MB per chip
Main Memory	30-60 ns	4-32 GB per system
Hard Disk	3,000,000-10,000,000 ns	over 1TB

# I/O and Disks and Files, Oh My!

## Categories of I/O Devices

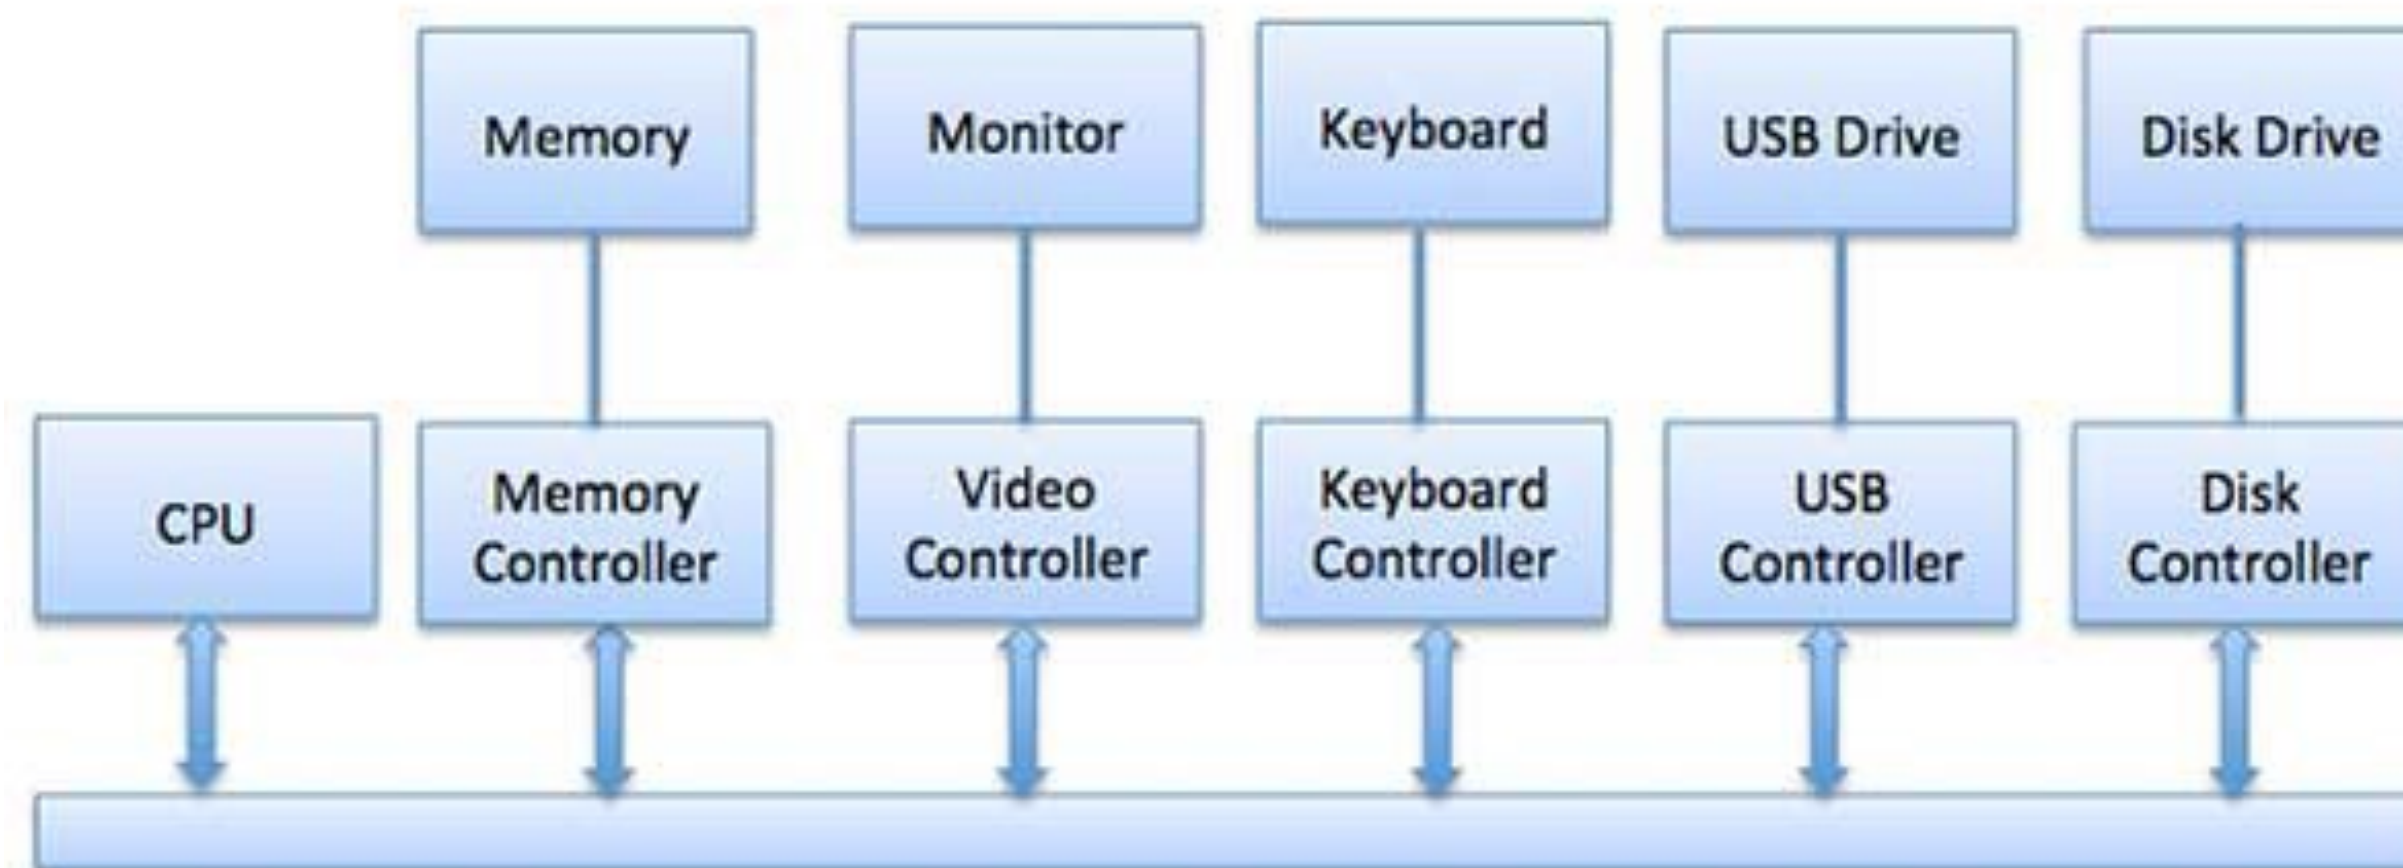
- Human Readable  
*Computer-User Interactions (e.g., printers, terminals, keyboards, mice)*
- Machine Readable  
*Computer-Device Interactions (E.g., disk drives, controllers, sensors, actuators)*
- Communication  
*Computer-Computer Interactions (E.g., network devices, modems)*

### Block-oriented devices

*Store information in  
(usually fixed-sized) blocks*

### Stream-oriented devices

*Transfer data in/out as a stream  
of bytes (no block structure)*



<https://medium.com/cracking-the-data-science-interview/the-10-operating-system-concepts-software-developers-need-to-remember-480d0734d710>

# I/O and Disks and Files, Oh My!

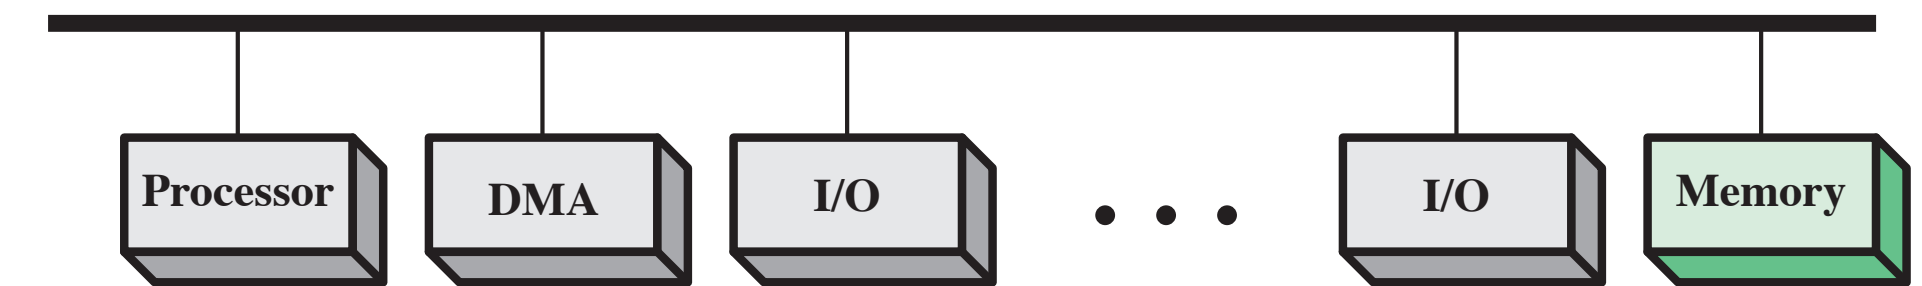
---

## I/O Approaches

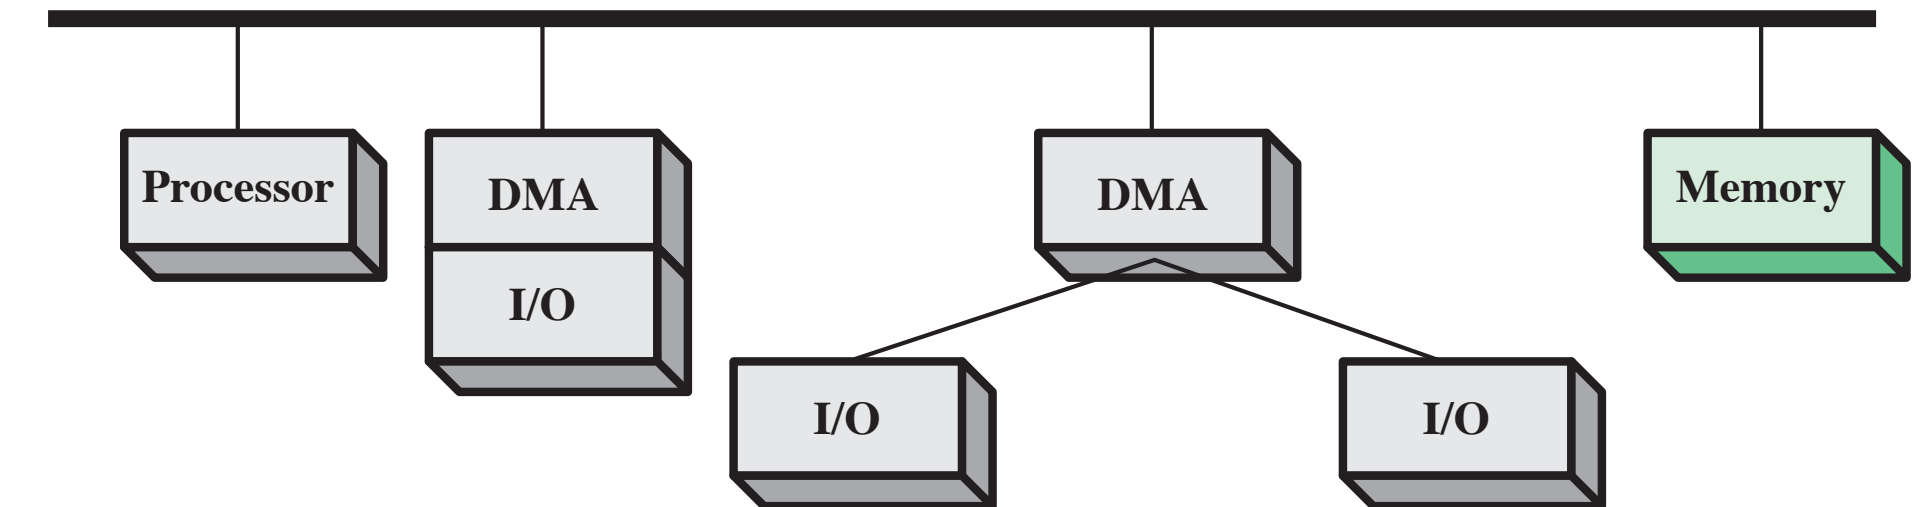
- Direct I/O  
*Processor directly controls I/O device(s)*
- Programmed I/O  
*processor issues command; busy wait until complete; no interrupts*
- Interrupt-Driven I/O  
*processor issues command; continue running (or block); interrupt when done; uses interrupts, obviously*
- Direct Memory Access (DMA)  
*control exchanges between MM and I/O devices w/ minimal involvement of the processor (only needed at beginning/end)*

# Direct Memory Access (DMA)


- **Role of DMA:**  
*transfer data to/from memory over the system bus*
- I/O processor\* (which realizes DMA) mimics the processor and its ability to take control over the system bus.
- A progression can show that more efficient implementations move all I/O off of the system bus onto a dedicated I/O bus



(a) Single-bus, detached DMA



(b) Single-bus, Integrated DMA-I/O



(c) I/O bus

\*I/O module more closely resembles a dedicated I/O processor for handling I/O



*What should I/O devices provide?*

# Overview of I/O, Files, & Storage Devices

## Overarching Objectives

### 1. Efficiency


*I/O is slowww compared to modern processors/memory systems*

*I/O is often the bottleneck*

*Want I/O to be as efficient as possible!*

Modern CPUs e.g., many BILLIONS of instructions per second  
 Modern memory systems e.g., 2-4GB/sec bandwidth

E.g., 10s-100s Mbits/sec



E.g., 9-10 keystrokes per second



# Overview of I/O, Files, & Storage Devices

## Overarching Objectives

### 1. Efficiency

*I/O is slowww compared to modern processors/memory systems*

*I/O is often the bottleneck*

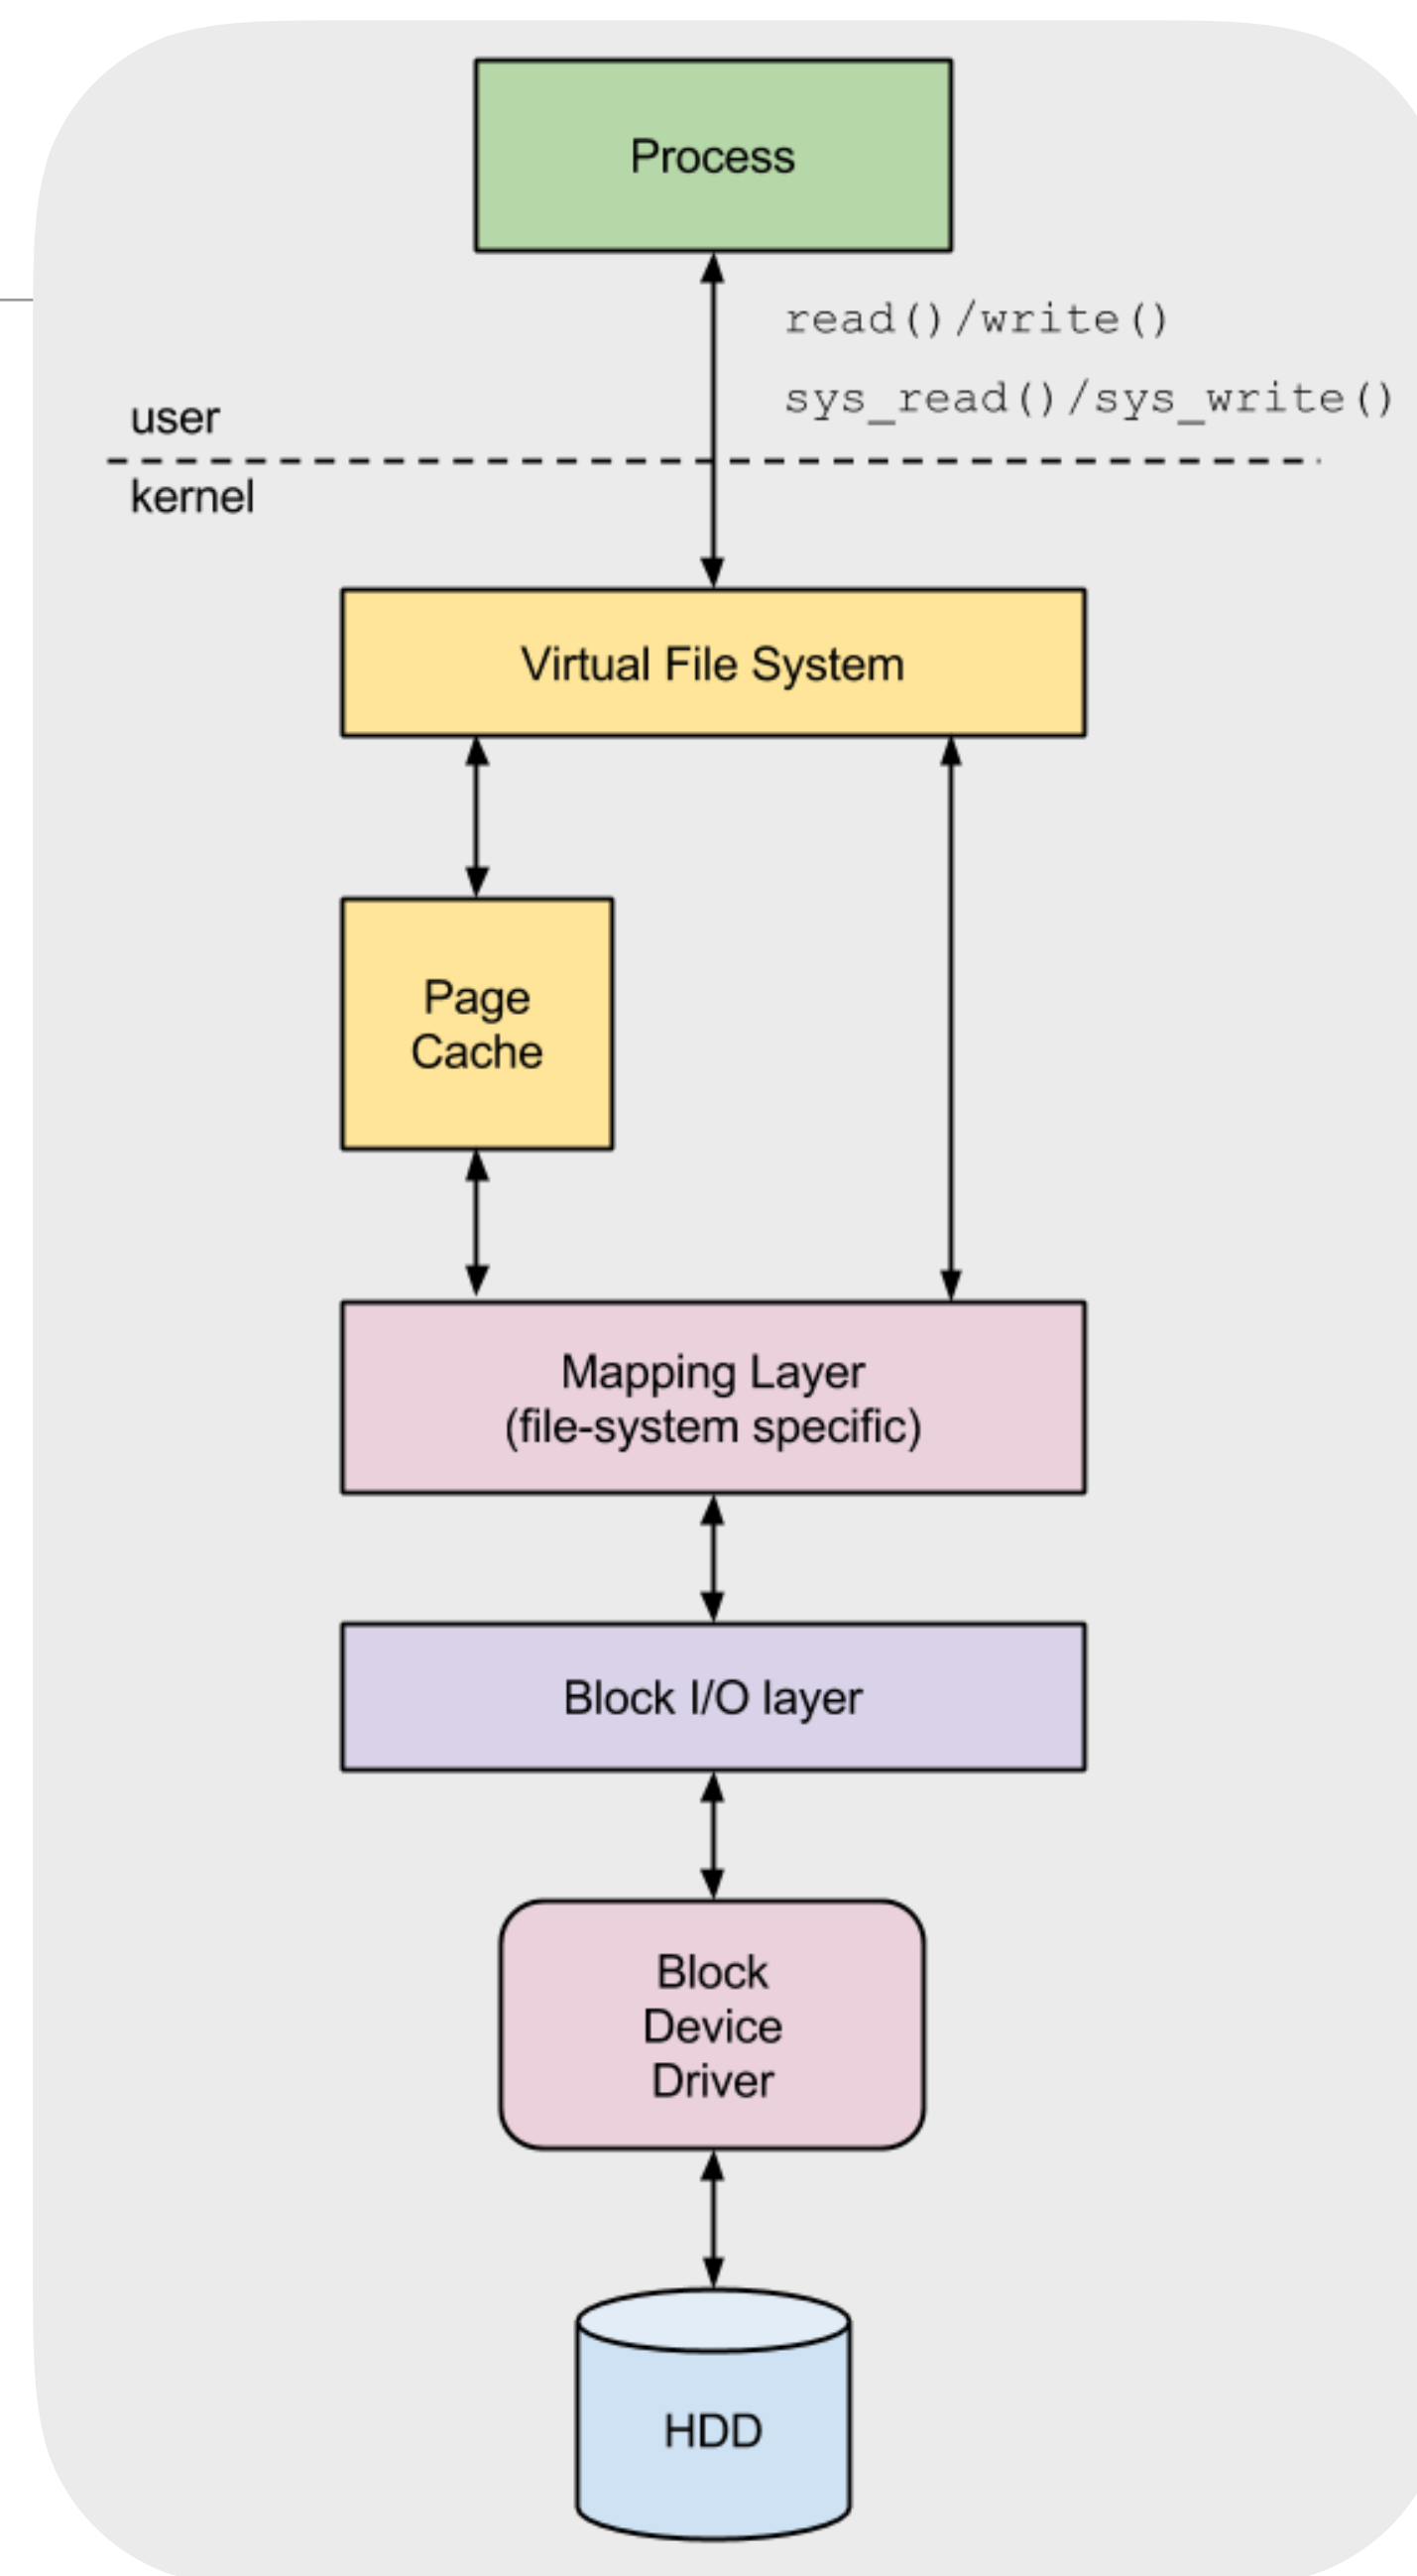
*Want I/O to be as efficient as possible!*

### 2. Generality

*Want simplicity and correctness*

*Want support for diverse range of devices*

→ Lots of abstractions — most things can be done with **read/write**, **open/close**, **lock/unlock**.



# *Disks & Disk Scheduling*

**What is a disk?**

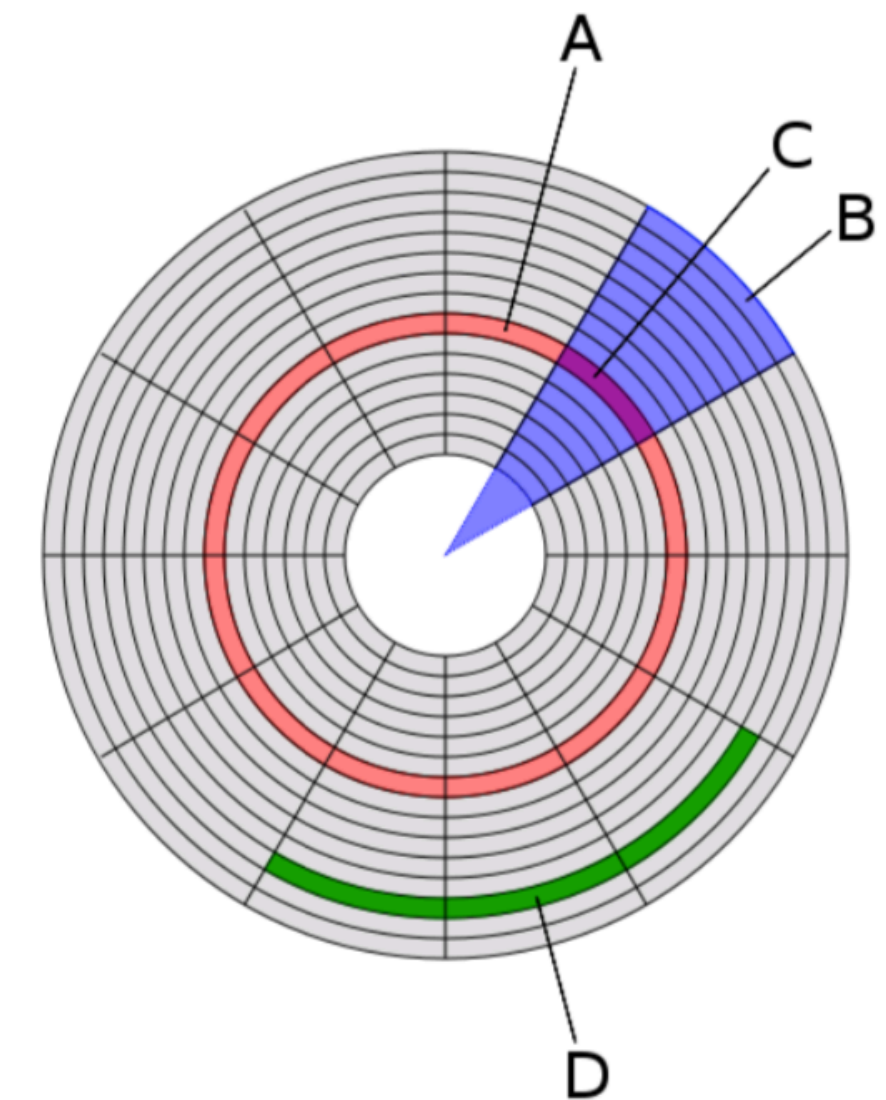
**How does a disk work?**

# Disk Scheduling

A = Track / B = Sector / C = Sector of Track / D = File

- **The Basics**

- Reading/writing data to/from a disk is done with a read/write head
- Disk rotates
- Head must move (**seek**) to appropriate position (**track** / **sector**) on the disk



## Disk Scheduling Decision

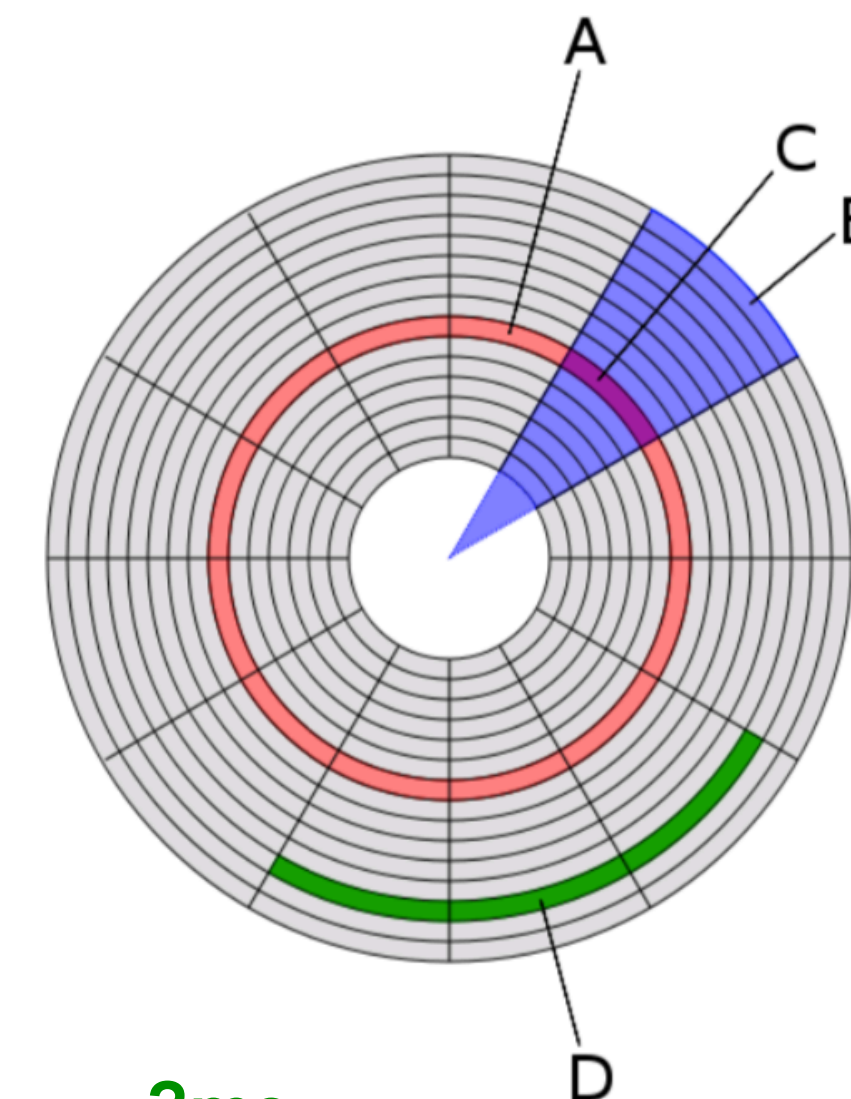
Given a series of access requests (reads/writes), on which track should the disk arm be placed next to maximize fairness, throughput, etc?

# Disk Scheduling

A = Track / B = Sector / C = Sector of Track / D = File

- **The Basics**

- Reading/writing data to/from a disk is done with a read/write head
- Disk rotates
- Head must move (**seek**) to appropriate position (**track / sector**) on the disk



- **Timing of Disk I/O Transfer**

- **Seek Time**

*time taken to move disk arm to a specified track*

- **Rotational Latency / Rotational Delay**

*time to wait for sector to reach the head*

- **Transfer Time**

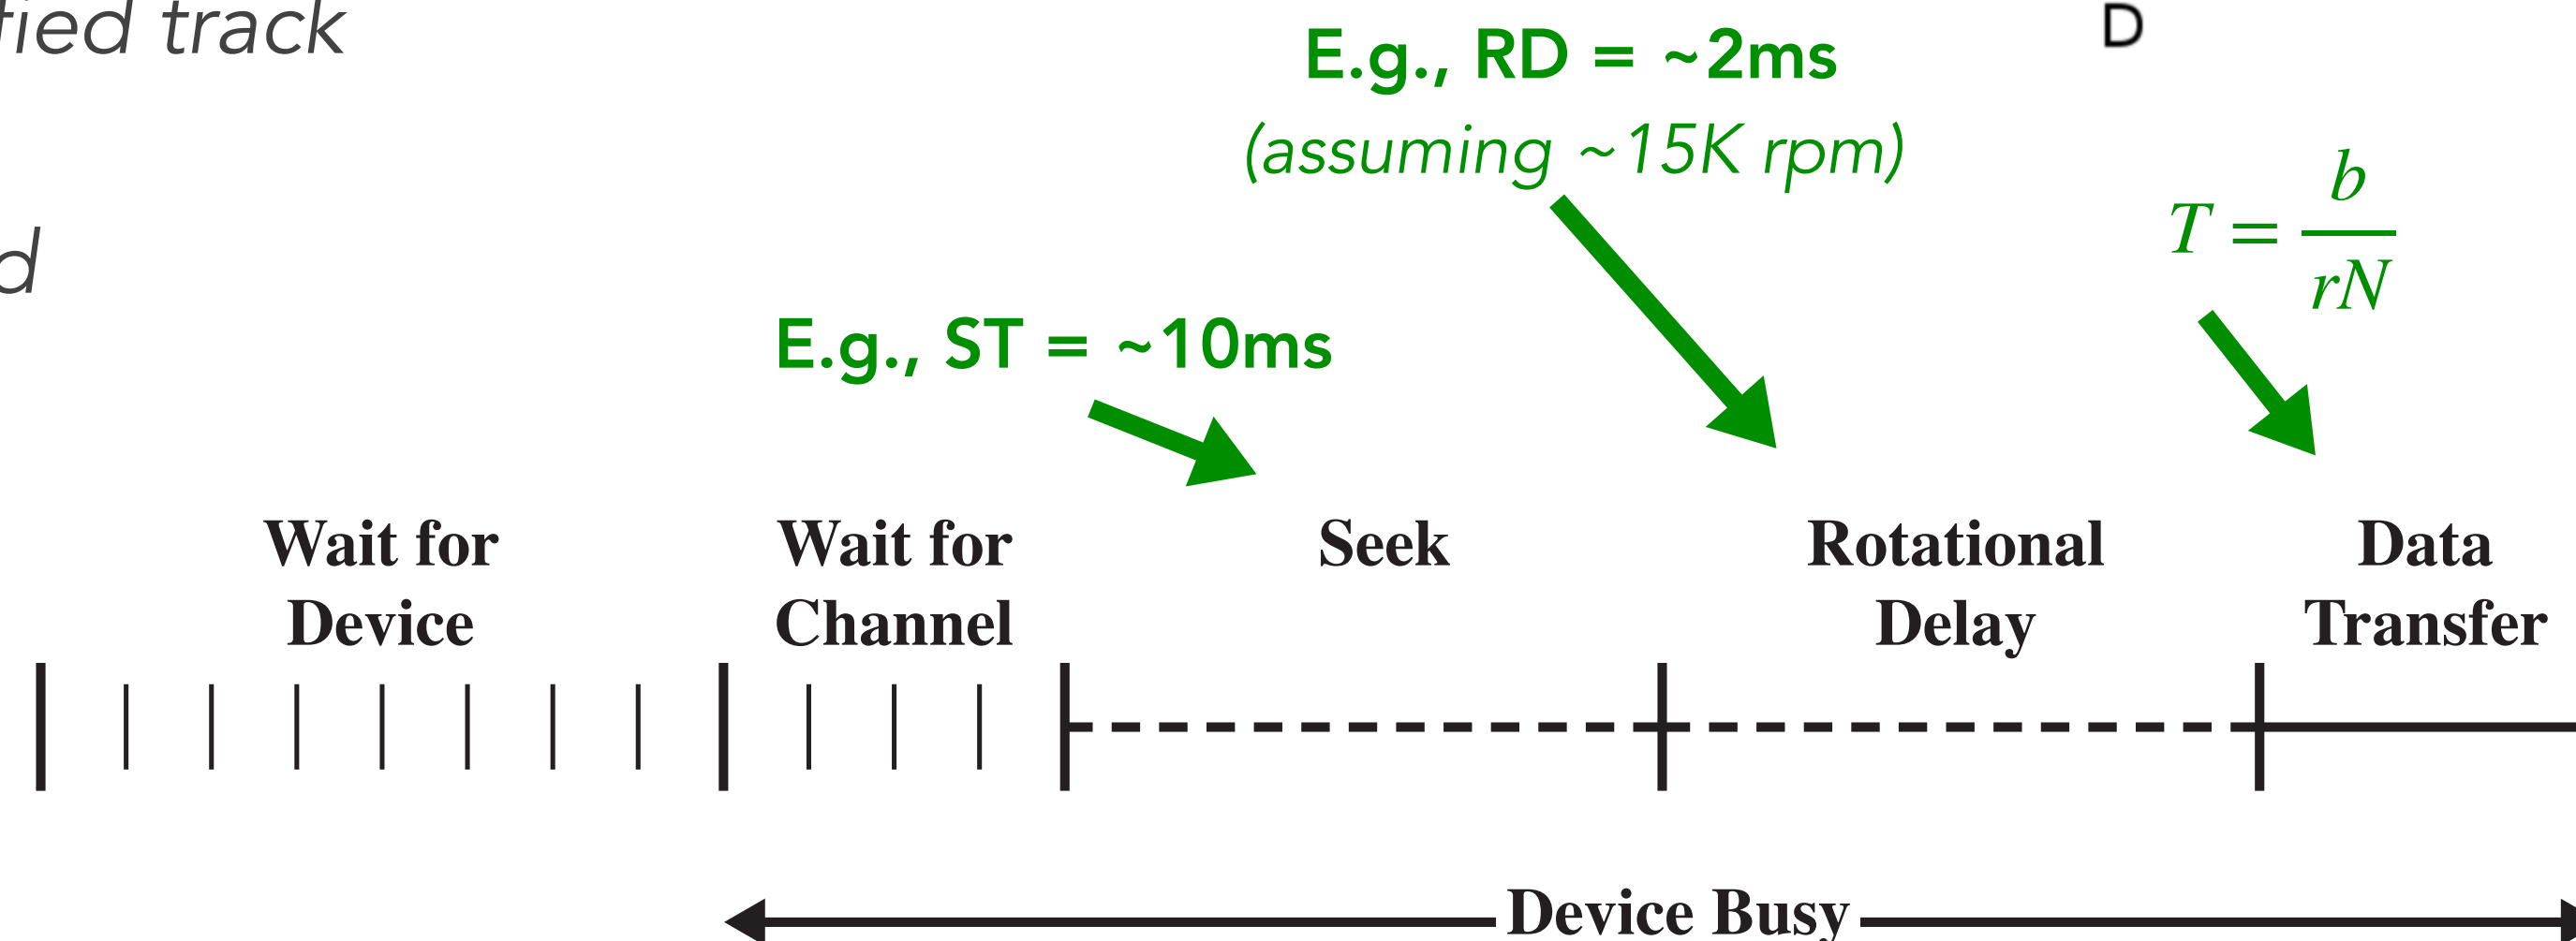
*time to transfer (read/write) data*

- **Access Time**

*= seek time*

*+ rotational latency*

*(+ transfer time)*



$$\rightarrow T_a = T_s + \frac{1}{2r} + \frac{b}{rN}$$




# Disk Scheduling Policies

*Order of ops matter! → seek time can have tremendous impact on access time*

Consider multiple processes queued up on various I/O requests...

*how to schedule I/O ops?*

- **Random Scheduling** — only really useful as a benchmark/baseline...
- **FIFO/FCFS** — approx. equivalent to random scheduling!
  - **Pros?** fairness among requests, requests handled in well-defined order
  - **Cons?** arrival may be on random spots on the disk (long seek times); wild “swings” can happen
  - **Analogy?** FCFS elevator scheduling?
- **Priority** — not targeted at optimal disk utilization (rather, more concerned w/ OS objectives)
- **LIFO** — always service the most recent process
  - **Why might this be good/bad?**
    - results in very little arm movement (*principle of locality!*)
    - could lead to starvation...



*Can we do better?*

*What if the scheduler knew about the current track / position of the head?*

# *Extras*



# Flash Memory *(1 Slide Only...)*

- **EEPROM** (Electrically Erasable Programmable Read Only Memory)
  - **Ex:** NAND Flash
  - READ performance
    - Random READ: 25 $\mu$ s, Sequential READ: 25ns
  - WRITE performance
    - PROGRAM PAGE: 25 $\mu$ s, BLOCK ERASE: 1.5ms
  - Endurance: 100,000 PROGRAM/ERASE cycles

- **In The News**

Tesla's cars use **Linux** and do **a huge amount of logging**. According to 057 Technology's Jason Hughes, "The information logged here is pretty much useless on production vehicles. Unless a developer has a specific reason for enabling it, it does the customer no good. These logs are also rarely downloaded by Tesla." That **excessive logging is causing the flash memory in the Media Control Unit (MCU) v1 to fail**, which causes the car to **lose touch screen functionality, which in a Tesla controls most everything**. V1 units are in model S and X Teslas made before 2018.

— <https://www.tomshardware.com/news/flash-memory-wear-killing-older-teslas-due-to-excessive-data-logging-report>

- **GANGRENE: Exploring the Mortality of Flash Memory**

<https://www.usenix.org/system/files/conference/hotsec12/hotsec12-final4.pdf>