

I/O, Files, & Storage Devices (Part II)

Professor Travis Peters
CSCI 460 Operating Systems
Fall 2019

Some slides & figures adapted from Stallings instructor resources.

Some slides adapted from Adam Bates's F'18 CS423 course @ UIUC
<https://courses.engr.illinois.edu/cs423/sp2018/schedule.html>

Some content adapted from the Disk Scheduling Algorithms tutorial
<http://www.cs.iit.edu/~cs561/cs450/disksched/disksched.html>

Today

Announcements

- Project Proposals Due **TONIGHT @ 10PM!**
- Exam #2 in class **NEXT FRIDAY!**
- PA2 Due **NEXT MONDAY (11/11)!** — *thank Parker for requesting this deadline be pushed...*
- Guest Lecture with Will Peteroy on Monday

Read before class!


- <https://msrc-blog.microsoft.com/2010/12/08/on-the-effectiveness-of-dep-and-aslr/>
- <https://msrc-blog.microsoft.com/2013/08/12/mitigating-the-ldrhotpatchroutine-depaslr-bypass-with-ms13-063/>
- <https://arstechnica.com/information-technology/2019/08/armed-with-ios-0days-hackers-indiscriminately-infected-iphones-for-two-years/>

Goals & Learning Objectives

- Understand key concepts behind I/O devices and functions
- Understand some of the key issues in the design of OS support for I/O
- Understand basics of secondary storage (emphasis on disks and disk scheduling)
- Understand basics behind files and file systems


Disk Scheduling Policies (cont.)

A = Track / B = Sector / C = Sector of Track / D = File




Shortest-Service-Time-First (SSTF)

- Select I/O ops that require the least amount of movement from disk arm's current position
- Choose next op that incurs minimum seek time
- **Pros?**
 - Try to minimize seek time...
- **Cons?**
 - Is SSTF optimal?
 - Are we worried about overhead of sorting?
 - Can we avoid starvation?




Disk Scheduling Policies (cont.)

A = Track / B = Sector / C = Sector of Track / D = File




SCAN (a.k.a. The Elevator Algorithm)

- The disk arm moves in one direction, servicing the closest request in the direction of travel
- Continue in that direction until...
 - last track is reached, or
 - there are no more requests to service in that direction (→ **LOOK policy**)
- **Pros?**
 - Bounded time for each request
- **Cons?**
 - Requests at the opposite end will take a while...
 - Does not exploit locality...
 - Which sectors have shorter wait times?
(How to fix this?)




Disk Scheduling Policies (cont.)

A = Track / B = Sector / C = Sector of Track / D = File



C-SCAN (Circular Scan)

- Similar to SCAN...
 - The disk arm moves in one direction only...
 - ...but *NEVER* actually changes directions!
 - When the last track is reached, reset the head ("wrap around") back to the opposite end of the disk and begin again
- **Pros?**
 - Uniform service time
 - Addresses issues with max delay that can occur in regular SCAN
- **Cons?**
 - Do nothing on the return




+ Many more... *N-Step-SCAN, FSCAN, C-LOOK, etc...*

Disks... a bit more on disks!!!

RAID (Redundant Array of Independent Disks)

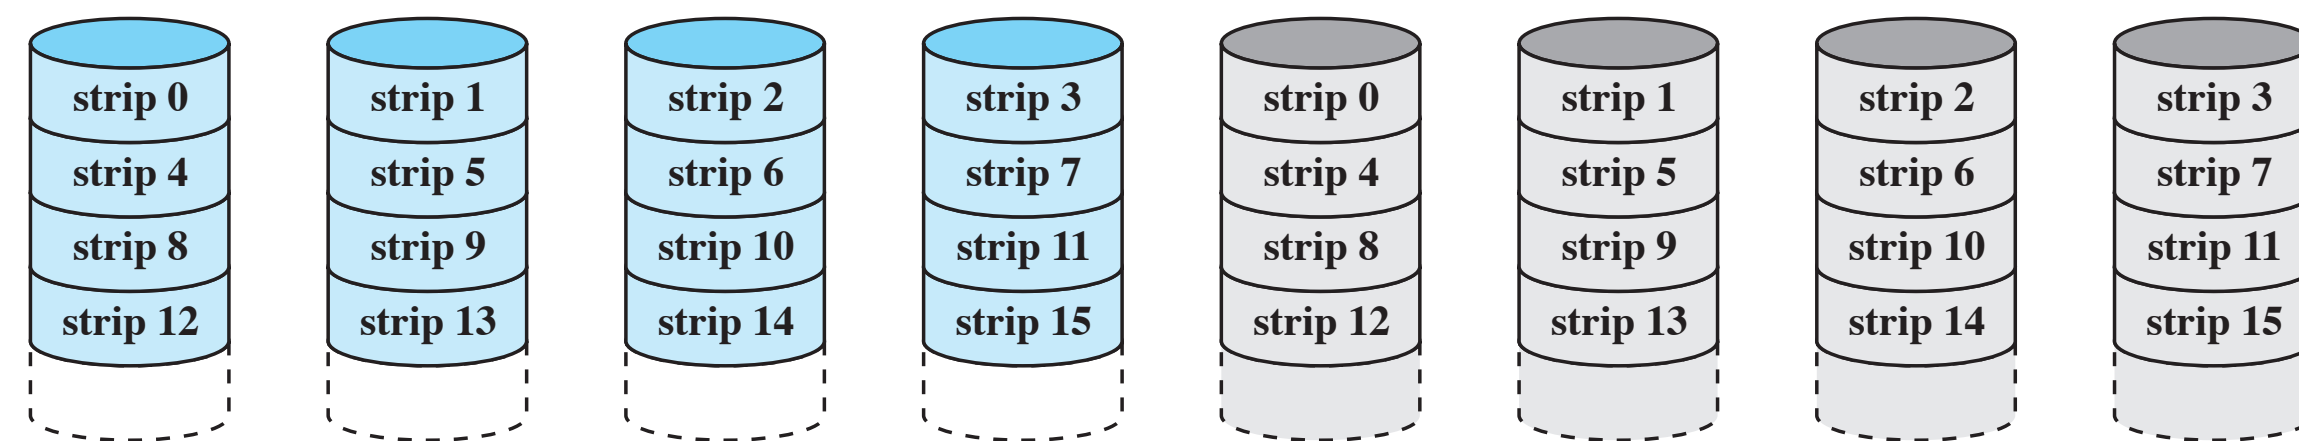
- Array of independent disks that work **cooperatively** and can be **used in parallel**
 - *Logically, treated as a single disk drive by the OS*
 - *Separate I/O requests can be handled in parallel*
 - *Data distributed/replicated across disks improves reliability (**striping**)*
- Various widely-accepted schemes (**RAID0, RAID1, RAID2, RAID3, RAID4, RAID5, RAID6**)
 - Primary contribution was aimed at addressing issues with redundancy
 - Only 4/7 levels are commonly used

RAID Levels (Redundant Array of Independent Disks)




(a) RAID 0 (non-redundant)

Non-redundant
Files are “striped” across disks
Improves performance
Data loss is possible though...




(b) RAID 1 (mirrored)


Improves performance + reliability
More expensive (high space overhead)...




(c) RAID 2 (redundancy through Hamming code)



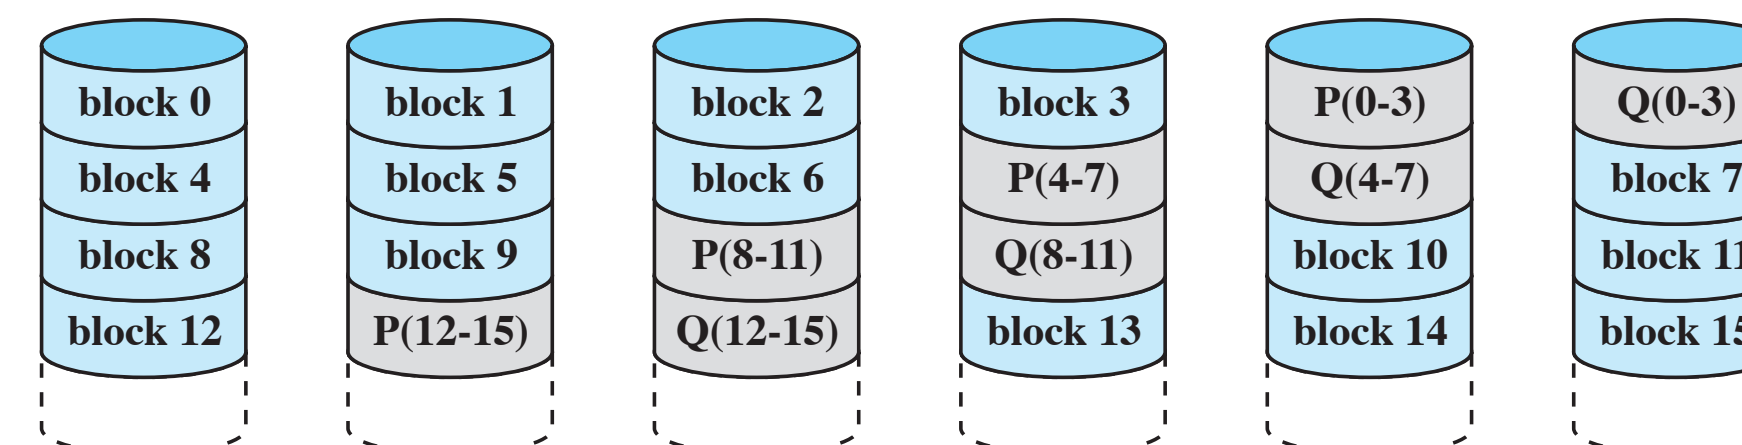
(d) RAID 3 (bit-interleaved parity)



(e) RAID 4 (block-level parity)



(f) RAID 5 (block-level distributed parity)



(g) RAID 6 (dual redundancy)

Less redundancy, less overhead, varying reliability...

File Allocation

When considering allocating space for a file...


- Is the maximum space required for the file allocated at once?
 - Preallocation vs. Dynamic Allocation
 - *Is it reasonable to assume we know how big a file can grow to be?*
- When allocating units of space, should the units be contiguous ("**portions**")?
And how big should those portions be?
 - Locality → better **performance**
 - Fewer portions → **smaller data structures** (e.g., tables) for managing allocations
 - Fixed-sized blocks → **simplifies** allocation space & allocation policies
 - **Variable-size / small fixed-size** portions → **minimizes waste**
- What sort of data structure(s) should be used to keep track of allocations?
 - Ex. File Allocation Table (FAT)
 - In general, 1 entry holds 1 file unit
 - Ex. File unit = entire file (1 contiguous set of blocks)
 - Ex. File unit = a fixed-size piece of the total file (blocks are scattered around)

Recall fragmentation issues from virtual memory
→ **Strategies: first-fit, best-fit, next-fit, etc.**

Example: Contiguous File Allocation


Allocate a single contiguous block to a file at the time of creation.

CFA is best from the perspective of a single sequential file!
...not so much for the system as a whole...



Before Compaction

File Allocation Table		
File Name	Start Block	Length
File A	2	3
File B	9	5
File C	18	8
File D	30	2
File E	26	3



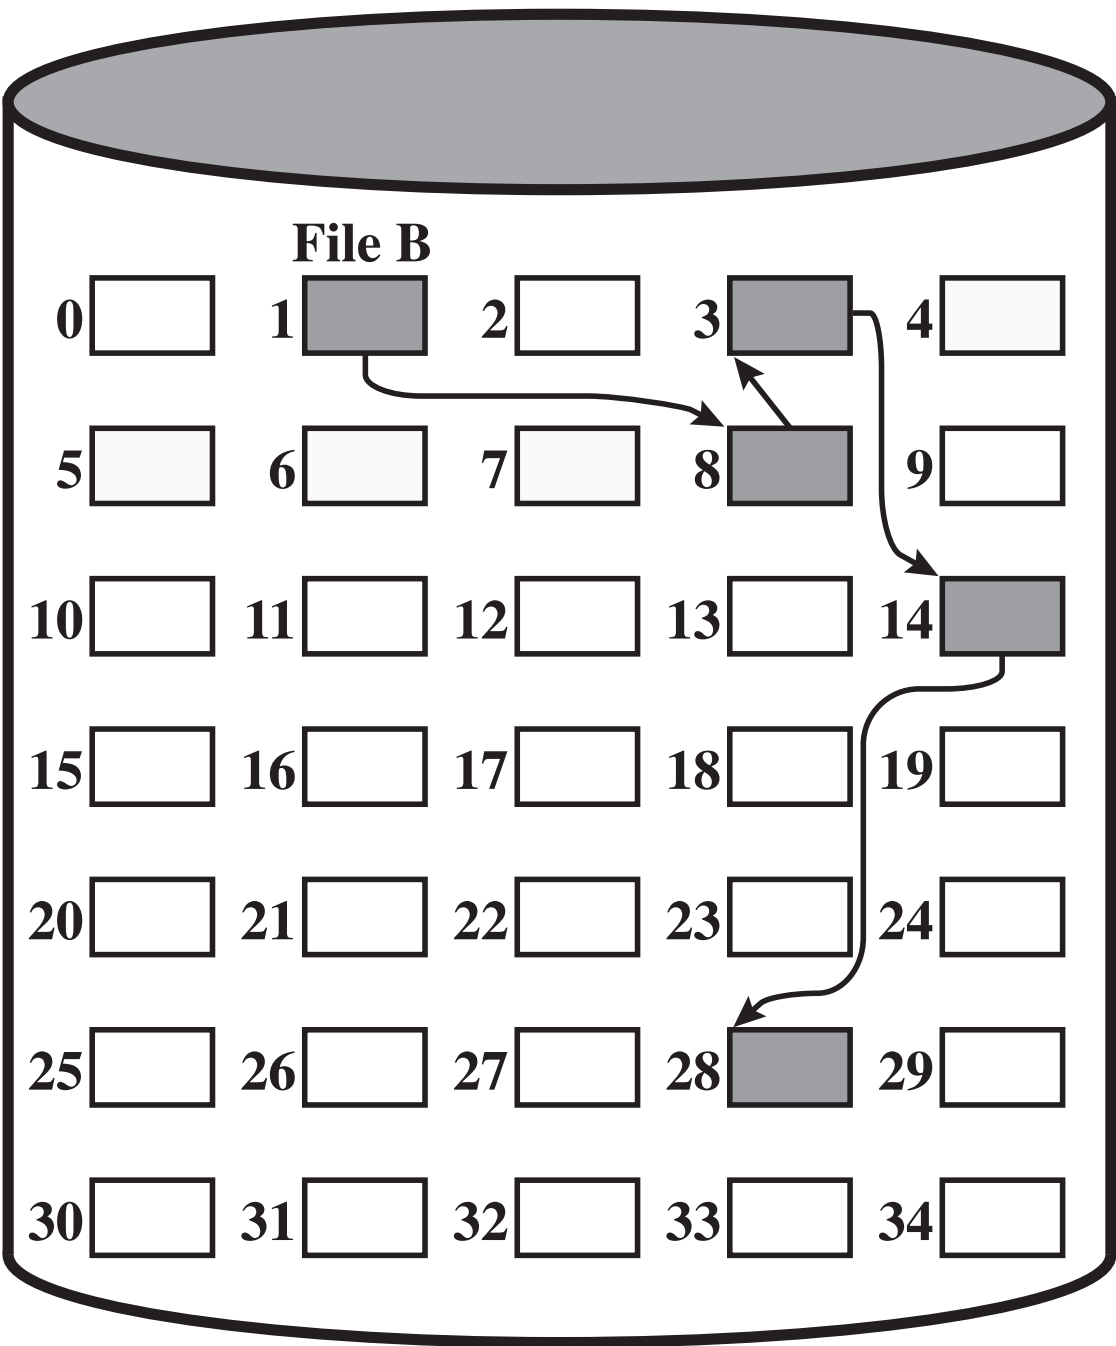
After Compaction

File Allocation Table		
File Name	Start Block	Length
File A	0	3
File B	3	5
File C	8	8
File D	19	2
File E	16	3

Example: Chained Allocation

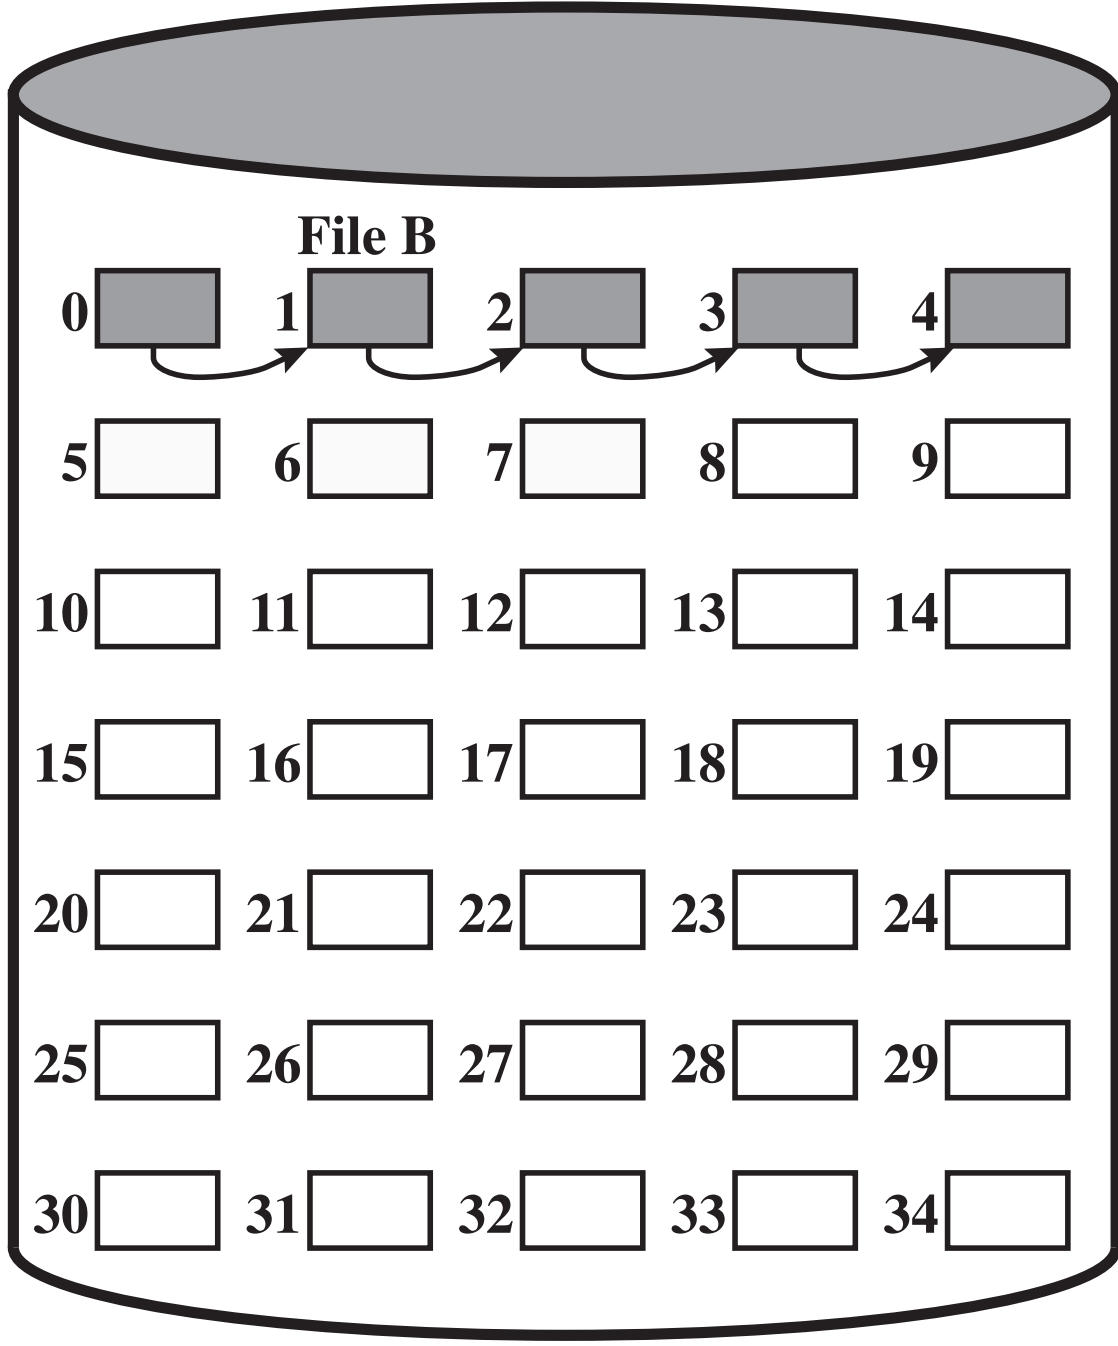
Allocate individual, fixed-size blocks until enough space is allocated for the file.
Recall: each block contains a pointer to the next.

CA is best for managing space for sequential files!
...no utilization of locality though...



Before Periodic Consolidation

File Allocation Table		
File Name	Start Block	Length
...
File B	1	5
...



After Periodic Consolidation

File Allocation Table		
File Name	Start Block	Length
...
File B	0	5
...

Example: Indexed Allocation

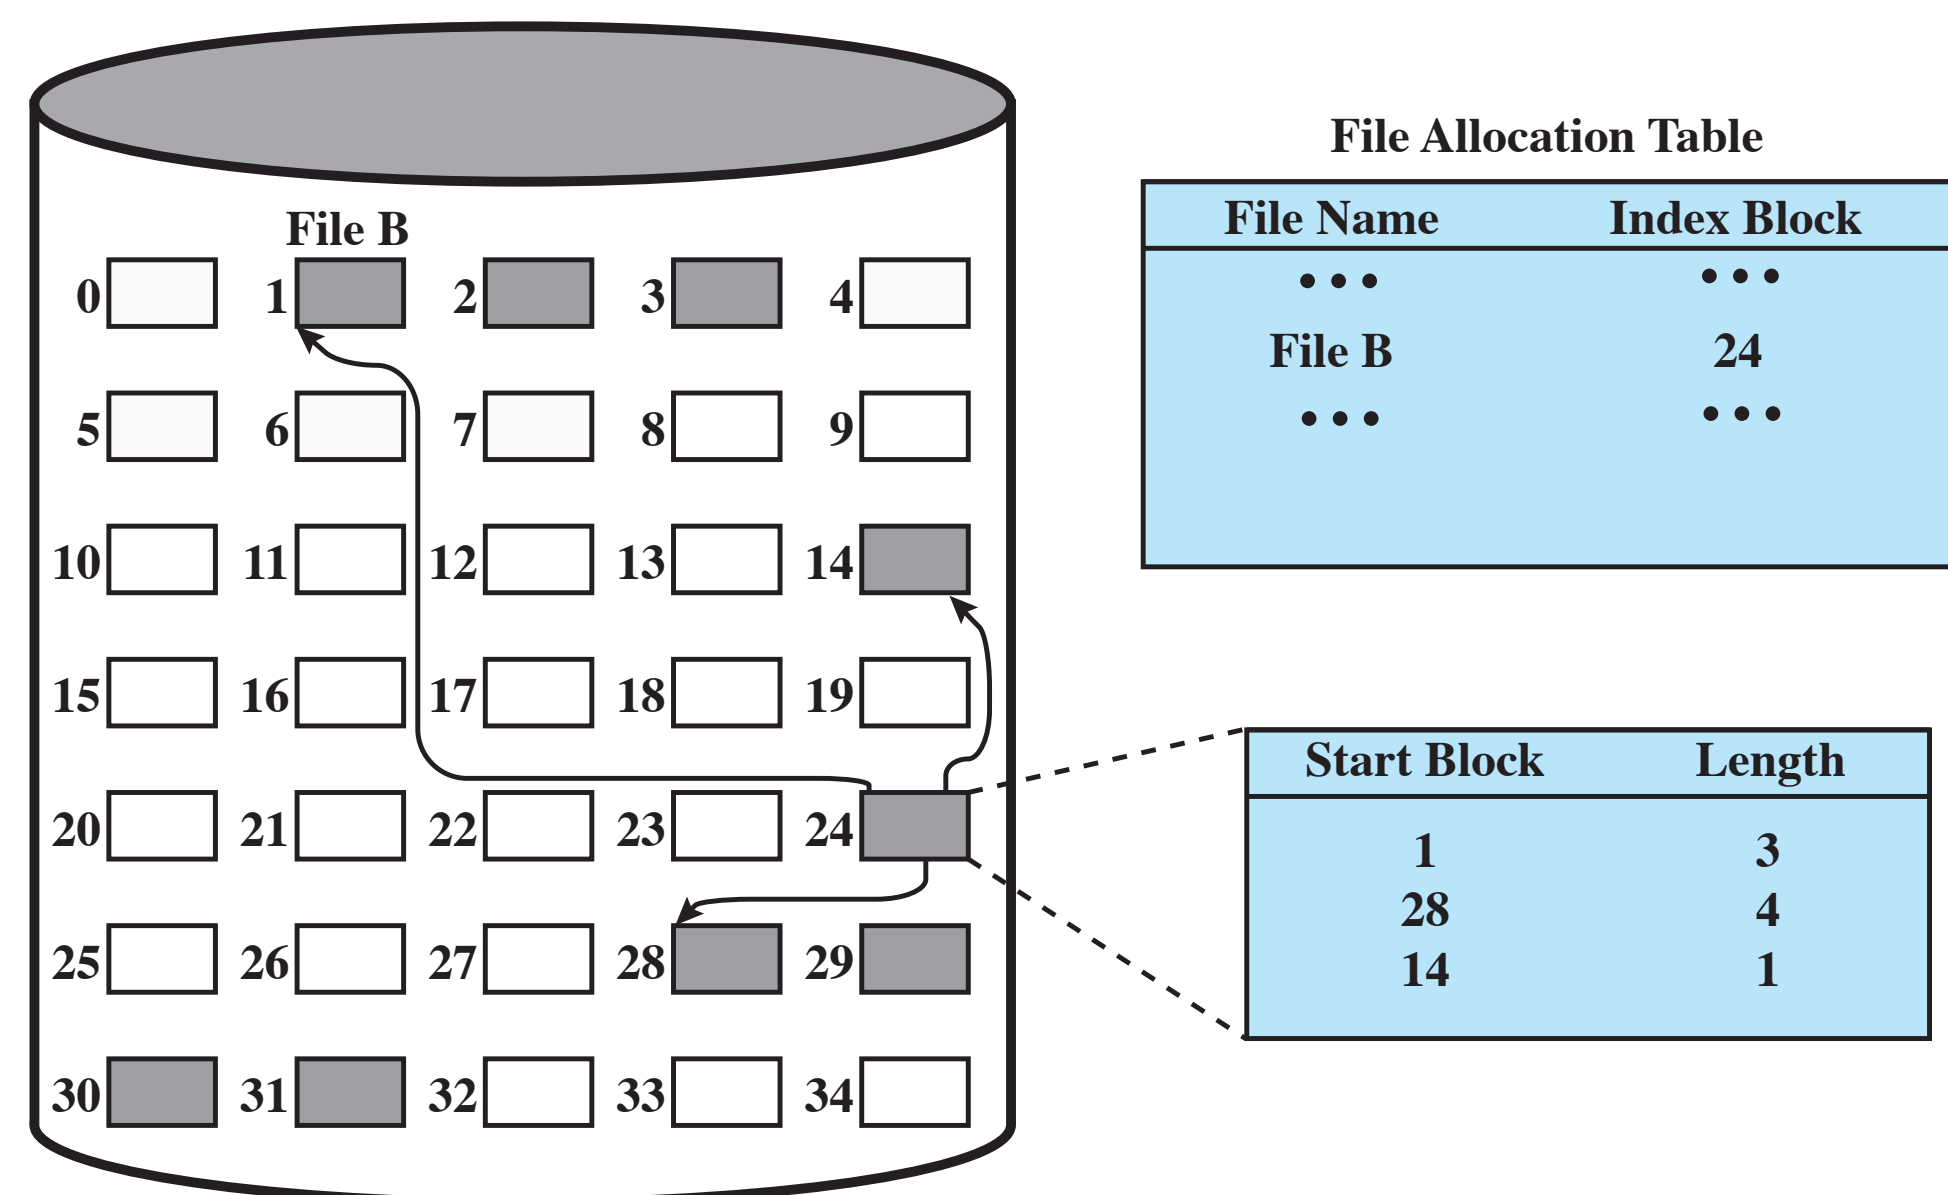
Allocate individual, fixed-size blocks until enough space is allocated for the file.

Maintain separate one-level index for each file in a separate block;

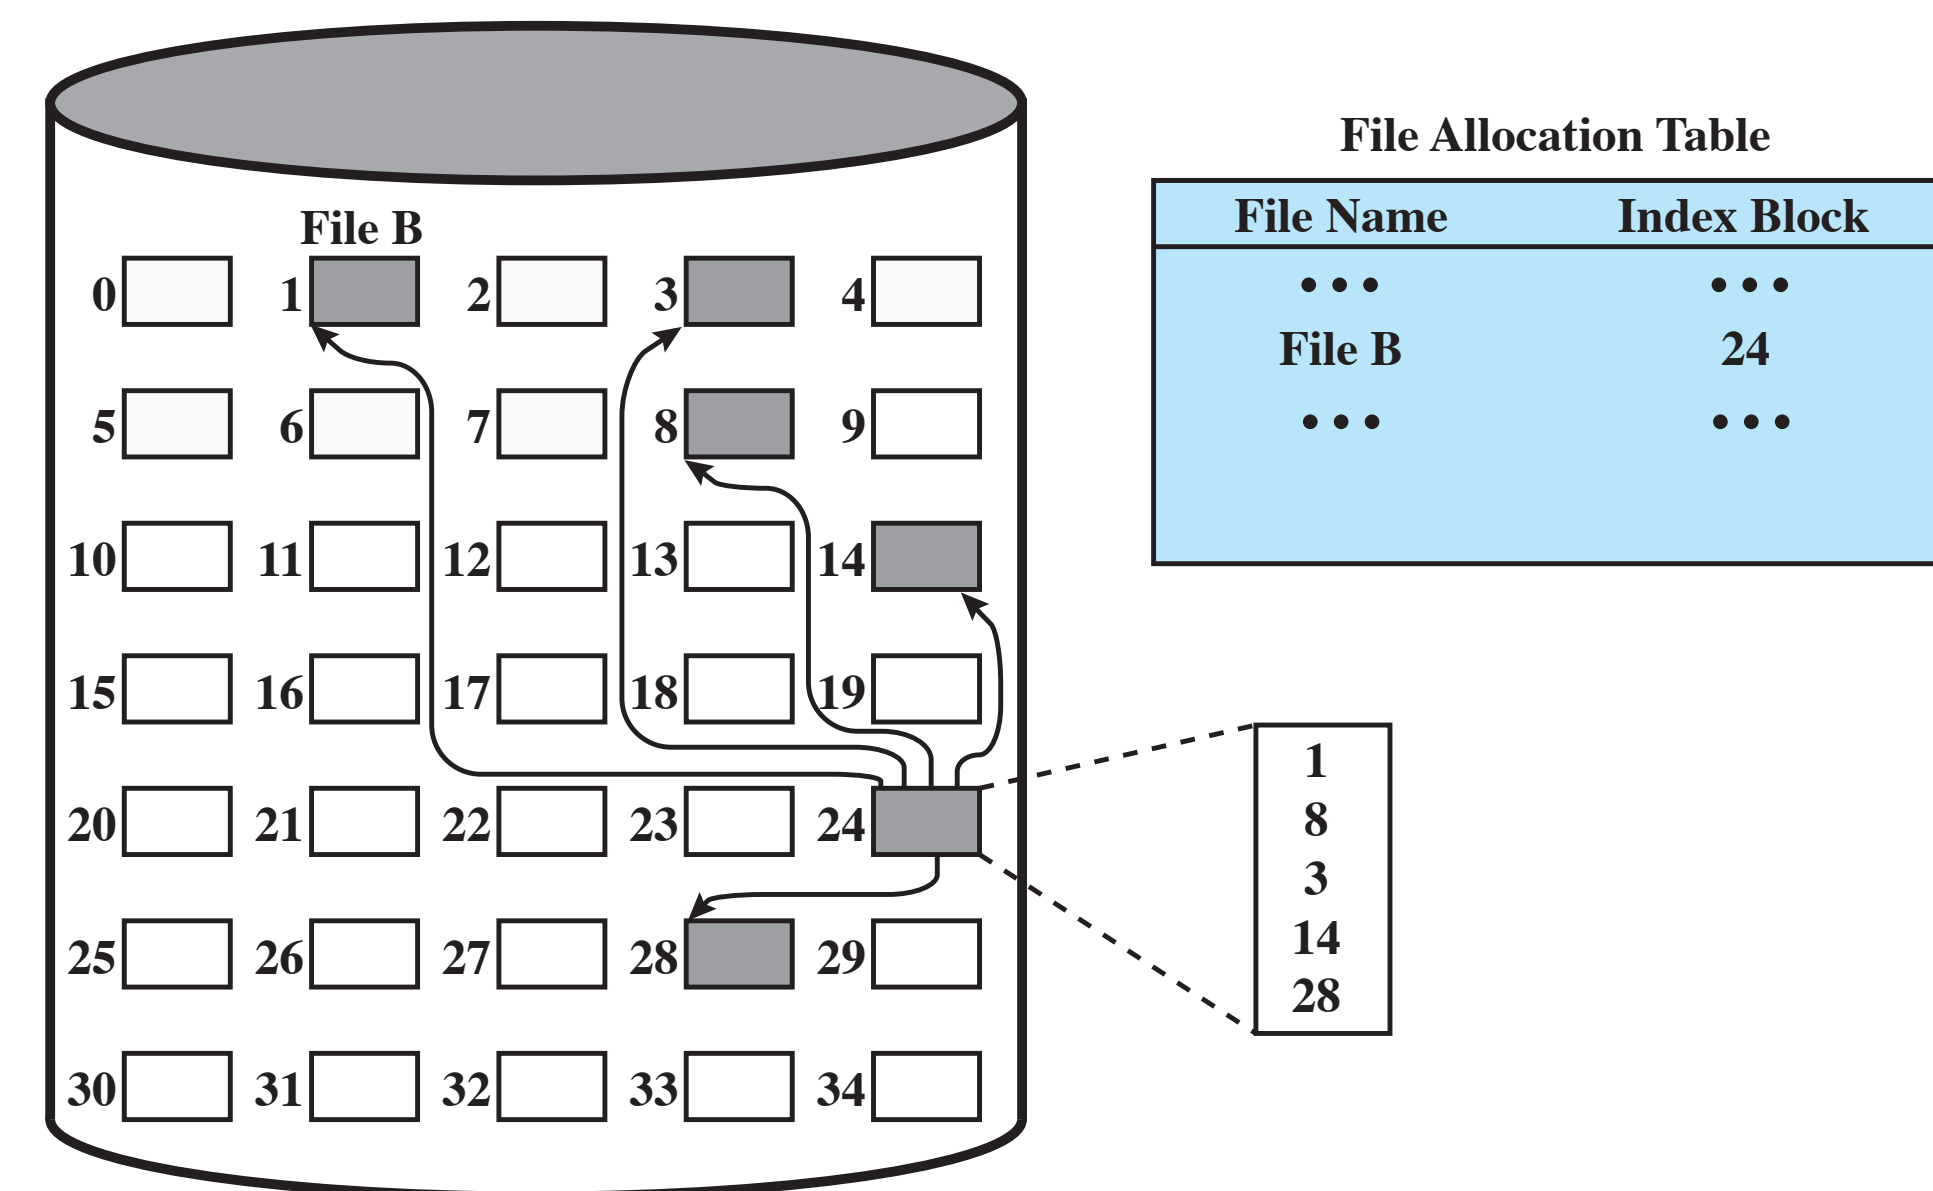
FAT points to block w/ the file's index;

Index has one entry per portion allocated to the file;

Addresses most problems from CFA and CA; still need to periodically run consultation.



IA / Block Portions



IA w/ Variable-Length Portions

Extras

Things alluded to in class but not fully covered

Example: Free Space Management

- We have to manage all of the free space as well...
→ *Disk Allocation Table (DAT)*

- Bit Tables

- Maintain vector of bits; 1 bit per block (0=free, 1=used)
- Easy to find a free block
- Easy to find a contiguous sequence of free blocks
- Compact representation for free space management

...001111000001110000111111111110110001....

$$\frac{\text{disk size in bytes}}{8 \times \text{file system block size}} = \frac{16\text{GB}}{8 \times 512} = 4\text{MB bit table}$$

With disk sizes today, complete bit tables can be large...

→ **create a summary table—break into subranges;
calculate # free blocks in subrange, max # of contiguous blocks, etc.**

See text for more examples of other strategies and their trade-offs...

Volumes — *A Logical Disk*

- A volume is nothing more than a logical representation for a physical disk. Indeed...
 - a volume can be a subset of a disk,
 - or represent a combination of disks,
 - Or even a subset/combination of another volume(s)
- Analogous to the relationship between physical memory & virtual memory