

# **Scheduling in Linux : CFS and it's implementation**

Fall 2019  
Montana State University

Ellen Marie Andersen  
Benjamin Cathelineau  
Olexandr Matveyev

<b>1 Brief history of CFS</b>	<b>3</b>
1.1 Reasons to change scheduler	3
1.2 Previous scheduler: O(1)	3
1.2.1 Shortcomings of the previous scheduler	4
<b>2 How CFS works : General presentation</b>	<b>4</b>
2.1 Fair share Scheduling and what is it?	4
2.2 The virtual runtime	5
2.3 Implementation of priorities	6
<b>3 Analysis of source code</b>	<b>7</b>
3.1 Data structure of CFS	7
3.2 Algorithm of CFS	7
<b>4 The user side : What you can do to configure your scheduler in a Linux system :</b>	<b>7</b>
4.1 Niceness	7
4.1.1 nice	8
4.1.2 renice	8
4.1.3 htop	8
4.2 Real time attribute	9
4.2.1 chrt command	9
4.2.2 Shell script with chrt command, to find which policies are the most common.	10
<b>5 Alternative to CFS : Other operating systems</b>	<b>11</b>
5.1 Windows	11
5.2 IOS	11
5.3 MACOS	11
5.4 Android	11
<b>6 Reference and source used</b>	<b>11</b>
6.1 Priority customisation in Linux	12
6.2 Source code of schedulers	12
6.3 Documentation of scheduler	12
6.4 Explanation of MACOS scheduler	12
6.5 Example of fair scheduling	12
6.6 vRuntime t increase, I/O bound vs CPU bound processes	12
6.7 group scheduling	13

# 1 Brief history of CFS

The Completely Fair Scheduler was introduced in 2007, and it was merged into kernel-2.6.23 [6]. First shift from  $O(1)$  happened by Con Kolivas' development, with his Rotating Staircase Deadline Scheduler (RSDL). Later, Ingo Molnar developed CFS scheduler, and he incorporated some of the ideas from Con Kolivas' RSDL scheduler. Molnar had also developed the previously used  $O(1)$  scheduler [5].

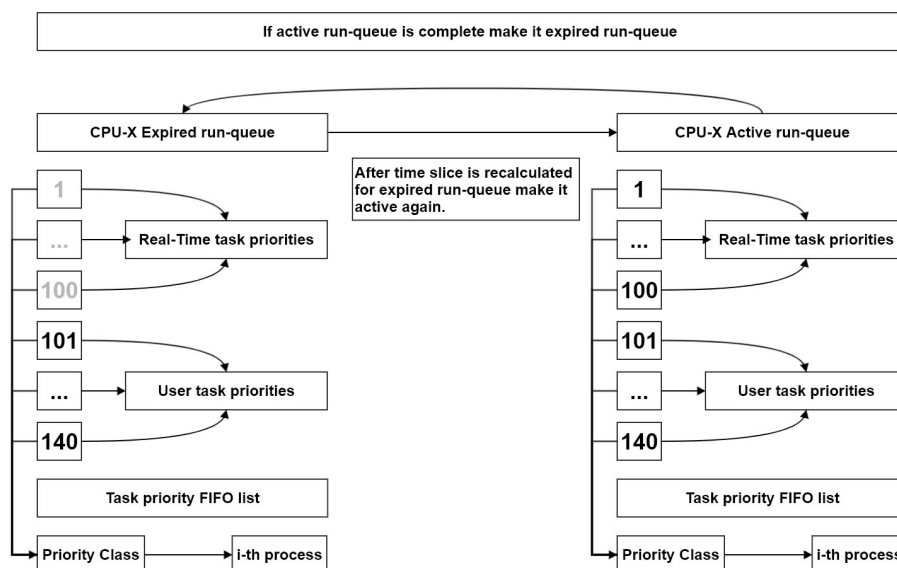
## 1.1 Reasons to change scheduler

It turned out that  $O(1)$  scheduler was very difficult to maintain due to the large amount of code and algorithm complexity ([3], pg: 462-463). The performance of interactivity with  $O(1)$  was poor, and it had low throughput for background jobs [2]. Interactive process requires short CPU bursts, and an example of interactive process is a text editor. The CFS provides and maintains balance by providing a fair amount of processor time to a task relative to other tasks [5].

## 1.2 Previous scheduler: $O(1)$

The  $O(1)$  scheduler is a constant time scheduler and it replaced  $O(n)$  scheduler. It was incorporated into Linux starting from kernel version 2.6.0 and till 2.6.22 version. In 2007 it was replaced by the Completely Fair Scheduler [1].

For the  $O(1)$  scheduling algorithm, in order to achieve constant scheduling time it must use active and expired arrays of processes. In this algorithm, each process gets a fixed time quantum. After time quantum expires for a task in the active array it will be placed into expired array, and array switch will take place. Due to that fact that arrays are accessed only via pointer, it makes array switching very fast  $O(1)$ . When algorithm performs task switching it makes the active array as the new empty expired array, while the expired array becomes the active array [1].



Real-Time Priority tasks from 1 to 100 and Normal Priority tasks from 101 to 140 [4]

### 1.2.1 Shortcomings of the previous scheduler

The design goals of the Linux kernel scheduler version 2.6 O(1) is to achieve fairness and boost interactivity performance, but because O(1) did try to achieve fairness that actually resulted decreasing interactive performance [12]. The main reason to replace O(1) scheduler was the large mass of code needed to calculate heuristics, to mark a task as interactive or non-interactive which is making it fundamentally difficult to increase interactivity performance [5].

## 2 How CFS works : General presentation

### 2.1 Fair share Scheduling and what is it?

Let's say we have 4 processes:

Process	Running time
Process A	3ms
Process B	2ms
Process C	4ms
Process D	3ms

How does Fair share scheduling work for these processes?

On the CPU, if only these four processes are to be executed and are ready, the processor will give each process a fair amount of CPU time. This means that the processes will alternate running on the CPU in fair time quantas while none of them are finished. This is calculated by N processes and  $(100/N)\%$  of CPU time. At time 0, N is 4 and each process will get a fair share of the processor. This is shown in example figure below:

	2ms quanta, N = 4	Remaining running time	2ms quanta, N =4	Remaining running time	2ms quanta, N = 4	Remaining running time
Process A	0.5ms	2.5ms	0.5ms	2ms	0.5ms	1.5ms
Process B	0.5ms	1.5ms	0.5ms	1ms	0.5ms	0.5ms
Process C	0.5ms	3.5ms	0.5ms	3ms	0.5ms	2.5ms

## DRAFT OF OUR TECHNICAL REPORT

Process D	0.5ms	2.5ms	0.5ms	2ms	0.5ms	1.5ms
-----------	-------	-------	-------	-----	-------	-------

continuing below:

	2ms quanta, N = 4	Remaining running time	2ms quanta, N = 3	Remaining running time	2ms quanta, N = 3	Remaining running time
Process A	0.5ms	1ms	0.6667ms	0.3333ms	0.6667ms	0ms
Process B	0.5ms	0ms	n/a	n/a	n/a	n/a
Process C	0.5ms	3ms	0.6667ms	2.3333ms	0.6667ms	1.6667ms
Process D	0.5ms	1ms	0.6667ms	0.3333ms	0.6667ms	0ms

continuing below:

	2ms quanta, N = 1	Remaining running time	2ms quanta, N = 1	Remaining running time
Process A	n/a	n/a	n/a	n/a
Process B	n/a	n/a	n/a	n/a
Process C	n/a	n/a	n/a	n/a
Process D	2ms	2ms	2ms	0ms

How does this differ if one of the processes have higher priority than another? Or if a process spawns many processes related to itself to capture more CPU resources? This is described in more detail below with virtual runtime and grouping scheduling [8].

## 2.2 The virtual runtime

The method of the scheduler is simple, CFS tries to run the process with the lowest virtual runtime. Virtual runtime is calculated by  $t$  (time spent running) multiplied with the weight of the process' niceness/priority (explained below). As a process runs on the processor,  $t$  increases accordingly.

The virtual runtime is simply incremented differently according to the priority of the process. Processes with low priority number get their virtual runtime increased **less**. Process with high priority number gets their virtual runtime increased **more**.

**I/O bound processes vs Processor bound processes:**

## DRAFT OF OUR TECHNICAL REPORT

Niceness/priorities set aside, due to I/O processes typically executing at a short time interval each time before being blocked, these processes' runtime will not increase a lot due to the short time bursts spent executing. If another process executes for double the time, the calculation will therefore be affected accordingly, and the process' place in the red black tree may likely be placed further to the right side of the tree. The further to the right side of the tree a process is placed, the less likely it is to get run next. The process on the leftmost and furthestmost down side of the tree is the process with least runtime and is picked to schedule next (see section 3.1 below). Starvation is combated in this model due to the 't' increase for each process, which ensures all processes are chosen to run. **how does this depend on niceness/priorities.**

In order for a process not to capture too much of the CPU by spawning new processes/tasks belonging to that process, Linux CFS can group certain processes/tasks that are related to each other and treat every group of tasks as one process. This is called group scheduling [9],[3], pg 464).

One can enable and disabling the grouping scheduling features in the scheduling algorithm in one's settings. Below shows that when 0 is returned, autogrouping is disabled and when 1 returns, autogrouping is enabled.

Here is an example of the autogrouping (grouping scheduling) disabled and enabled.

0

```
/proc/sys/kernel/sched autogroup enabled (END)
```

1

```
/proc/sys/kernel/sched autogroup enabled (END)
```

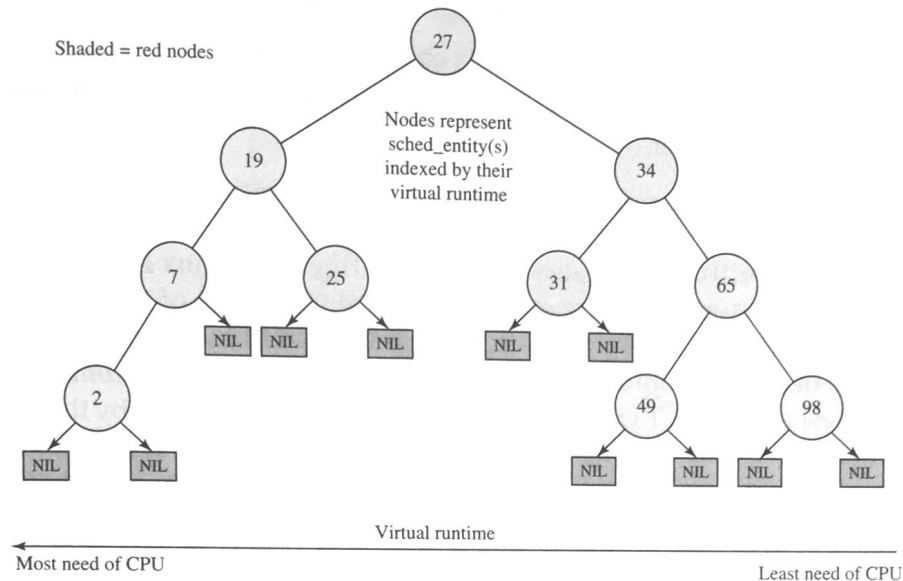
## 2.3 Implementation of priorities

Priorities range from 0 to 139. The lower the number the higher the priority.

The user can customize priorities as explained [later in this report](#)

## 3 Analysis of source code

### 3.1 Data structure of CFS



**Figure 10.11 Example of Red Black Tree for CFS**

RB Tree ([3]: pg 463)

The leftmost task is always picked by the scheduler. The RBTree is ordered by vruntime.

### 3.2 Algorithm of CFS

Based on the CFS documentation we can see that CFS has about 10000 lines of code [10].

## 4 The user side : What you can do to configure your scheduler in a Linux system :

### 4.1 Niceness

On Linux, the user can customize the priority of a process, to a certain degree. This customisation is made through the use of niceness. The niceness of a process can range from -20 to 19. The lower the niceness, the higher the priority.

Only positive (lower) priority number are accessible to non root user. To set a negative priority you need to be root. However, like many things in Linux, this can be changed in the `/etc/security/limits.conf` file, to allow non privileged user to set negative niceness.[\[7\]](#)

## DRAFT OF OUR TECHNICAL REPORT

Still, it is important to keep in mind that these priority adjustment only concern non-real time process, as explained [earlier in this report](#)

### 4.1.1 nice

Nice is the standard utility to start a program with a different niceness than the default one (0).

```
benjamin@debian-benjamin-laptop:~$ nice -n 10 ls
Android  Dev      Downloads  mbox    opt      Public  Templates virt-manager
Desktop  Documents  github    Music   Pictures Steam   Videos
benjamin@debian-benjamin-laptop:~$
```

*Example usage of nice*

### 4.1.2 renice

Renice is the utility that allows to change the niceness of an already started process.

```
benjamin@debian-benjamin-laptop: ~
File Edit View Search Terminal Help
benjamin@debian-benjamin-laptop:~$ ping debian.org -i 20 &
[1] 16010
benjamin@debian-benjamin-laptop:~$ PING debian.org (130.89.148.77) 56(84) bytes
of data.
64 bytes from klecker-misc.debian.org (130.89.148.77): icmp_seq=1 ttl=49 time=16
9 ms
64 bytes from klecker-misc.debian.org (130.89.148.77): icmp_seq=2 ttl=49 time=14
5 ms

benjamin@debian-benjamin-laptop:~$ man renice
benjamin@debian-benjamin-laptop:~$ sudo renice -n 19 16010
16010 (process ID) old priority 0, new priority 19
benjamin@debian-benjamin-laptop:~$
```

There is also a ionice program but it is outside the scope of this report because it's related to IO.

### 4.1.3 htop

HTOP is a general purpose command line task manager, inspired by top but better.



## DRAFT OF OUR TECHNICAL REPORT

```
1  [|||||] 17.1%] Tasks: 155, 830 thr; 2 running
2  [||||] 9.3%] Load average: 1.10 0.94 0.93
3  [|||||] 17.2%] Uptime: 1 day, 21:42:21
4  [||||] 8.6%]
Mem[|||||] 6.36G/7.69G]
Swp[||] 361M/7.46G]
```

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
2766	benjamin	20	0	1263M	170M	27256	S	0.0	2.2	0:00.00	/usr/bin/gnome-software -
2803	benjamin	20	0	4150M	849M	348M	S	4.6	10.8	1h02:36	/usr/lib/firefox-esr/fire
2807	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:00.09	/usr/lib/firefox-esr/fire
2808	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:00.05	/usr/lib/firefox-esr/fire
2810	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	5:24.94	/usr/lib/firefox-esr/fire
2811	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:01.29	/usr/lib/firefox-esr/fire
2812	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:08.52	/usr/lib/firefox-esr/fire
2813	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:08.99	/usr/lib/firefox-esr/fire
2814	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:08.61	/usr/lib/firefox-esr/fire
2815	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:08.98	/usr/lib/firefox-esr/fire
2816	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:20.52	/usr/lib/firefox-esr/fire
2817	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:00.01	/usr/lib/firefox-esr/fire
2818	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	1:20.13	/usr/lib/firefox-esr/fire
2819	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:10.09	/usr/lib/firefox-esr/fire
2820	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:04.33	/usr/lib/firefox-esr/fire
2821	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:35.17	/usr/lib/firefox-esr/fire
2822	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:22.26	/usr/lib/firefox-esr/fire
2826	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:30.36	/usr/lib/firefox-esr/fire
2827	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:00.00	/usr/lib/firefox-esr/fire
2830	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:00.00	/usr/lib/firefox-esr/fire
2831	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:00.00	/usr/lib/firefox-esr/fire
2832	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	3:07.13	/usr/lib/firefox-esr/fire
2833	benjamin	20	0	4150M	849M	348M	S	0.7	10.8	19:01.11	/usr/lib/firefox-esr/fire
2834	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:00.13	/usr/lib/firefox-esr/fire
2835	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:00.07	/usr/lib/firefox-esr/fire
2840	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:00.00	/usr/lib/firefox-esr/fire
2841	benjamin	20	0	4150M	849M	348M	S	0.7	10.8	3:09.27	/usr/lib/firefox-esr/fire
2842	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:07.23	/usr/lib/firefox-esr/fire
2850	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:01.58	/usr/lib/firefox-esr/fire
2851	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:01.63	/usr/lib/firefox-esr/fire
2852	benjamin	20	0	4150M	849M	348M	S	0.0	10.8	0:01.61	/usr/lib/firefox-esr/fire

```
1Help F2Setup F3Search F4Filter F5Tree F6SortBy F7Nice F8Nice +F9Kill F10Quit
```

The priority can be increased (f7 decrease niceness) or decreased (f8 increase niceness). This tool is easier to use than renice because it displays the name of the program.

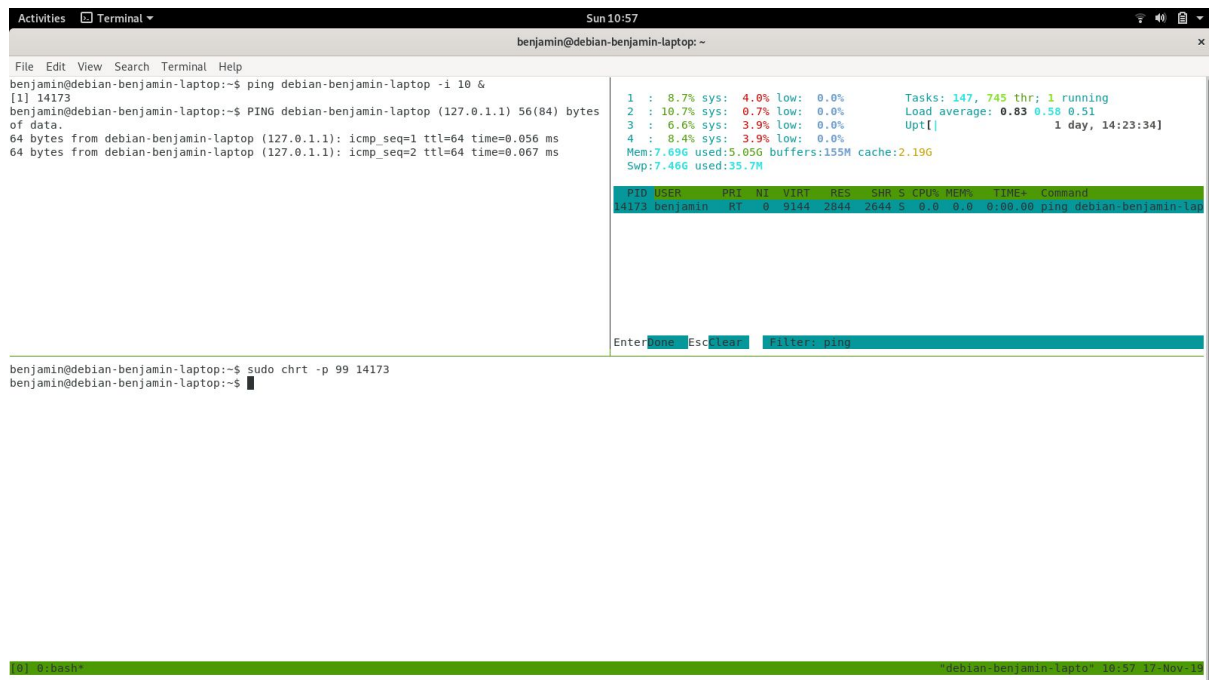
## 4.2 Real time attribute

### 4.2.1 chrt command

The default scheduling policies for processes is SCHED\_OTHER, or SCHED\_NORMAL (same thing). Processes under this policies are scheduled with CFS, and they use niceness. However they are other scheduling policies, such as SCHED\_RR and SCHED\_FIFO. Theses policies don't use CFS and niceness. Instead they use static priorities. [\[14\]](#) With chrt you can change the policy of your process, and it's static priority. Here is an example usage :

```
sudo chrt -p 99 10540
```

## DRAFT OF OUR TECHNICAL REPORT



The screenshot shows a terminal window titled "benjamin@debian-benjamin-laptop: ~". The terminal output includes a ping command and system statistics. The ping command is: `benjamin@debian-benjamin-laptop:~$ ping debian-benjamin-laptop -i 10 &`. The output shows: `[1] 14173`, `benjamin@debian-benjamin-laptop:~$ PING debian-benjamin-laptop (127.0.1.1) 56(84) bytes of data.`, `64 bytes from debian-benjamin-laptop (127.0.1.1): icmp_seq=1 ttl=64 time=0.056 ms`, and `64 bytes from debian-benjamin-laptop (127.0.1.1): icmp_seq=2 ttl=64 time=0.067 ms`. The system statistics show: `1 : 8.7% sys: 4.0% low: 0.0% Tasks: 147, 745 thr: 1 running`, `2 : 10.7% sys: 0.7% low: 0.0% Load average: 0.83 0.58 0.51`, `3 : 6.6% sys: 3.9% low: 0.0% Upt[ 1 day, 14:23:34]`, `4 : 8.4% sys: 3.9% low: 0.0%`, `Mem: 7.69G used: 5.05G buffers: 155M cache: 2.19G`, and `Swp: 7.46G used: 35.7M`. A table of processes is also shown with columns: `PID USER PRI NI VIRT RES SHR S CPU% MEM% TIME+ Command`. The first row is: `14173 benjamin RT 0 9144 2844 2644 S 0.0 0.0 0:00.00 ping debian-benjamin-lap`. The terminal prompt is `benjamin@debian-benjamin-laptop:~$ sudo chrt -p 99 14173` and `benjamin@debian-benjamin-laptop:~$`.

Here I did a ping and I changed it's static priority to 99( highest possible priority).

4.2.2 Shell script with chrt command, to find which policies are the most common.

This shell script is heavily inspired and slightly modified, from[\[13\]](#)

```
#!/bin/bash
ps -aux |grep [0-9]|awk '{print $2}' > test.txt
cat test.txt |while read line
do
chrt -p $line 2> /dev/null
done
```

```
benjamin@debian-benjamin-laptop:~/Documents/montana/
CSCI460/project$ ./shellscript |grep OTHER |wc
      218      1308      10833

benjamin@debian-benjamin-laptop:~/Documents/monta
CSCI460/project$ ./shellscript |grep FIFO |wc
       12        72        572

benjamin@debian-benjamin-laptop:~/Documents/monta
CSCI460/project$ ./shellscript |grep IDLE|wc
        1         6         49
```

We see that most processes are running as SCHED\_OTHER on my system. Some are running as SCHED\_FIFO. And only one is using SCHED\_IDLE, which is the lowest priority policy.

## 5 Alternative to CFS : Other operating systems

### 5.1 Windows

In the current version of windows multilevel level feedback queue are used. Each of theses are scheduled with round robin.

### 5.2 IOS

Just a brief description

### 5.3 MACOS

Just a brief description

### 5.4 Android

Probably identical to the Linux one

## 6 Reference and source used

- [1]: [https://en.wikipedia.org/wiki/O\(1\)\\_scheduler](https://en.wikipedia.org/wiki/O(1)_scheduler)
- [2]: J. Jose, O. Sujisha, M. Gilesh and T. Bindima, "On the Fairness of Linux O(1) Scheduler," *2014 5th International Conference on Intelligent Systems, Modelling and Simulation*, Langkawi, 2014, pp. 668-674. doi: 10.1109/ISMS.2014.120  
URL: <https://ieeexplore.ieee.org/document/7280991>
- [3]: Stallings, W. (2018). *Operating systems: internals and design principles*. Harlow, Essex: Pearson.
- [4]: Jones, T. (2006, June 30). Inside the Linux scheduler. Retrieved November 10, 2019, from <https://www.ibm.com/developerworks/library/l-scheduler/index.html>.
- [5]: Jones, . (2009, December 15). Inside the Linux 2.6 Completely Fair Scheduler. November 10, 2019, Retrieved from <https://developer.ibm.com/tutorials/l-completely-fair-scheduler/>.
- [6]: [https://en.wikipedia.org/wiki/Completely\\_Fair\\_Scheduler](https://en.wikipedia.org/wiki/Completely_Fair_Scheduler)
- [7]: Silberschatz, A., Galvin, P. B., & Gagne, G. (2018). *Operating system concepts*. Hoboken, NJ: Wiley.
- [8]: Operating System #22 Completely Fair Scheduling (CFS). Retrieved from <https://www.youtube.com/watch?v=scfDOof9pww>.
- [9]: Retrieved from [https://www.eit.lth.se/fileadmin/eit/courses/eitf60/Rapporter/Ludwig\\_Hellgren\\_Winblad\\_4547\\_assignsubmission\\_file\\_ludwig.hellgren.winblad.pdf](https://www.eit.lth.se/fileadmin/eit/courses/eitf60/Rapporter/Ludwig_Hellgren_Winblad_4547_assignsubmission_file_ludwig.hellgren.winblad.pdf)
- [10]: <https://www.kernel.org/doc/Documentation/scheduler/sched-design-CFS.txt>
- [11]: Priority customisation in Linux. Retrieved from [https://en.wikipedia.org/wiki/Nice\\_\(Unix\)](https://en.wikipedia.org/wiki/Nice_(Unix)) and <https://askubuntu.com/questions/656771/process-niceness-vs-priority>

[12]: Wong, C., Tan, I., Kumari, R., Lam, J., & Fun, W. (2008). Fairness and interactive performance of O(1) and CFS Linux kernel schedulers. *2008 International Symposium on Information Technology*. doi: 10.1109/itsim.2008.4631872

[13]: <https://unix.stackexchange.com/questions/407497/how-to-find-scheduling-policy-and-active-processes-priority>

[14]: <http://man7.org/linux/man-pages/man7/sched.7.html>

## 6.1 Priority customisation in Linux

[https://en.wikipedia.org/wiki/Nice\\_\(Unix\)](https://en.wikipedia.org/wiki/Nice_(Unix))

<https://askubuntu.com/questions/656771/process-niceness-vs-priority>

## 6.2 Source code of schedulers

Current linux version : CA10000 lines of code

<https://elixir.bootlin.com/linux/latest/source/kernel/sched/fair.c>

First version introduced : 1200 lines of code

[https://elixir.bootlin.com/linux/v2.6.23/source/kernel/sched\\_fair.c](https://elixir.bootlin.com/linux/v2.6.23/source/kernel/sched_fair.c)

O(1) scheduler : 7000 lines of code

<https://elixir.bootlin.com/linux/v2.6.22.9/source/kernel/sched.c>

## 6.3 Documentation of scheduler

<https://www.kernel.org/doc/html/latest/scheduler/index.html>

## 6.4 Explanation of MACOS scheduler

<https://developer.apple.com/library/archive/documentation/Darwin/Conceptual/KernelProgramming/scheduler/scheduler.html>

## 6.5 Example of fair scheduling

<https://www.youtube.com/watch?v=scfDOof9pww>

## 6.6 vRuntime t increase, I/O bound vs CPU bound processes

<https://www.youtube.com/watch?v=scfDOof9pww>

<https://www.youtube.com/watch?v=scfDOof9pww>

## 6.7 group scheduling

<http://man7.org/linux/man-pages/man7/sched.7.html>

<http://www.informit.com/articles/article.aspx?p=397655&seqNum=6>

### 7.8 other sources:

<http://www.informit.com/articles/printerfriendly/101760>

<https://trepo.tuni.fi/bitstream/handle/10024/96864/GRADU-1428493916.pdf>

<https://developer.ibm.com/tutorials/l-completely-fair-scheduler/#fig2>

MAYBE WILL WORK ON THIS

Alternativ data structure that can be used for processor schedulers:

B-Trees, Suffix Trees, Suffix Arrays, FM-Index

By representing a process via character we can build a sequence of characters, and we can apply Suffix Trees, Suffix Arrays, and FM-Index to manage processes.