

Process Management

Jakob Frank

Dillon Tice

Linux

- Processes are managed in memory as pointers to structs known as `task_structs`
- Each struct contains a large amount of information for each process, such as process id, state, name of executable, and links to other processes such as children, parents, and siblings
- `fork()` and `exec()` are primary used for the creation of new processes

Structure

- task_struct
 - Differs by architecture
- Init
 - Statically allocated task_struct
 - /init/init_task.c

```
56 struct task_struct init_task
57 #ifdef CONFIG_ARCH_TASK_STRUCT_ON_STACK
58     __init_task_data
59 #endif
60 = {
61     #ifdef CONFIG_THREAD_INFO_IN_TASK
62     .thread_info = INIT_THREAD_INFO(init_task),
63     .stack_refcount = REFCOUNT_INIT(1),
64     #endif
65     .state = 0,
66     .stack = init_stack,
67     .usage = REFCOUNT_INIT(2),
68     .flags = PF_KTHREAD,
69     .prio = MAX_PRIO - 20,
70     .static_prio = MAX_PRIO - 20,
71     .normal_prio = MAX_PRIO - 20,
72     .policy = SCHED_NORMAL,
73     .cpus_ptr = &init_task.cpus_mask,
74     .cpus_mask = CPU_MASK_ALL,
75     .nr_cpus_allowed = NR_CPUS,
76     .mm = NULL,
77     .active_mm = &init_mm,
78     .restart_block = {
79         .fn = do_no_restart_syscall,
80     },
81     .se = {
82         .group_node = LIST_HEAD_INIT(init_task.se.group_node),
83     },
84     .rt = {
85         .run_list = LIST_HEAD_INIT(init_task.rt.run_list),
86         .time_slice = RR_TIMESLICE,
87     },
88     .tasks = LIST_HEAD_INIT(init_task.tasks),
```

task_struct fields

- state
 - A set of bits indicating process state
 - Running, stopped, interruptible, uninterruptible
- flags
 - Is a process being created? Exiting? Allocating?
- comm
 - Name of the executable without the path
- static_prio
 - Priority
 - Lower priority > higher priority
 - *Actual* priority determined dynamically

```
629 struct task_struct {
630     #ifdef CONFIG_THREAD_INFO_IN_TASK
631         /*
632          * For reasons of header soup (see current_thread_info()), this
633          * must be the first element of task_struct.
634          */
635         struct thread_info      thread_info;
636     #endif
637     /* -1 unrunnable, 0 runnable, >0 stopped: */
638     volatile long              state;
639
640     /*
641      * This begins the randomizable portion of task_struct. Only
642      * scheduling-critical items should be added above here.
643      */
644     randomized_struct_fields_start
645
646     void                        *stack;
647     refcount_t                  usage;
648     /* Per task flags (PF_*), defined further below: */
649     unsigned int                flags;
650     unsigned int                ptrace;
651
652     #ifdef CONFIG_SMP
653         struct llist_node        wake_entry;
654         int                      on_cpu;
655     #endif CONFIG_THREAD_INFO_IN_TASK
656     /* Current CPU: */
657     unsigned int                cpu;
658     #endif
659     unsigned int                wakee_flips;
660     unsigned long               wakee_flip_decay_ts;
661     struct task_struct          *last_wakee;
662 }
```

task_struct fields

- mm and active_mm
 - Process memory descriptors
 - Active is the previous process descriptors
 - Context switching
- tasks
 - Used for linked list representation
 - Init is the head
 - Prev and next pointer
- thread_struct
 - Context switch storage
 - Registers, program counter, etc.

```
#ifndef CONFIG_SMP
    struct plist_node    pushable_tasks;
    struct rb_node        pushable_dl_tasks;
#endif

    struct mm_struct      *mm;
    struct mm_struct      *active_mm;

    /* Per-thread vma caching: */
    struct vmacache        vmacache;

#ifdef SPLIT_RSS_COUNTING
    struct task_rss_stat    rss_stat;
#endif
```

```
1276    /*
1277     * New fields for task_struct should be added above here, so that
1278     * they are included in the randomized portion of task_struct.
1279     */
1280    randomized_struct_fields_end
1281
1282    /* CPU-specific state of this task: */
1283    struct thread_struct    thread;
1284
1285    /*
1286     * WARNING: on x86, 'thread_struct' contains a variable-sized
1287     * structure. It *MUST* be at the end of 'task_struct'.
1288     *
1289     * Do not put anything below here!
1290     */
```

Overview of Linux Process Management

- With regards to creating and managing process, linux contains the following commands:
 - Commands for creating processes via `fork()` and `exec()`
 - Can also use `System()` but it is slow and a security risk
 - Commands for killing processes such as `kill`, `exit`, `pkill`, and `killall`
 - Commands for viewing process information such as `pgrep()`, `ps()`, and `top()`

Process States

- Possible process states:
 - `TASK_RUNNING`
 - Running or in run queue
 - `TASK_INTERRUPTIBLE`
 - Sleeping
 - `TASK_UNINTERRUPTIBLE`
 - Sleeping but unable to be awoken
 - `TASK_STOPPED`
 - Task becomes a zombie (stopped but still contains an entry in the process table)

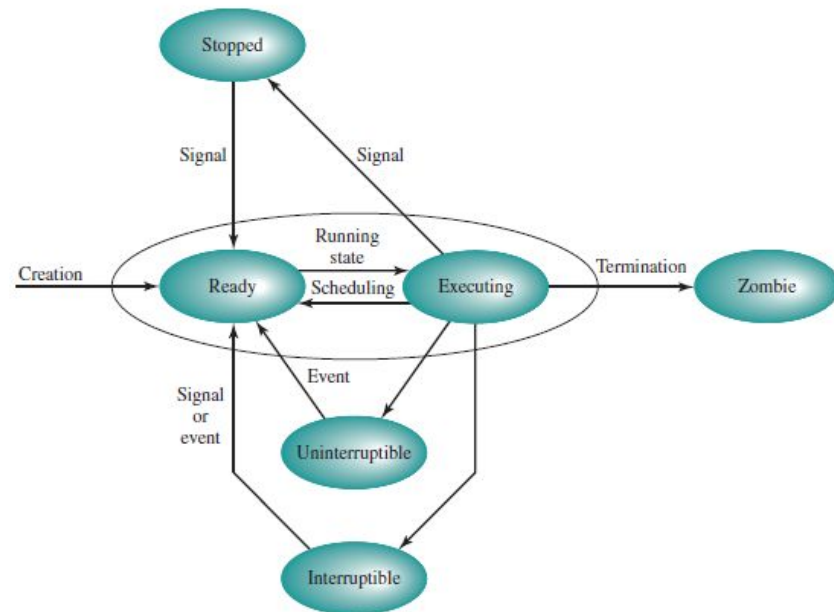



Figure 4.15 Linux Process/Thread Model

Creation

- fork() -> clone() -> do_fork()
- do_fork() - clone_flags
 - copy_process(..., args)
- copy_process()
 - Security_task_create
 - Further SELinux functions
 - Return value == 0 ?
 - dup_task_struct()
 - copy_creds()
 - CPU time reset
 - CPU time sharing setup

```
1 long _do_fork(struct kernel_clone_args *args)
2 {
3     u64 clone_flags = args->flags;
4     struct completion vfork;
5     struct pid *pid;
6     struct task_struct *p;
7     int trace = 0;
8     long nr;
9
10    /*
11     * Determine whether and which event to report to ptracer. When
12     * called from kernel_thread or CLONE_UNTRACED is explicitly
13     * requested, no event is reported; otherwise, report if the event
14     * for the type of forking is enabled.
15     */
16    if (!(clone_flags & CLONE_UNTRACED)) {
17        if (clone_flags & CLONE_VFORK)
18            trace = PTRACE_EVENT_VFORK;
19        else if (args->exit_signal != SIGCHLD)
20            trace = PTRACE_EVENT_CLONE;
21        else
22            trace = PTRACE_EVENT_FORK;
23
24        if (likely(!ptrace_event_enabled(current, trace)))
25            trace = 0;
26    }
27
28    p = copy_process(NULL, trace, NUMA_NO_NODE, args);
29    add_latent_entropy();
30
31    if (IS_ERR(p))
32        return PTR_ERR(p);
33}
```

 _do_fork

Destruction

POSIX defines two ways regarding how a process can be terminated

- A process can terminate itself, either by calling `exit()`, `_exit()`, returning from `main()`, or the last thread of the process terminates
- A process can be killed by a signal, possibly sent by the kernel, another process, or the process itself

Alternative Ideas/Strategies

- Different built in Linux scheduling policies
 - SCHED_FIFO
 - SCHED_RR
 - SCHED_OTHER
- Different task manager implementations
 - htop
 - conky
 - pstree
 - GNOME System Monitor