

# Spectre and Meltdown

---

By Max Weimer

# Overview

- Two distinct attacks
- Spectre
  - User ↔ User memory breakdown
  - Sandbox breakout
  - Less serious, harder to mitigate
- Meltdown
  - User ↔ Kernel memory breakdown
  - Access to kernel memory
  - More serious, easy to mitigate

# Severity

- BAD

## Vulnerability Note VU#584653

### CPU hardware vulnerable to side-channel attacks

Original Release date: 03 Jan 2018 | Last revised: 03 Jan 2018



#### Overview

CPU hardware implementations are vulnerable to side-channel attacks. These vulnerabilities are referred to as [Meltdown](#) and [Spectre](#).

#### Description

CPU hardware implementations are vulnerable to side-channel attacks referred to as [Meltdown](#) and [Spectre](#) (also KAISER and KPTI). These attacks are described in detail by [Google Project Zero](#) and the Institute of Applied Information Processing and Communications (IAIK) at Graz University of Technology (TU Graz).

#### Impact

An attacker able to execute code with user privileges can achieve various impacts, such as reading otherwise protected kernel memory and bypassing KASLR.

#### Solution

Replace CPU hardware

The underlying vulnerability is primarily caused by CPU architecture design choices. Fully removing the vulnerability requires replacing vulnerable CPU hardware.

# Solution

# Replace CPU hardware


# Background: CPU Pipelining

How do we execute an instruction?

Fetch → Decode → Execute

What are we running? → How do we run it? → Run it.

Optimize?



# Spectre: Branch Prediction

```
int read(int i) {  
    int result = -1;  
    if (i < array_size)  
        result = array[i];  
    return result;  
}
```

- Resolving condition takes time
- Where is `array_size` in memory?
  - Uncached is beneficial
- Are we going to run the code in the branch?
  - Do we wait for the condition?
  - Do we guess?

Of course we guess.

# Spectre: Branch Prediction

```
int read(int i) {  
    int result = -1;  
    if (i < array_size)  
        result = array[i];  
    return result;  
}
```

- What's doing the predicting?
  - Branch predictor in hardware
- Evil mistraining
- What if we're wrong?
  - Rewind state
  - *Can't rewind cache*

Why does it matter?

We evaluate cache, read secrets

Break out of sandbox

# Spectre: Mitigations

- No fix-all solution
- Javascript engines have been updated
  - Chrome uses individual processes to bolster sandboxing
  - Webkit
    - Index masking to prevent the bounds from getting *too* out of hand
    - Pointers are XORed with a pseudo random “poison” value, poisoned pointers can’t be used to read arbitrary memory unless the poison is known

# Meltdown Background: $\mu$ OPs & OOE

- What does decode do?
- Breaks down a command like `add` into micro operations ( $\mu$ OPs)
- Ex: `add ax, bx` breaks down to the following  $\mu$ OPs roughly:
  1. Put contents of `ax` on bus, read bus into `ALU input a`
  2. Put contents of `bx` on bus, read bus into `ALU input b`
  3. Put `ALU output` on system bus, read bus into `acc`
- Independent operations may be executed in a different order than written
  - Exception handling gets muddy
  - Worse yet,  $\mu$ OPs may run out of order

KEY: Errors are handled at the **end** of the pipeline



# Meltdown Background: Out of Order Execution

```
1 ; rcx = kernel address, rbx = probe array
2 xor rax, rax
3 retry:
4   mov al, byte [rcx]           ; yes officer, this line right here
5   shl rax, 0xC
6   jz retry                    ; Still alive? Go again
7 mov rbx, qword [rbx + rax]    ; OOE: quick throw data somewhere!
```

- Meltdown causes an exception on purpose by attempting to read kernel memory
  - Kernel memory is read into register before exception is raised
  - With dying breath the program passes the register value out through a side-channel, run out of order
  - Rinse, repeat, dump kernel memory out at ~500KB/s

# Meltdown Mitigation: Kernel PTI

- For performance, modern OSes map kernel memory into user space
  - I/O calls, etc are faster
- Solution? Kernel Page Table Isolation (KPTI)
  - Don't map kernel memory into user space *on affected CPUs*
    - Speed decrease, but only really viable method
    - ~10% performance hit on Intel CPUs
      - Workload dependent

Fin